



Fundamentals of Istio

Logistics



- **Class Hours:**
- Start time is 8:30am
- End time is 3:30pm
- Class times may vary slightly for specific classes
- Breaks mid-morning and afternoon (10 minutes)

- **Lunch:**
- Lunch is 11:30am to 1pm
- Yes, 1 hour and 30 minutes
- Extra time for email, phone calls, or simply a walk.



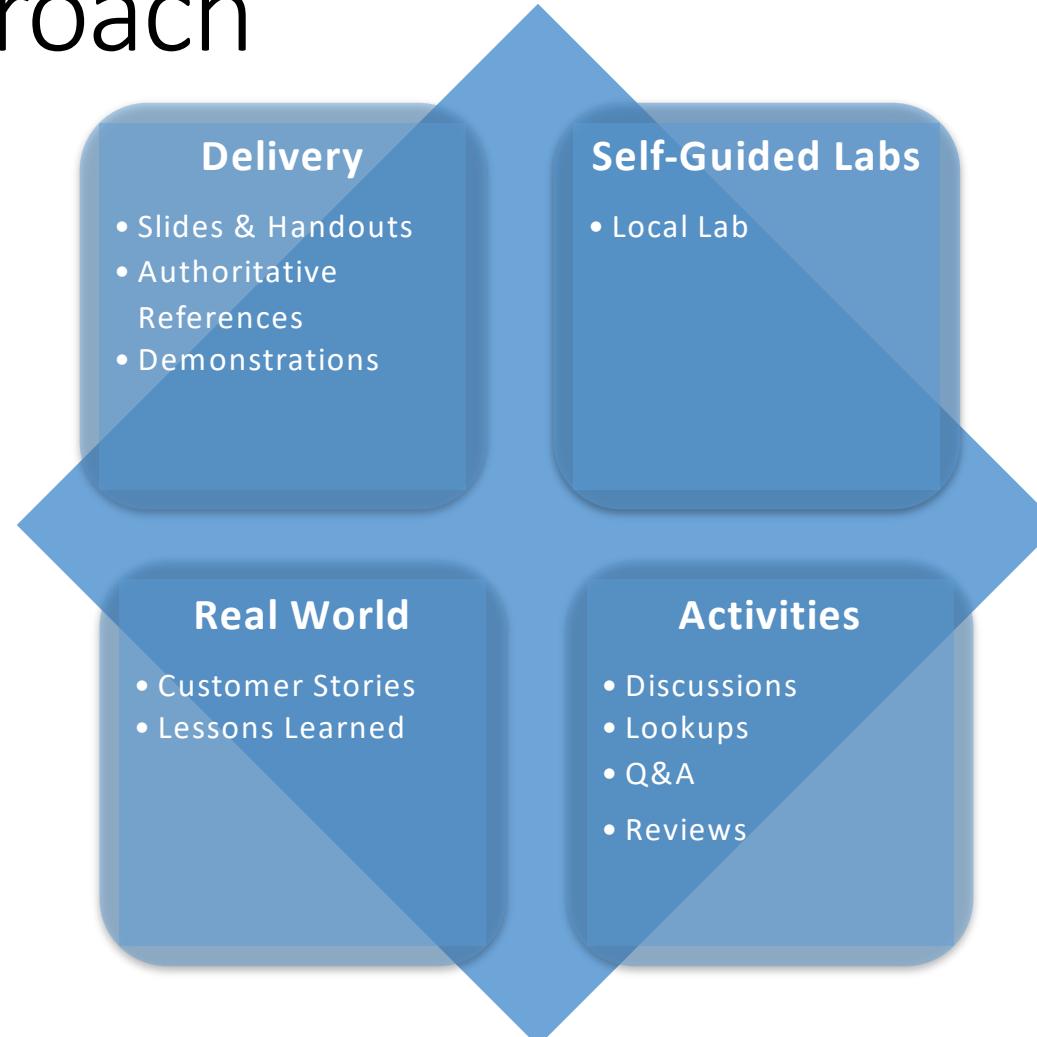
- **Telecommunication:**
- Turn off or set electronic devices to vibrate
- Reading or attending to devices can be distracting to other students

- **Miscellaneous**
- Courseware
- Bathroom

Course Objectives

- Monolithic vs Microservices
- Challenges of microservices
- Overview of Kubernetes components
 - Pods
 - Services
 - Deployments
- Install and configure Kubernetes
- How Istio simplifies microservice communications
- Install and configure Istio
- Analyze microservice communications

Course Approach



End of Day Office Hours - 1:1 Assistance

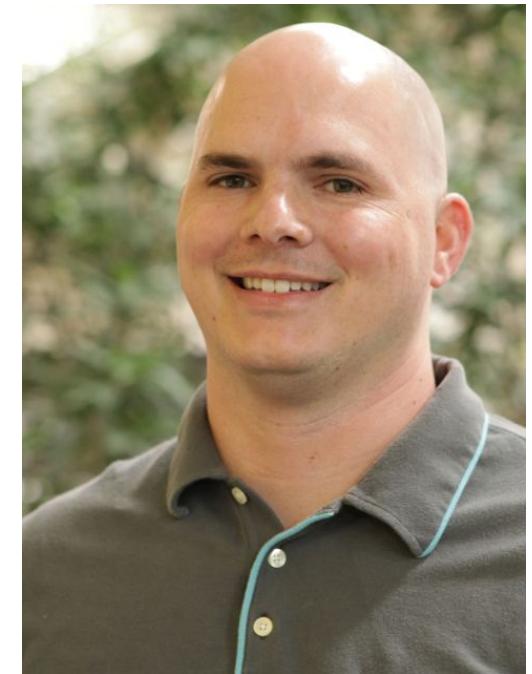
The Training Dilemma



Meet the Instructor

Jason Smith

Cloud Consultant with a Linux sysadmin background.
Focused on cloud-native technologies: automation,
containers & orchestration



Twitter
[@iruels](https://twitter.com/iruels)

github
<https://github.com/iruels>

mail
jason@innovationinsoftware.com

Expertise

- Cloud
- Automation
- CICD
- Docker
- Kubernetes

Introductions

- Name
- Job Role
- Which statement best describes your Istio/Kubernetes experience?
 - a. I am ***currently working*** with Istio/Kubernetes on a project/initiative
 - b. I ***expect to work*** with Istio/Kubernetes on a project/initiative in the future
 - c. I am ***here to learn*** about Istio/Kubernetes outside of any specific work related project/initiative
- Expectations for course (please be specific)

What are Monoliths?

Monolithic

- Simple to develop
 - IDEs designed for building one application
- Simple to test
 - Selenium
- Easy to deploy
 - Copy packaged application onto server
- Easy to scale by running multiple copies behind Load Balancer
- Works well in the beginning

Monolithic

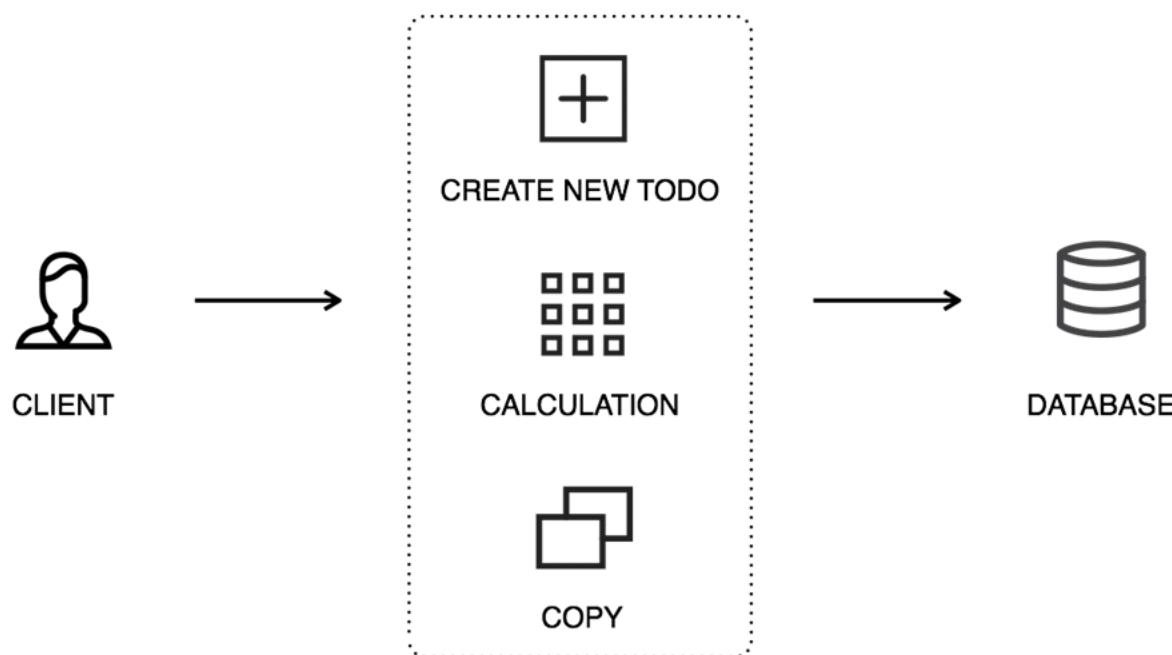
- Designed to be self contained
- Inter-connected services
- Tightly coupled components
- Must test all components if anything is changed
- Slower deploys
- Slower scaling

Monolithic

- Single development stack
- Changes are difficult
- Afraid to update

Monolithic

The Monolithic Architecture

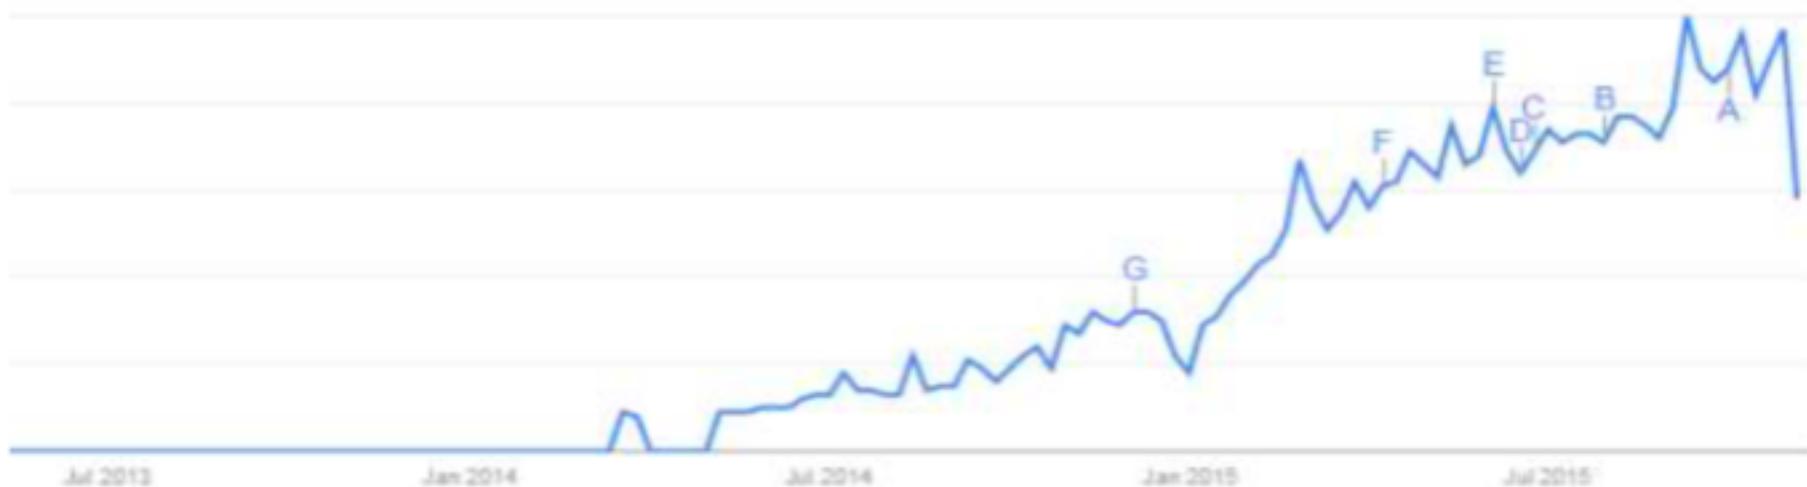
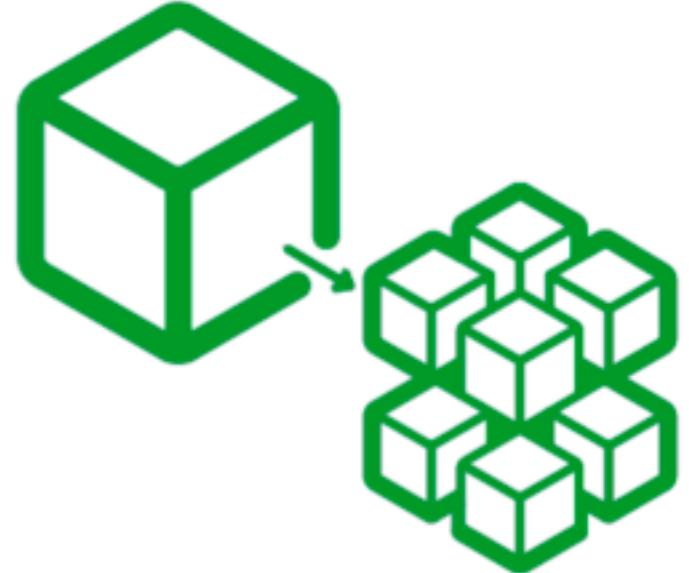


All services combined into one build,
written in the same language and application framework

What are Microservices?

What are Microservices?

The term microservices was coined in 2012 by a group of software architects. It started gaining in popularity in 2014 when Martin Fowler wrote a blog post describing it.



What are Microservices?

- Services communicate with each other over the network.
- Update services independently.
 - No need to change other services.
- Self-contained.
 - You can update the code without knowing anything about internals of other microservices.

What are Microservices?



adrian cockcroft
@adrianco

Follow

Replies to @kellabyte

@kellabyte @mamund I used to call what we did "fine grain SOA". So microservices is SOA with emphasis on small ephemeral components

5:16 PM - 10 Dec 2014

3 Retweets 5 Likes



0 7 12 3 5

Functional decomposition of systems into manageable and independently deployable components,

Microservice Architectures by Dr. Andreas Schroeder

“Service-oriented architecture composed of loosely coupled elements that have bounded contexts” (VP Cloud Architecture at AWS) - Adrian Cockcroft

What are Microservices?

“In short, the microservice architectural style is an approach to developing a single application as a suite of small services, each running in its own process and communicating with lightweight mechanisms, often an HTTP resource API.”

- Martin Fowler

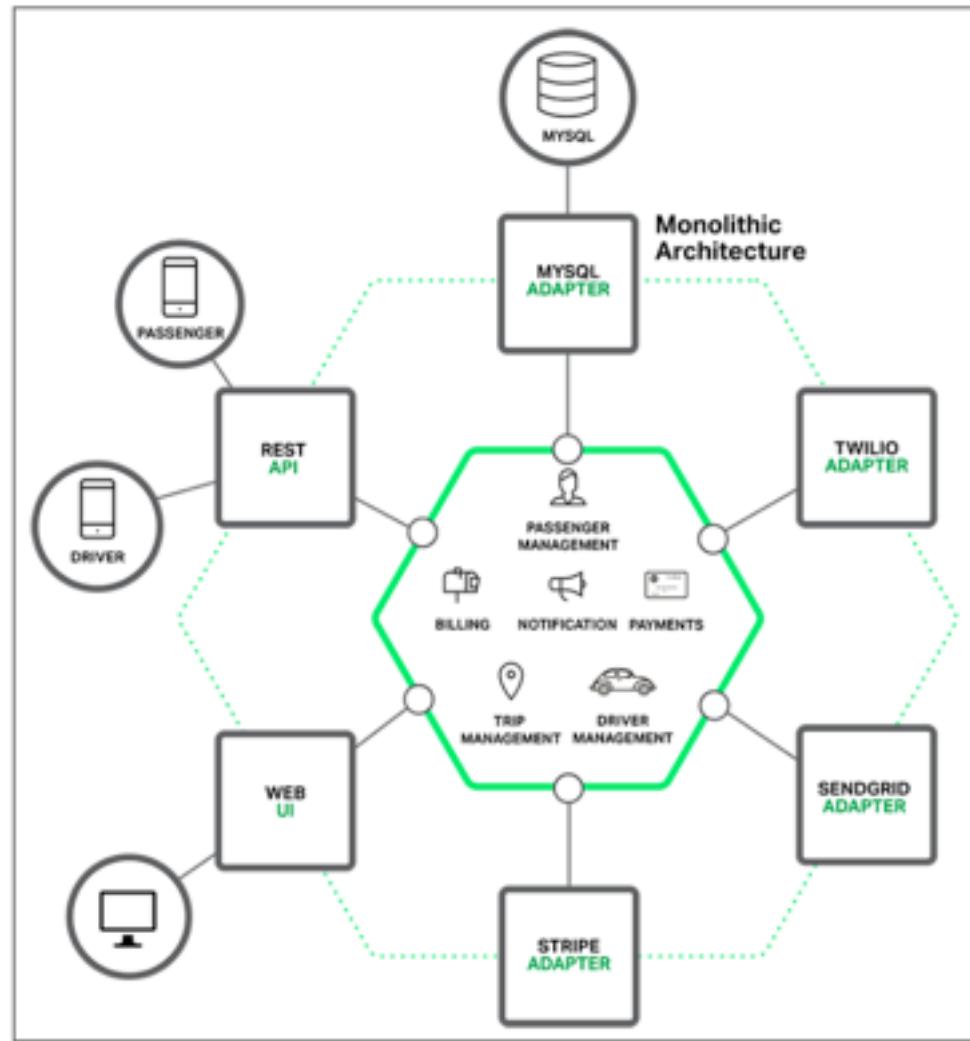
Guidelines for Microservices

- Make each one focused
 - “2 pizza teams”
- Treat each service as one app
 - Have one SCM repo, own pipeline, etc...
- Each service should be independent of other services
- Service owners are responsible for entire lifecycle
- Use lightweight protocols
- Use the right tool for the job

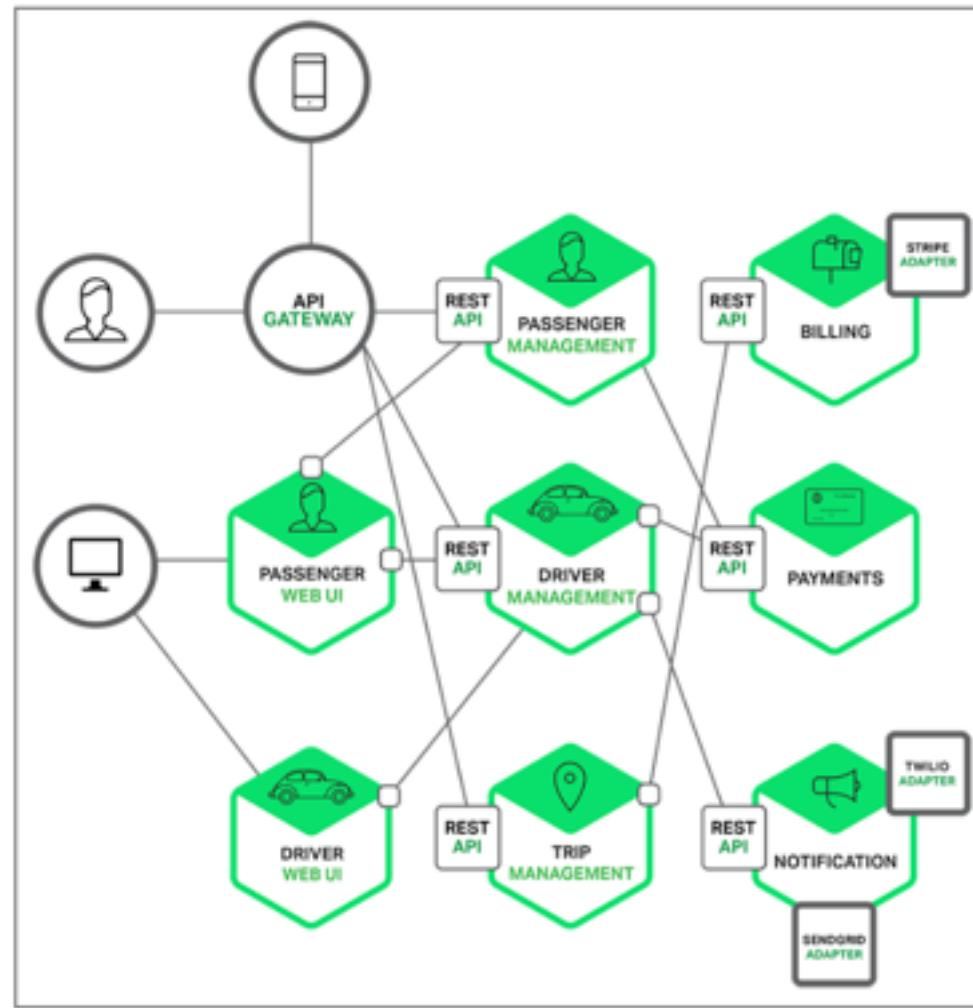
Guidelines for Microservices

- Each service should have it's own datastore
- Failure will happen, design for it
- Automate everything!

Monolithic Architecture

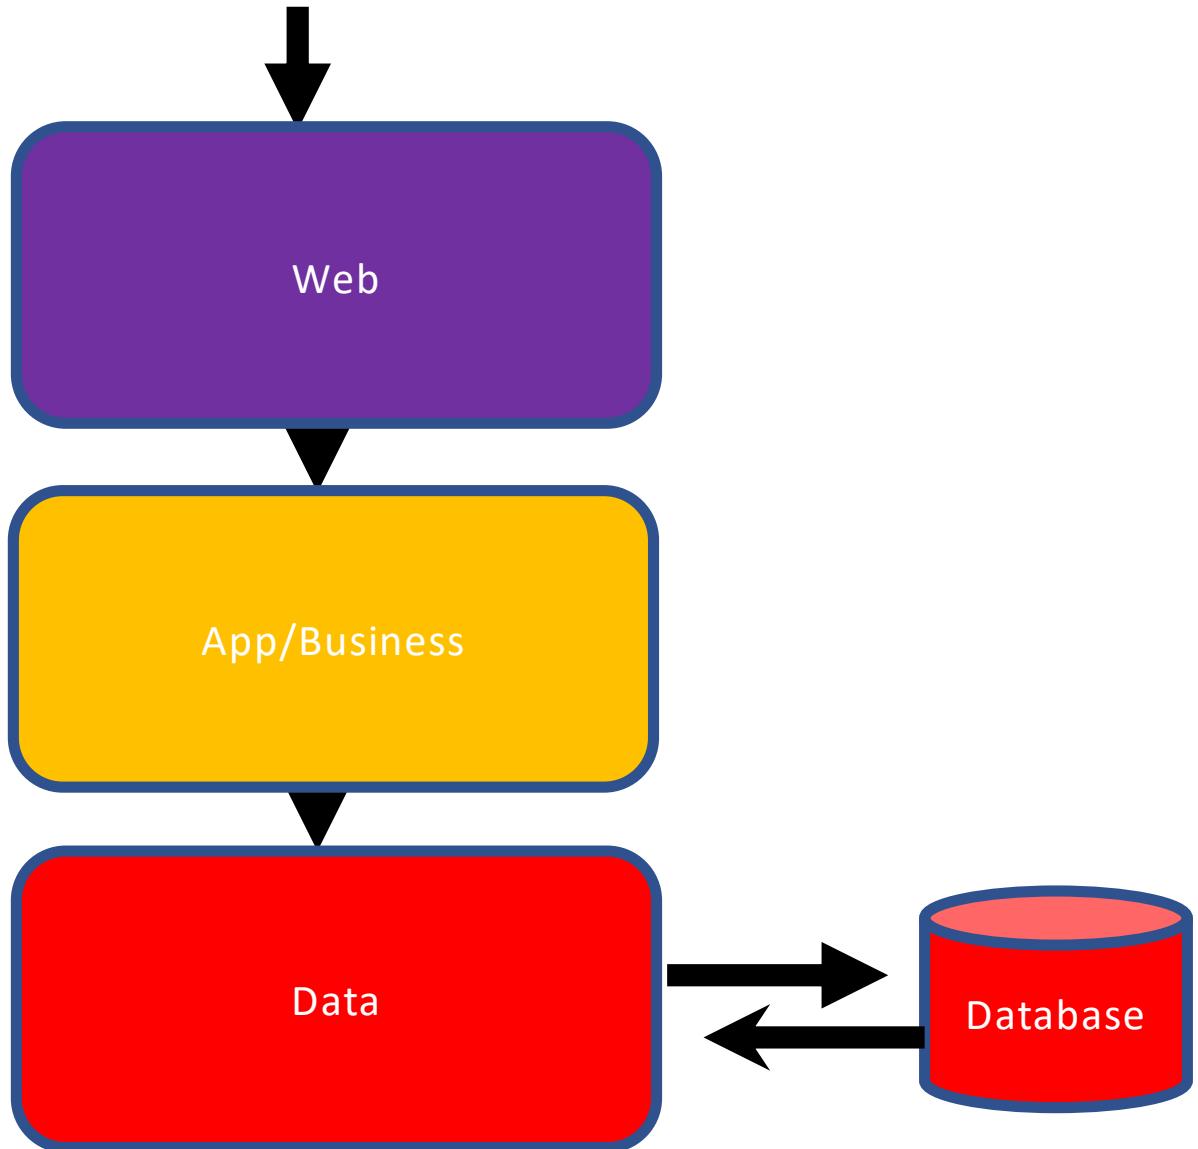


Microservice Architecture



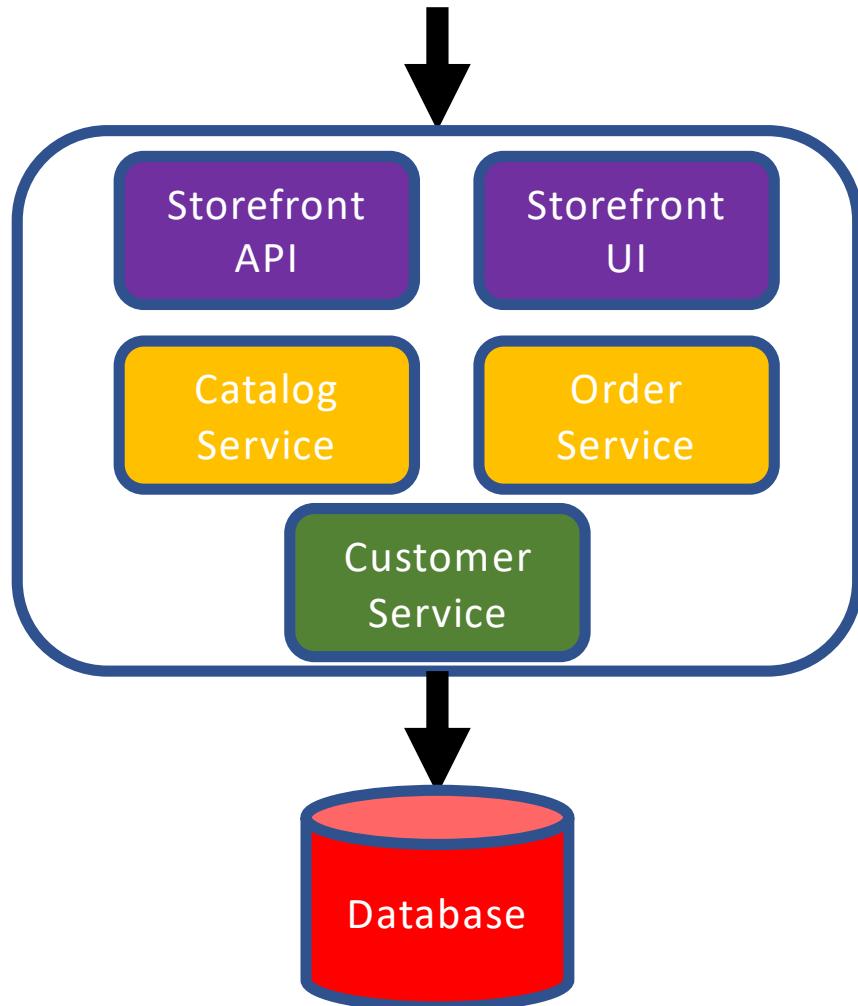
Monolithic Architecture

- Classic “three-tier” architecture
- Deploy all three tiers as single entity
 - Java WAR/EAR file



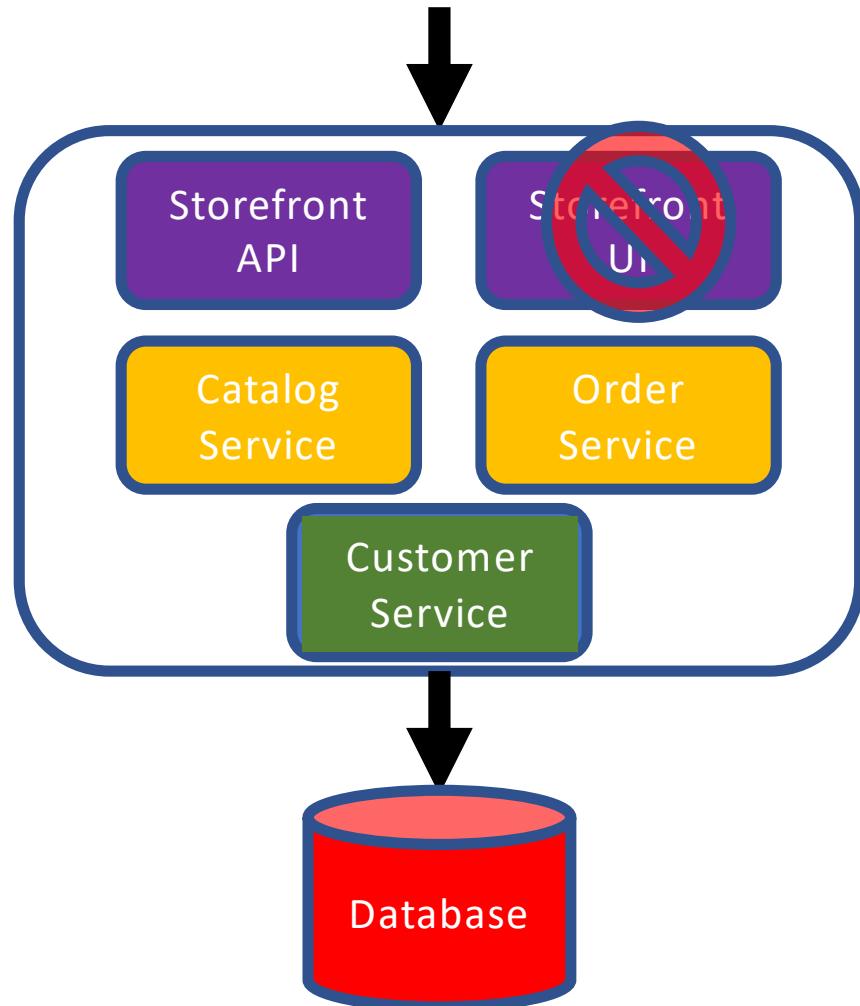
Monolithic Architecture

- All services built into one application

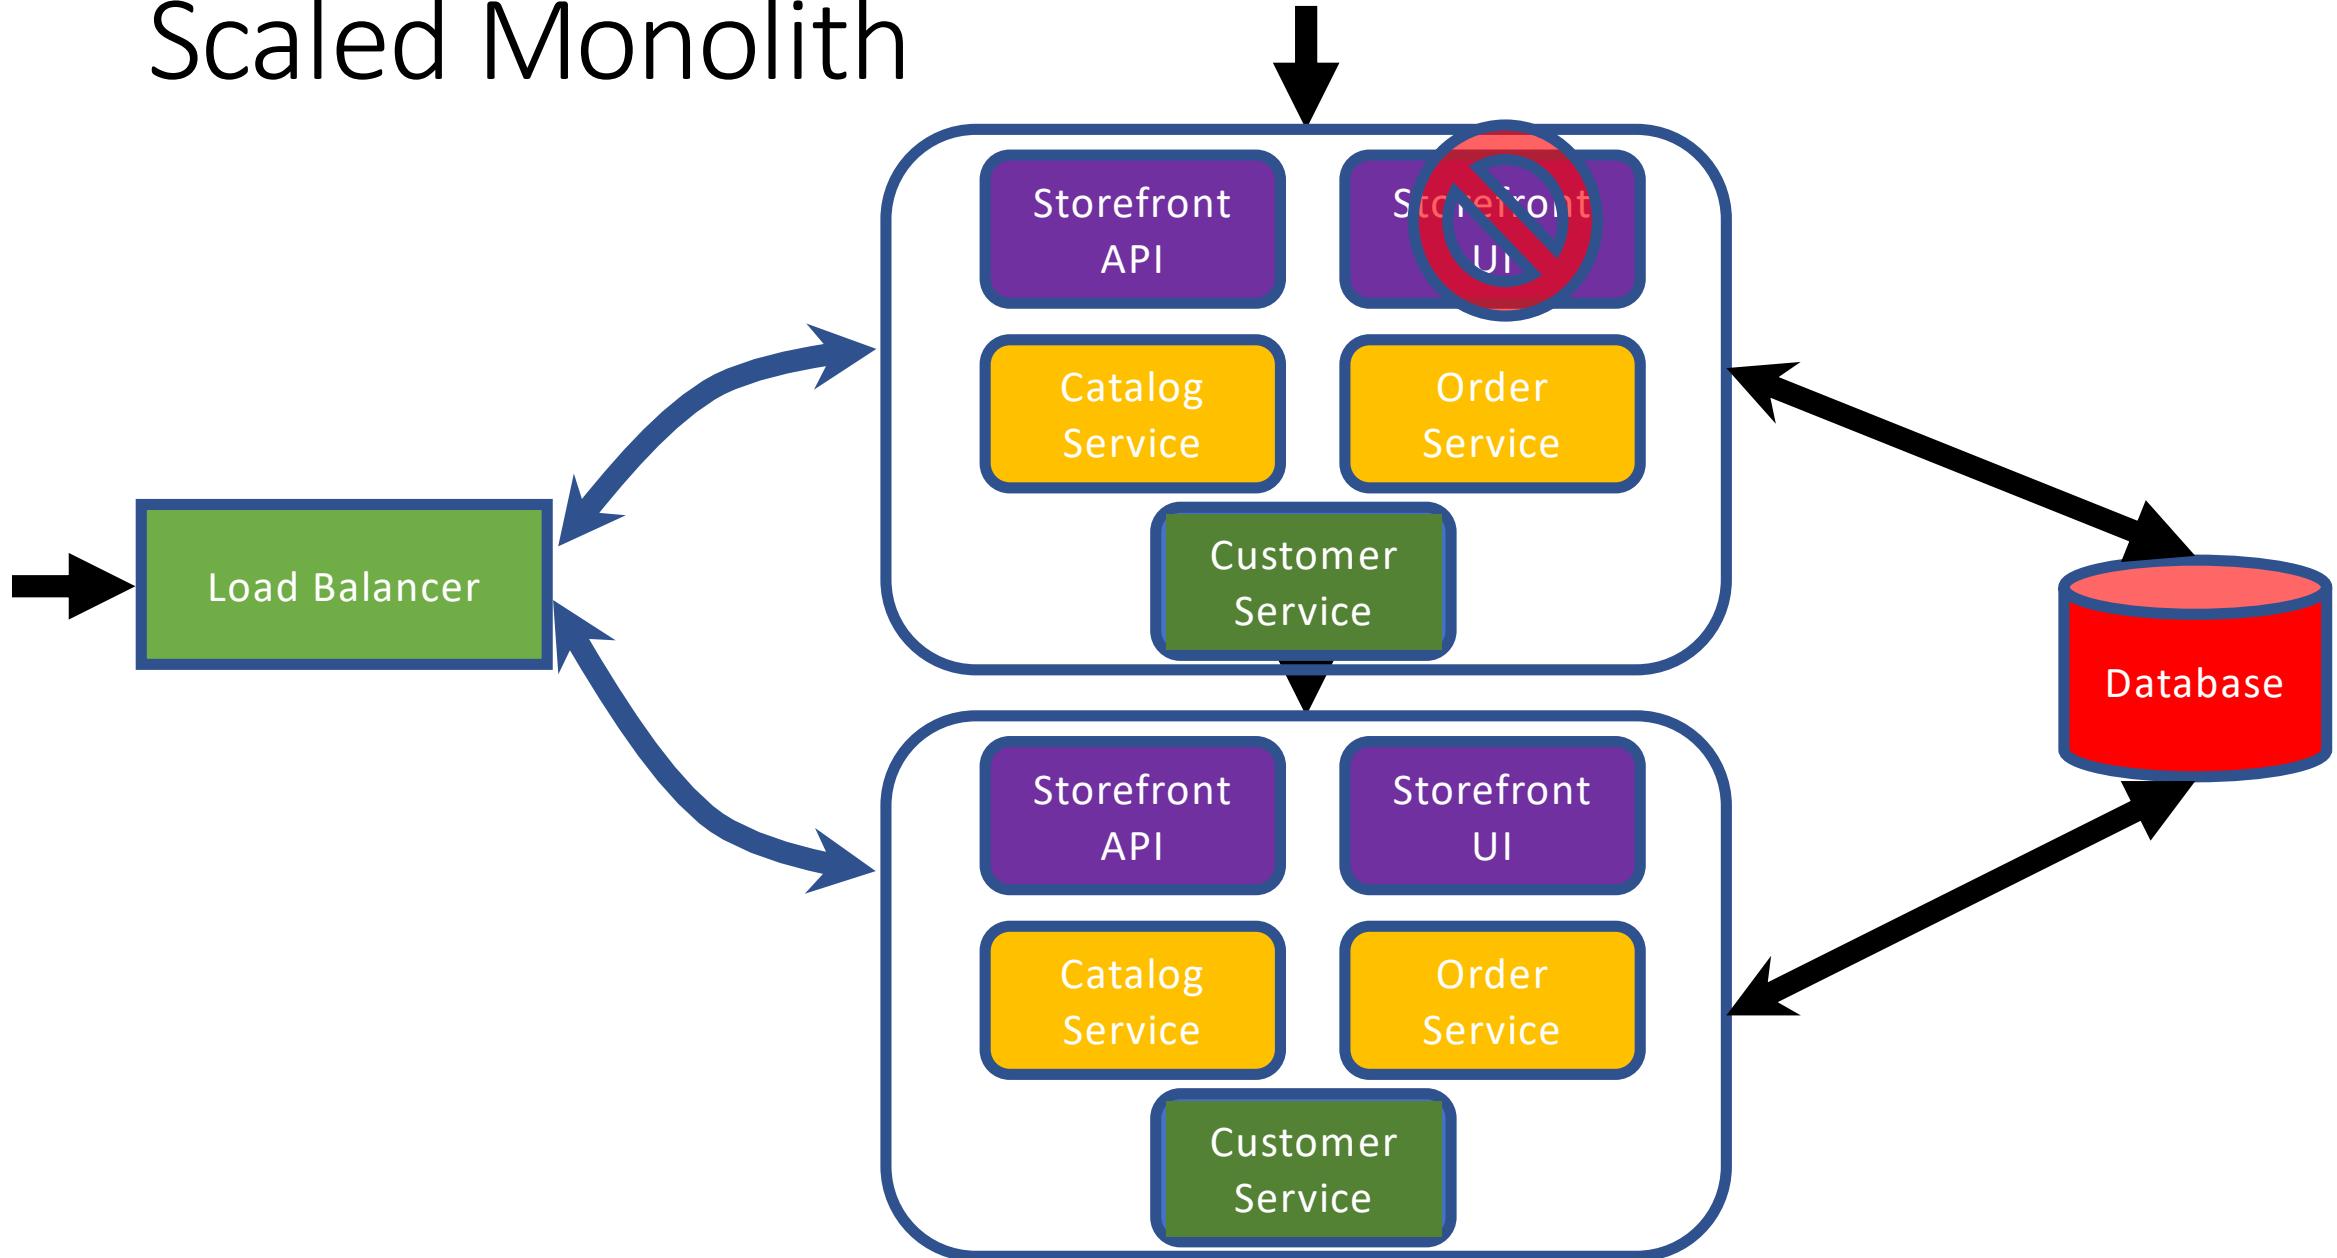


Monolithic Architecture

- Failure in monolith

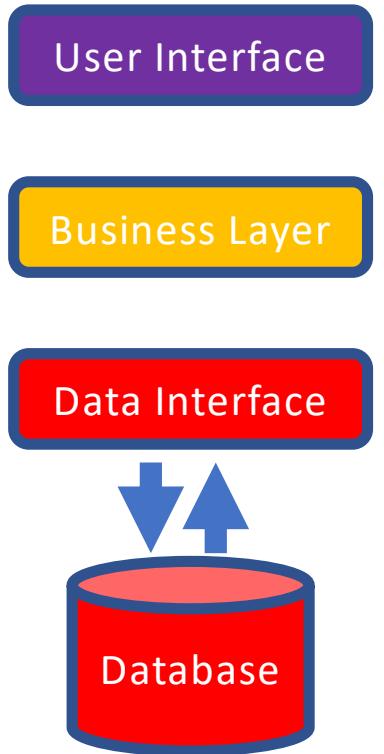


Scaled Monolith

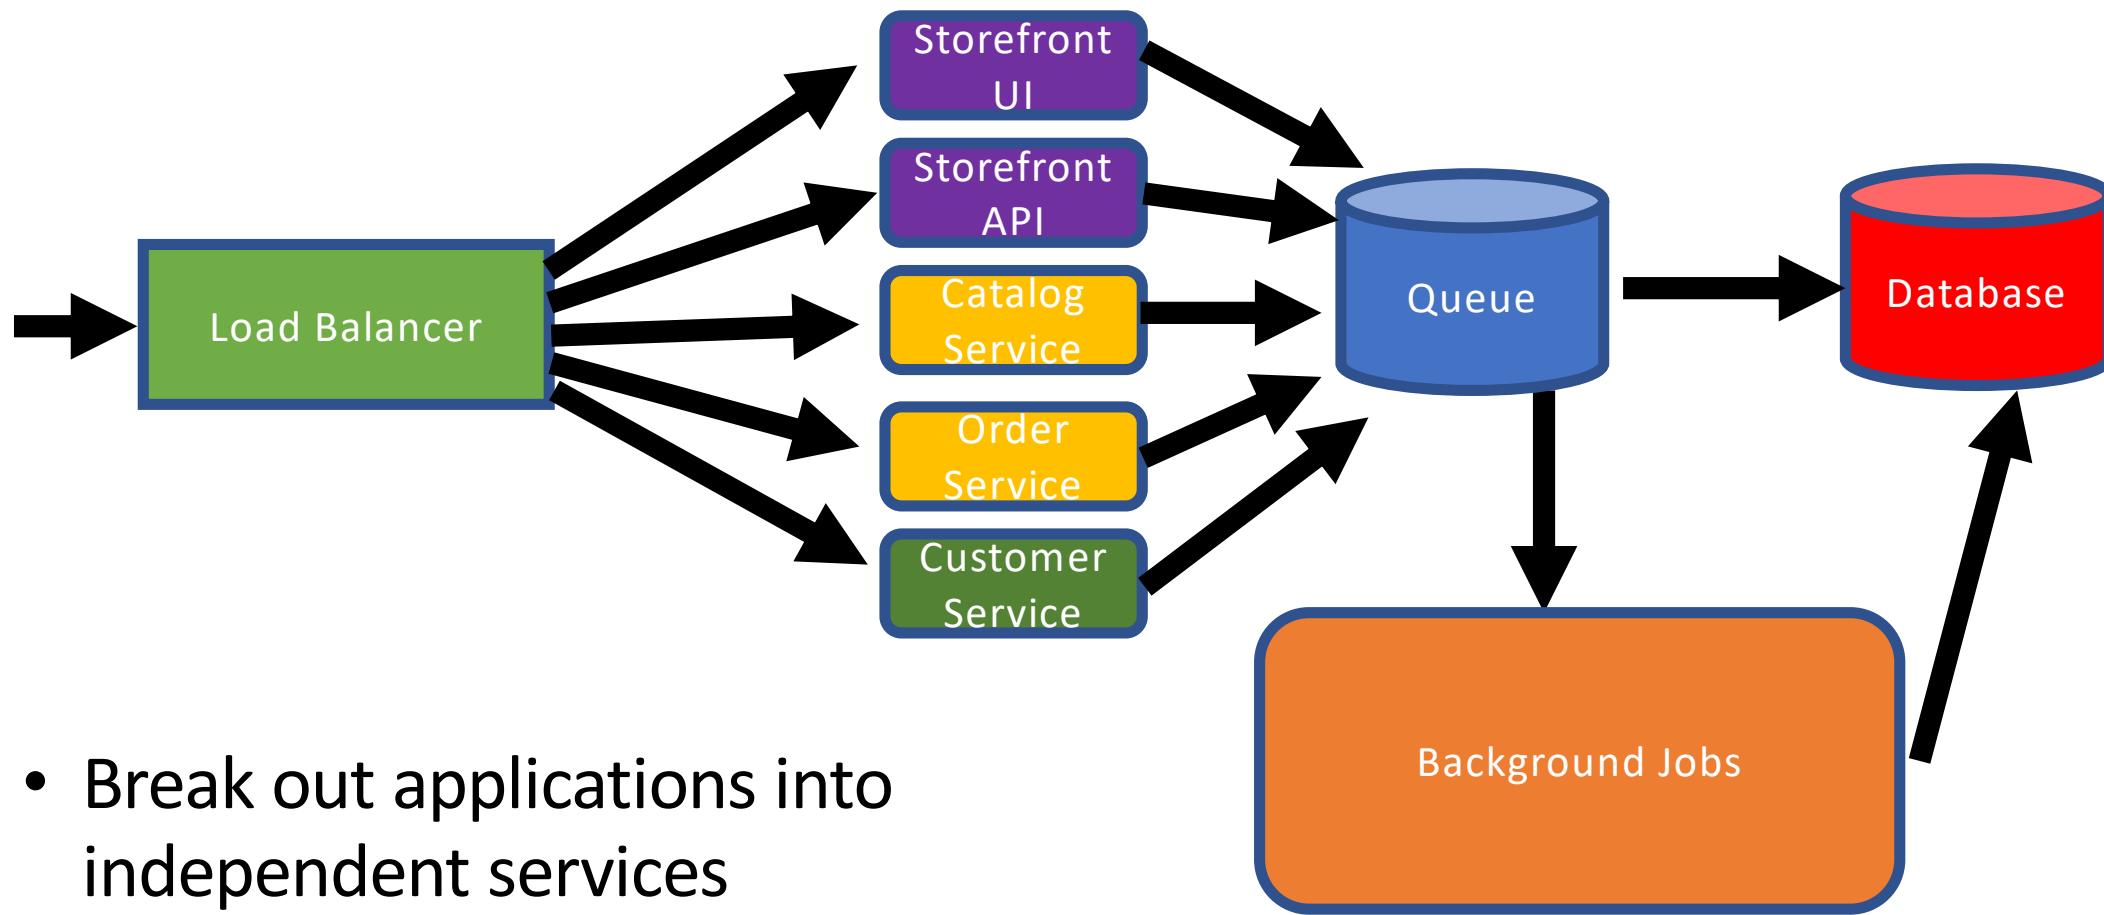


Monoliths = Siloed Functional Teams

- Components are not aligned to business functionality, silos are created.
- Example: new feature requires three cross functional teams
 - UI, Middleware, DBAs
- Communication breakdown
- Anti-DevOps

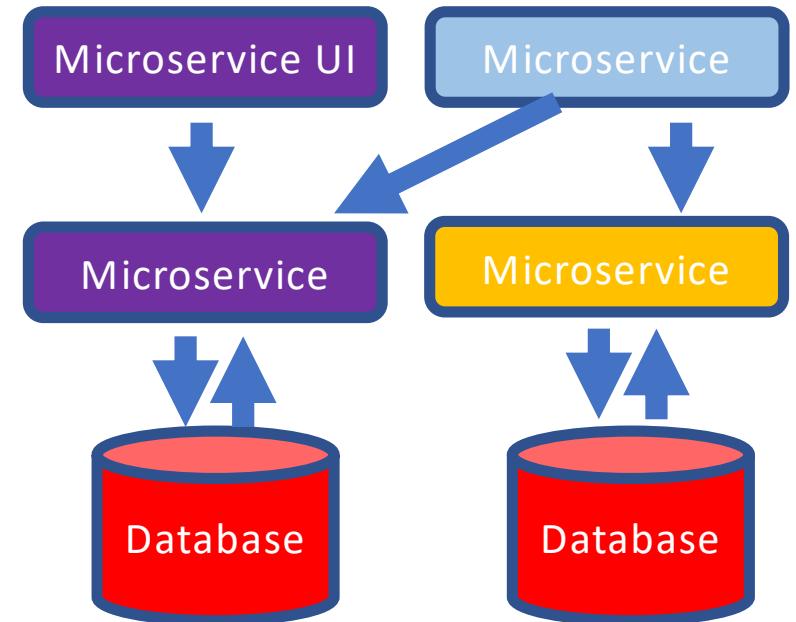


Microservices Architecture



Microservices Application

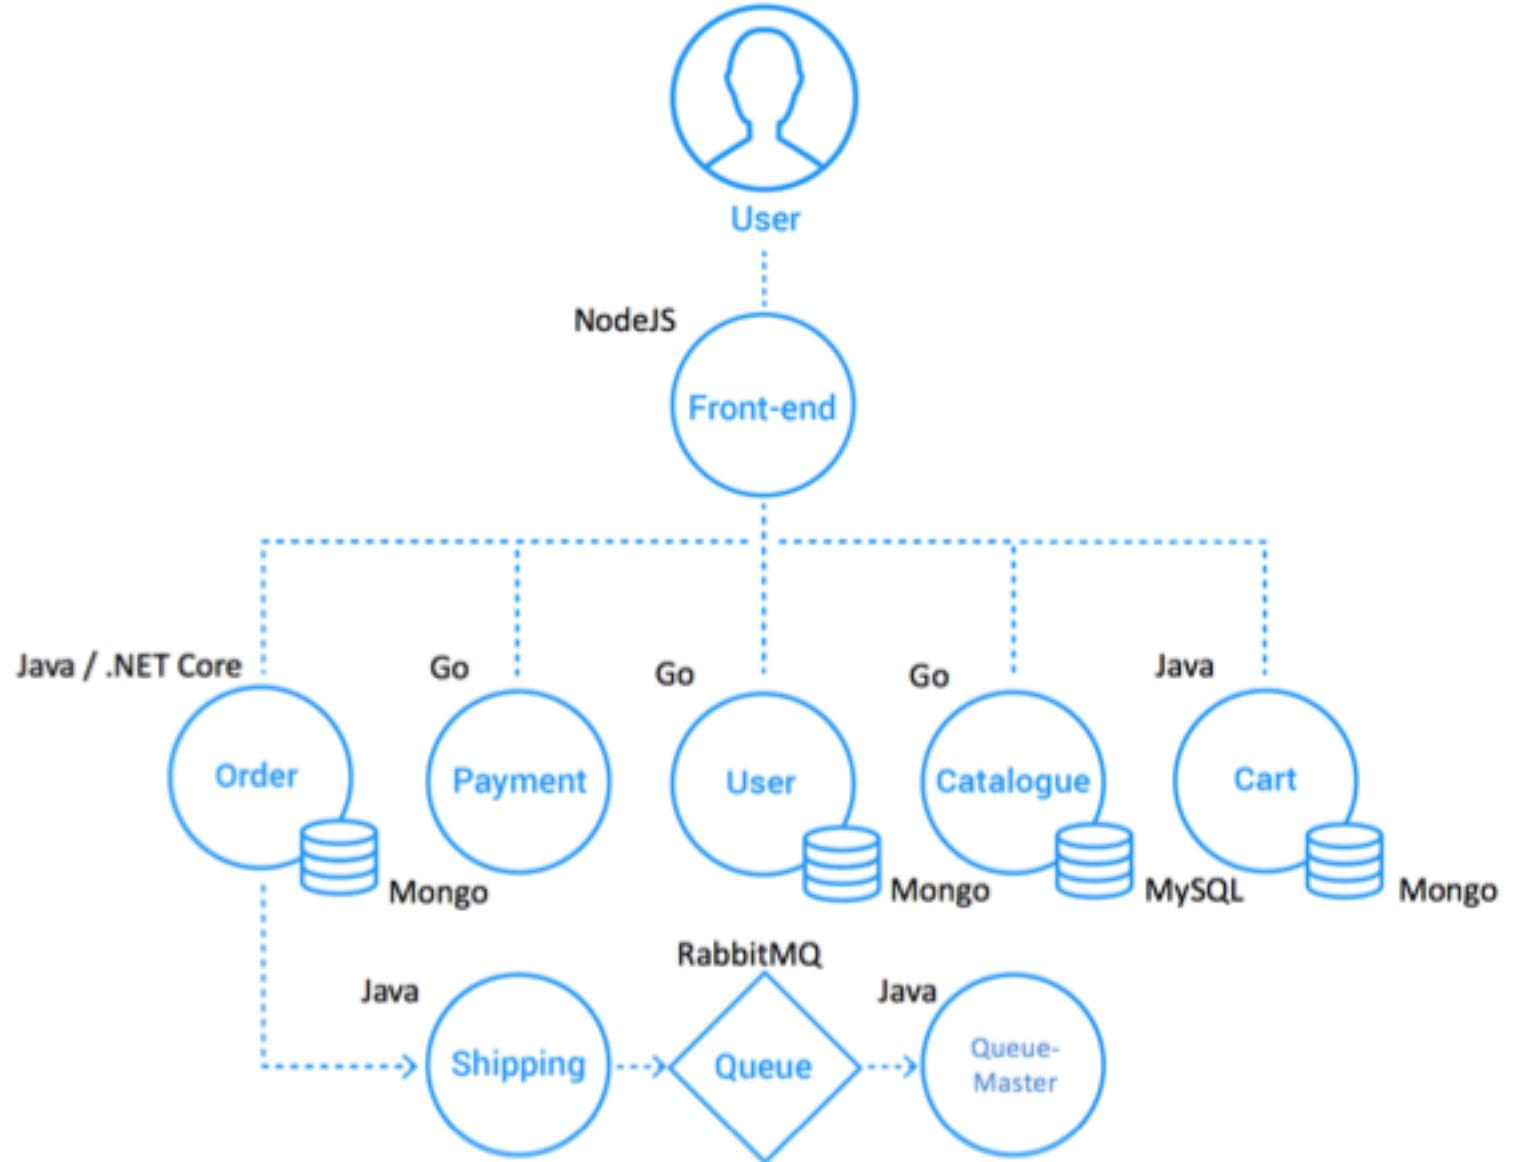
- Components developed (as services)
AND deployed separately
- Self-contained
- Adding new code much easier
- Less chance for inadvertent bugs



Demo Microservices Applications

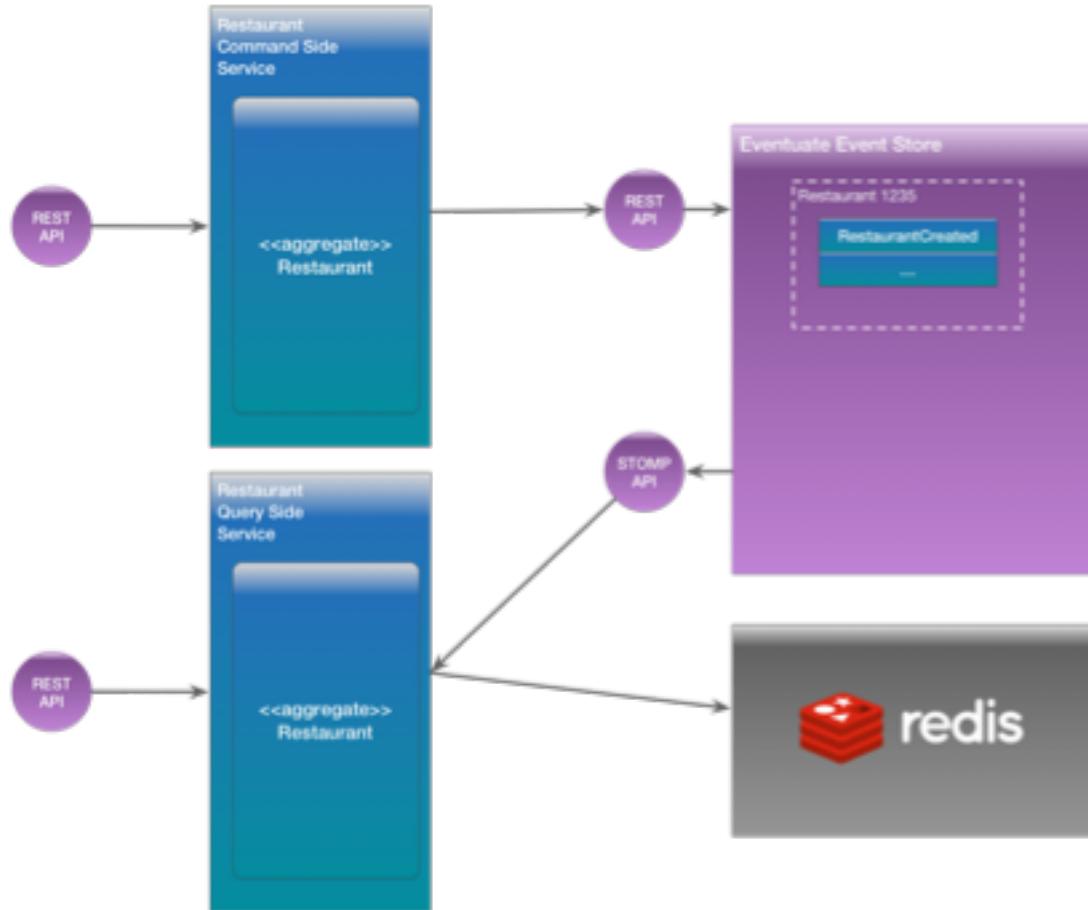
Sock Shop

- Go
- Spring Boot
- Node.js
- RabbitMQ
- MySQL
- Mongo



Restaurant Manager

- Spring Boot
- MongoDB
- SwaggerUI
- Gradle
- Zookeeper
- Kafka
- MySQL
- Redis



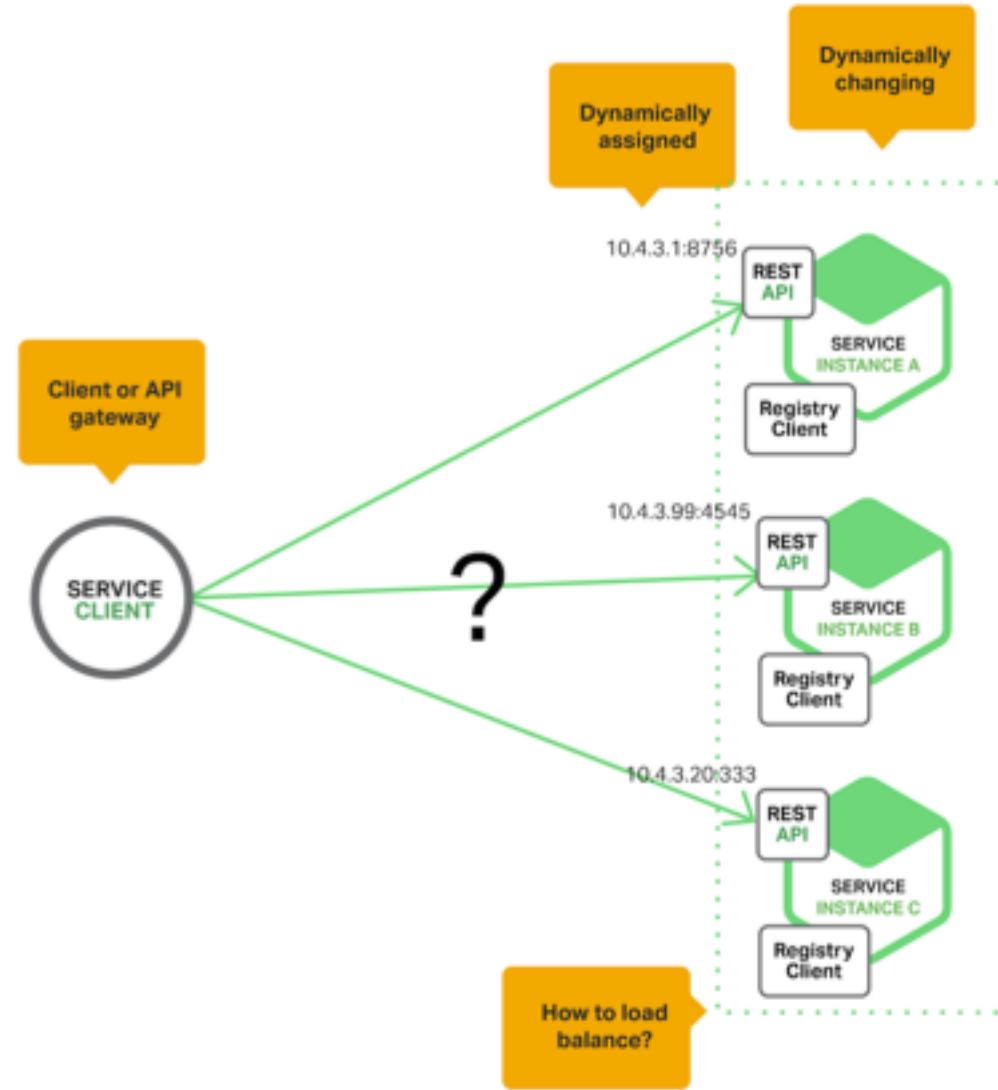
Microservices Challenges

Microservices Challenges

- Service discovery & request routing
- Security and access control
- Efficiently scale up/down
 - connectivity
 - service resiliency
- Logging/Monitoring
- Canary/AB testing/rollouts/rollbacks

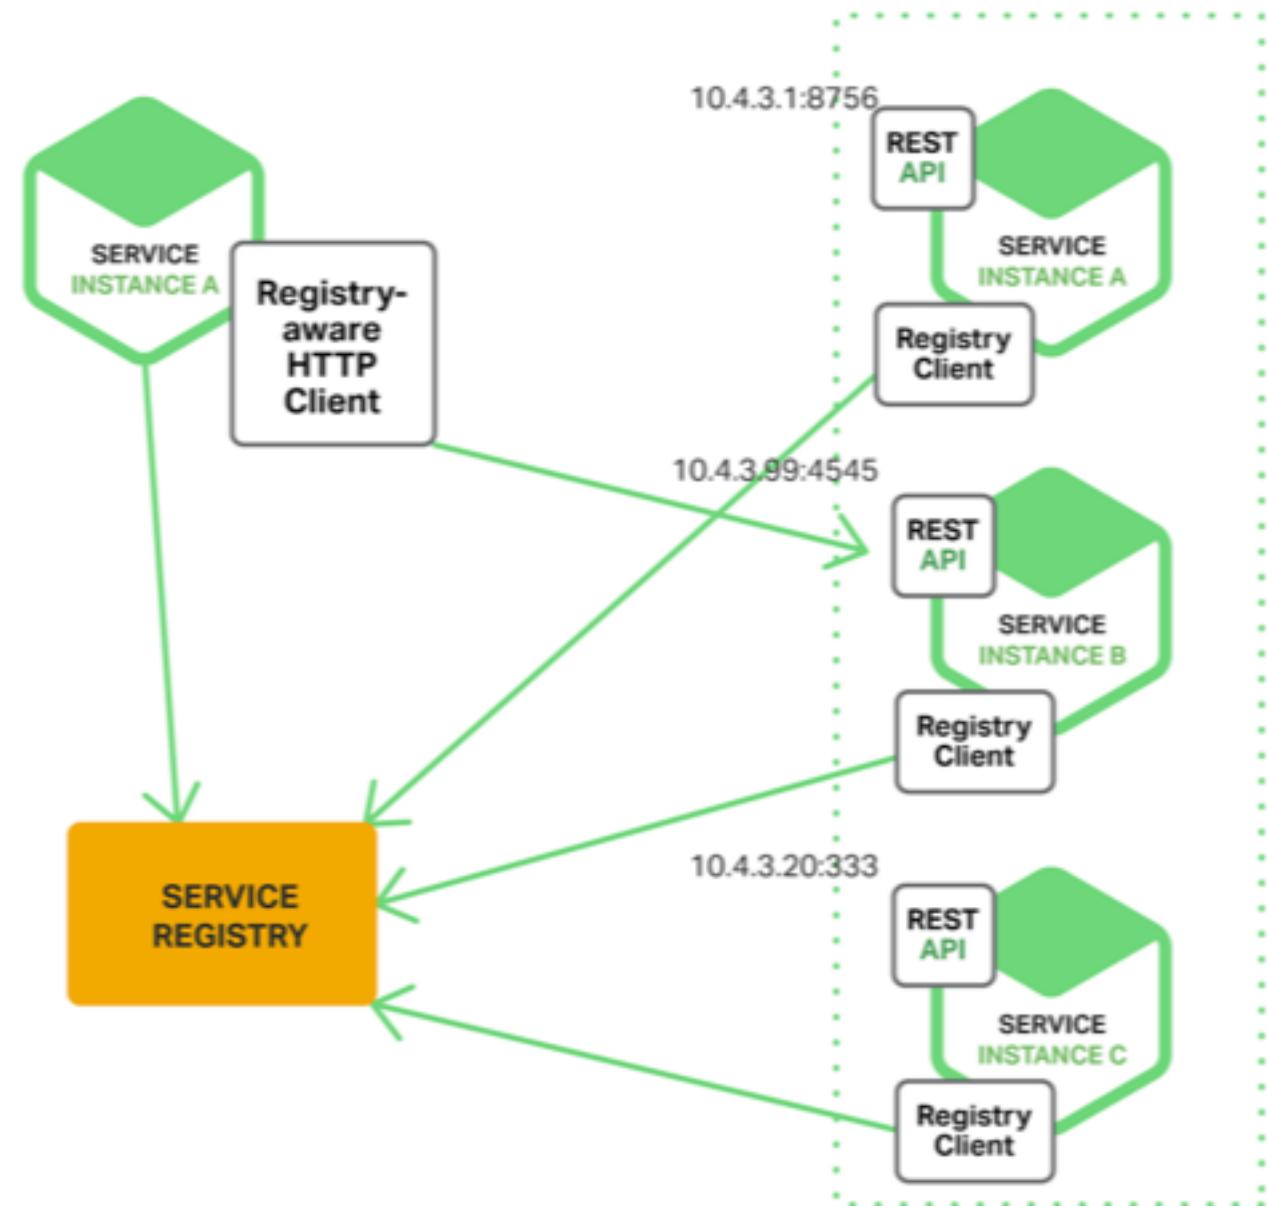
Service discovery

- Service discovery
 - Client side?
 - Server side?



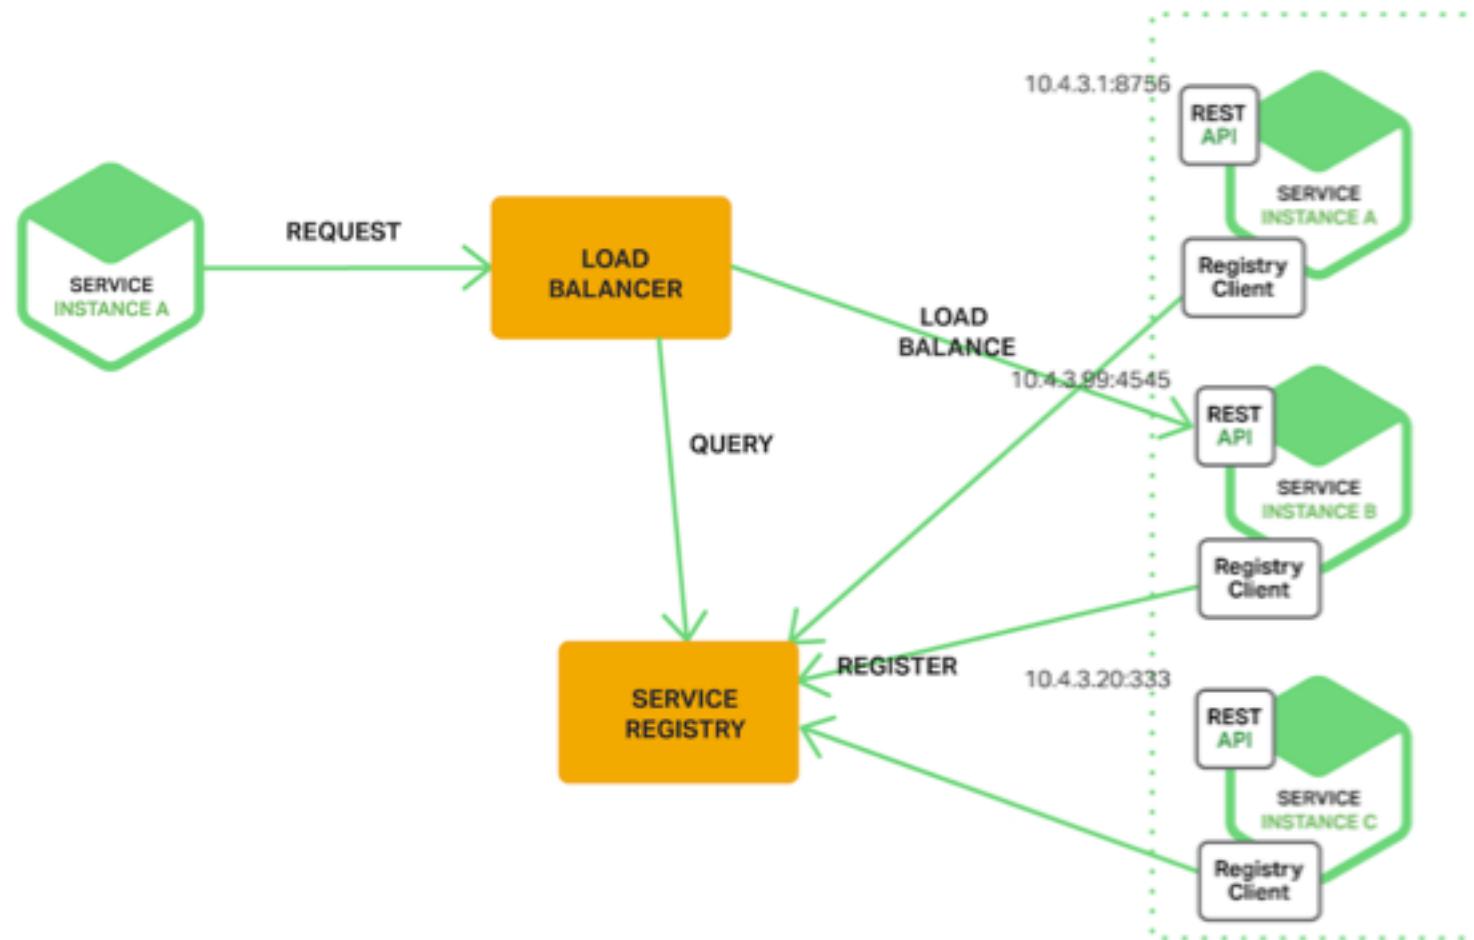
Service discovery

- Client side
 - Client responsible for finding services
 - Couples client and services requiring code changes



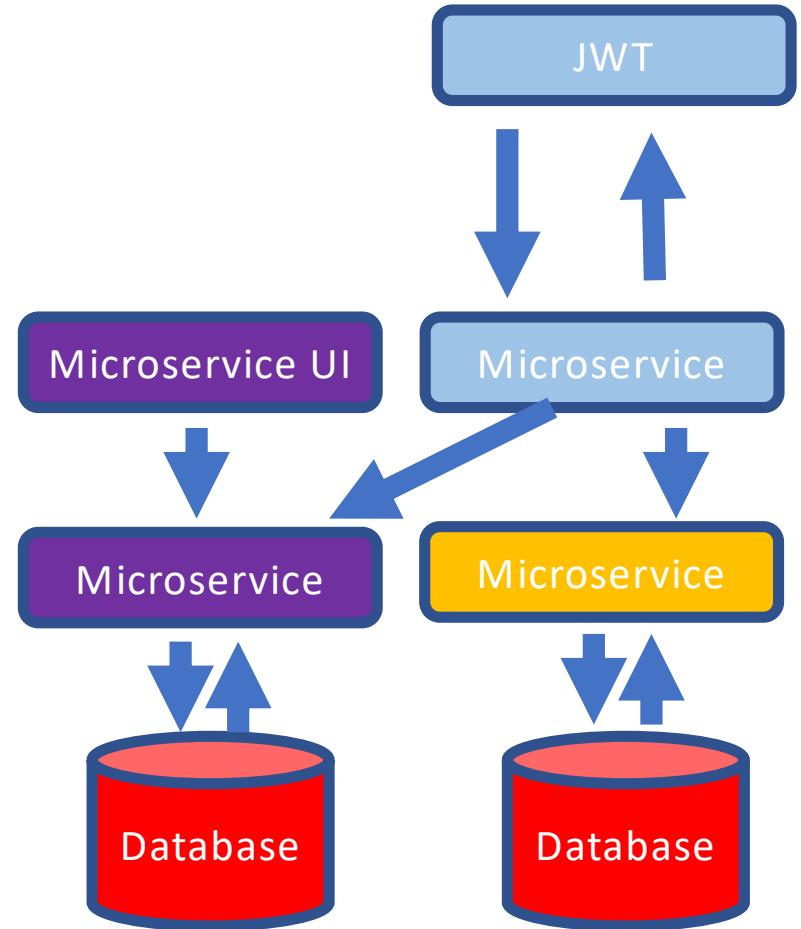
Service discovery

- Server side
 - ELB
 - Nginx & Consul
 - included in Kubernetes
 - Discovery abstractected away from client

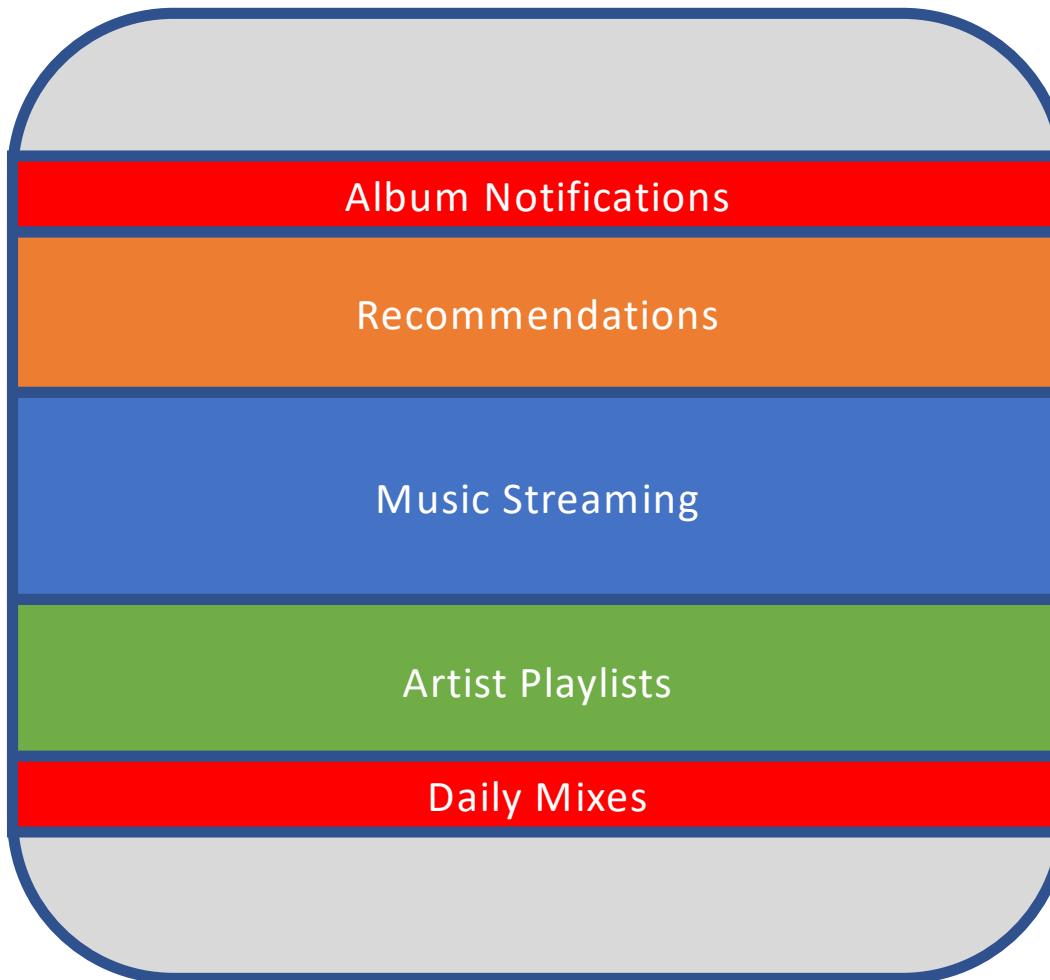


Security & Access Control

- Stateless
- JSON Web Token (JWT)
- API Gateway usually handles auth



Scaling - Monolithic

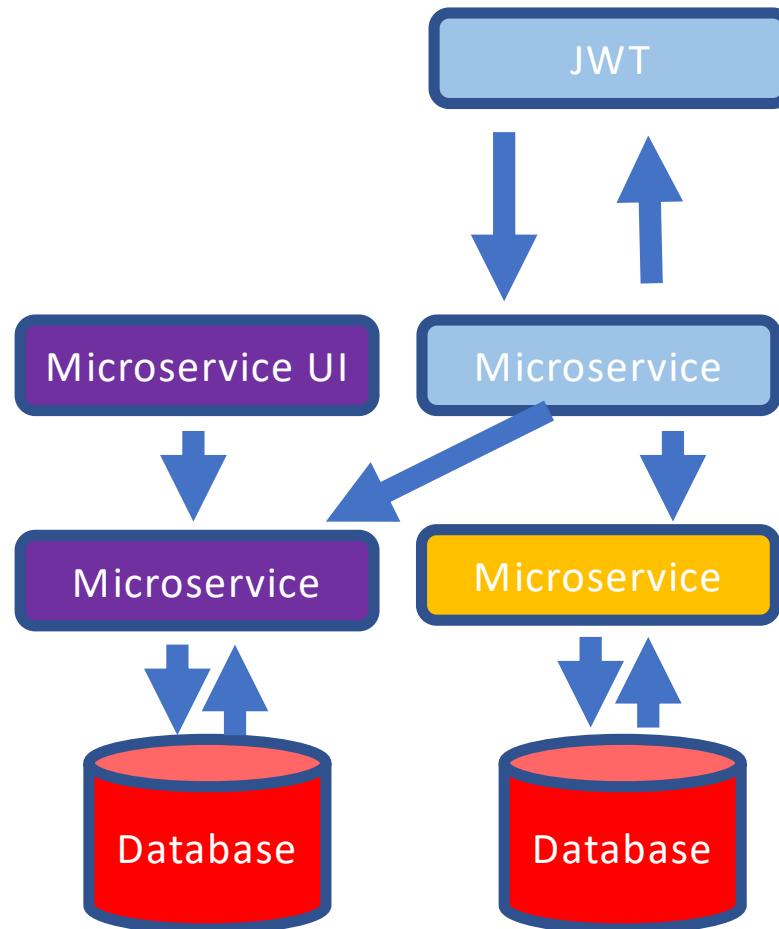


Scaling - Monolithic



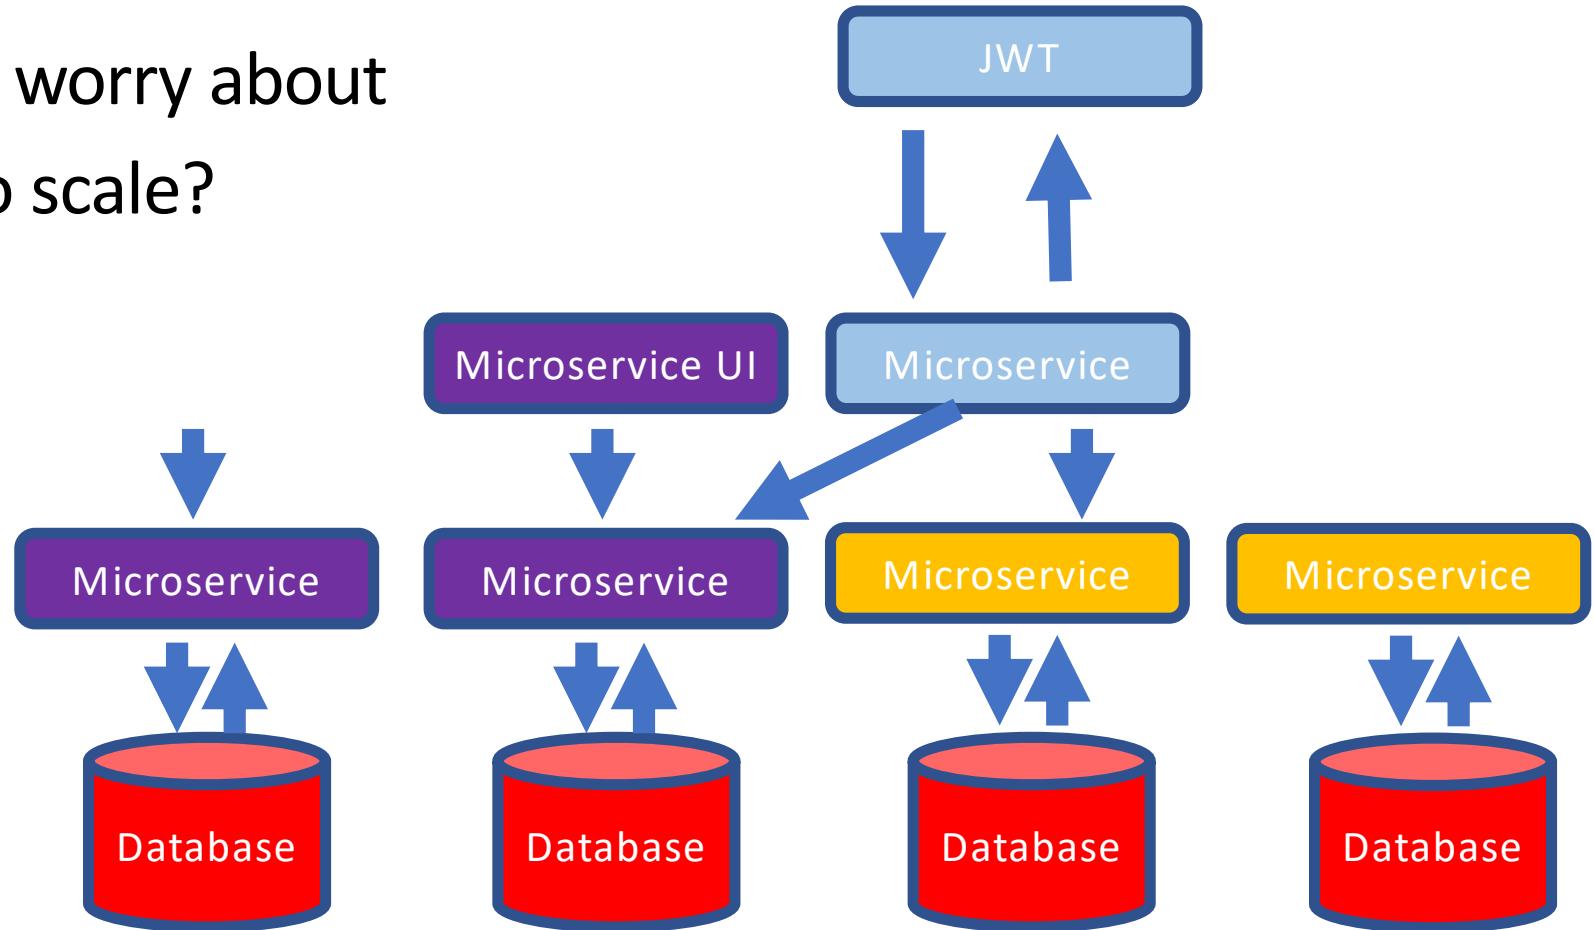
Scaling - Microservices

- Multiple components to worry about
- Which service(s) need to scale?



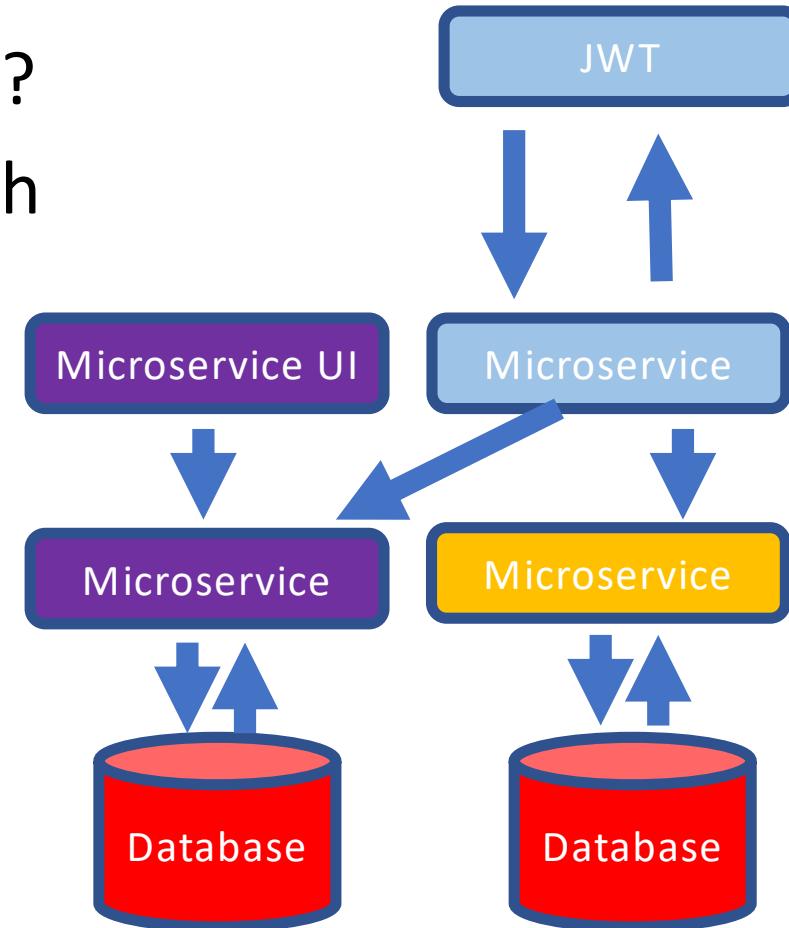
Scaling - Microservices

- Multiple components to worry about
- Which service(s) need to scale?



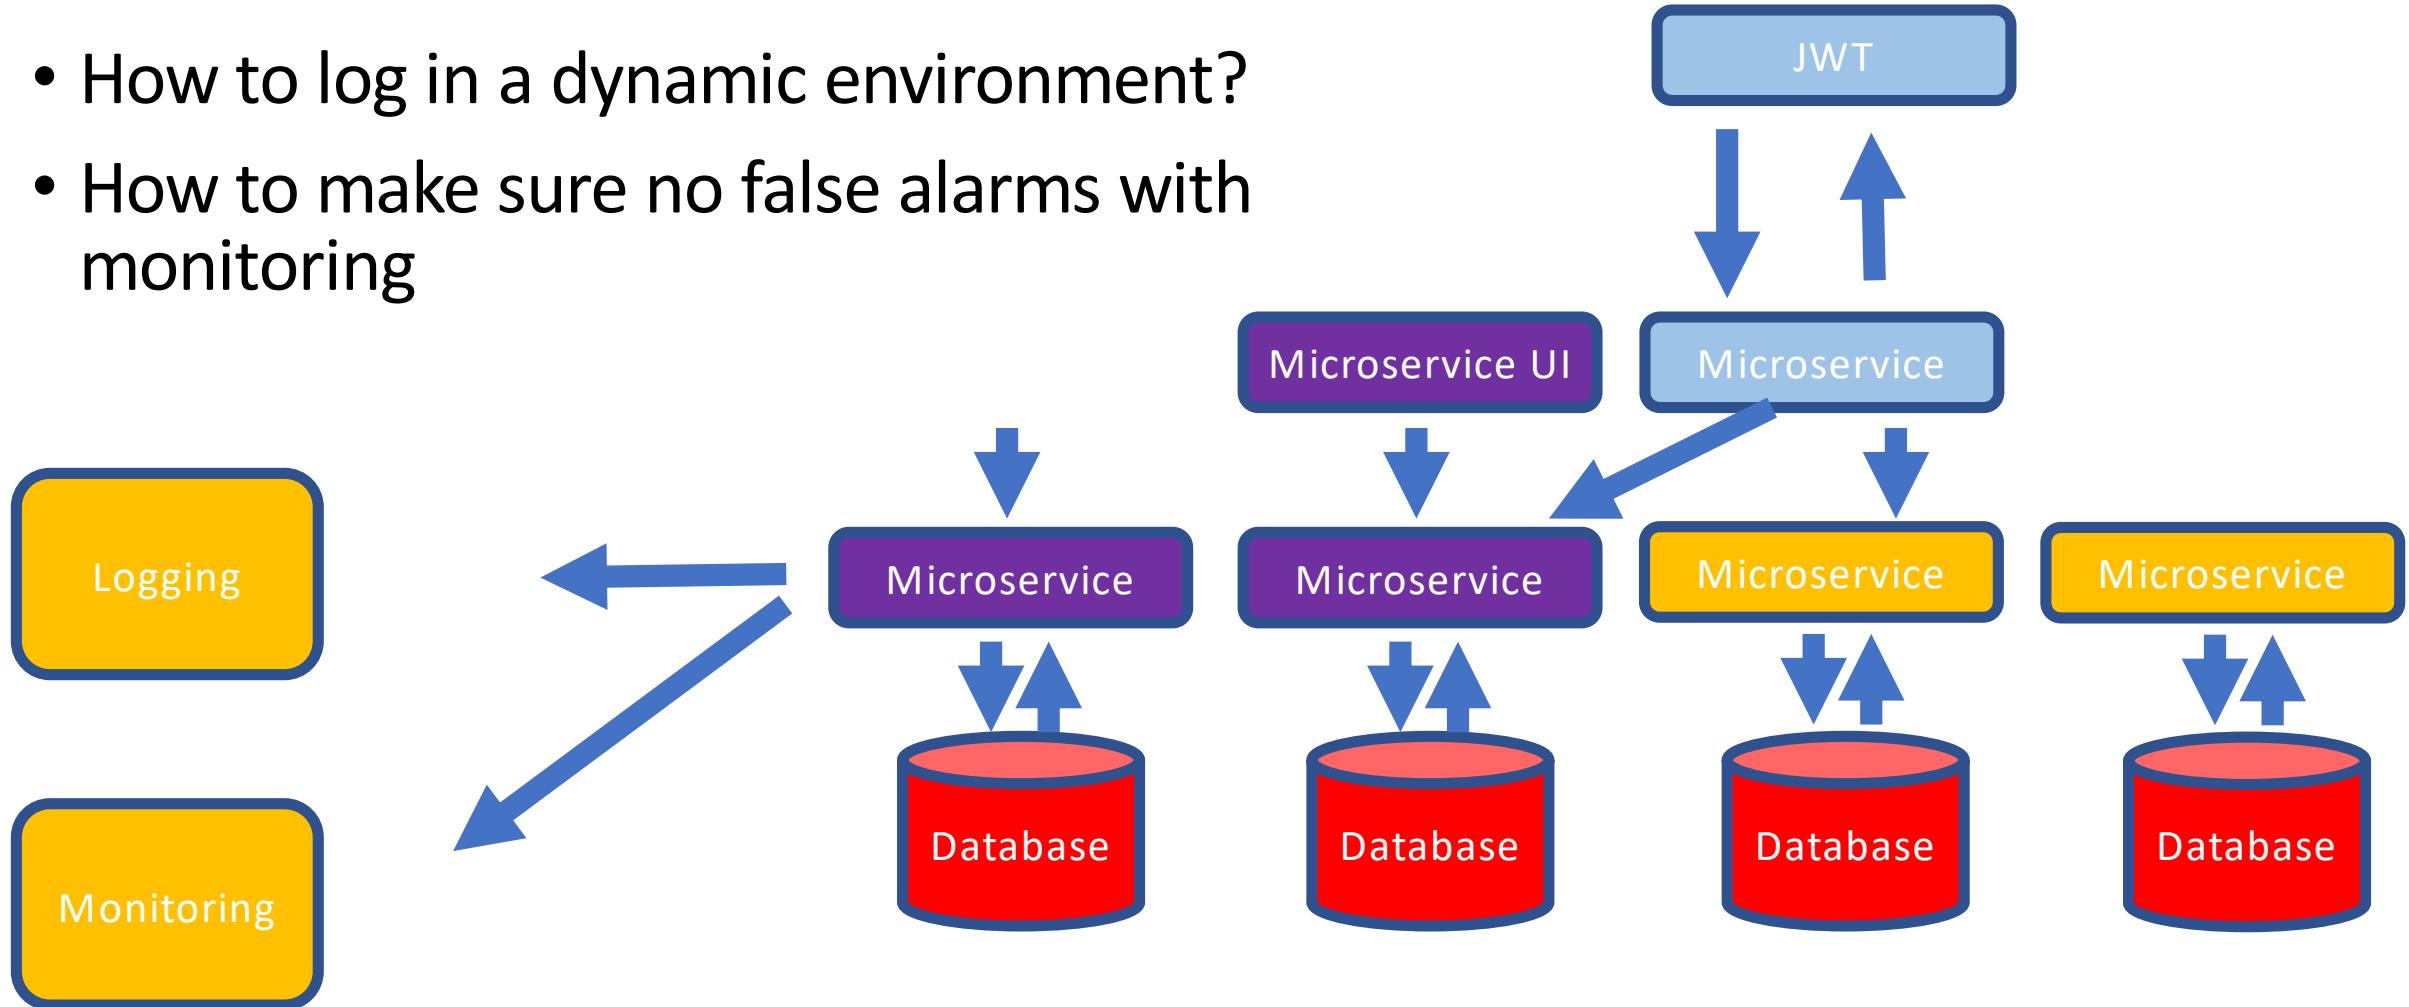
Logging and Monitoring

- How to log in a dynamic environment?
- How to make sure no false alarms with monitoring



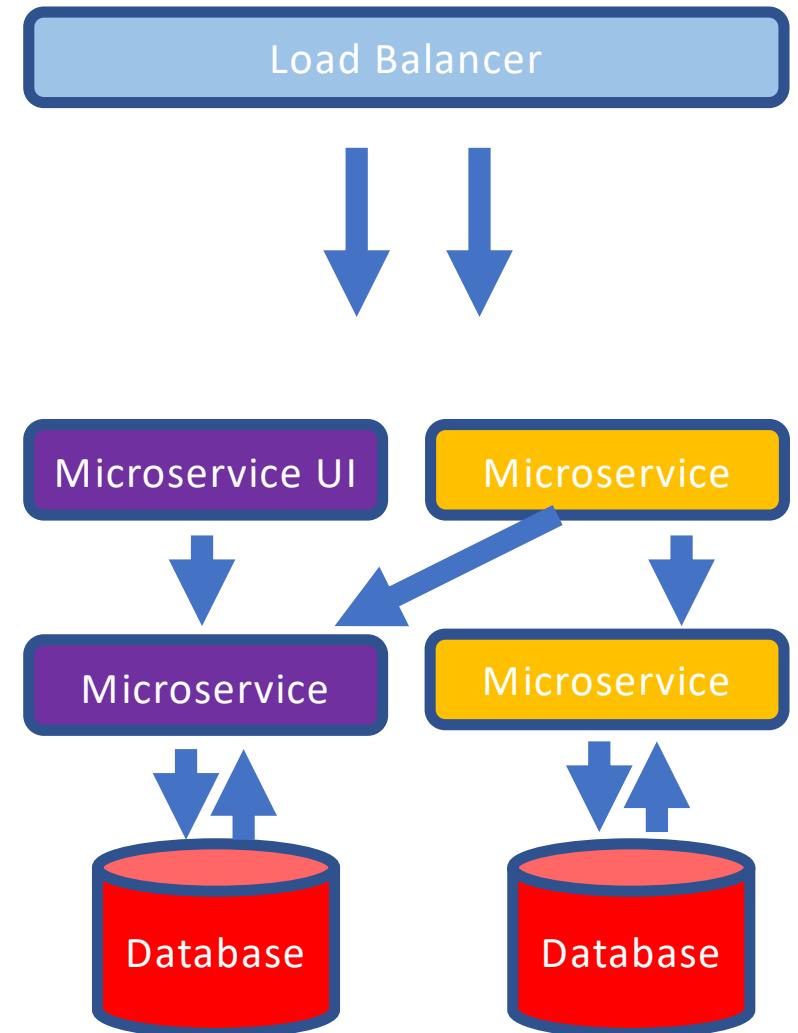
Logging and Monitoring

- How to log in a dynamic environment?
- How to make sure no false alarms with monitoring



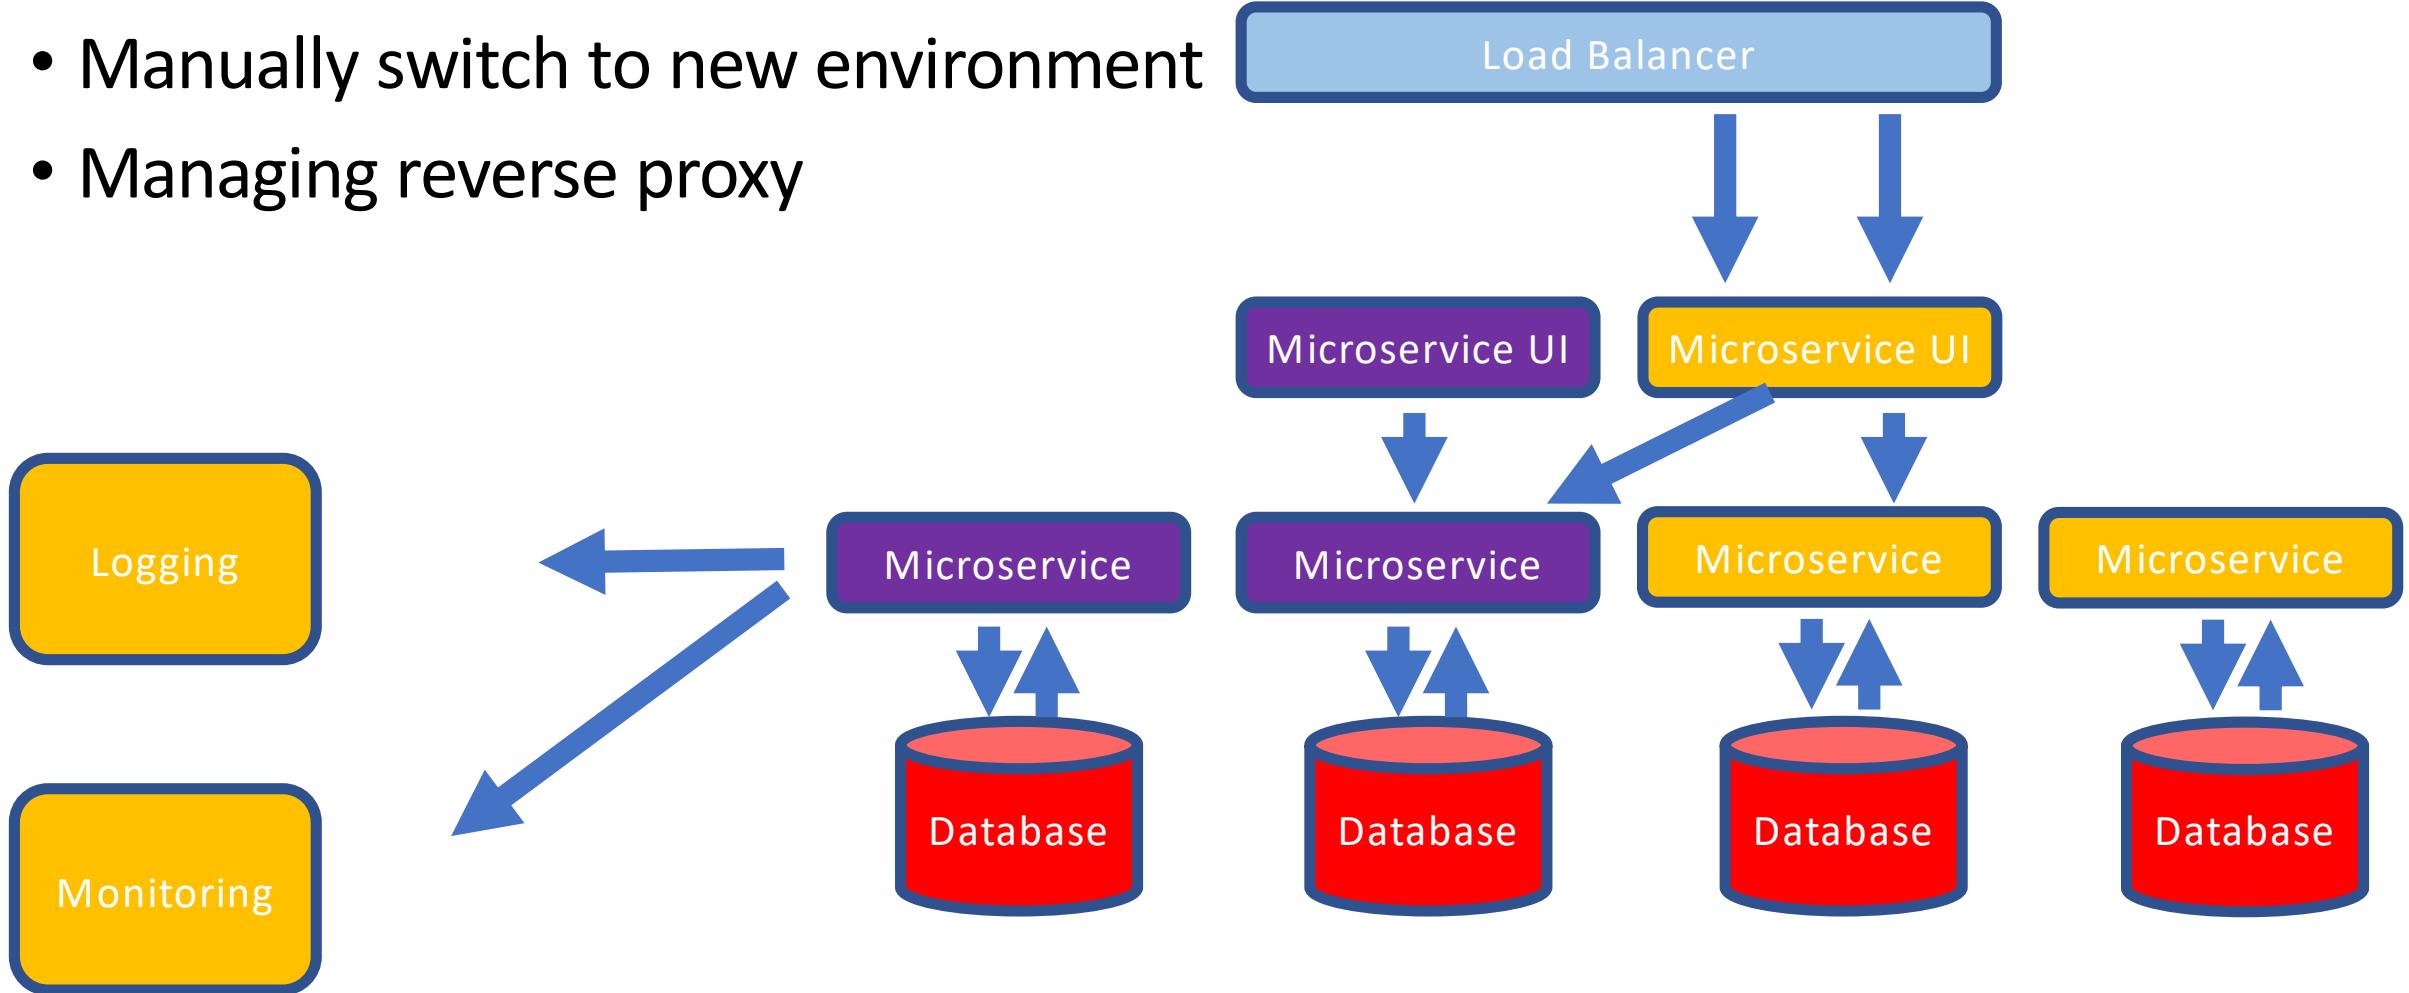
Canary and A/B Testing

- Manually switch to new environment
- Managing reverse proxy



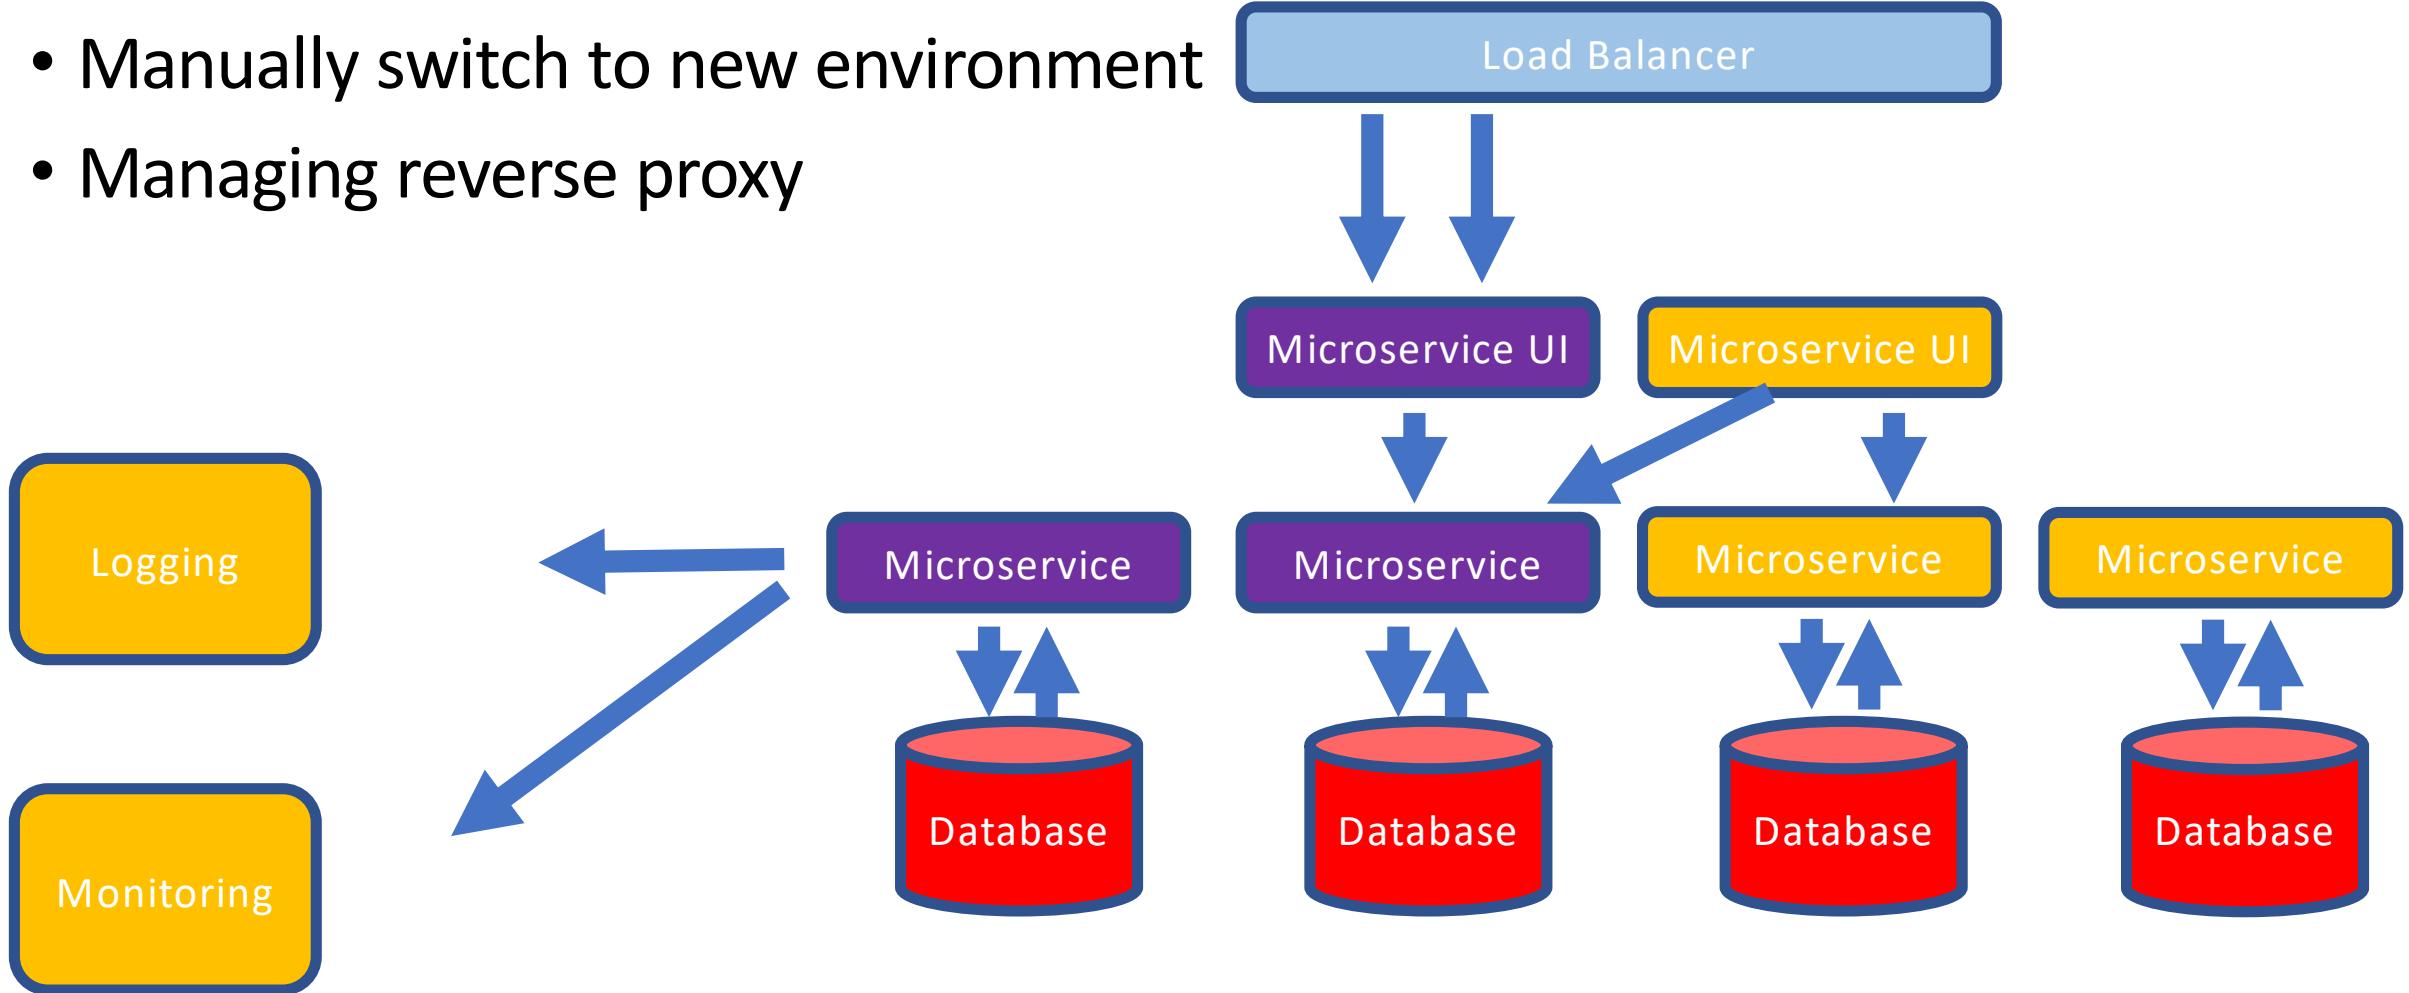
Canary and A/B Testing

- Manually switch to new environment
- Managing reverse proxy



Canary and A/B Testing

- Manually switch to new environment
- Managing reverse proxy



The Docker Platform Components



Docker Engine

Docker Engine

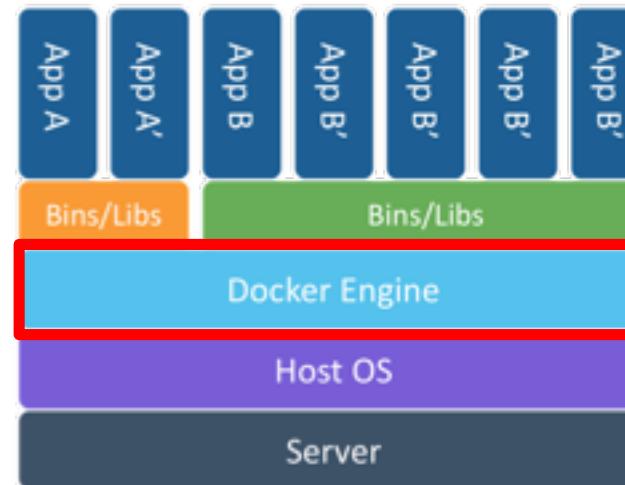
Docker Registry

Docker Compose

Docker Swarm

Lightweight runtime program to build, ship, and run Docker containers

- Also known as **Docker Daemon**
- Uses Linux Kernel namespaces and control groups
 - **Linux Kernel (>= 3.10)**
- Namespaces provide an isolated workspace



Docker Engine

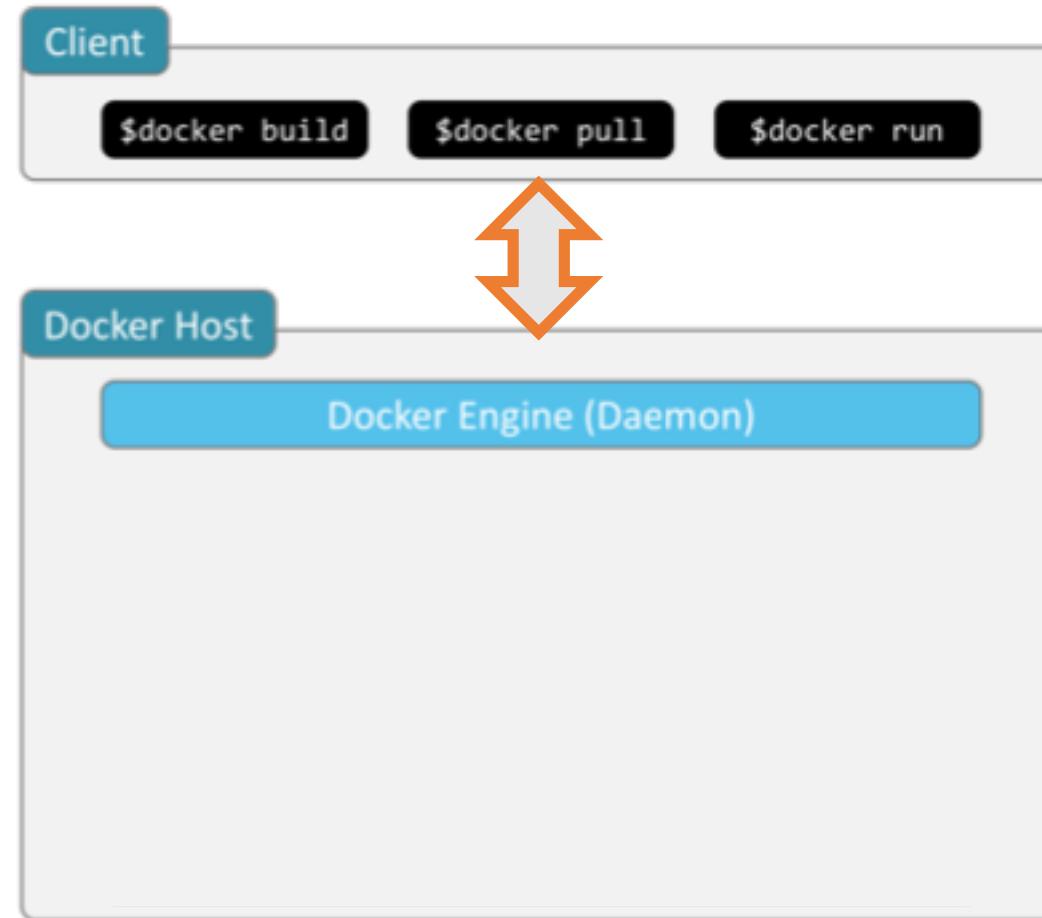
Docker Engine

Docker Registry

Docker Compose

Docker Swarm

- The Docker Client is the `docker` binary
 - Primary interface to the Docker Host
 - Accepts commands and communicates with the Docker Engine (Daemon)



Docker Engine

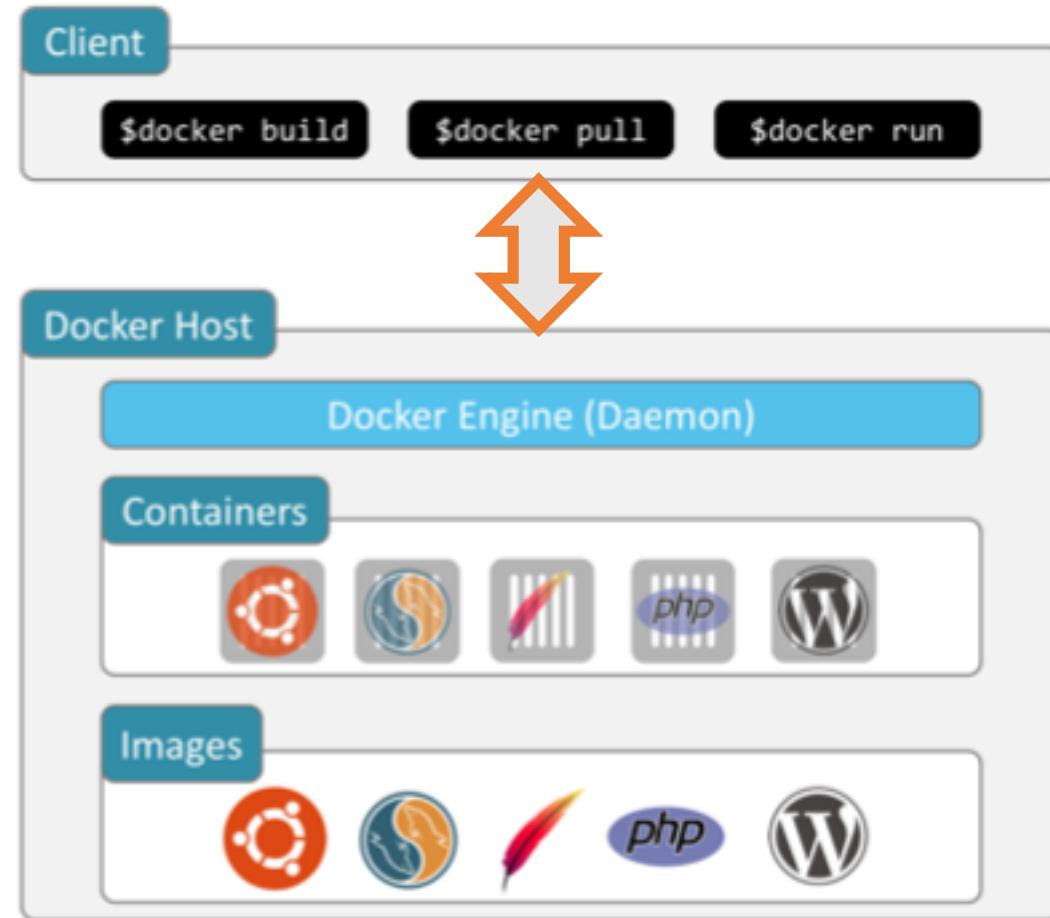
Docker Engine

Docker Registry

Docker Compose

Docker Swarm

- Lives on a Docker host
- Creates and manages containers on the host



Docker Registry

Docker Engine

Docker Registry

Docker Compose

Docker Swarm

Image Storage & Retrieval System

- Docker Engine **Pushes** Images to a Registry
- Version Control
- Docker Engine **Pulls** Images to Run



Docker Registry

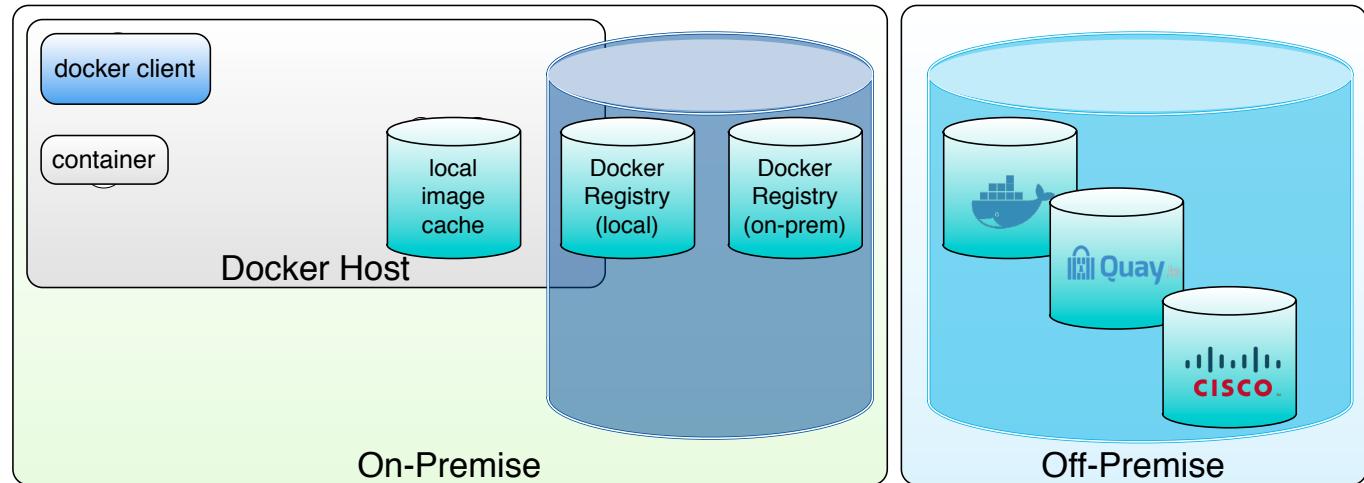
Docker Engine

Docker Registry

Docker Compose

Docker Swarm

- **Types of Docker Registries**
 - Local Docker Registry (On Docker Host)
 - Remote Docker Registry (On-Premise/Off-Premise)
 - Docker Hub (Off-Premise)



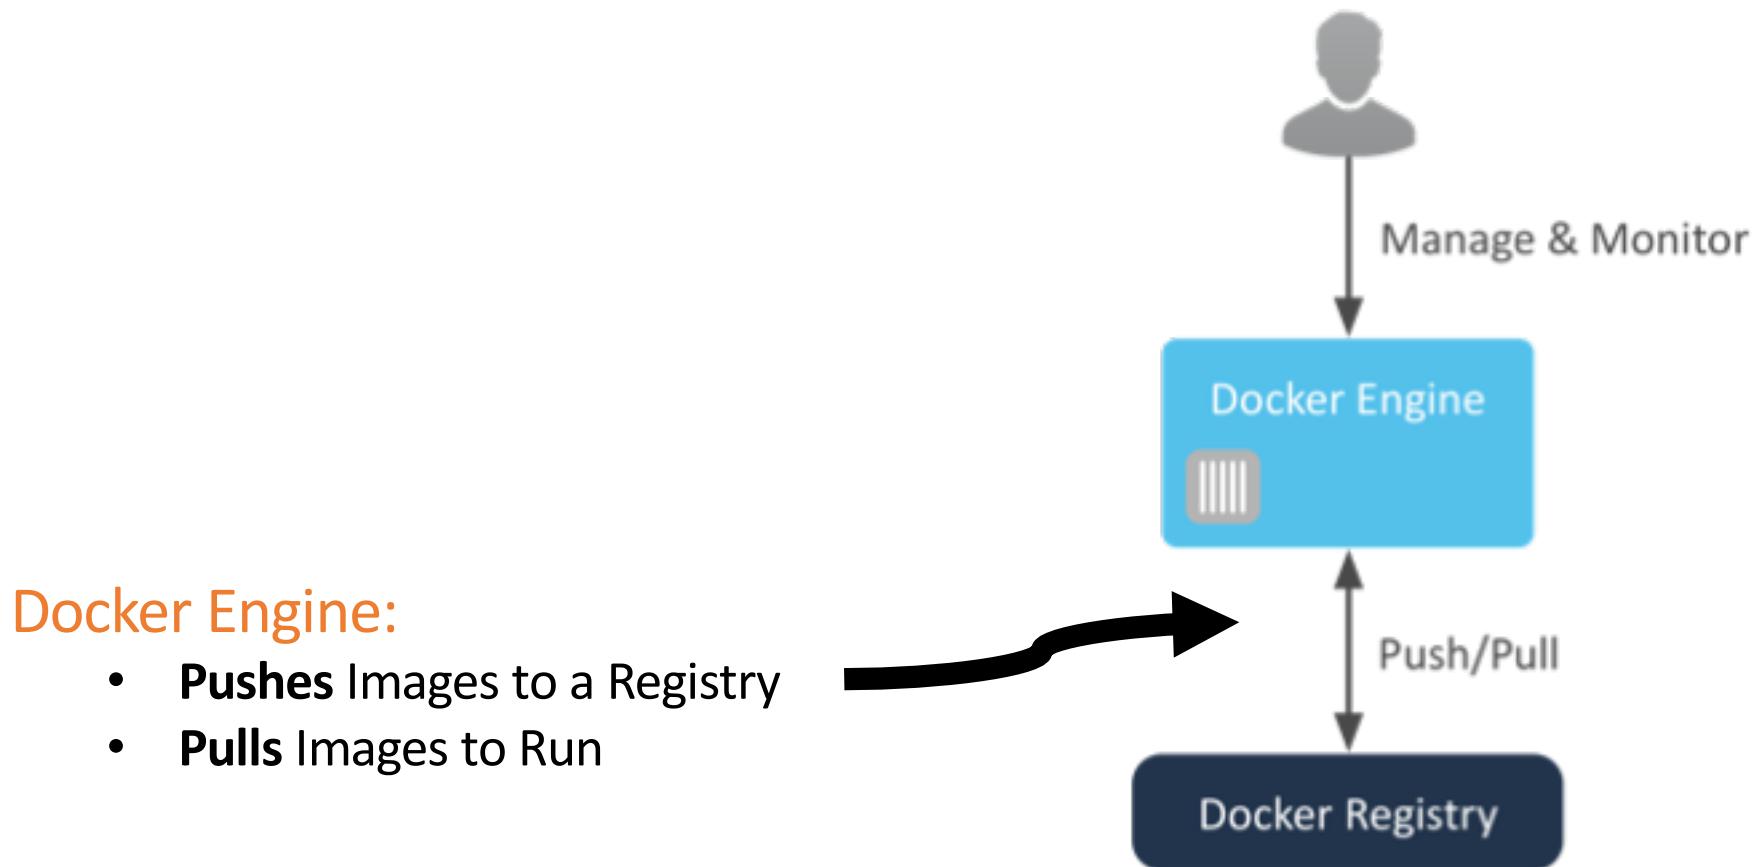
Docker Engine and Registry

Docker Engine

Docker Registry

Docker Compose

Docker Swarm



Docker Registry

Docker Engine

Docker Registry

Docker Compose

Docker Swarm

The registry and engine both present APIs

- All of Docker's functionality will utilize these APIs
- RESTFUL API
- Commands presented with Docker's CLI tools can also be used with curl and other tools

Docker Compose

Docker Engine

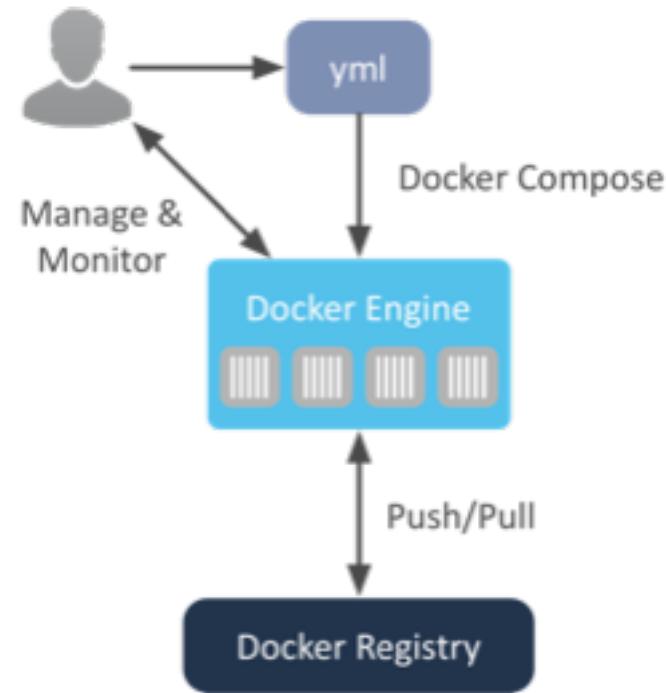
Docker Registry

Docker Compose

Docker Swarm

Tool to create and manage multi-container applications

- Applications defined in a single file:
- **docker-compose.yml**
- Transforms applications into individual containers that are linked together
- Compose will start all containers in a single command



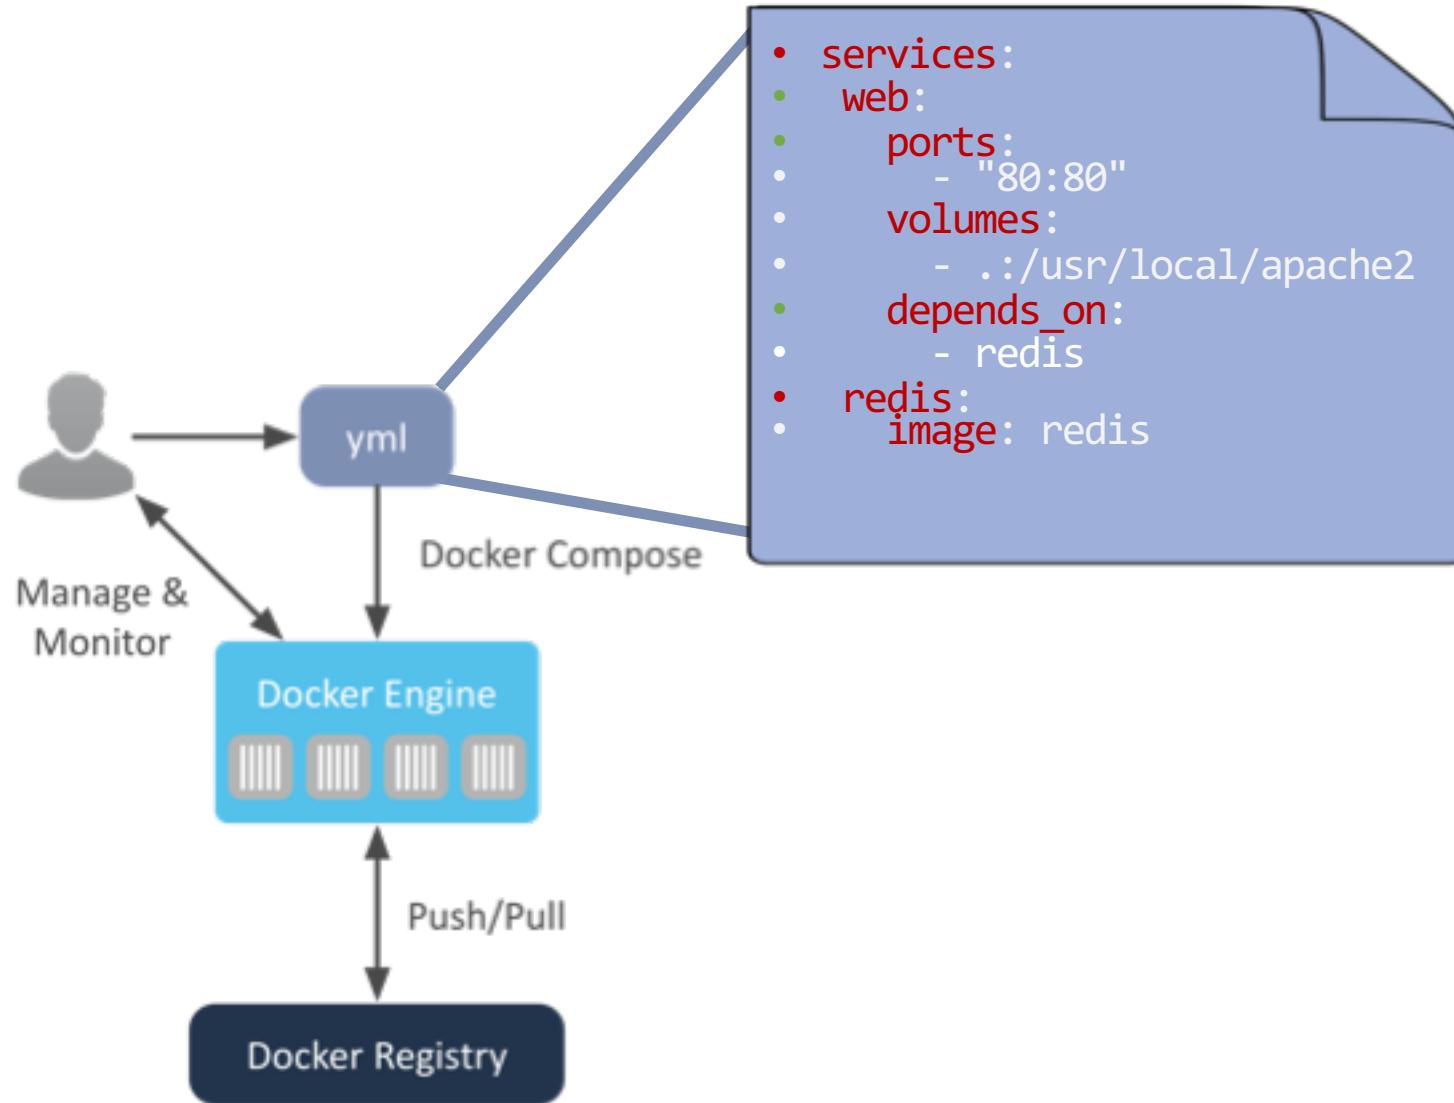
Docker Compose

Docker Engine

Docker Registry

Docker Compose

Docker Swarm



Docker Swarm

Docker Engine

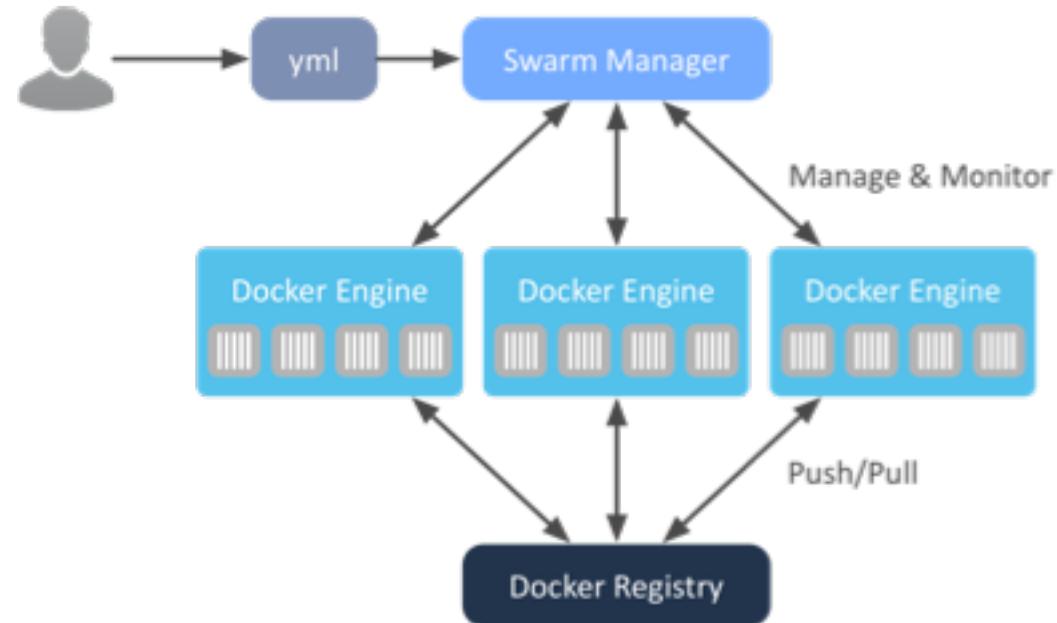
Docker Registry

Docker Compose

Docker Swarm

Clusters Docker hosts and schedules containers

- Native Clustering for Docker
 - Turn a pool of Docker hosts into a single, virtual host
 - Serves the standard Docker API



Labs:

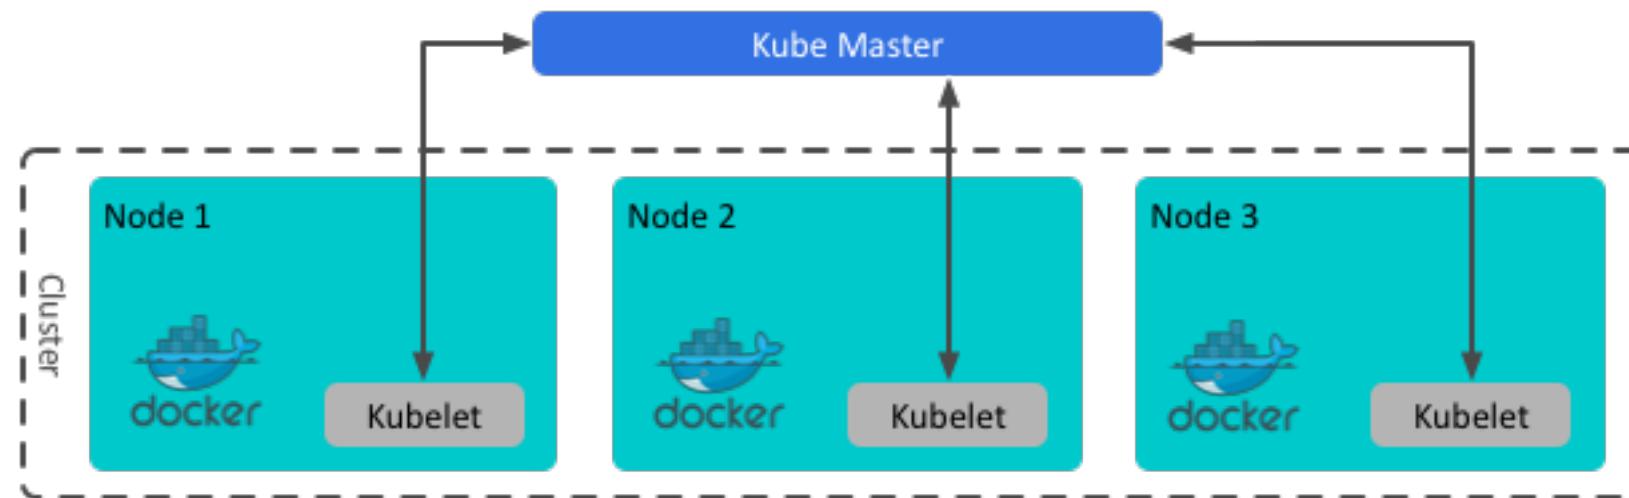
1. Manually build Docker images
2. Package Microservices in containers

Overview of Kubernetes Clusters



Kubernetes Clusters

- Kubernetes deployed on a set of physical or virtual hosts – K8s nodes
- Hosts run host OS that supports Linux containers, e.g. Docker or rkt hosts
- Kubernetes runs well in both private and public IaaS environments
- Users and admins control Kubernetes resources via REST API on K8s master



Kubernetes Components: Master

Main Components:

Master Node:

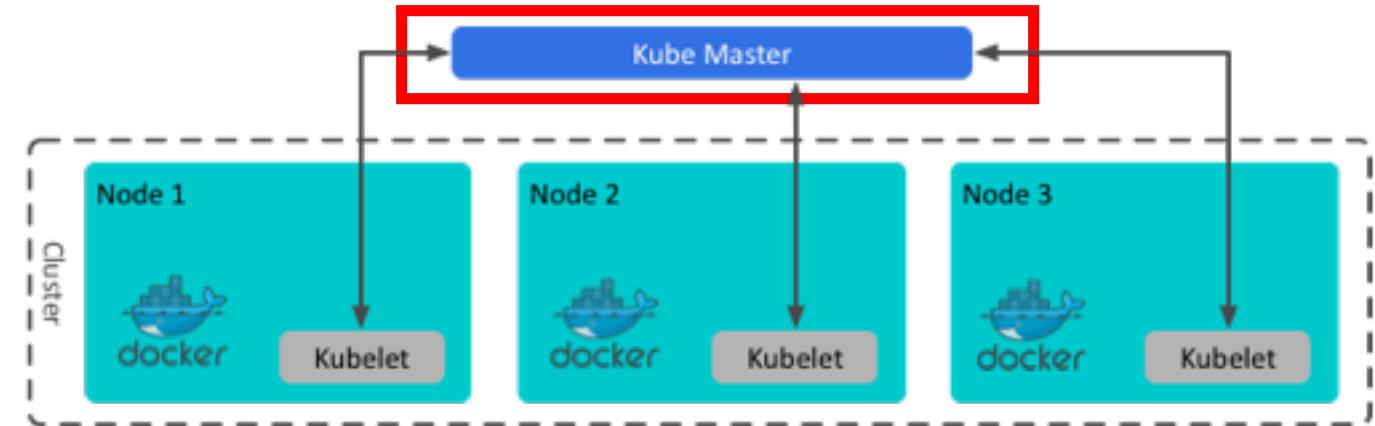
- This node manages and oversees other nodes and hosts the K8s API, scheduler, and controllers

Worker Nodes:

- Run user workloads as directed by the K8s master
- The master may also serve as a worker node

Clusters:

- A collection of nodes bound to a Master and managed as a single logical unit of capacity



Kubernetes Components: Nodes

Main Components:

Master Node:

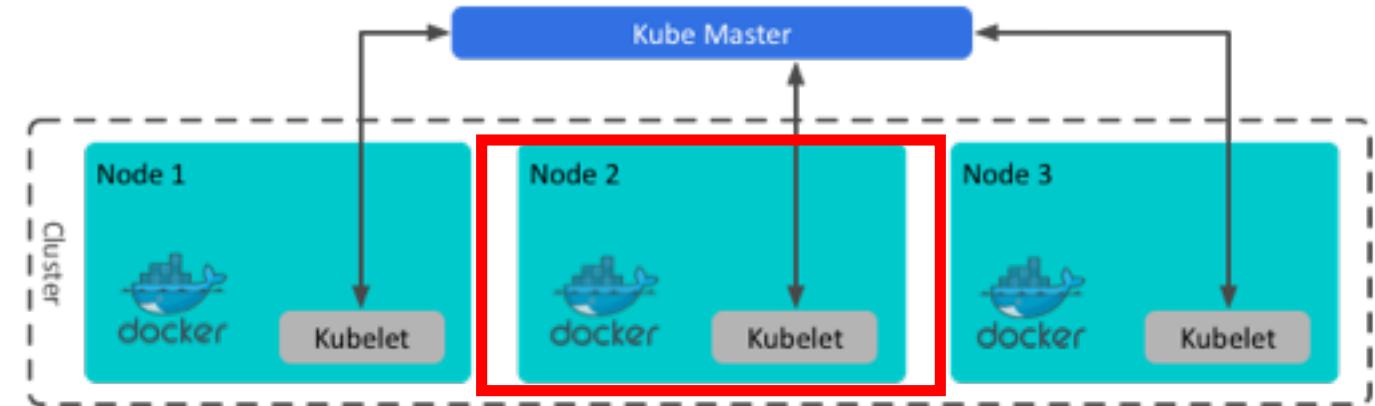
- This node manages and oversees other nodes and hosts the K8s API, scheduler, and controllers

Worker Nodes:

- Run user workloads as directed by the K8s master
- The master may also serve as a worker node

Clusters:

- A collection of nodes bound to a Master and managed as a single logical unit of capacity



Kubernetes Cluster Components

Main Components:

Master Node:

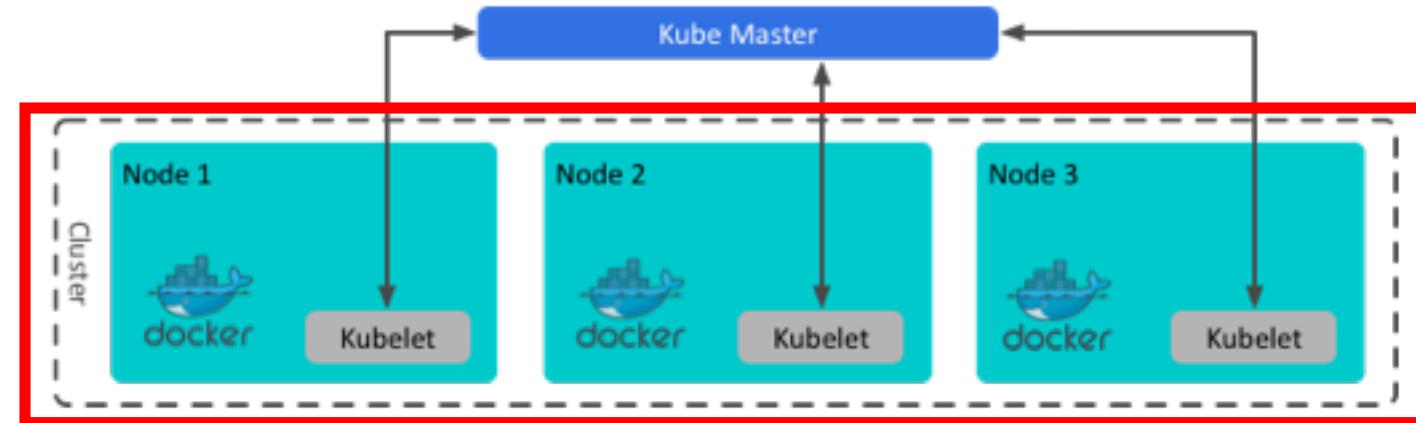
- This node manages and oversees other nodes and hosts the K8s API, scheduler, and controllers

Worker Nodes:

- Run user workloads as directed by the K8s master
- The master may also serve as a worker node

Clusters:

- A collection of nodes bound to a Master and managed as a single logical unit of capacity

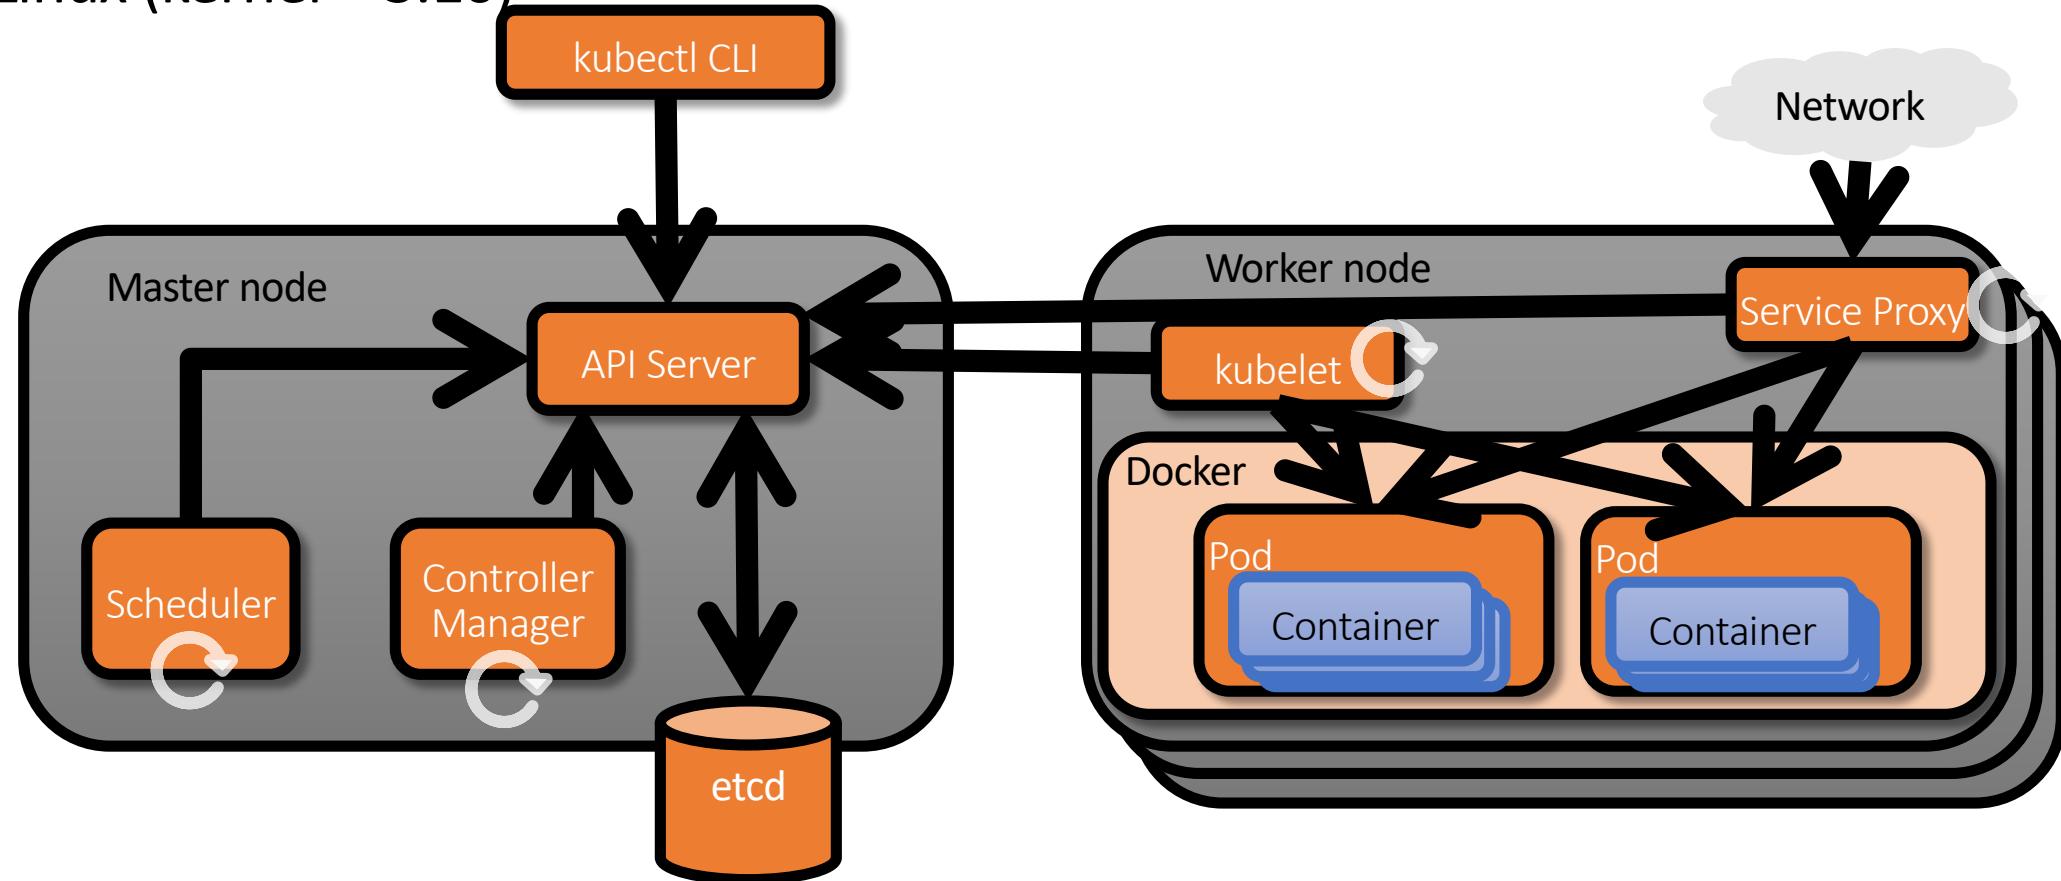


Kubernetes Architecture

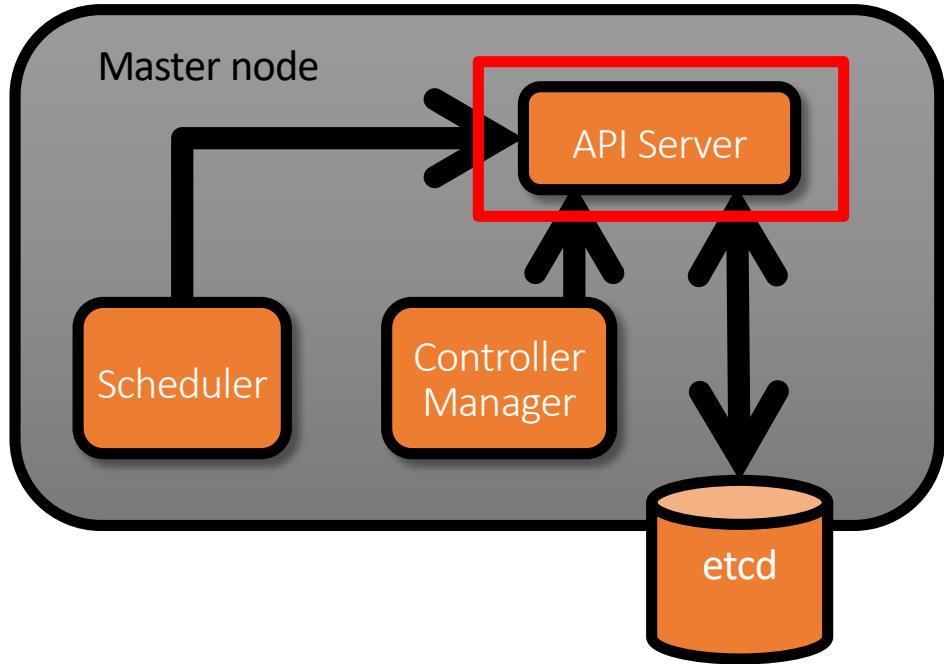


Kubernetes Cluster Architecture

- Kubernetes nodes can be physical hosts or VM's running a container-friendly Linux (kernel > 3.10)



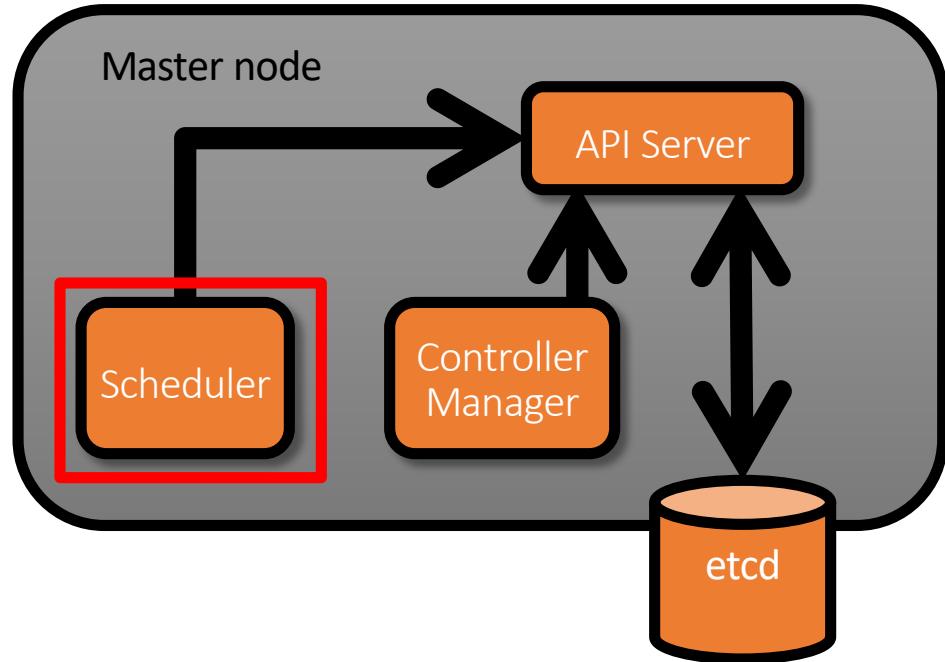
Kubernetes Master Node Components



K8s components
written in Go
(golang.org)

- **API Server (`kube-apiserver`)**: exposes the Kubernetes REST API, and can be scaled horizontally
- **Scheduler (`kube-scheduler`)**: selects nodes for newly created pods to run on
- **Controller manager (`kube-controller-manager`)**: runs background controller processes for the system to enforce declared object states, e.g. Node Controller, Replication Controller, ...
- **Persistent data store (`etcd`)**: all K8s system data is stored in a distributed, reliable key-value store. etcd may run on separate nodes from the master

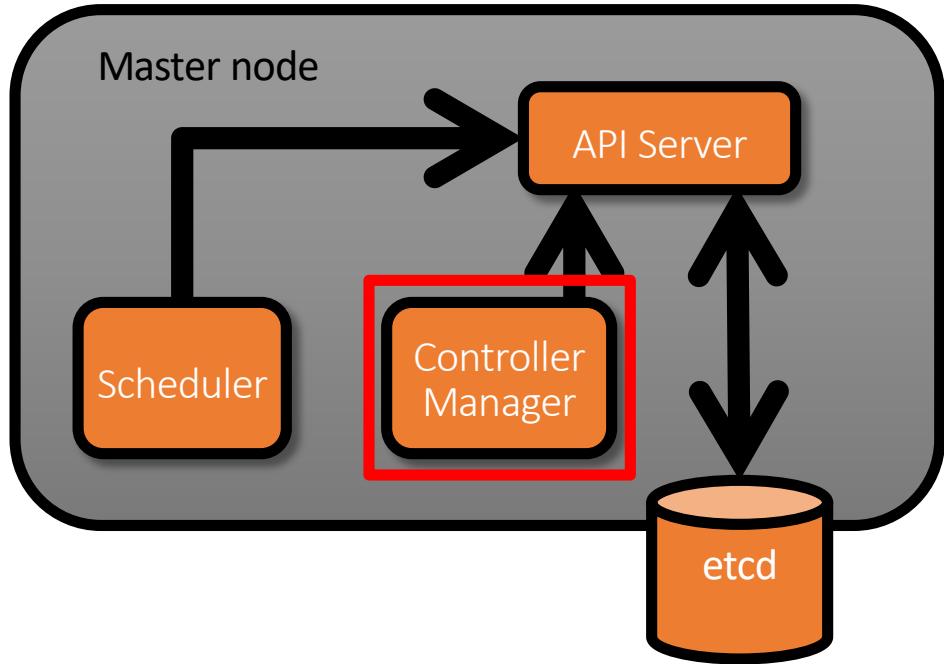
Kubernetes Master Node Components



K8s components
written in Go
(golang.org)

- **API Server (`kube-apiserver`)**: exposes the Kubernetes REST API, and can be scaled horizontally
- **Scheduler (`kube-scheduler`)**: selects nodes for newly created pods to run on
- **Controller manager (`kube-controller-manager`)**: runs background controller processes for the system to enforce declared object states, e.g. Node Controller, Replication Controller, ...
- **Persistent data store (`etcd`)**: all K8s system data is stored in a distributed, reliable key-value store. etcd may run on separate nodes from the master

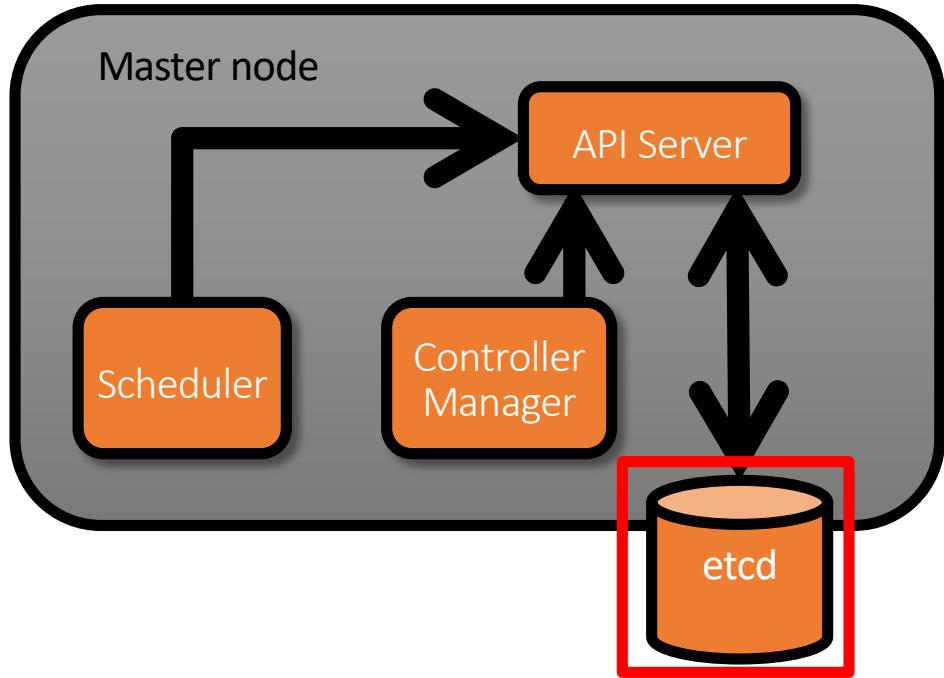
Kubernetes Master Node Components



K8s components
written in Go
(golang.org)

- **API Server (`kube-apiserver`)**: exposes the Kubernetes REST API, and can be scaled horizontally
- **Scheduler (`kube-scheduler`)**: selects nodes for newly created pods to run on
- **Controller manager (`kube-controller-manager`)**: runs background controller processes for the system to enforce declared object states, e.g. Node Controller, Replication Controller, ...
- **Persistent data store (`etcd`)**: all K8s system data is stored in a distributed, reliable key-value store. etcd may run on separate nodes from the master

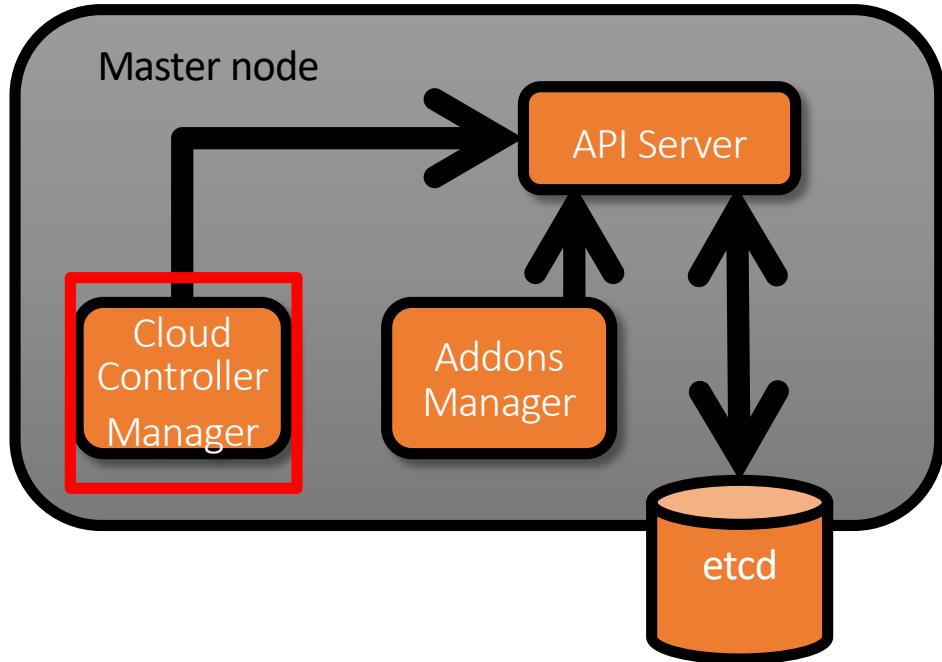
Kubernetes Master Node Components



K8s components
written in Go
(golang.org)

- **API Server (`kube-apiserver`)**: exposes the Kubernetes REST API, and can be scaled horizontally
- **Scheduler (`kube-scheduler`)**: selects nodes for newly created pods to run on
- **Controller manager (`kube-controller-manager`)**: runs background controller processes for the system to enforce declared object states, e.g. Node Controller, Replication Controller, ...
- **Persistent data store (`etcd`)**: all K8s system data is stored in a distributed, reliable key-value store. etcd may run on separate nodes from the master

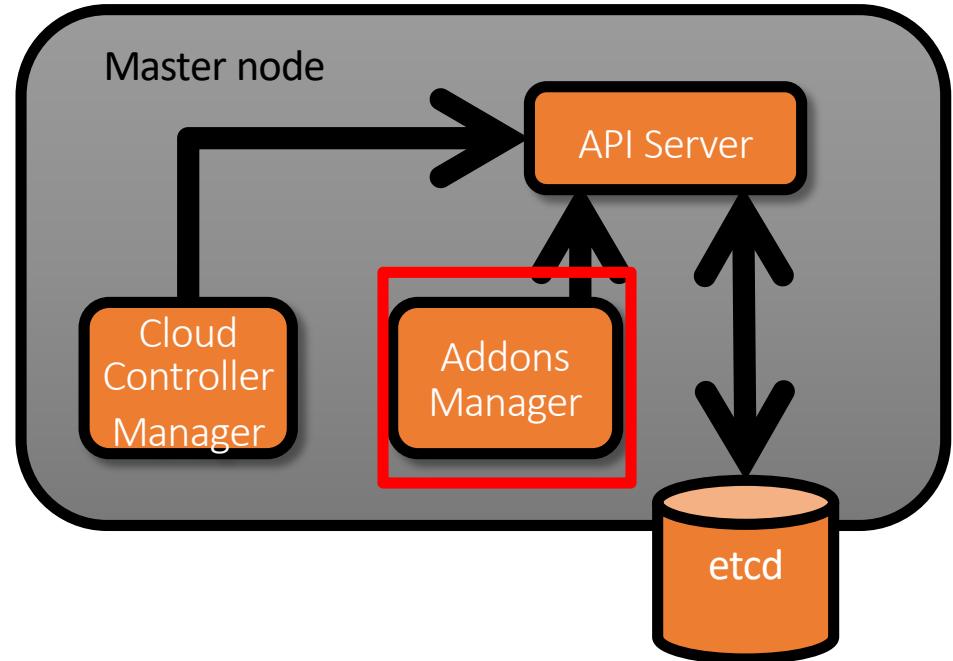
Master Node Additional Components



K8s components
written in Go
(golang.org)

- **cloud-controller-manager**: runs controllers interacting with underlying IaaS providers – alpha feature
 - Allows cloud vendor-specific code to be separate from main K8s system components
- **addons-manager**: creates and maintains cluster addon resources in ‘kube-system’ namespace, e.g.
- **Kubernetes Dashboard**: general web UI for application and cluster management
- **kube-dns**: serves DNS records for K8s services and resources
- Container resource monitoring and cluster-level logging

Master Node Additional Components

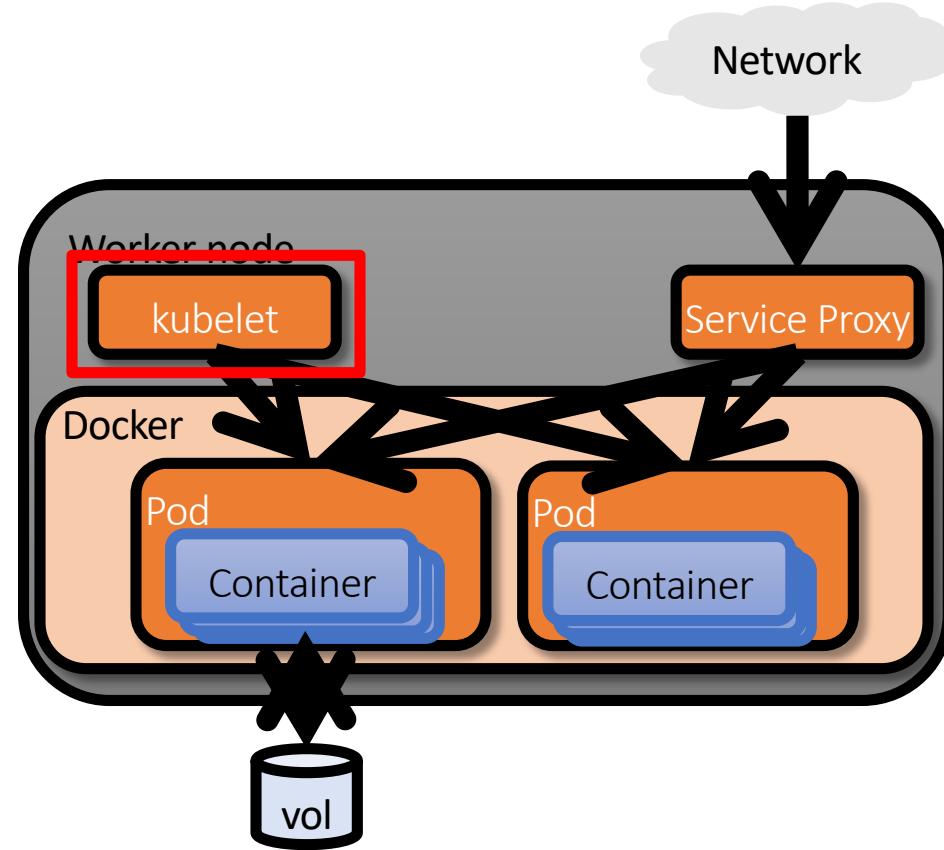


K8s components
written in Go
(golang.org)

- **cloud-controller-manager**: runs controllers interacting with underlying IaaS providers – alpha feature
 - Allows cloud vendor-specific code to be separate from main K8s system components
- **addons-manager**: creates and maintains cluster addon resources in ‘kube-system’ namespace, e.g.
 - **Kubernetes Dashboard**: general web UI for application and cluster management
 - **kube-dns**: serves DNS records for K8s services and resources
 - Container resource monitoring and cluster-level logging

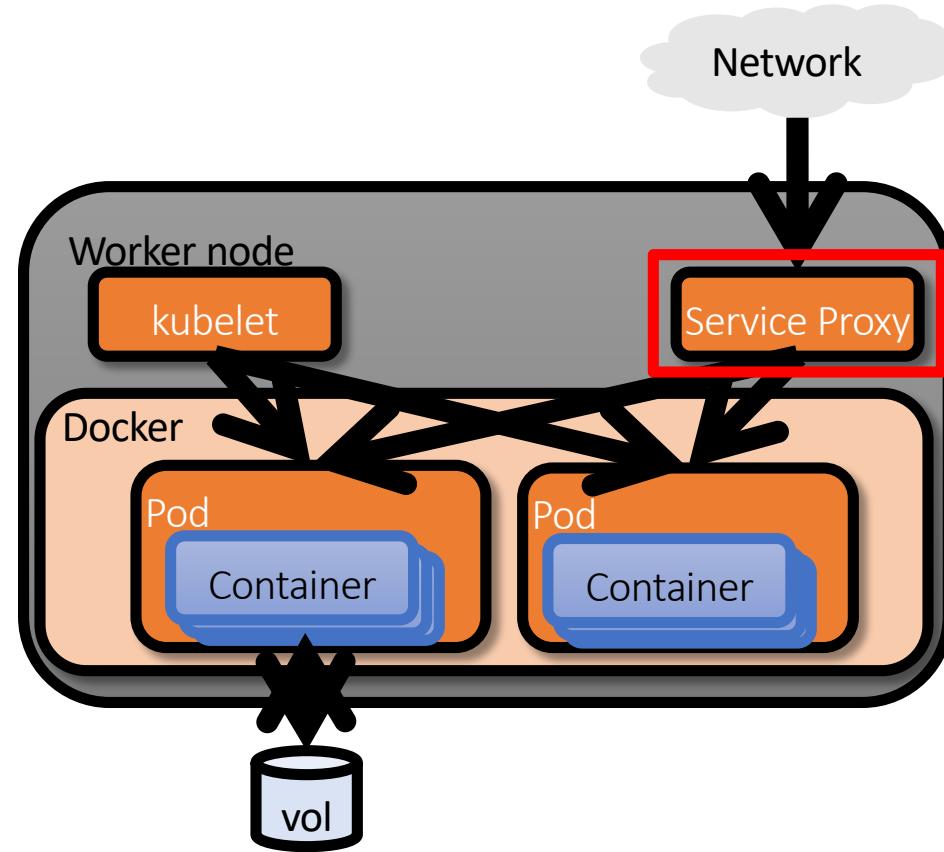
Kubernetes Worker Node Components

- **kubelet**: local K8s agent that is responsible for operations on the node, including
 - Watching for pod assignments
 - Mounting pod required volumes
 - Running a pod's containers
 - Executing container liveness probes
 - Reporting pod status to system
 - Reporting node status to system
- **Service proxy (kube-proxy)**: enables K8s service abstractions by maintaining host network rules and forwarding connections
- **Docker**: runs the containers



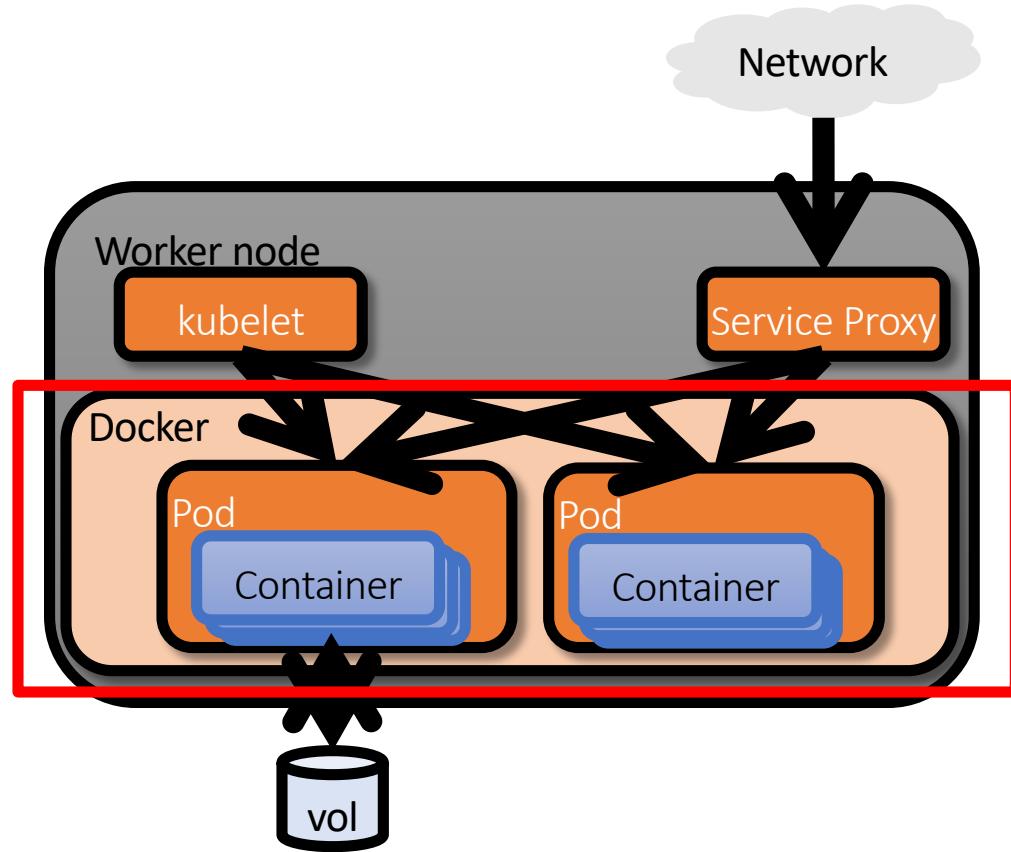
Kubernetes Worker Node Components

- **kubelet**: local K8s agent that is responsible for operations on the node, including
 - Watching for pod assignments
 - Mounting pod required volumes
 - Running a pod's containers
 - Executing container liveness probes
 - Reporting pod status to system
 - Reporting node status to system
- **Service proxy (kube-proxy)**: enables K8s service abstractions by maintaining host network rules and forwarding connections
- **Docker**: runs the containers



Kubernetes Worker Node Components

- **kubelet:** local K8s agent that is responsible for operations on the node, including
 - Watching for pod assignments
 - Mounting pod required volumes
 - Running a pod's containers
 - Executing container liveness probes
 - Reporting pod status to system
 - Reporting node status to system
- **Service proxy (kube-proxy):** enables K8s service abstractions by maintaining host network rules and forwarding connections
- **Docker:** container runtime

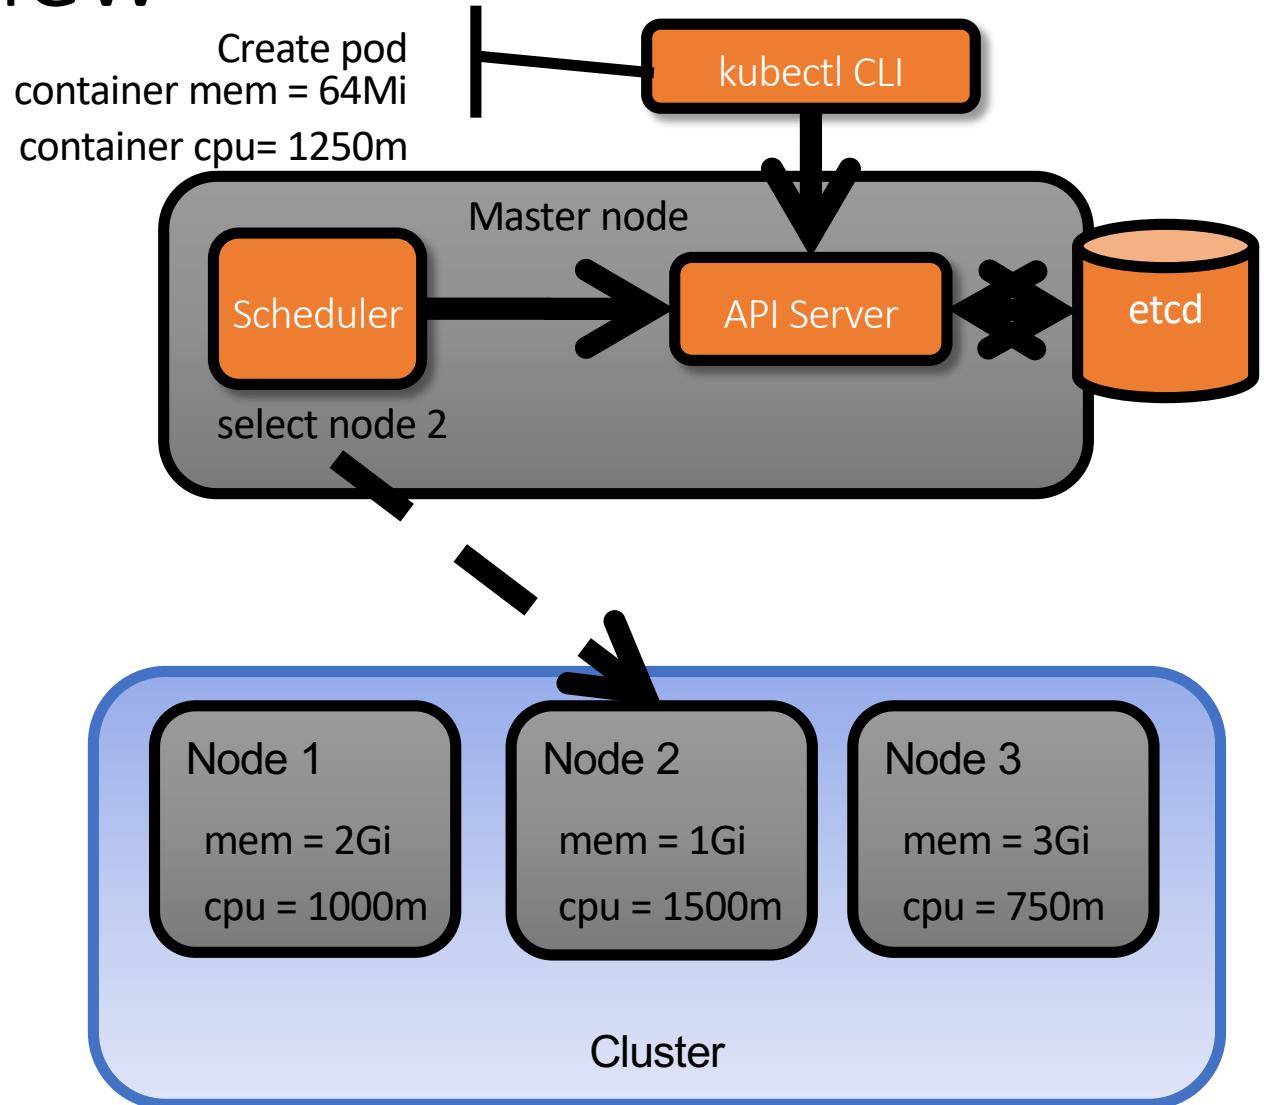


Kubernetes Pod Scheduling



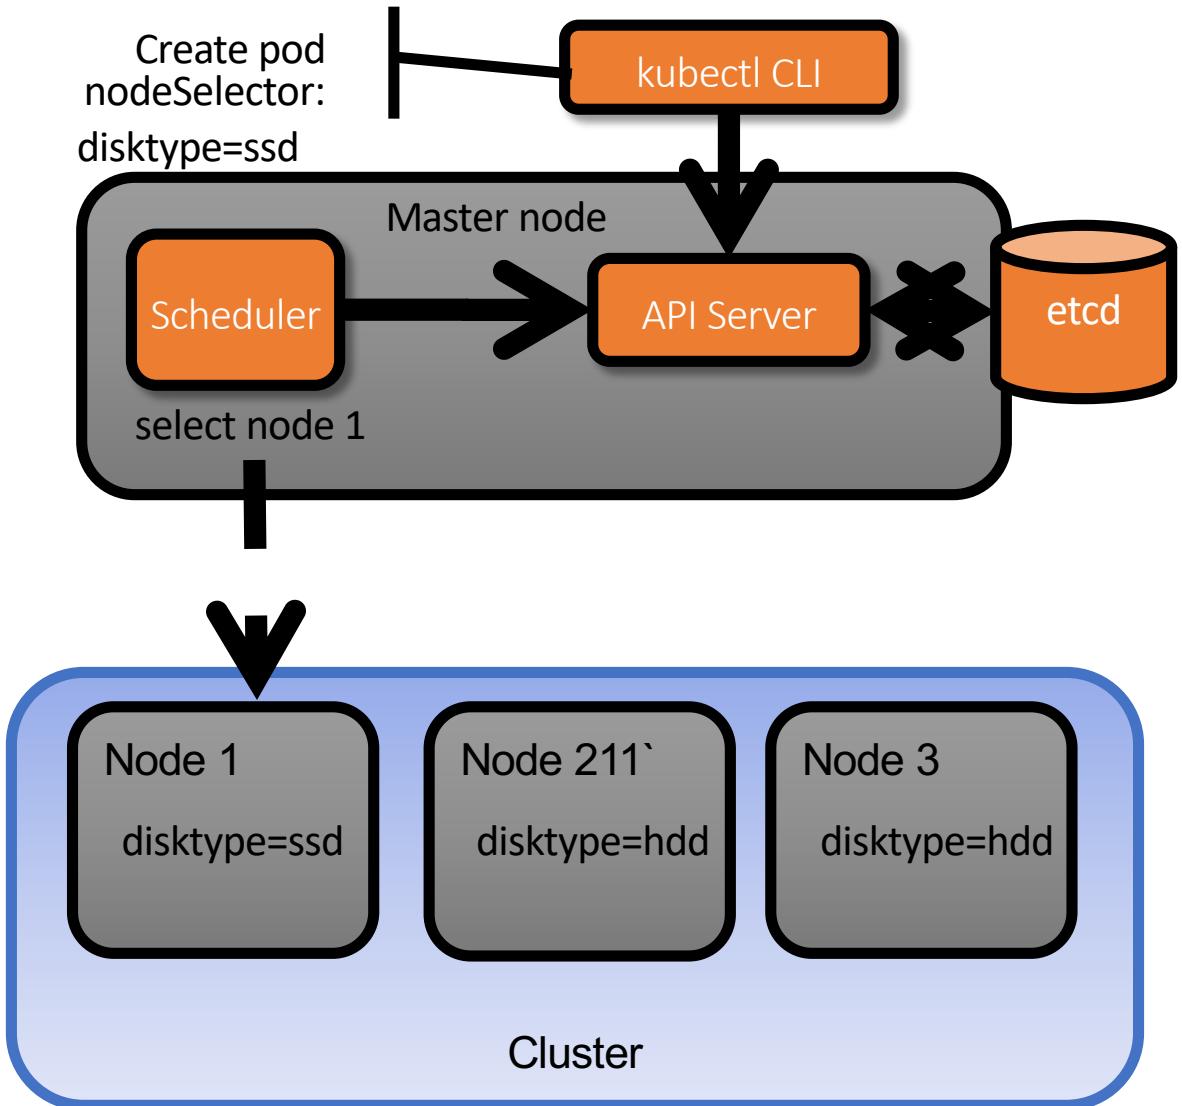
Pod Scheduling Overview

- kube-scheduler on the master node assigns new pod requests to appropriate worker nodes
- Default scheduler takes account of
 - Available node CPU/RAM
 - Resource requests from new pod – sum of resource requests of pod containers
- Default scheduler will automatically
 - Schedule pod on node with sufficient free resources
 - Spread pods across nodes in the cluster
- Can specify custom schedulers for pods



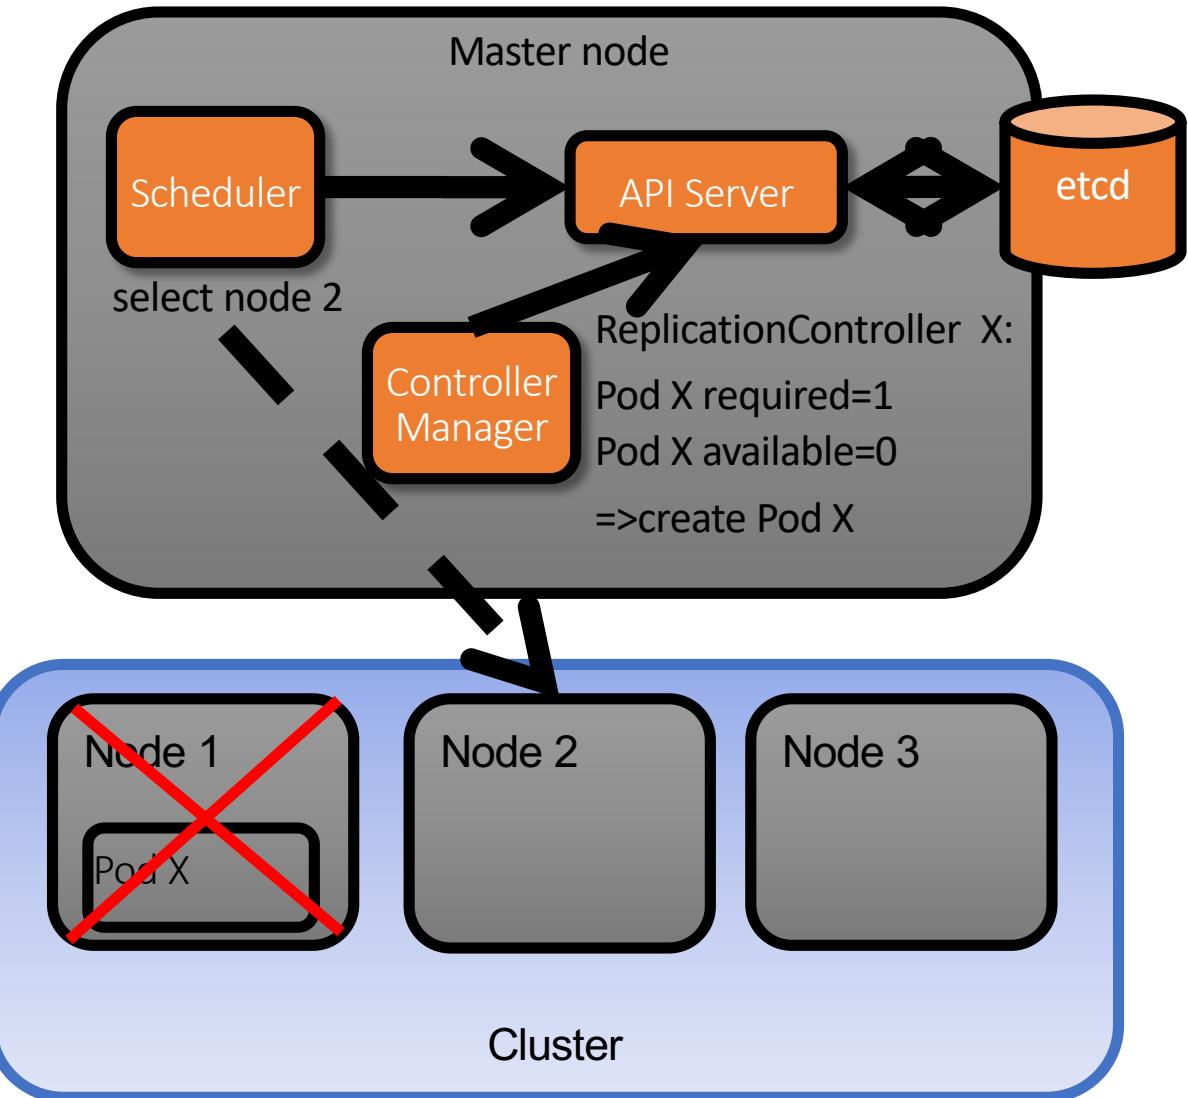
Pod Scheduling with Default Scheduler

- Nodes carry labels indicating topology and other resource notes
- Users can require pods to be scheduled on nodes with specific label(s) via a nodeSelector in container spec
- K8s 1.6 adds more advanced scheduling options, including
 - Node affinity/anti-affinity generalizing the nodeSelector feature
 - Node taints that prevent scheduling of pods to nodes unless pod ‘tolerates’ the taint
 - Pod affinity/anti-affinity to control relative placement of pods



Pod Re-creation Driven by Control Loops

- kube-scheduler performs same node selection operation when new pod created due to e.g. node loss
- kube-controller-manager runs controllers like ReplicationController managing number of pod instances available
- kube-controller-manager will initiate request for new pods as needed, which will be scheduled by kube-scheduler per pod/container spec





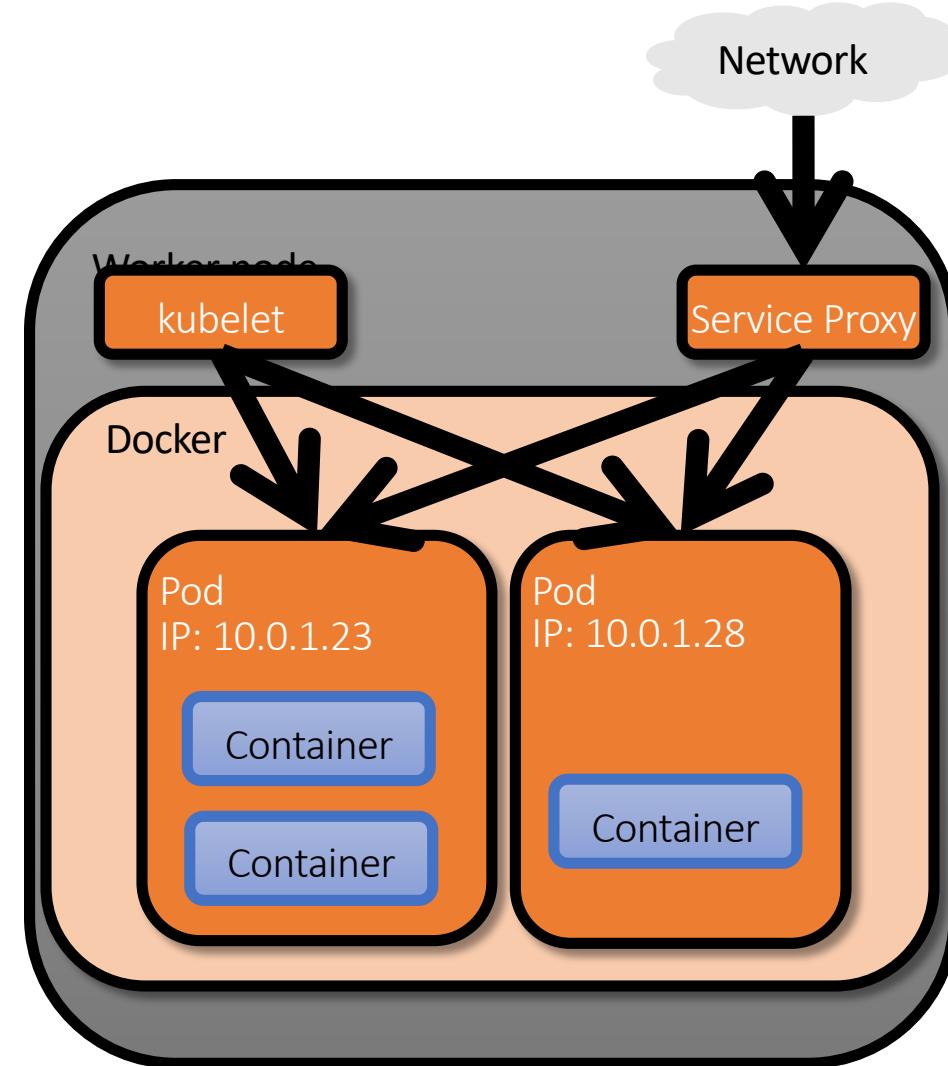
Questions

Kubernetes Network Models



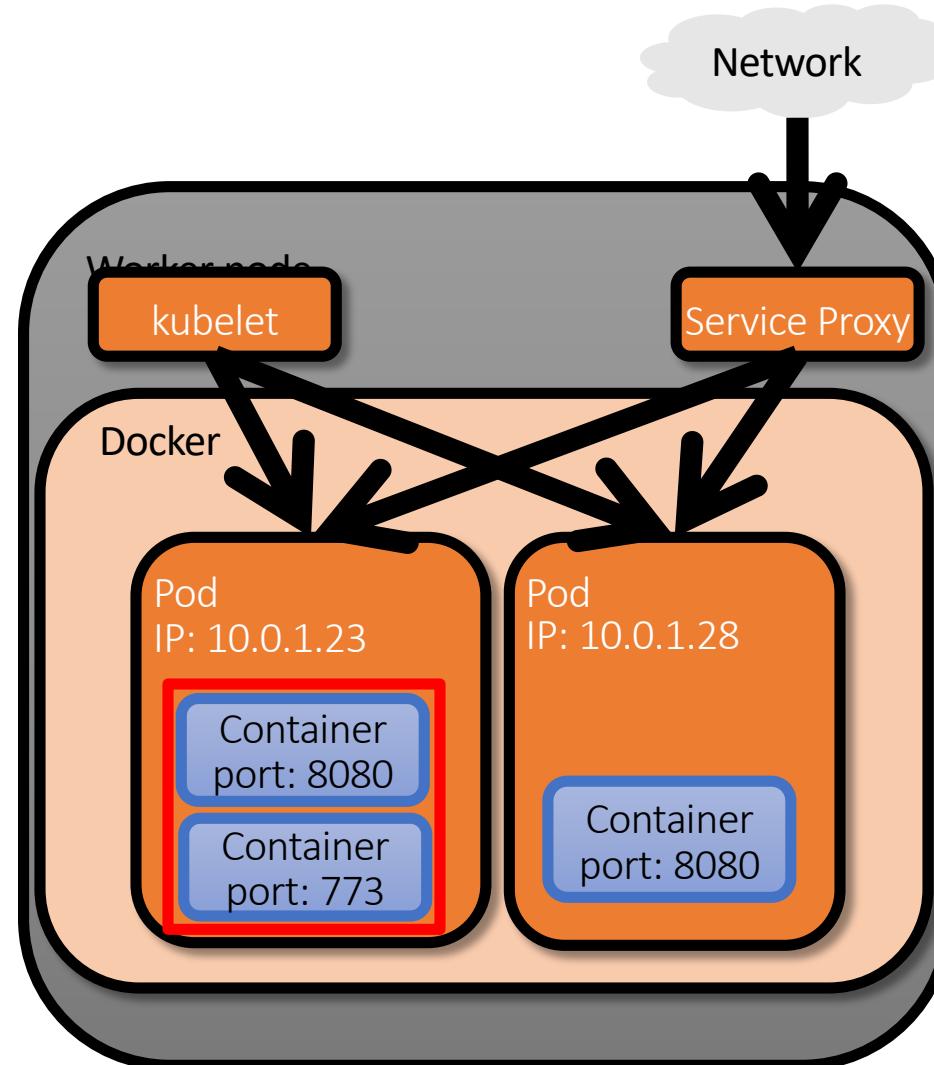
Networking in Kubernetes – One IP per Pod

- Kubernetes design goal is to simplify configuration – tracking lots of port mappings is hard!
- K8s networking principles:
 - All containers can communicate with other containers without NAT
 - All nodes can communicate with containers without NAT
 - Container sees its own IP exactly as other containers and nodes do
- K8s assigns IP to every **pod** in the cluster as part of flat shared network



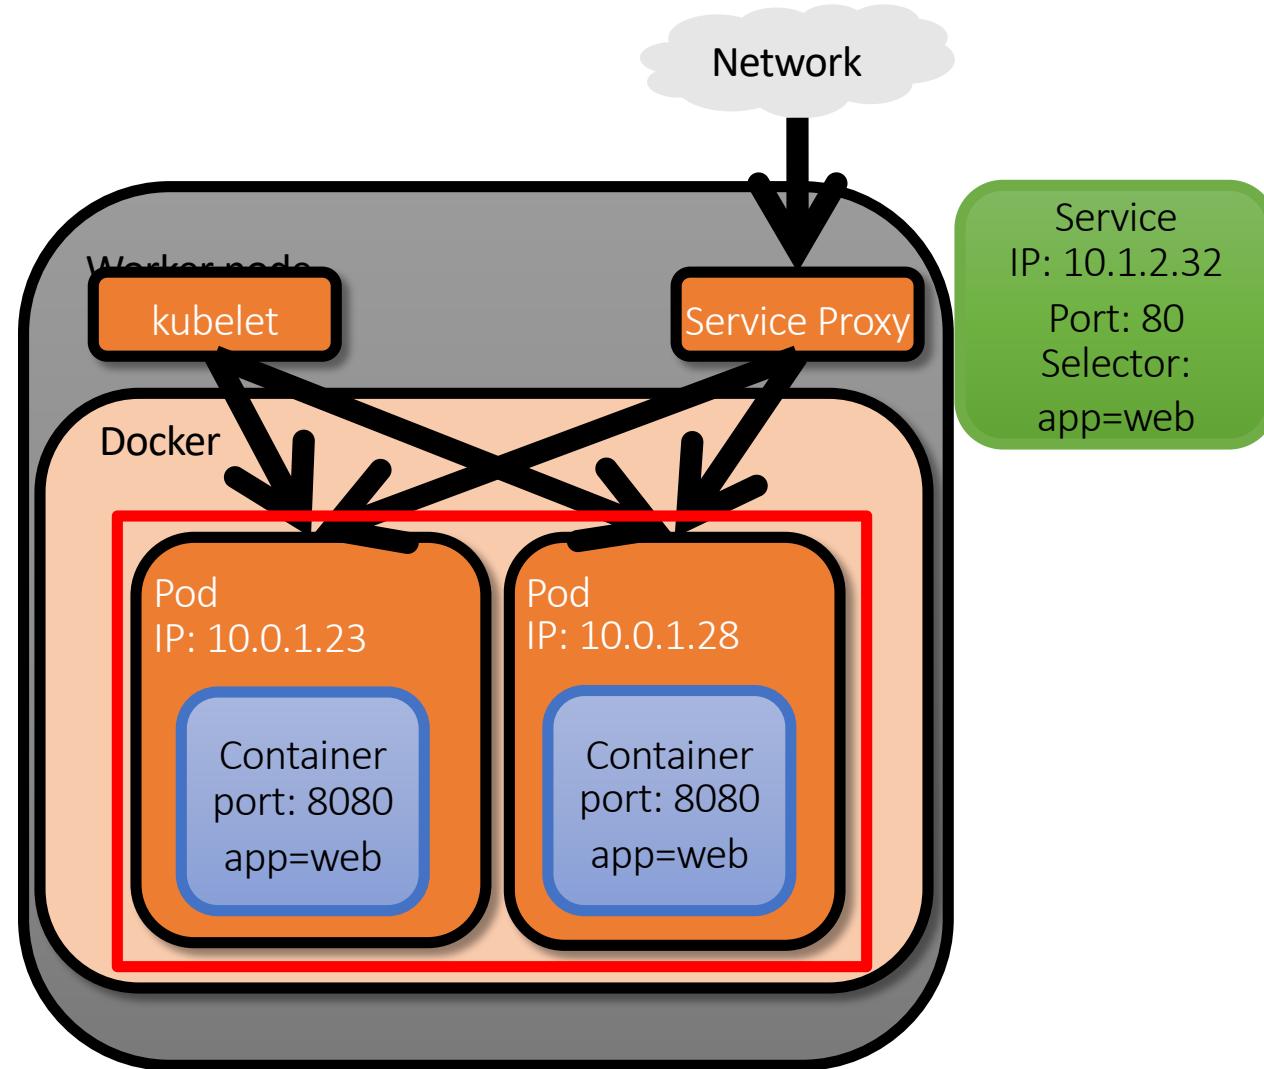
Container Communications – Same Pod

- Containers in the same pod share the same network namespaces and IP address
- Containers in same pod can reach each other's port on 'localhost'
- Need to manage ports used by co-located containers to avoid conflicts, but done in same pod spec
- Other containers in the cluster can connect to the pod's containers via their ports on the pod's IP



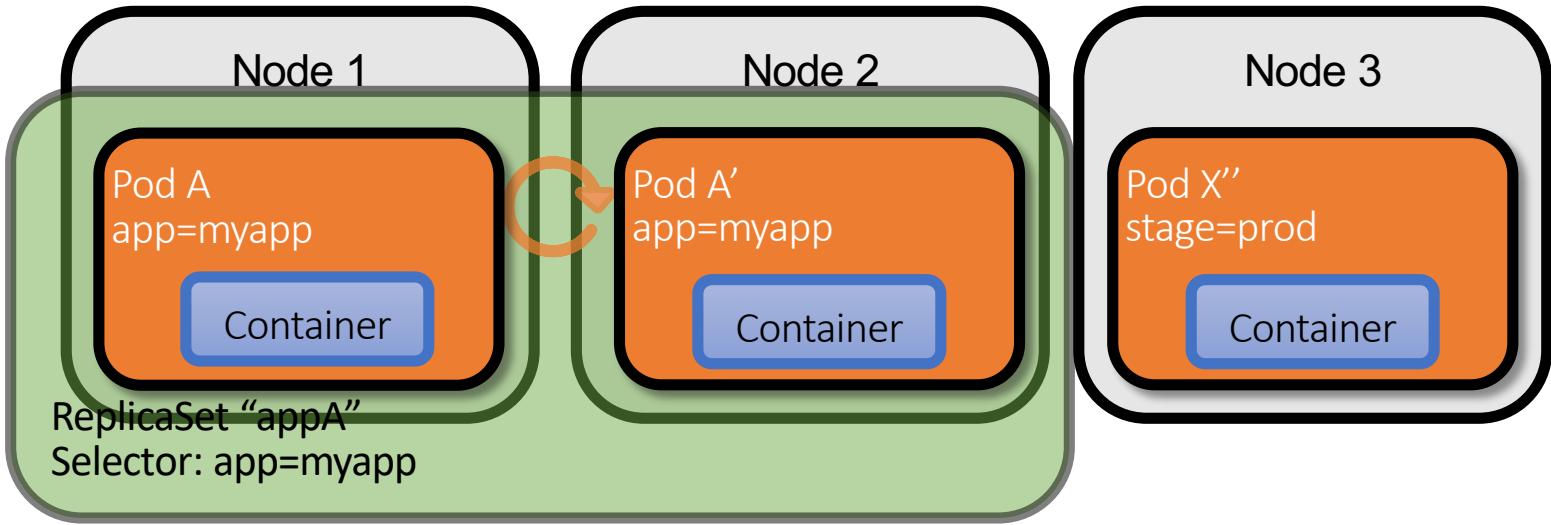
Abstraction Layer for Applications

- Pod IP's are simple but tricky to track across replicas, restarted pods, etc.
- Service objects permit allocation of a stable IP that forwards to multiple pods matching the service's selector
- kube-proxy maintains iptables rules for service objects on each worker node
 - NATs requests to a backend pod
 - Maps each service port to backend pod port

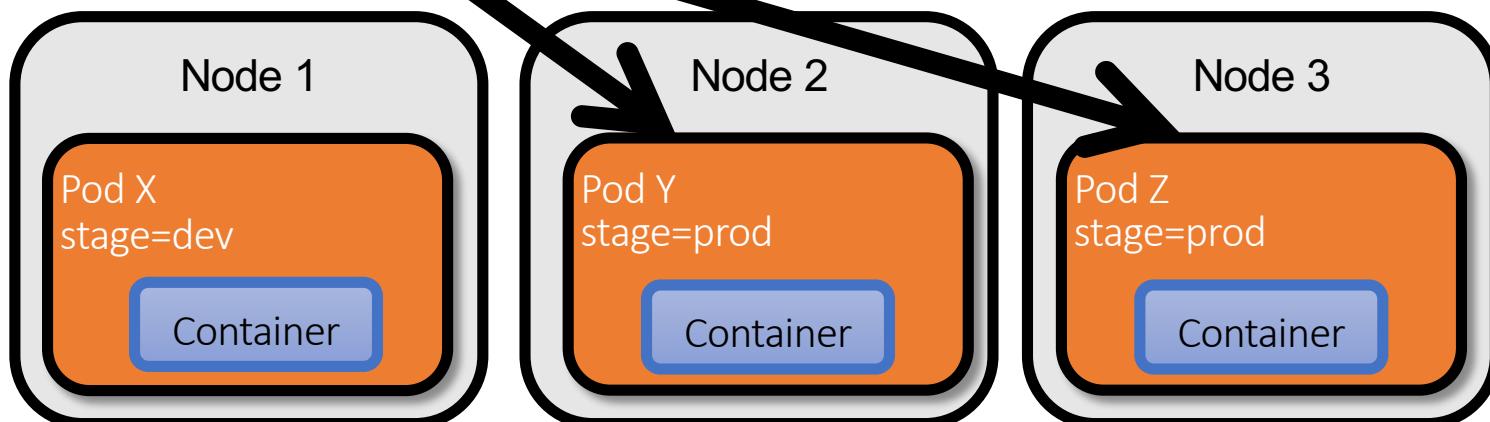


Object Grouping and Selection via Labels

Service label selector
targets Pod endpoints



Controller label selector
targets managed Pods



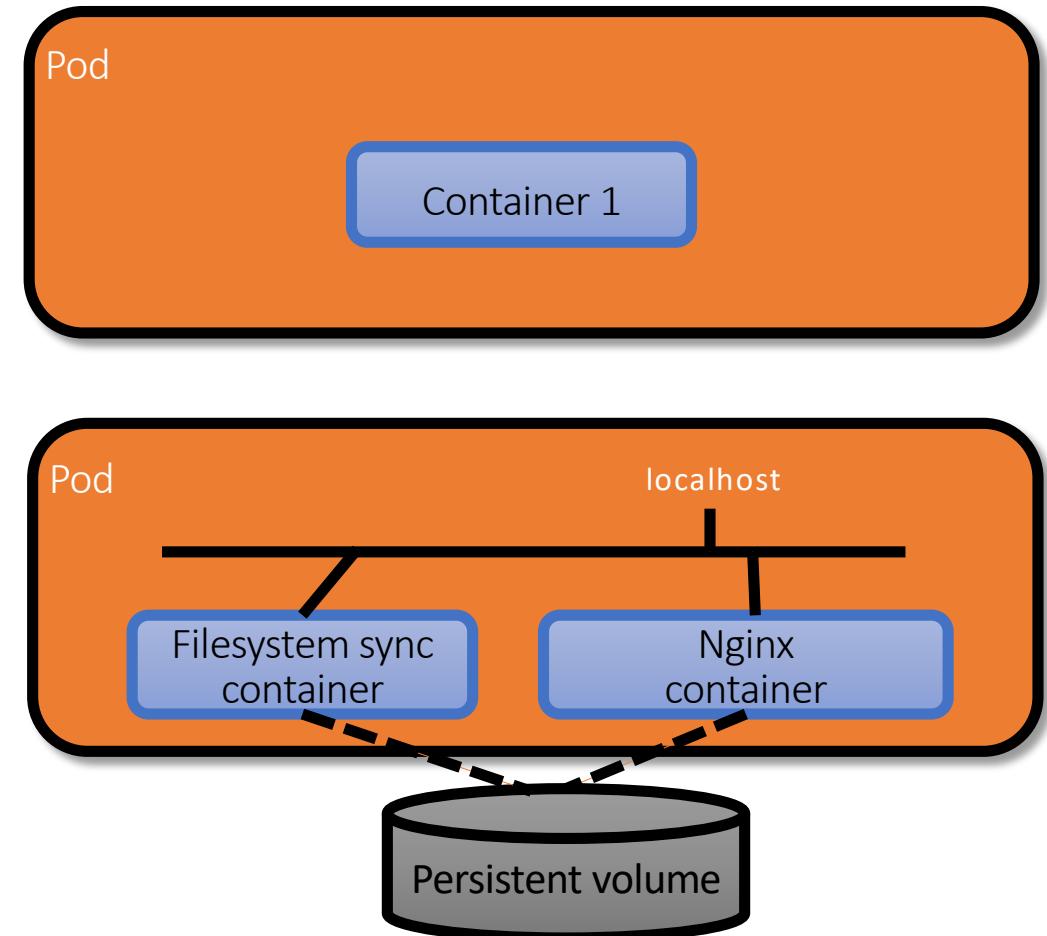
Kubernetes Pods



What is a Kubernetes Pod?

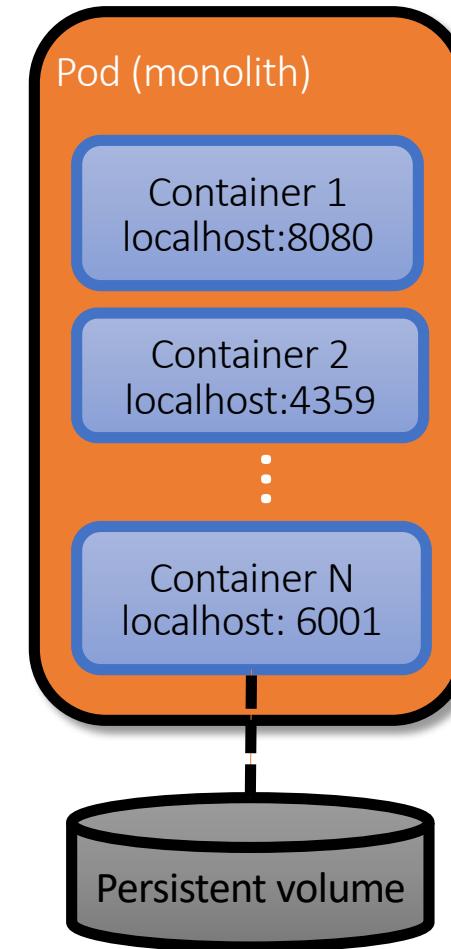
Kubernetes design intention: Pod == application instance

- Basic unit of deployment is the **pod**, a set of co-scheduled containers and shared resources
- Pods can include more than one container, for tightly-coupled application components, e.g.
 - Sidecar containers : nginx + filesystem synchronizer to update www from git
 - Content adapter: transform data to common output standard
- Containers in a pod share network namespaces and mounted volumes



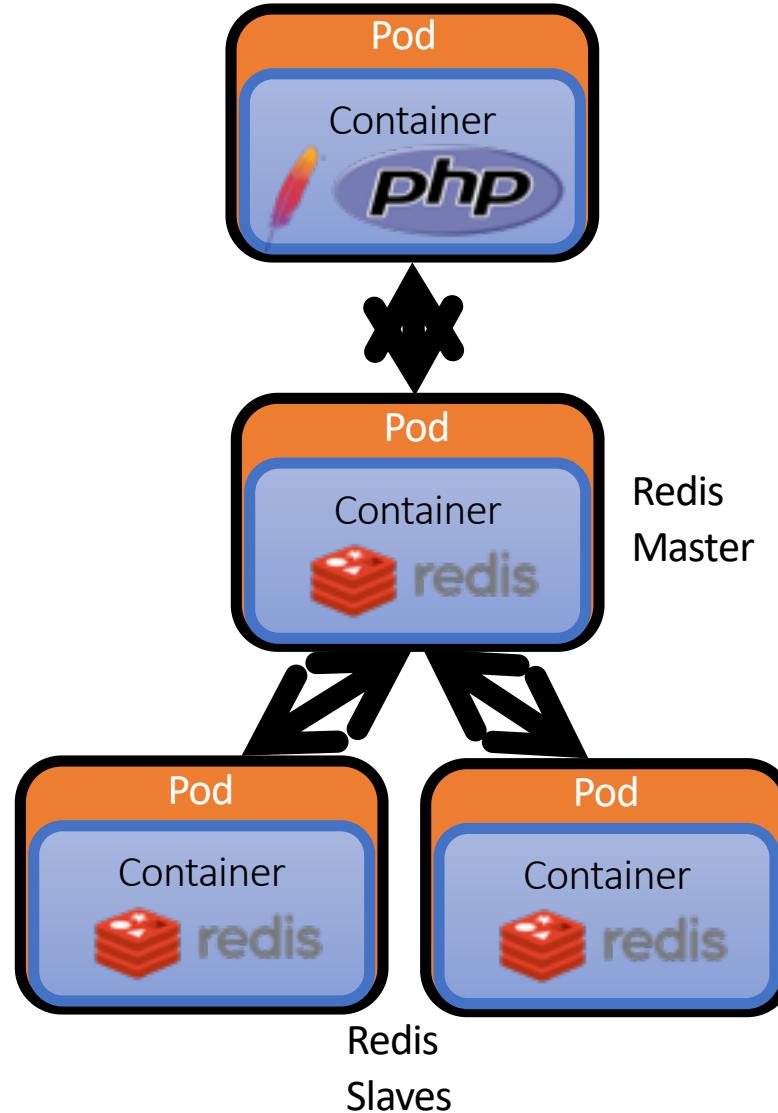
Pods Enable Deployment Flexibility

- Possible to use a single pod to run a monolithic application
 - Each application process can be built as a container
 - All containers can access each other's ports on localhost
- More advanced features of the K8s system available if application built instead from assemblages of pods, e.g.
 - Web tier: Apache pods
 - Data tier: Redis master/slave pods
- Pods provide scale and elasticity via replication – not possible in the monolith
- Best practice: assume every pod is mortal



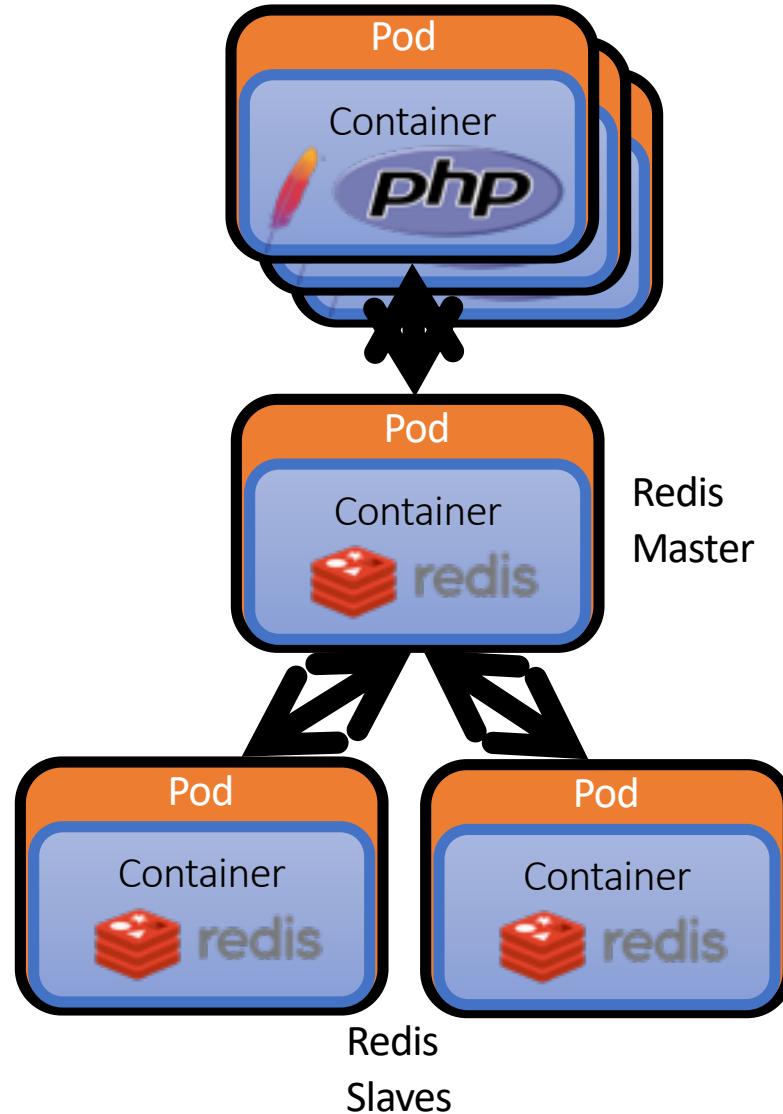
Pods Enable Deployment Flexibility

- Possible to use a single pod to run a monolithic application
 - Each application process can be built as a container
 - All containers can access each other's ports on localhost
- More advanced features of the K8s system available if application built instead from assemblages of pods, e.g.
 - Web tier: Apache pods
 - Data tier: Redis master/slave pods
- Pods provide scale and elasticity via replication – not possible in the monolith
- Best practice: assume every pod is mortal



Pods Enable Deployment Flexibility

- Possible to use a single pod to run a monolithic application
 - Each application process can be built as a container
 - All containers can access each other's ports on localhost
- More advanced features of the K8s system available if application built instead from assemblages of pods, e.g.
 - Web tier: Apache pods
 - Data tier: Redis master/slave pods
- Pods provide scale and elasticity via replication – not possible in the monolith
- Best practice: assume every pod is mortal



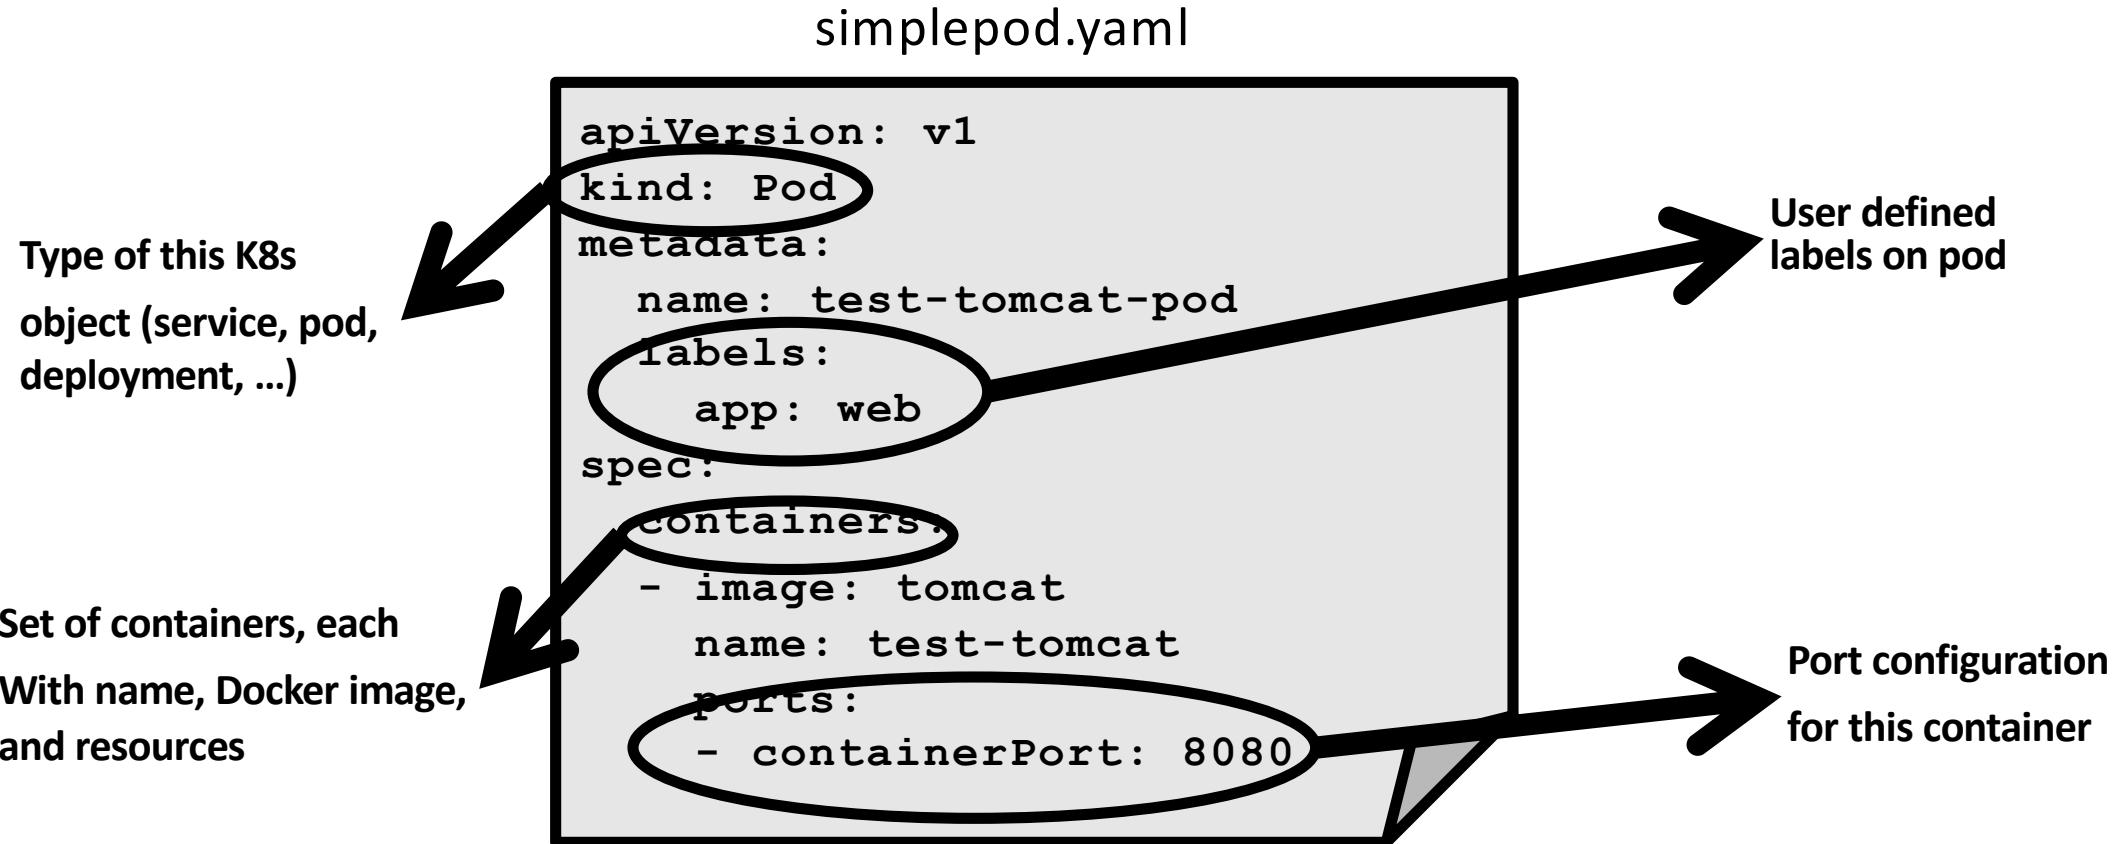
Defining a Pod via a Manifest File

Like other K8s objects, pods can be defined in YAML or JSON files

- K8s API accepts object definitions in JSON, but manifests often in YAML
- YAML format used by a variety of other tools, e.g. Docker Compose, Ansible, etc.
- **kind** field value is ‘Pod’
- **metadata** includes
 - **name** to assign to pod
 - **label** values
- **spec** includes specifics of container images, ports, and other resources

```
apiVersion: v1
kind: Pod
metadata:
  name: test-tomcat-pod
  labels:
    app: web
spec:
  containers:
    - image: tomcat
      name: test-tomcat
      ports:
        - containerPort: 8080
```

Looking at a Pod Manifest File



- Configuration options similar to creating Docker container directly

Defining a Pod with Multiple Containers

```
apiVersion: v1
kind: Pod
metadata:
  name: test-tomcat-pod
spec:
  containers:
    - image: tomcat
      name: test-tomcat
      ports:
        - containerPort: 8080
    - image: mysql
      name: test-mysql
      ports:
        - containerPort: 3306
```



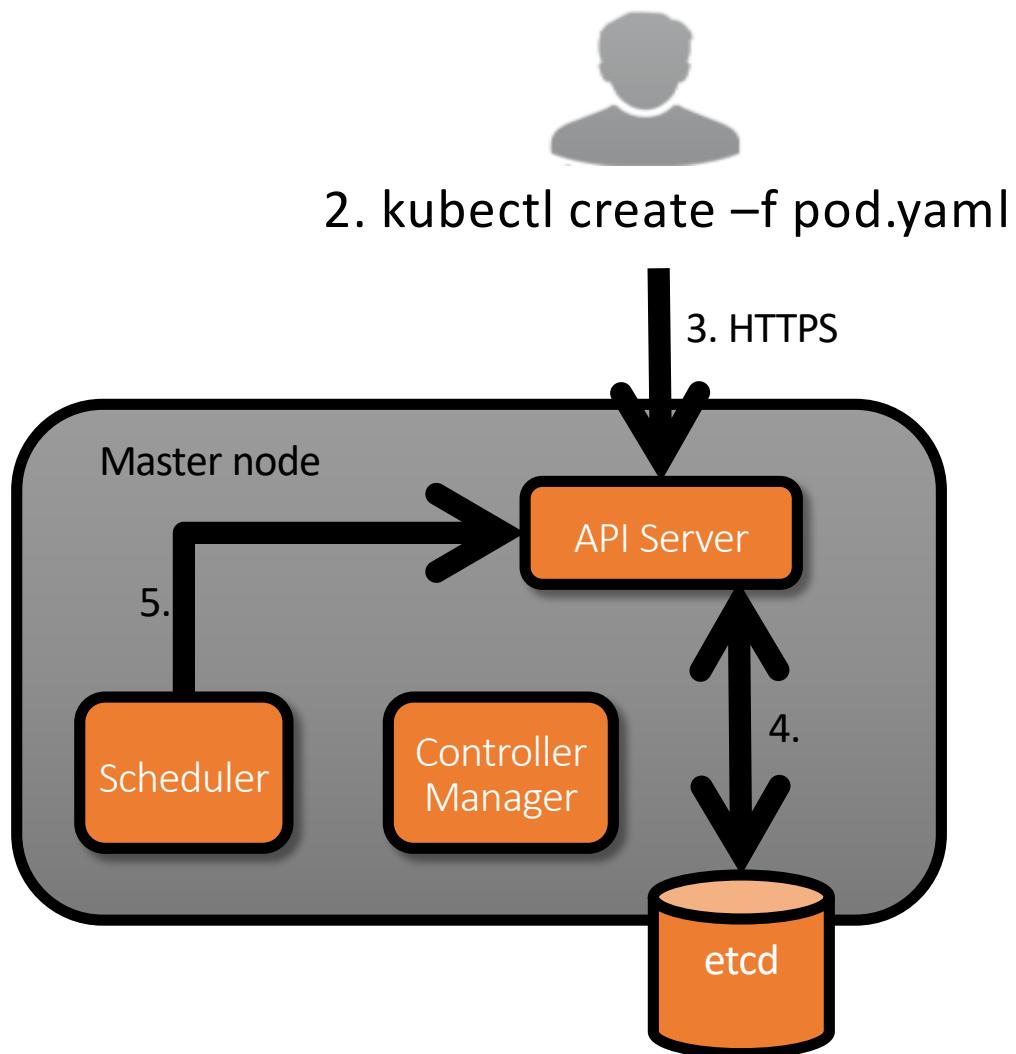
multipod.yaml

- Pod spec can contain multiple containers from different images
- Containers in pod share local network context and cluster IP for pod

Pod Creation and Management

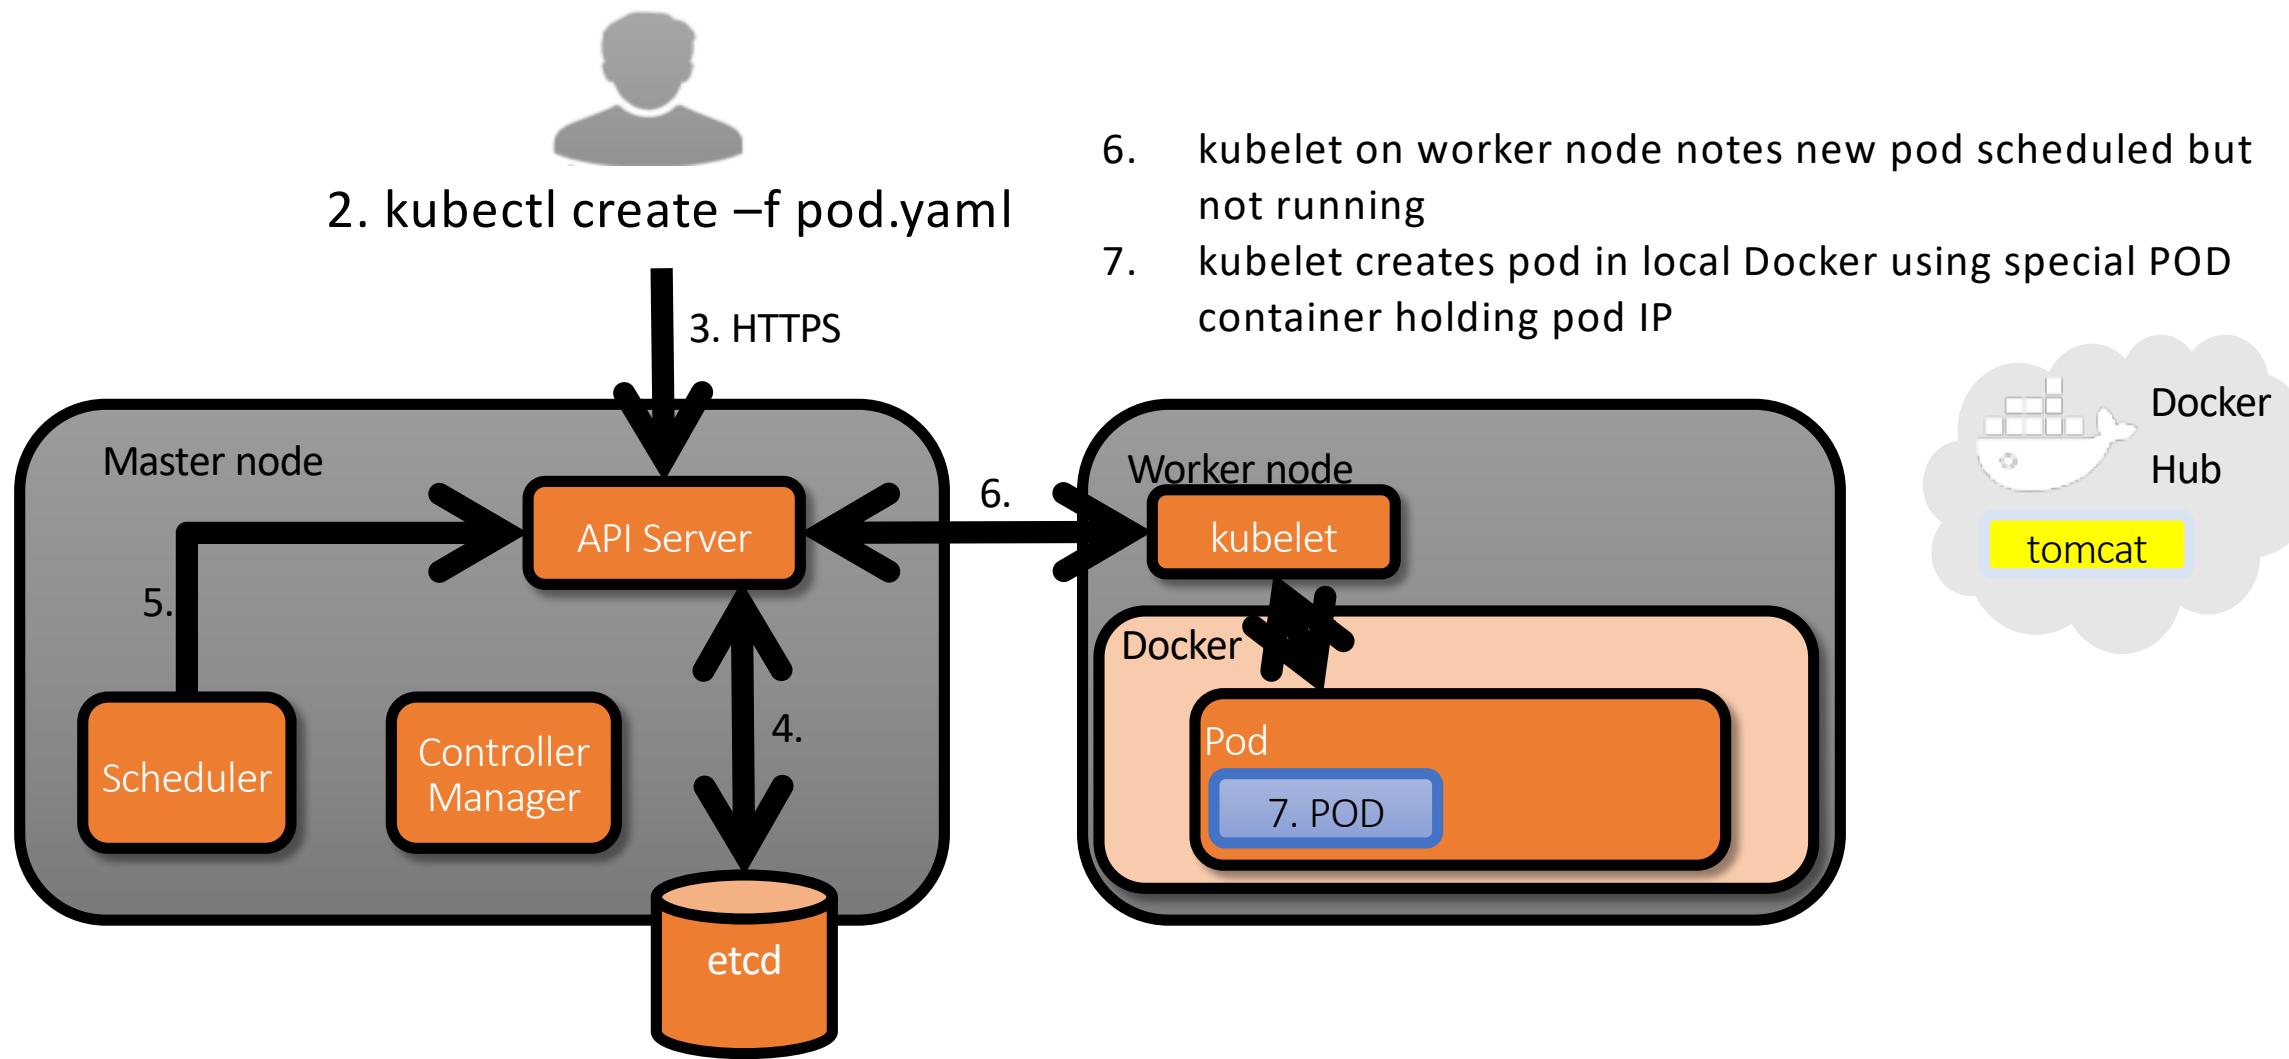


Pod Creation Process

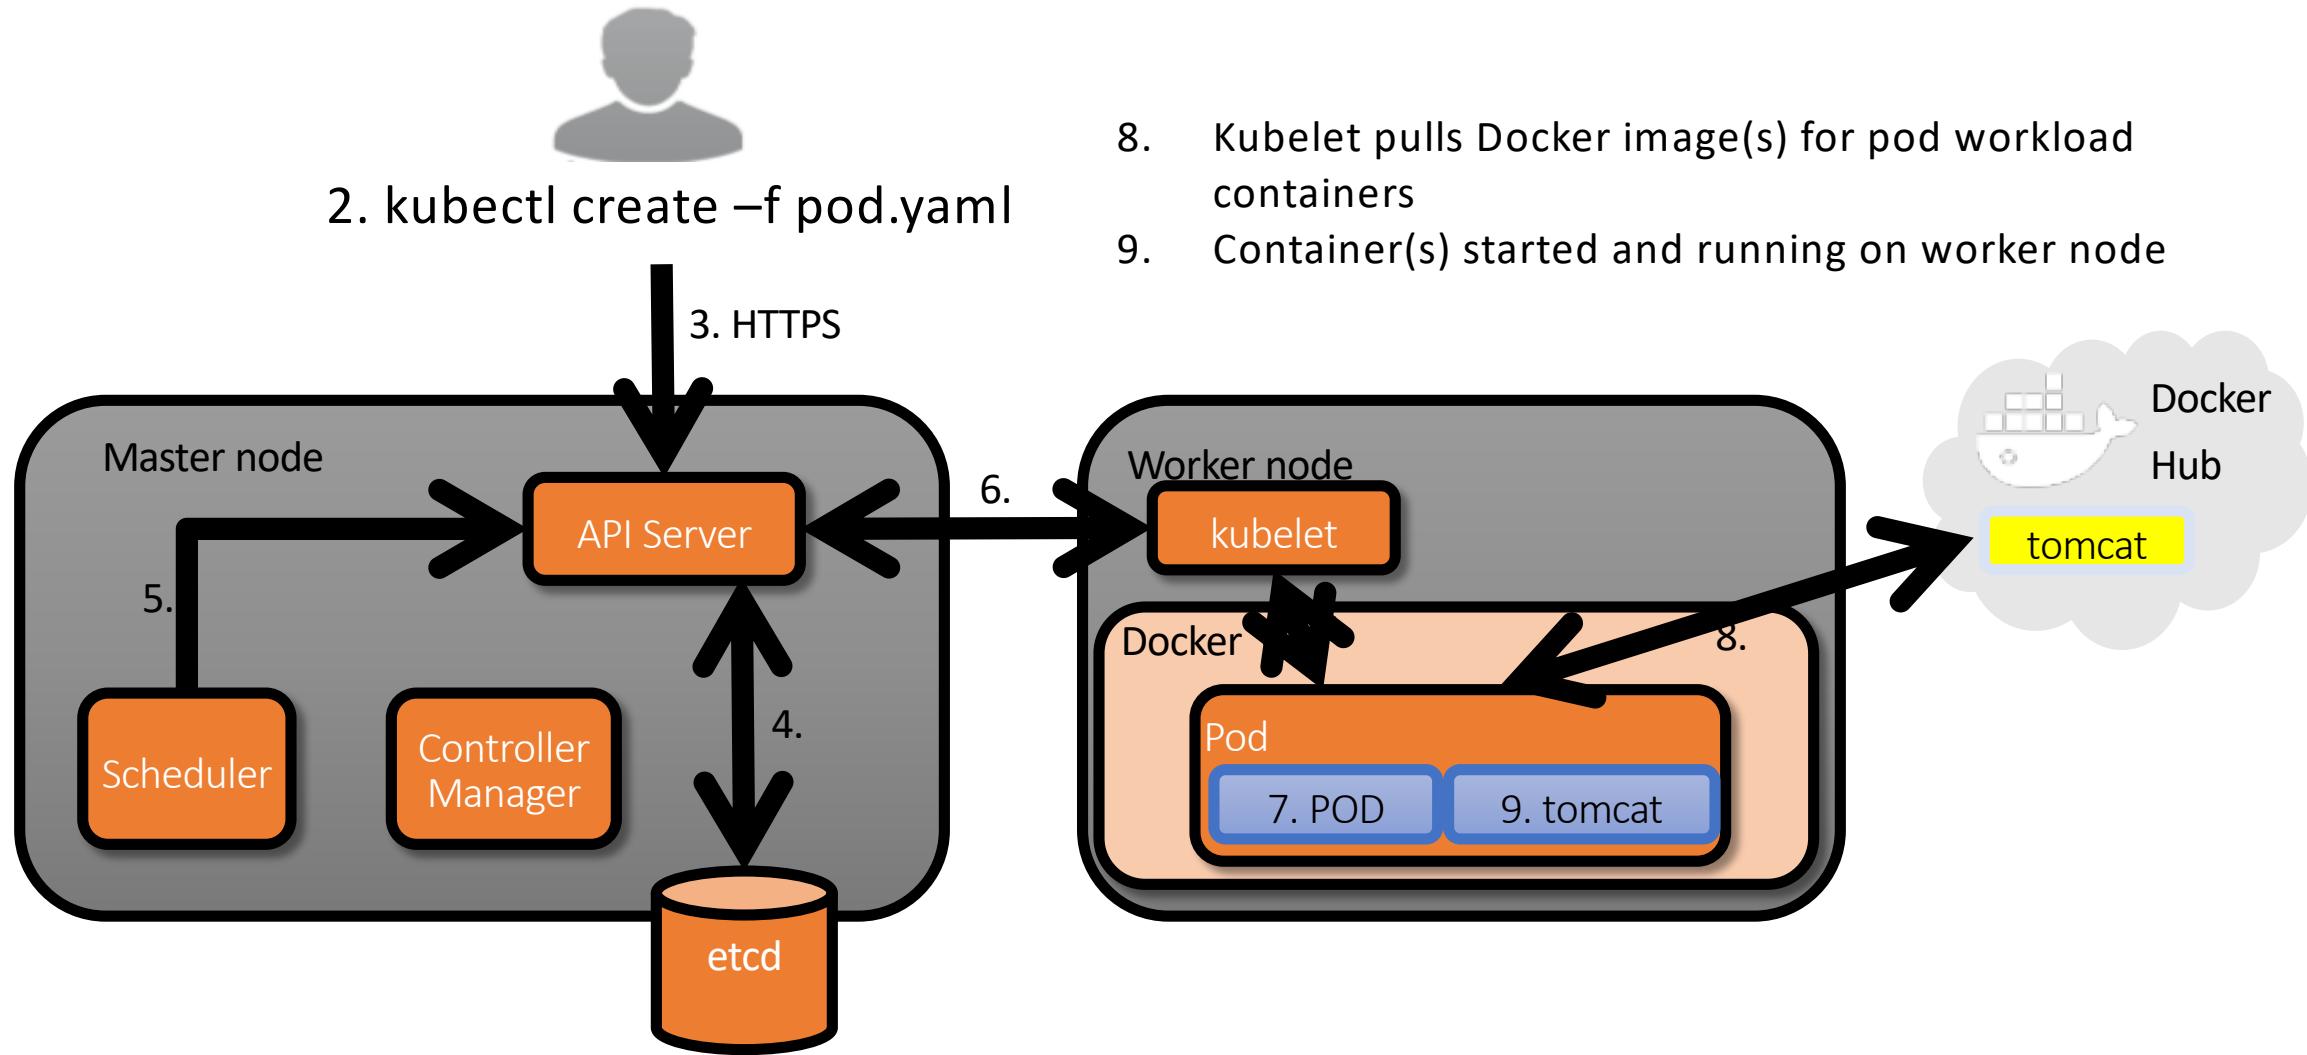


1. User writes a pod manifest file
2. User requests creation of pod from manifest via CLI
3. CLI tool marshals parameters into K8s RESTful API request (HTTP POST)
4. kube-apiserver creates new pod object record in etcd, with no node assignment
5. kube-scheduler notes new pod via API
 - a. Selects node for pod to run on
 - b. Updates pod record via API with node assignment

Pod Creation Process



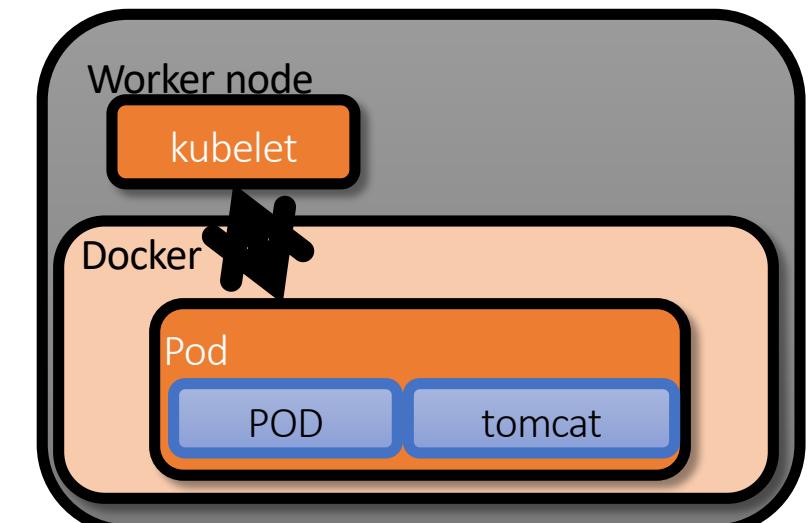
Pod Creation Process



Pod Lifecycles

- By default, K8s Pods have an indefinite lifetime, which is not immortality
 - **restartPolicy** of Always by default
 - **restartPolicy** of Never or OnFailure also available for terminating jobs
- Node's kubelet will create and keep running containers for pods assigned to node, per the pod specs
- If a Pod container fails to start, or unexpectedly exits, kubelet will restart it
 - Can see container lifecycle events via 'kubectl describe pod <PODNAME>'
- If node is lost, its Pods are also lost – K8s will not rebind Pods to another node

```
apiVersion: v1
kind: Pod
metadata:
  name: test-tomcat-pod
  labels:
    app: web
spec:
  containers:
  - image: tomcat
    name: test-tomcat
    ports:
    - containerPort: 8080
```



Modifying a Pod

Change the container version

- You can make changes to the desired state of a pod via updating the manifest file
- Changes can then be applied to the pod via the command
 - `kubectl apply -f <manifest.yaml>`
- Changing a container image as shown will result in K8s automatically killing and recreating the pod's workload container

```
$ vi simplepod.yaml
```

```
apiVersion: v1
kind: Pod
metadata:
  name: test-tomcat-pod
spec:
  containers:
    - image: tomcat:8.5.5
      name: test-tomcat
    ports:
      - containerPort: 8080
```

New image version



Modifying a Pod

```
$ kubectl apply -f simplepod.yaml
pod "test-tomcat-pod" configured

$ kubectl describe pod test-tomcat-pod
Name:           test-tomcat-pod
Namespace:      default
...
Labels:         tier=frontend
Status:         Running
...
Containers:
  test-tomcat:
    Image:          Tomcat:8.0.5
    Image ID:       ...
    Port:           8080/TCP
    State:          Running
...

```

New version running



8080/TCP

Labeling Pods

User-defined labels help organize K8s resources

- Labels are key/value pairs that users can assign and update on any K8s resources, including pods
- Other K8s objects, like controllers, use labels to select pods to govern
- Labels can also be used to filter data queries with *kubectl*, e.g.
 - `kubectl get pods -l <label=value>`
- Labels can be used to distinguish pods on any criteria, such as
 - Application, application tier, version, environment state, etc.
- K8s system does not require specific labels to be used – all user-defined

Labeling a Pod

```
$ kubectl label pod test-tomcat-pod tier=frontend  
pod "test-tomcat-pod" labeled
```

```
$ kubectl describe pods test-tomcat-pod  
Name:           test-tomcat-pod  
...  
Labels:         tier=frontend
```

```
$ kubectl get pods -l tier=frontend  
NAME          READY     STATUS    RESTARTS   AGE  
test-tomcat-pod  1/1      Running   0          1d
```

Reviewing Labels on Pods

Changing kubectl output

- You can display pod labels via a flag on the *kubectl* command

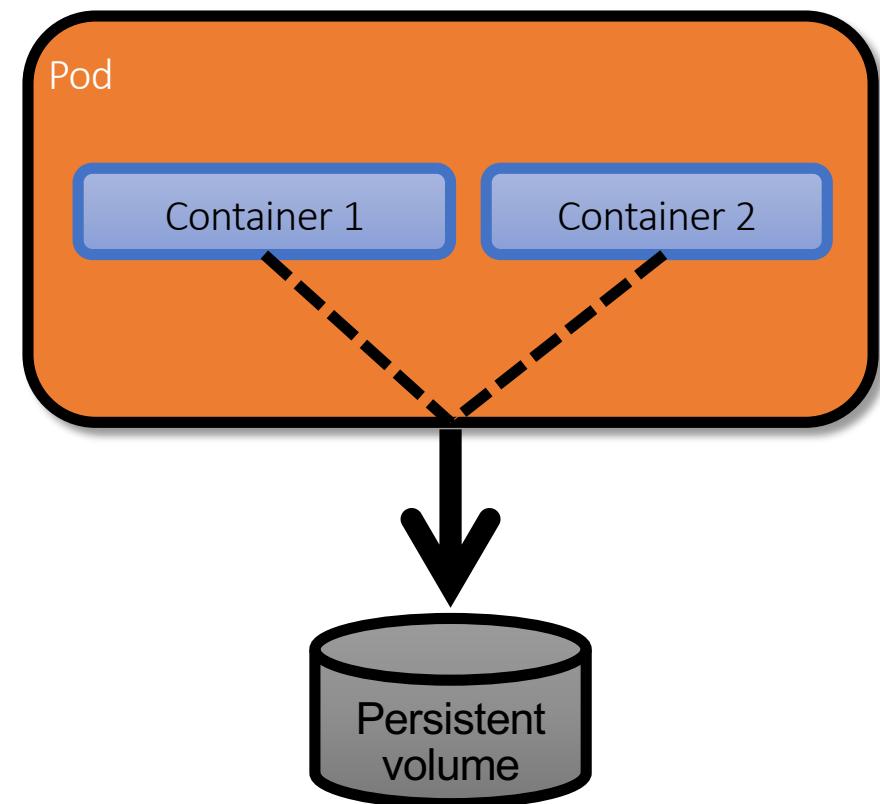
```
$ kubectl get pods --show-labels
NAME          READY   STATUS    RESTARTS   AGE   LABELS
test-tomcat-pod 1/1     Running   1          1d    tier=frontend

$ kubectl get pods --show-labels --namespace=kube-system
NAME          READY   STATUS    RESTARTS   AGE   LABELS
kube-addon-manager-minikube 1/1     Running   3          8d    component=kube-addon-
manager,kubernetes.io/minikube-addons=addon-manager,version=v6.4-alpha.1
kube-dns-v20-mm0zl           3/3     Running   9          8d    k8s-app=kube-
dns,version=v20
kubernetes-dashboard-kc9rk   1/1     Running   3          8d    app=kubernetes-
dashboard,kubernetes.io/cluster-service=true,version=v1.6.0
```

Deleting Pods

Pod deletion will discard all local pod resources

- When deleting a Pod, its containers will be removed and pod IP relinquished
- If an application needs to persist data, its pods must be configured to use persistent volumes for storage
- If a node dies, its local pods are also gone
- Best practice: use controller resources instead of managing pods directly
- Best practice: use service resources to build reliable abstraction layers for clients



Lab: Create Pods



Kubernetes Services Overview



Kubernetes Service Objects: Microservices

Services provide abstraction
between layers of an application

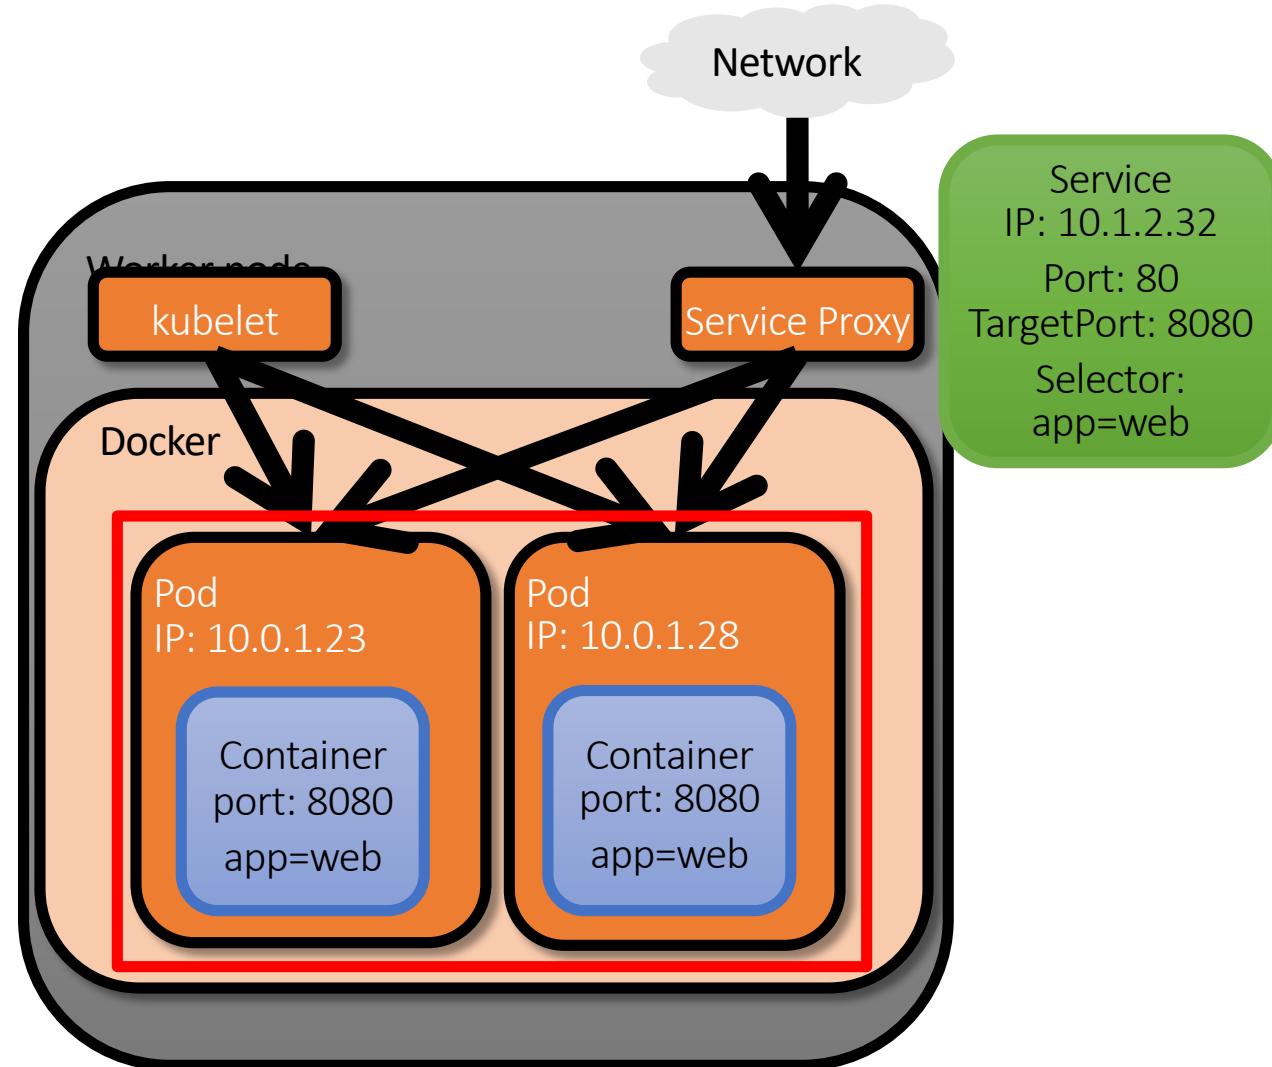
- **Service** object provides a stable IP for a collection of Pods
- Services use a label **selector** to target a specific set of pods as endpoints to receive proxied traffic
- Clients can reliably connect via the service IP and port(s), even as individual endpoint pods are dynamically created & destroyed
- Can model other types of backends using services without selectors

```
apiVersion: v1
kind: Service
metadata:
  name: wordpress
  labels:
    app: wordpress
spec:
  ports:
    - port: 80
      targetPort: 8080
  selector:
    app: wordpress
    tier: frontend
  type: LoadBalancer
```

sampleservice.yaml

Services Provide Abstraction Layer for Applications

- Services can be used for communications between application tiers
- Services can also be used to expose applications outside the K8s cluster
- Services distribute requests over the set of Pods matching the service's selector
 - Service functions as TCP and UDP proxy for traffic to Pods
 - Service maps its defined ports to listening ports on Pod endpoints



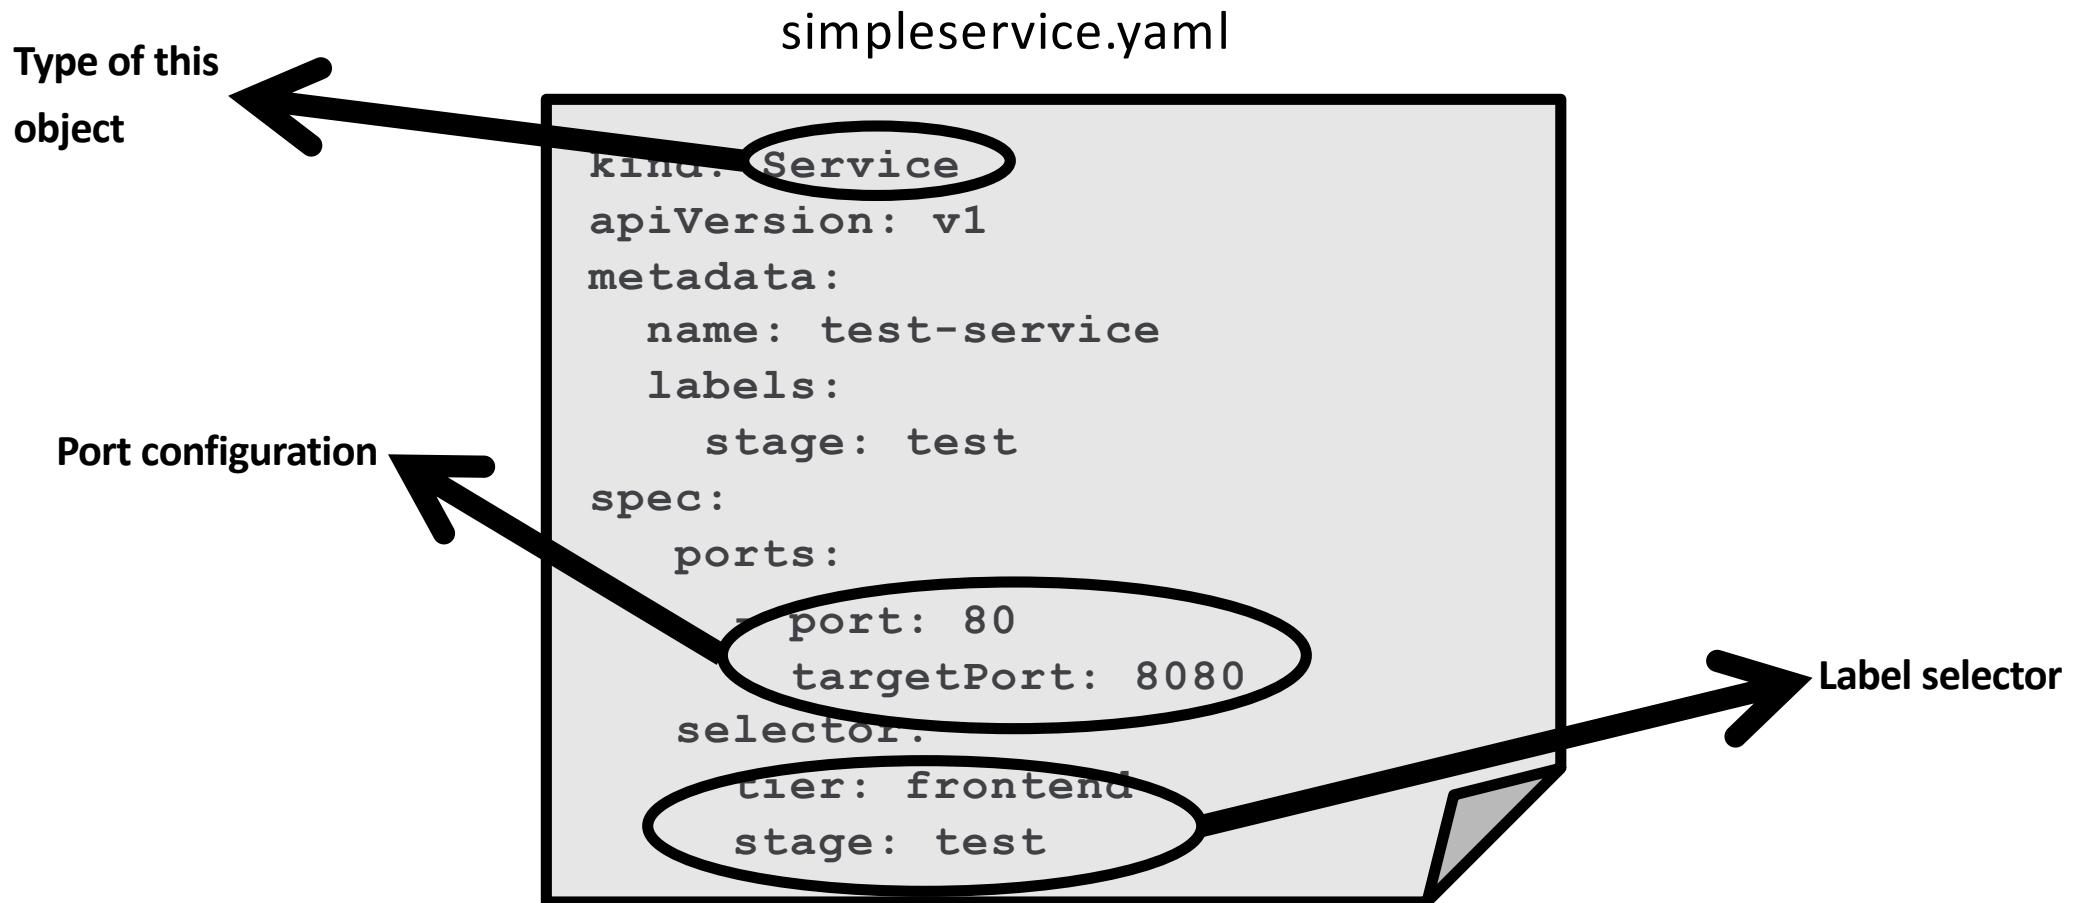
Defining a Service Using a Manifest File

Services can be defined in YAML or JSON, like other K8s resources

- **kind** field value is ‘Service’
- **metadata** includes
 - **name** to assign to Service
- **spec** includes the ports associated with the Service
 - **port** is the Service’s port value
 - **targetPort** is connection port on selected pods (default: **port** value)
- **selector** specifies a set of label KV pairs to identify the endpoint pods for the Service

```
kind: Service
apiVersion: v1
metadata:
  name: test-service
  labels:
    stage: test
spec:
  ports:
    - port: 80
      targetPort: 8080
  selector:
    tier: frontend
    stage: test
```

Reviewing a Service Manifest File



ServiceTypes and Exposing Applications

- By default, a Service is assigned a cluster-internal IP – good for back-ends
 - **ServiceType ClusterIP**
- To make a front-end Service accessible outside the cluster, there are other ServiceTypes available
 - **ServiceType NodePort** exposes Service on each Node's IP on a static port
 - **ServiceType LoadBalancer** exposes the Service externally using cloud provider's load balancer
- Can also use a Service to expose an external resource via **ServiceType ExternalName**

```
kind: Service
apiVersion: v1
metadata:
  name: test-service
  labels:
    stage: test
spec:
  ports:
    - port: 80
      targetPort: 8080
      nodePort: 30080
  selector:
    tier: frontend
  type: NodePort
```

Exposing Services at L7 through Ingresses

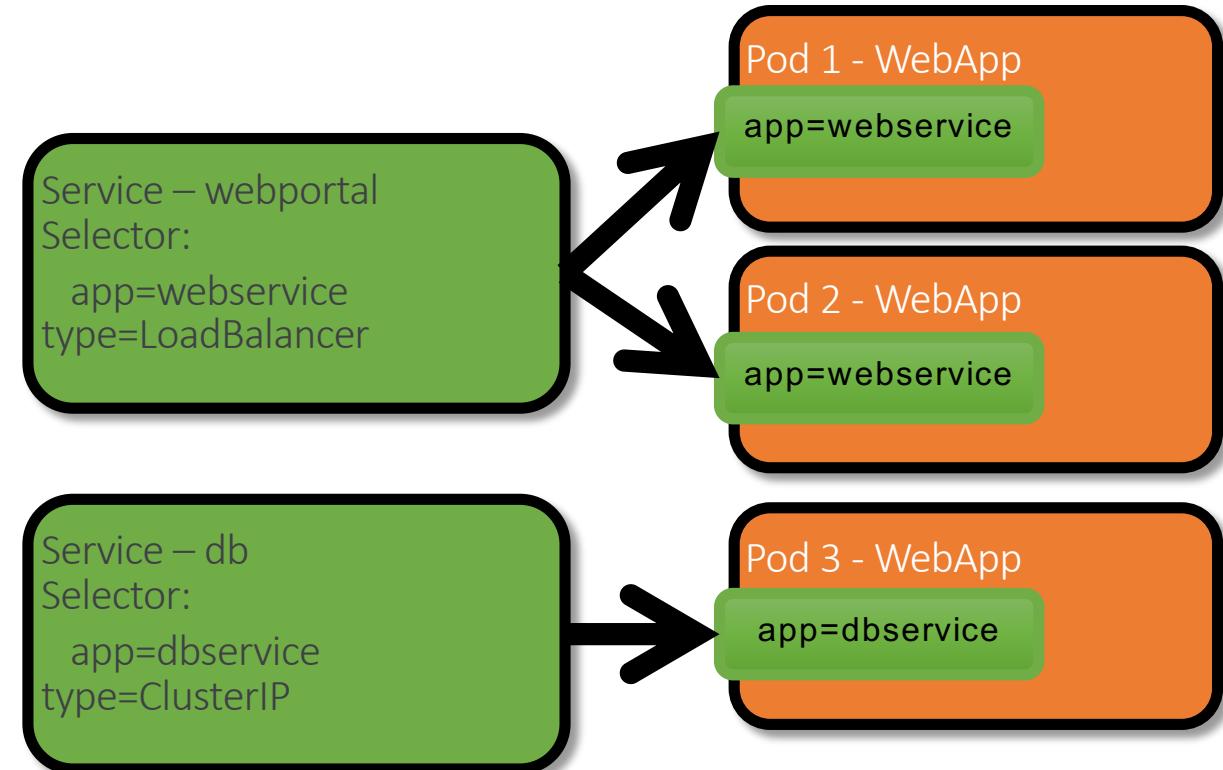
- Kubernetes also provides the facility to define **Ingress** resource to configure an external loadbalancer at L7
- Spec of Ingress resource is a set of rules matching HTTP host/url paths to specific Service backends
- Ingresses require the cluster to be running an appropriately configured Ingress controller to function (e.g. nginx)
- Useful for implementing fanout, Service backends for virtual hosts, etc.

```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: test-ingress
spec:
  rules:
  - host: bar.foo.com
    http:
      paths:
      - path: /first
        backend:
          serviceName:
            firstservice
            servicePort: 80
      - path: /second
        backend:
          serviceName:
            secondservice
            servicePort: 80
```

Selecting Pods as Service Endpoints

Service's pod selector based on labels

- Multiple pods can have the same label, unlike pod names which are unique in the namespace
- K8s system re-evaluates Service's selector continuously
- K8s maintains **endpoints** object of same name with list of pod IP:port's matching Service's selector



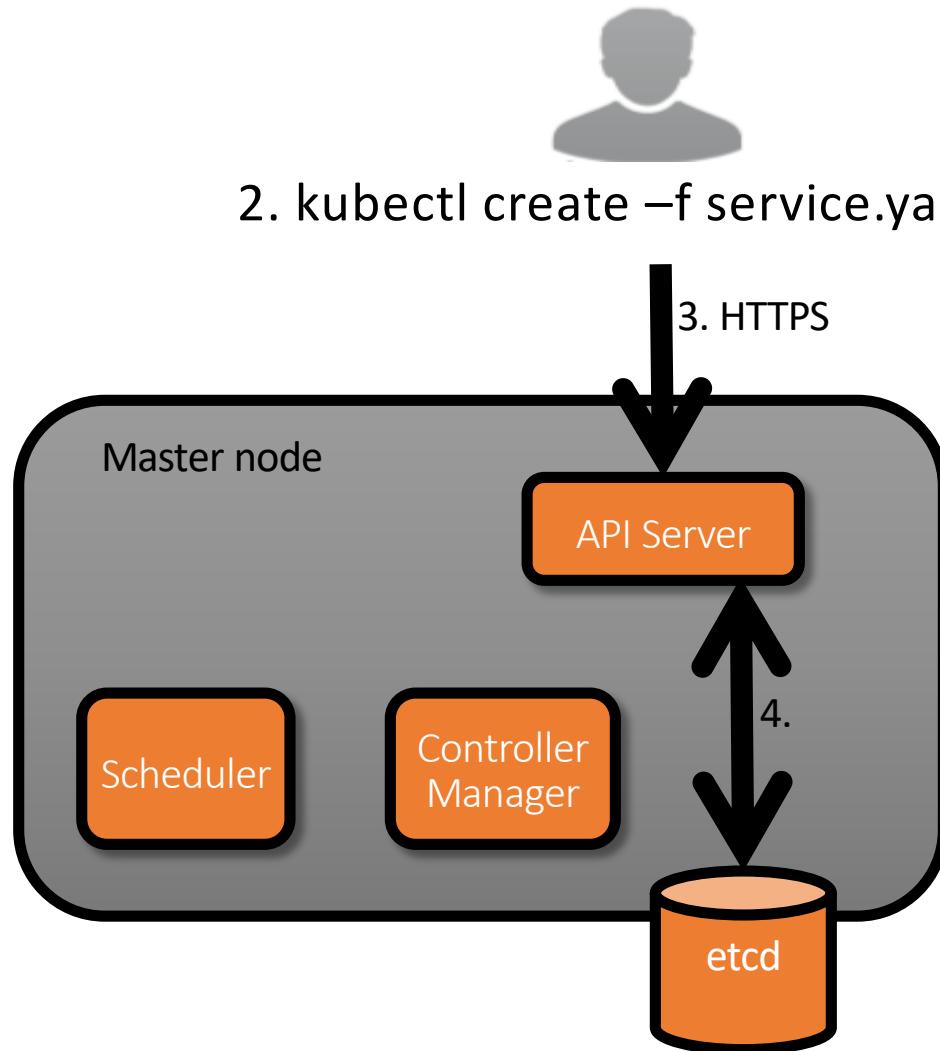


Questions

Services Management

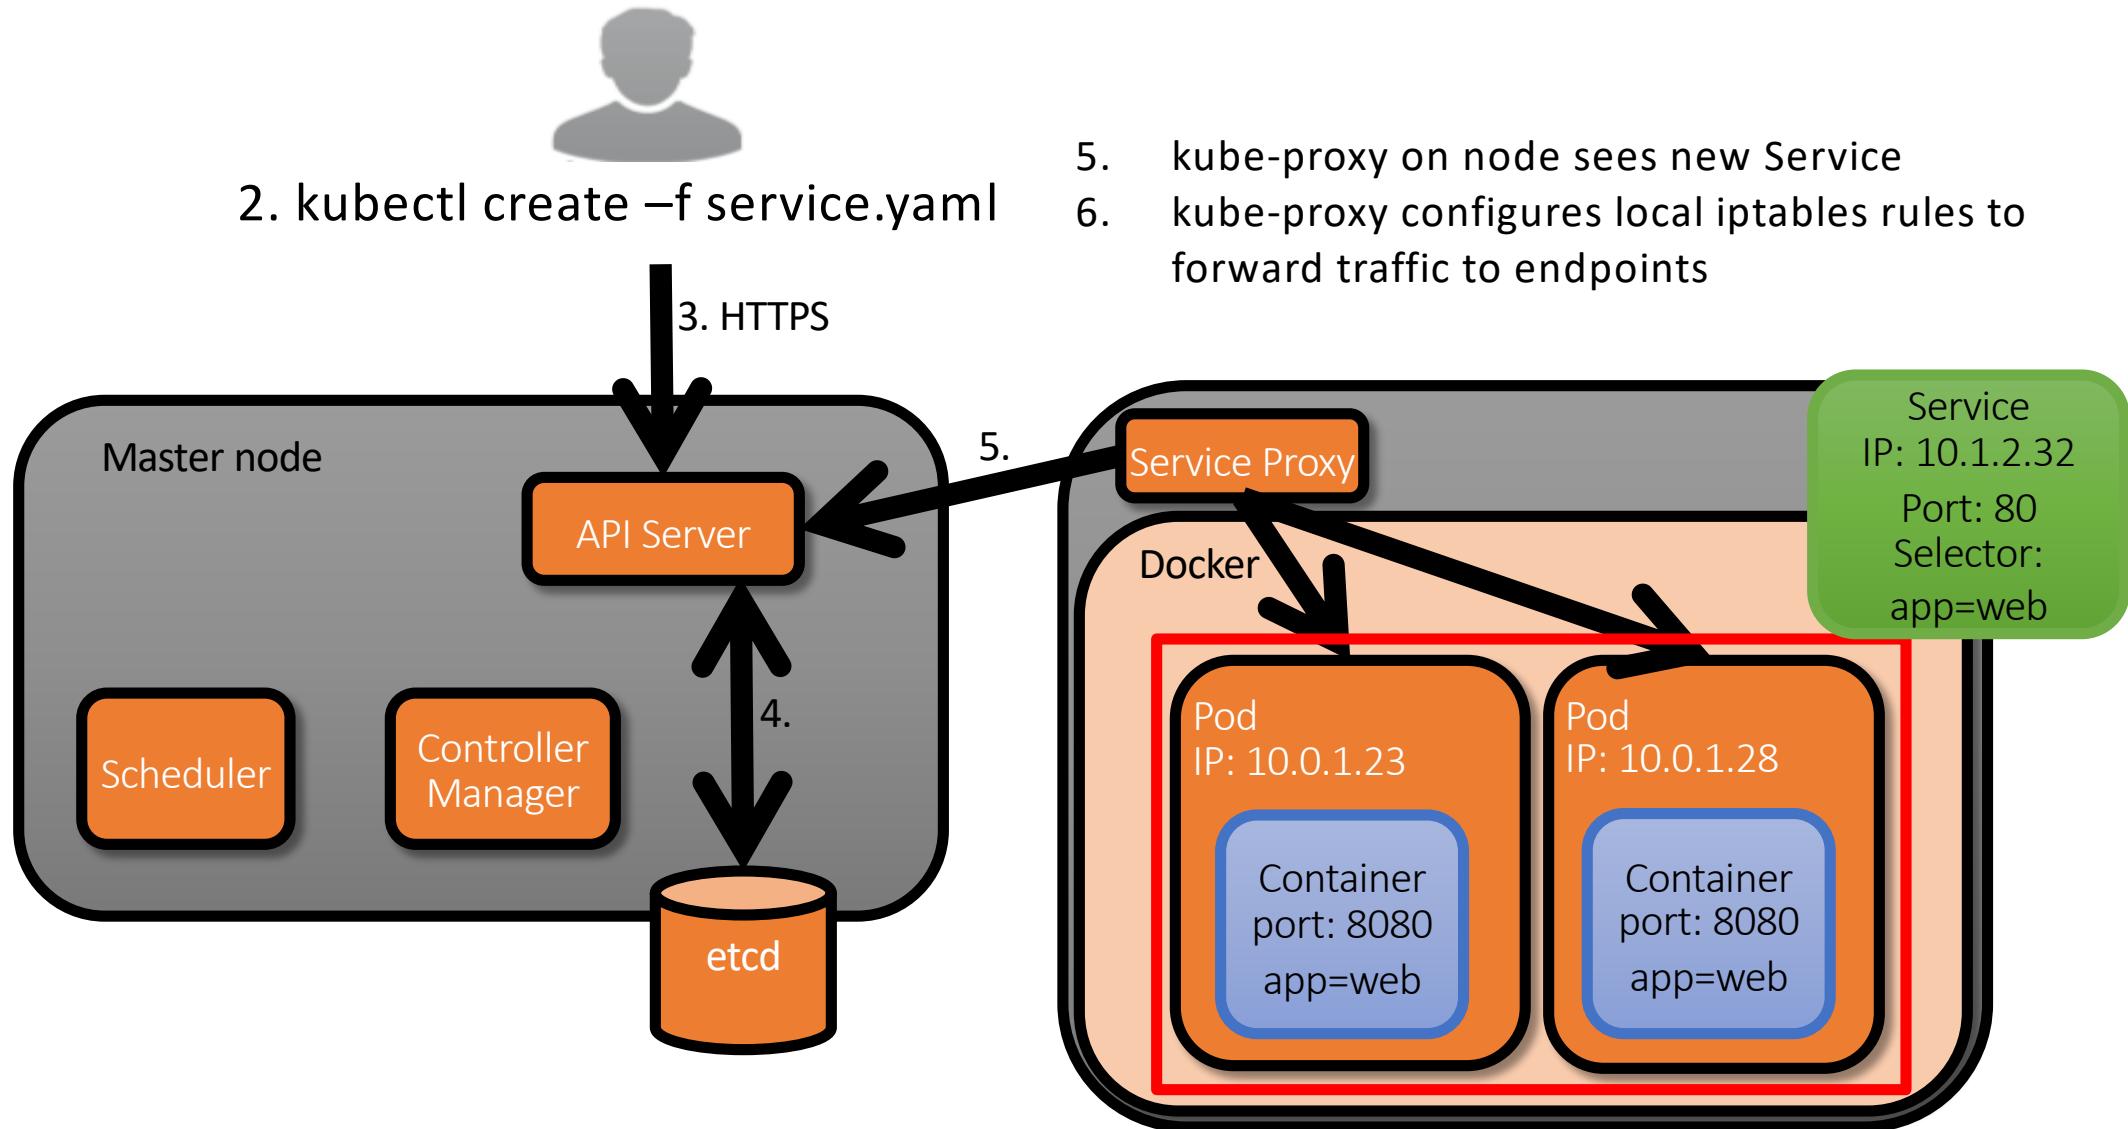


Service Creation Process



1. User writes a Service manifest file
2. User requests creation of Service from manifest via CLI
3. CLI tool marshals parameters into K8s RESTful API request (HTTP POST)
4. kube-apiserver creates new Service object record in etcd

Service Creation Process



Accessing Services Externally

Checking the service external IP

- Configured as type LoadBalancer, a configured cluster will provide an externally accessible IP for your Service

```
$ kubectl get services test-service
NAME      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
test-service  10.3.247.123  104.154.105.198  8080/TCP  4h
```

Deleting Services

Services can be deleted anytime

- Service deletion does not affect pods targeted by the Service's selector
- Can be done referencing the Service name or a manifest for the resource

```
$ kubectl delete -f simpleservice.yaml
```

```
$ kubectl delete service test-service
```

Service Discovery via DNS

Kubernetes advertises services via cluster DNS

- Kubernetes uses a cluster-internal DNS addon to create and manage records for all Services in the cluster
- Pod's DNS search list includes its own namespace and cluster default domain by default

```
$ kubectl get svc
NAME           CLUSTER-IP      EXTERNAL-IP    PORT(S)        AGE
kubernetes     10.0.0.1        <none>        443/TCP       11d
test-service   10.0.0.102      <nodes>       8080:30464/TCP 2d
```

```
kubectl exec -ti busybox1 -- nslookup test-service.default
Server:  10.0.0.10
Address 1: 10.0.0.10 kube-dns.kube-system.svc.cluster.local
```

```
Name:      test-service.default
Address 1: 10.0.0.102 test-service.default.svc.cluster.local
```

Service Discovery via Environment Variables

Kubernetes advertises services via local environment

- When a Pod is started on a node, kubelet will create environment variables for Services in the Pod's namespace and system namespace

```
$ kubectl exec -ti busybox -- printenv | grep SERVICE
KUBERNETES_SERVICE_HOST=10.0.0.1
KUBERNETES_SERVICE_PORT=443
KUBERNETES_SERVICE_PORT_HTTPS=443
TEST_SERVICE_SERVICE_HOST=10.0.0.102
TEST_SERVICE_PORT_8080_TCP_PROTO=tcp
TEST_SERVICE_PORT_8080_TCP_PORT=8080
TEST_SERVICE_PORT=tcp://10.0.0.102:8080
TEST_SERVICE_PORT_8080_TCP_ADDR=10.0.0.102
TEST_SERVICE_SERVICE_PORT=8080
TEST_SERVICE_PORT_8080_TCP=tcp://10.0.0.102:8080
```

Lab: Multi-container deployment

