



# Fundamentals of Istio

# What is Istio?

- An open platform to connect, manage, and secure microservices.
- Supports:
  - load balancing
  - service-to-service authentication
  - monitoring
  - more
- No changes to service code!

# What is Istio?

- Traffic management
  - Control the flow of traffic and API calls between services, make calls more reliable, and make the network more robust in the face of adverse conditions.
- Observability
  - Gain understanding of the dependencies between services and the nature and flow of traffic between them, providing the ability to quickly identify issues.

# What is Istio?

- Policy Enforcement
  - Apply organizational policy to the interaction between services, ensure access policies are enforced and resources are fairly distributed among consumers. Policy changes are made by configuring the mesh, not by changing application code.
- Service Identity and Security
  - Provide services in the mesh with a verifiable identity and provide the ability to protect service traffic as it flows over networks of varying degrees of trustability.

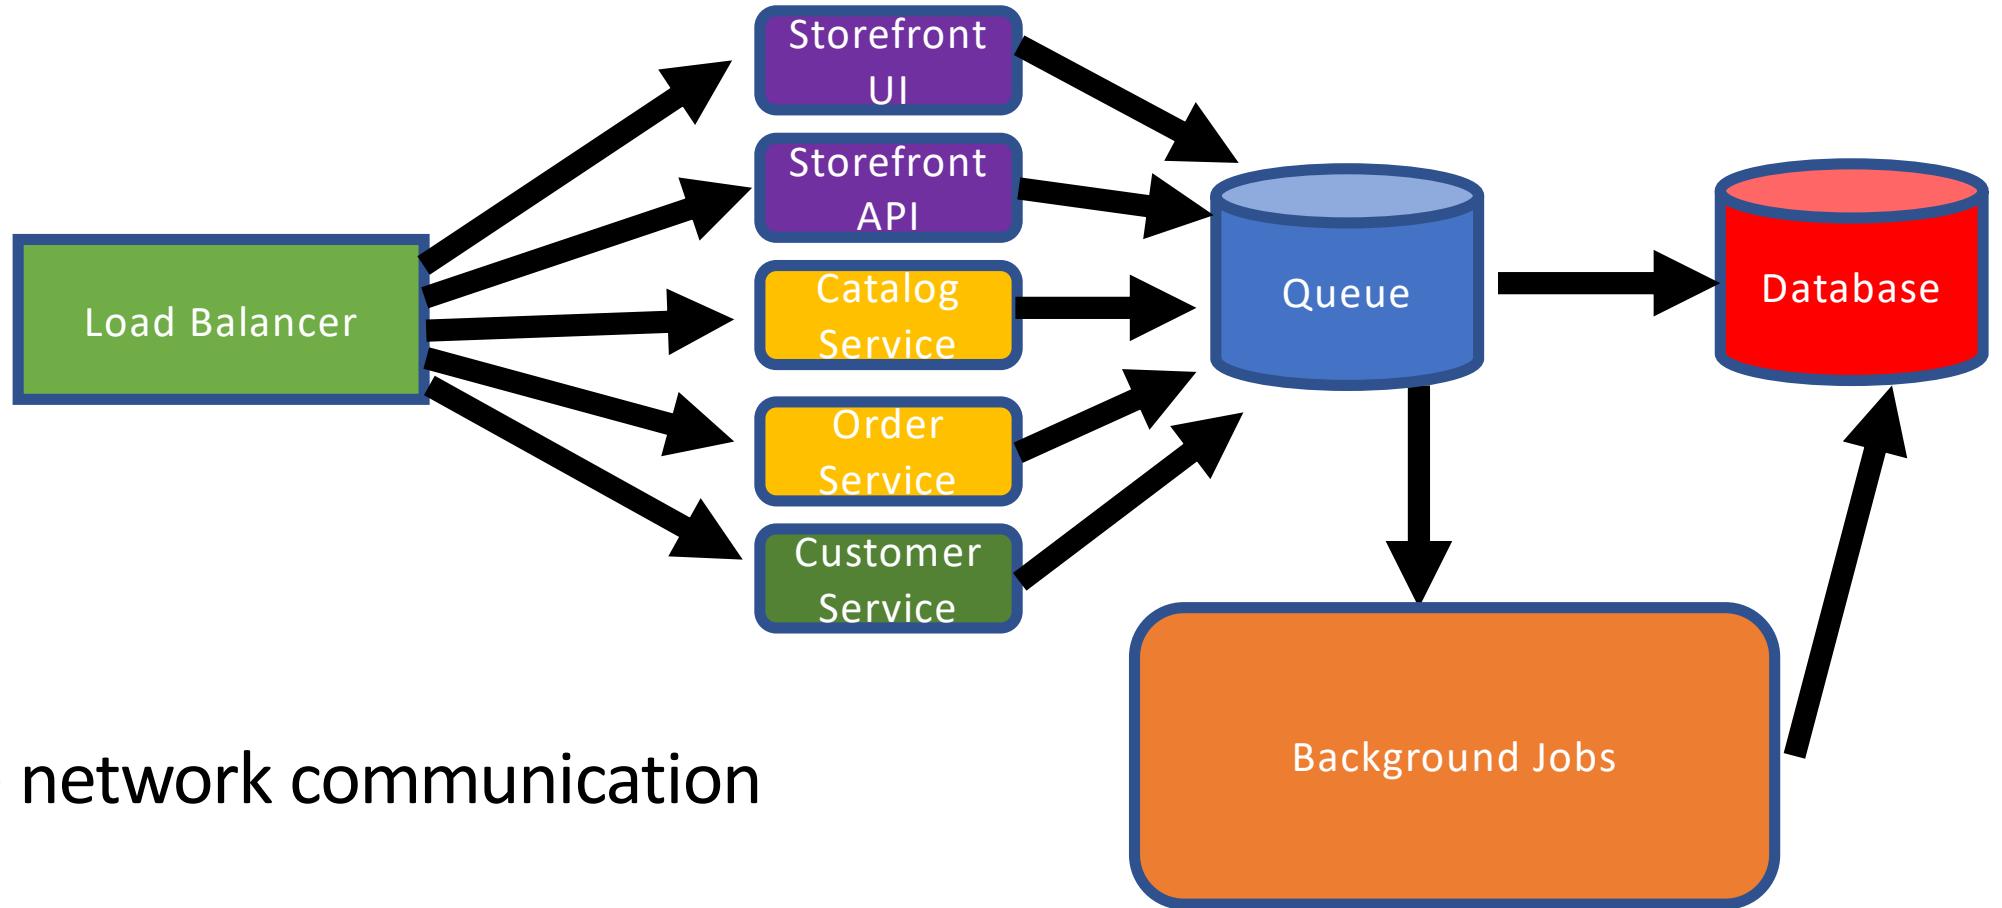
# Microservices Challenges Review

# Monolithic Architecture

- Components communicate seamlessly
- Easy to monitor
- Easy to determine where error is
- Communication inside monolith is secured



# Microservices Architecture

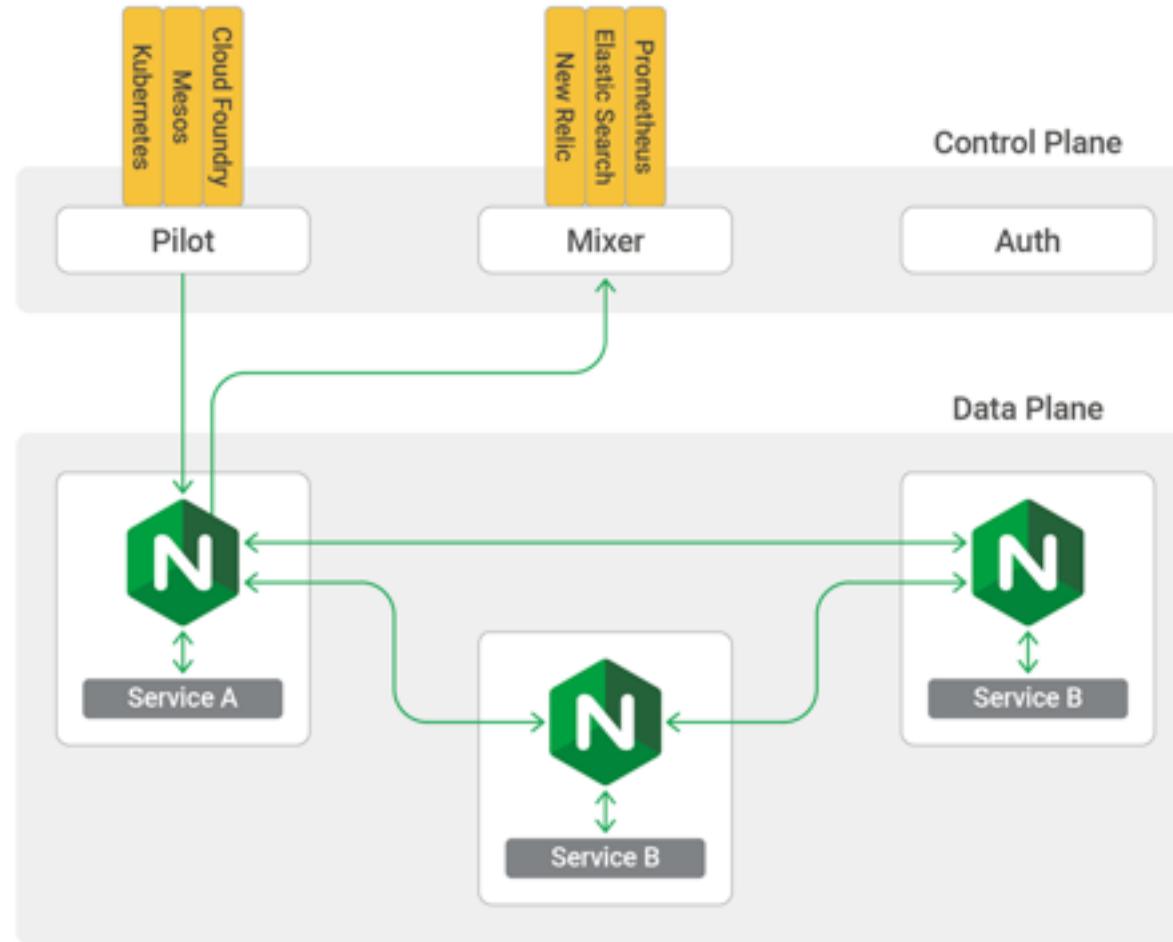


- Insecure network communication
- Visibility
- Policy enforcement

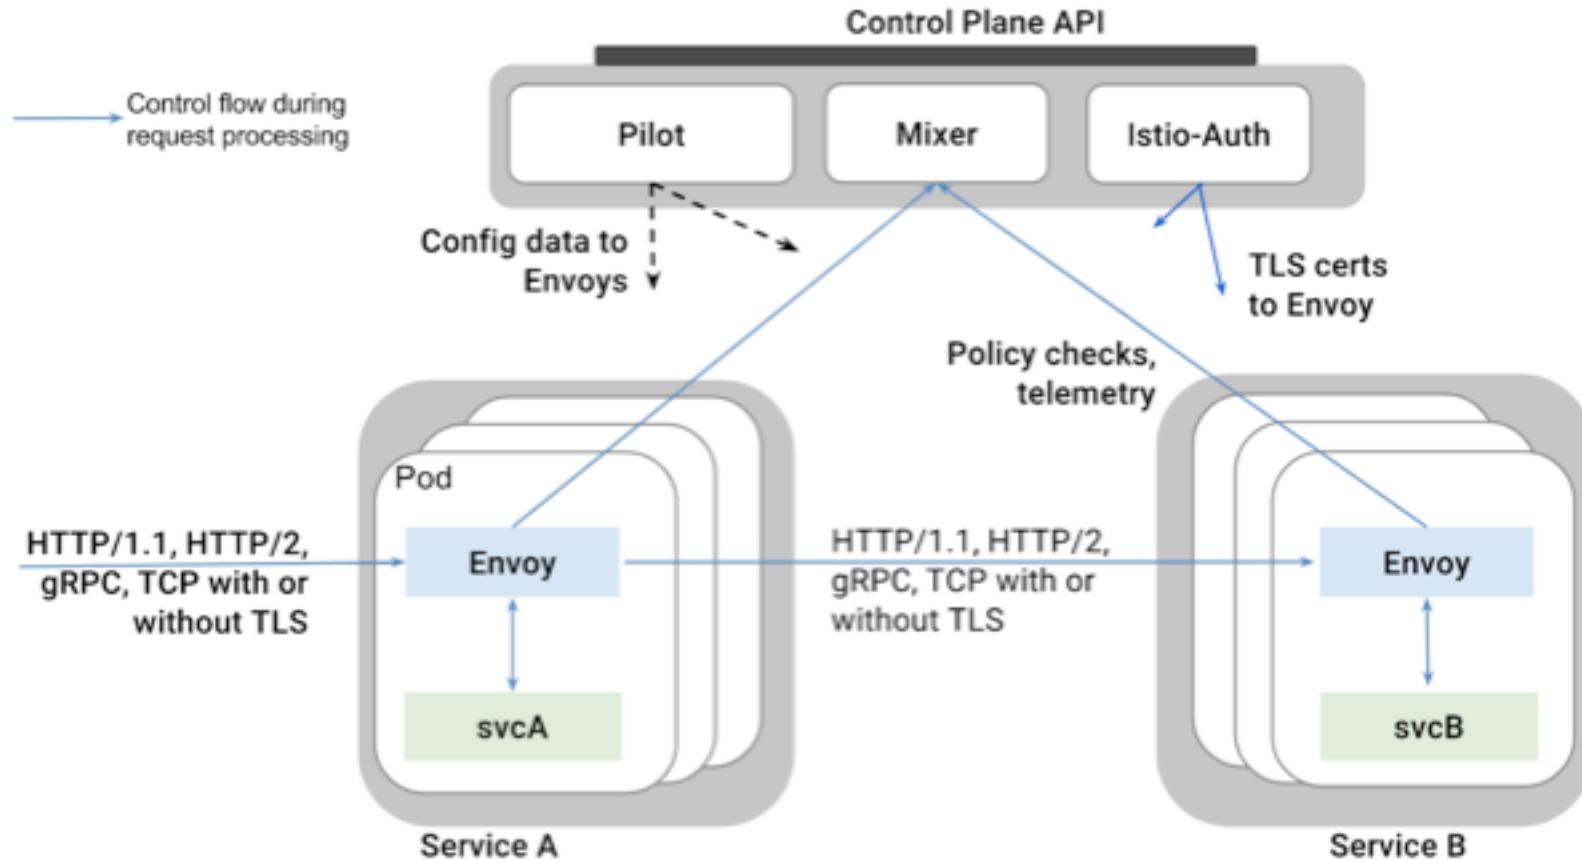
# Service Mesh

- Dedicated infrastructure layer for making service-to-service communication safe, fast, and reliable.
- Array of lightweight network proxies
  - Deployed alongside application code
- Encryption
  - Developers no longer have to add encrypt/decrypt code to their applications
- Circuit breaker pattern support
  - Service signals it's "unhealthy", service mesh removes it, then when fixed adds it back in to pool.

# Service Mesh



# Service Mesh

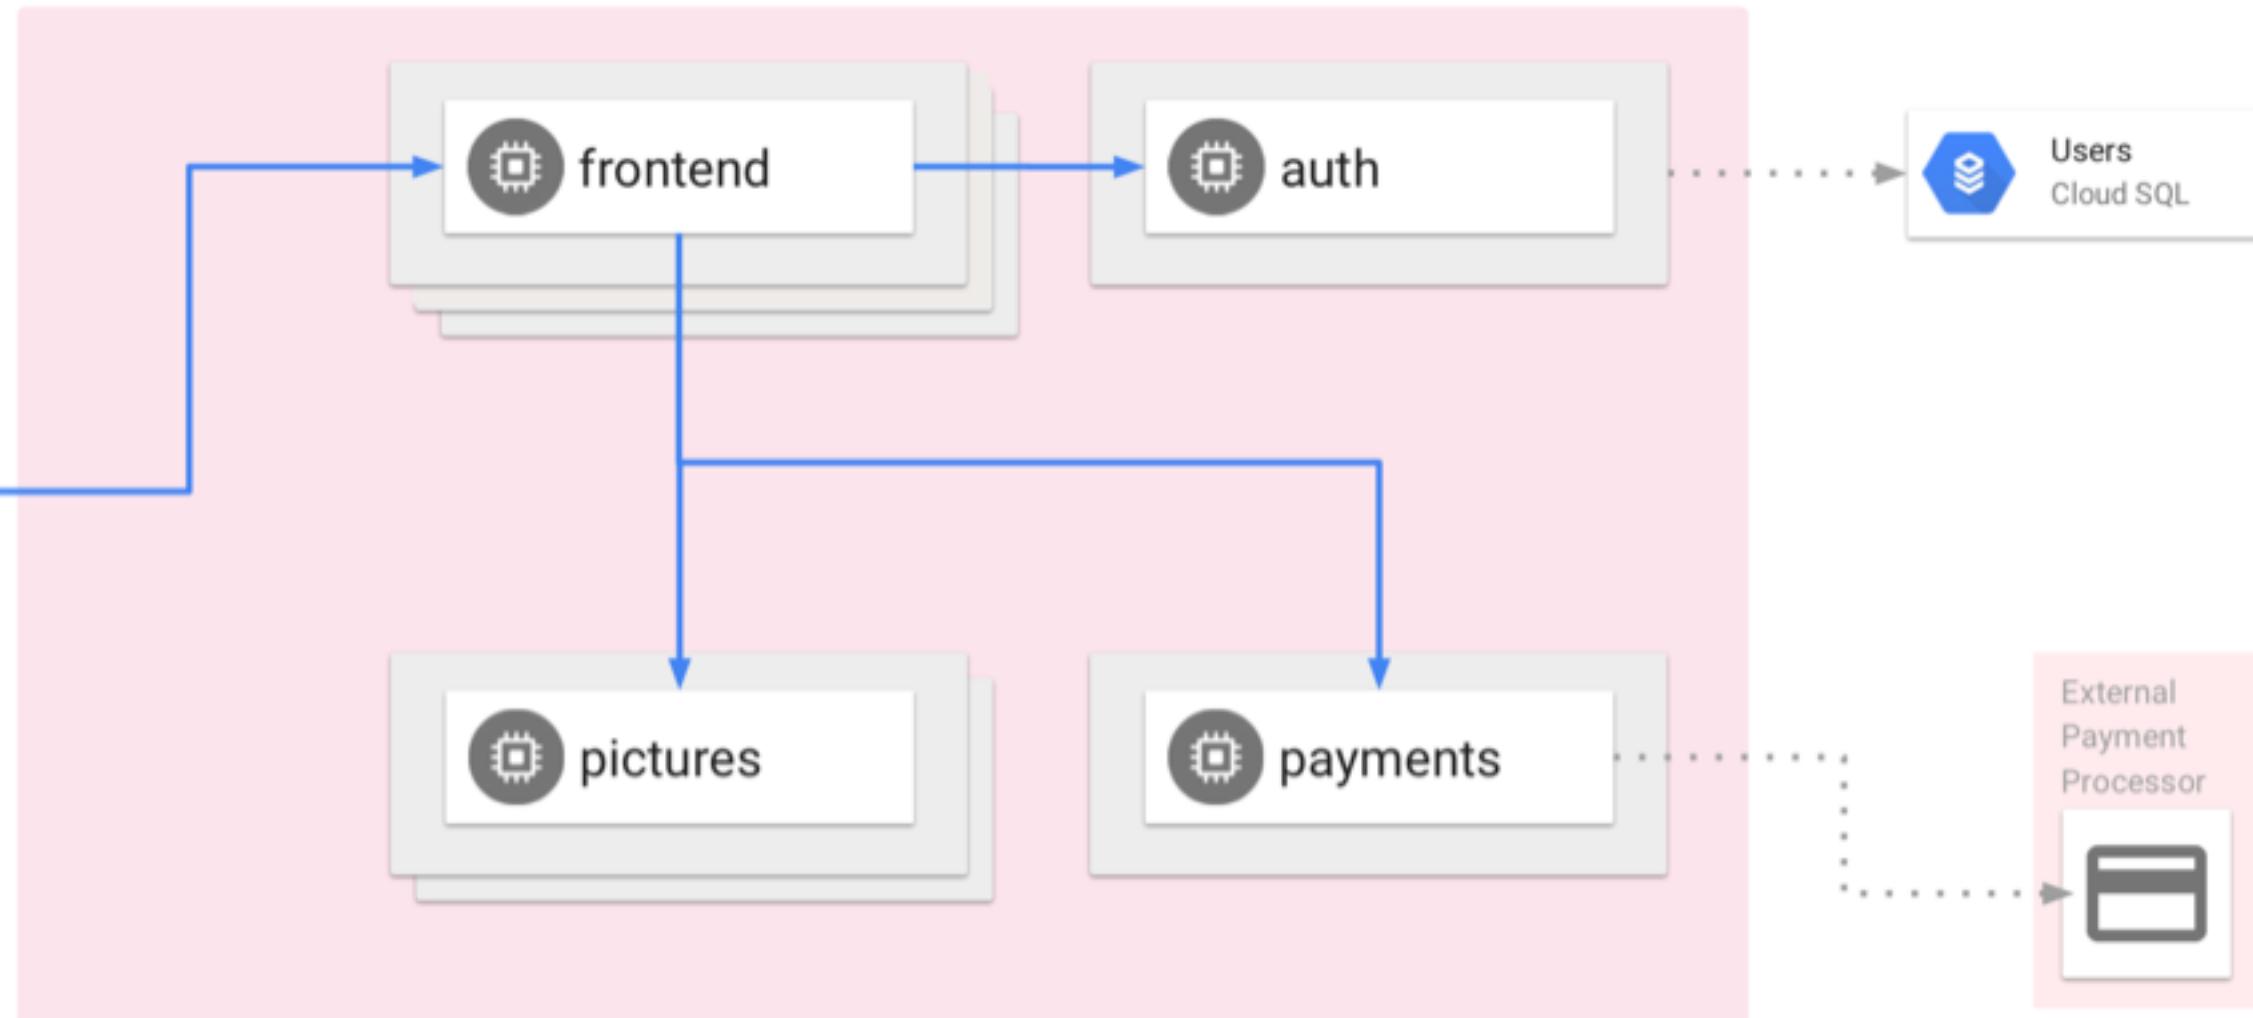


Istio Architecture

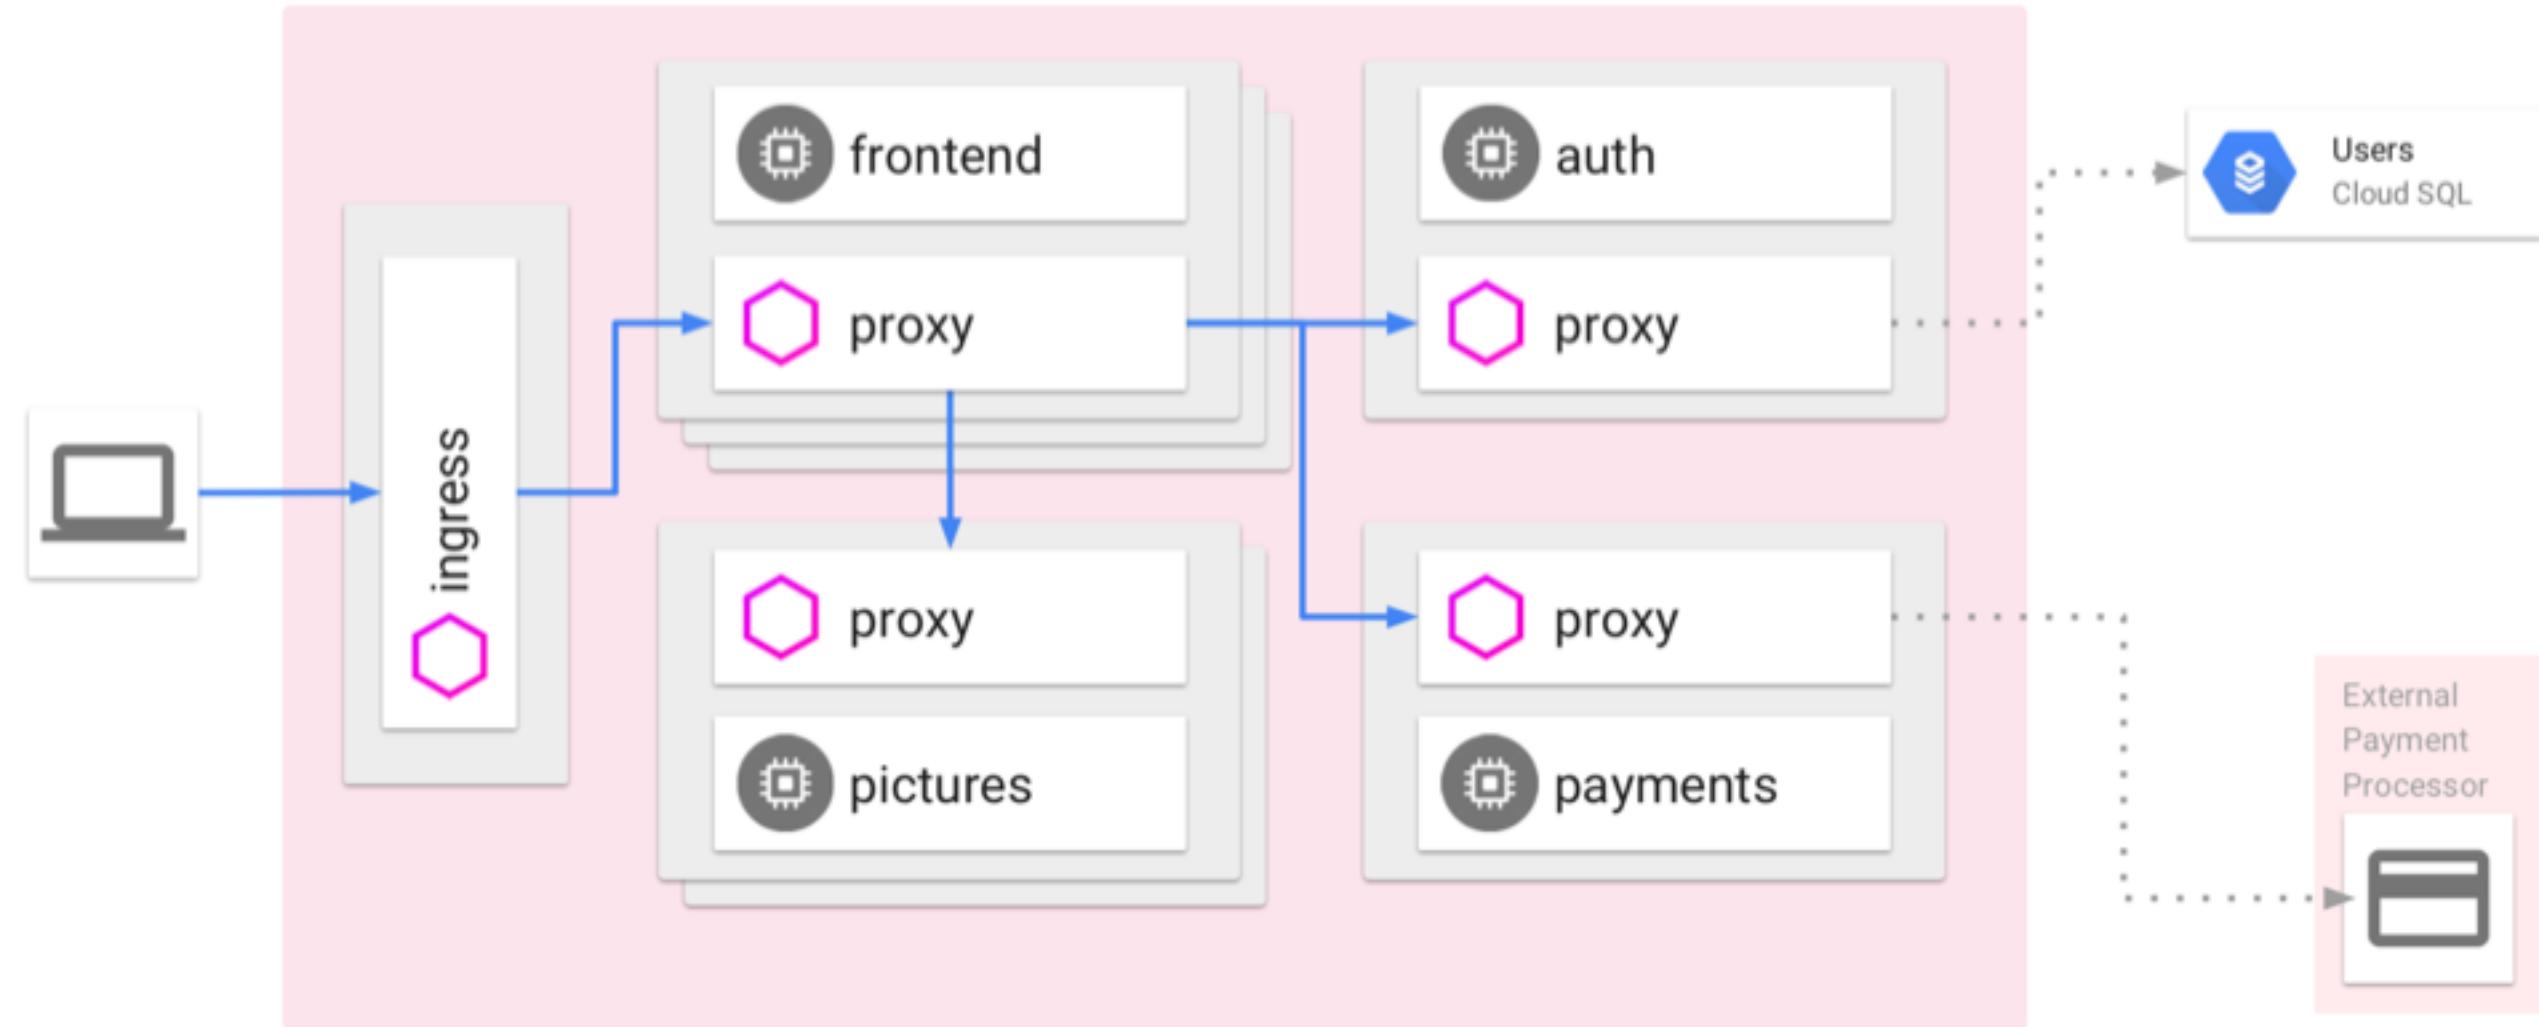
# Service Mesh



# Service Mesh



# Service Mesh

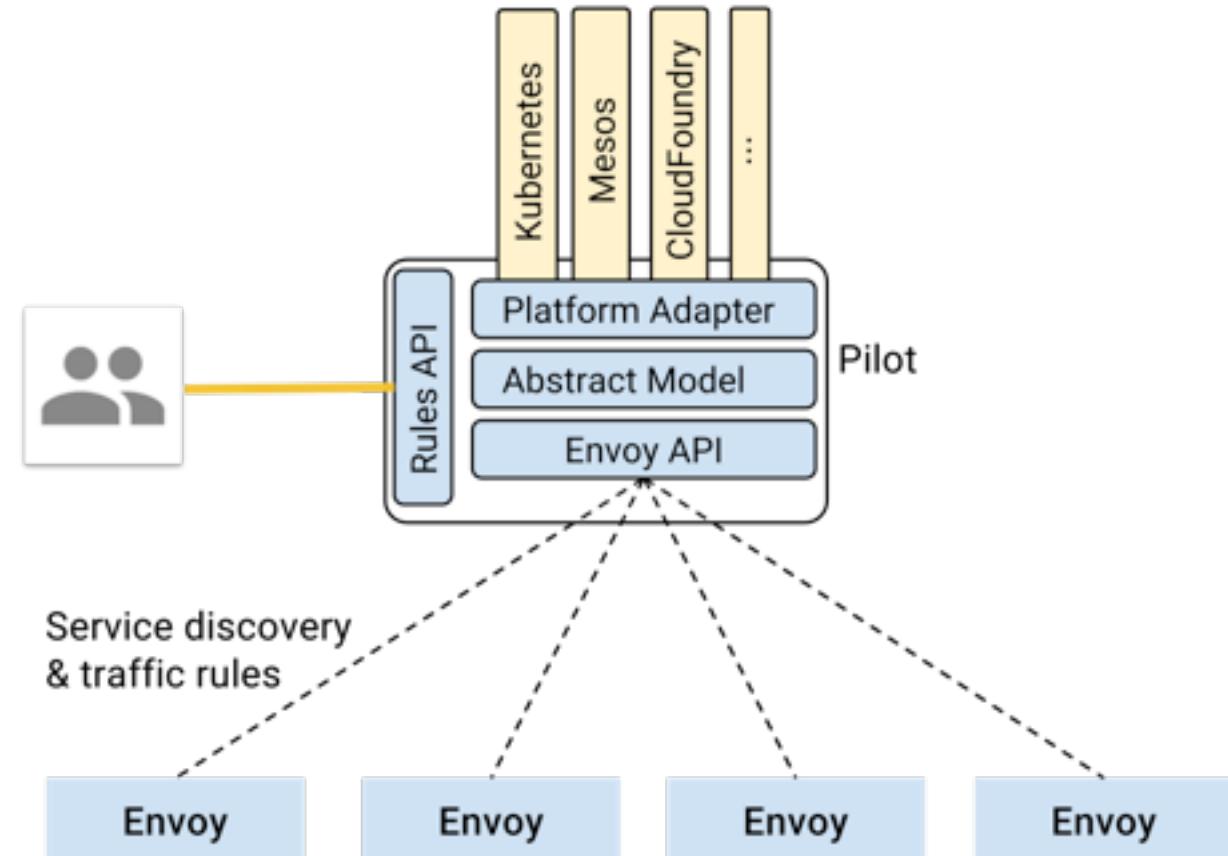


# Istio Components

- Pilot
  - Manages Envoy
- Mixer
  - Input data collector
    - telemetry
    - policy
- Auth/Citadel (0.8)
  - Envoy security
- Envoy
  - Proxy

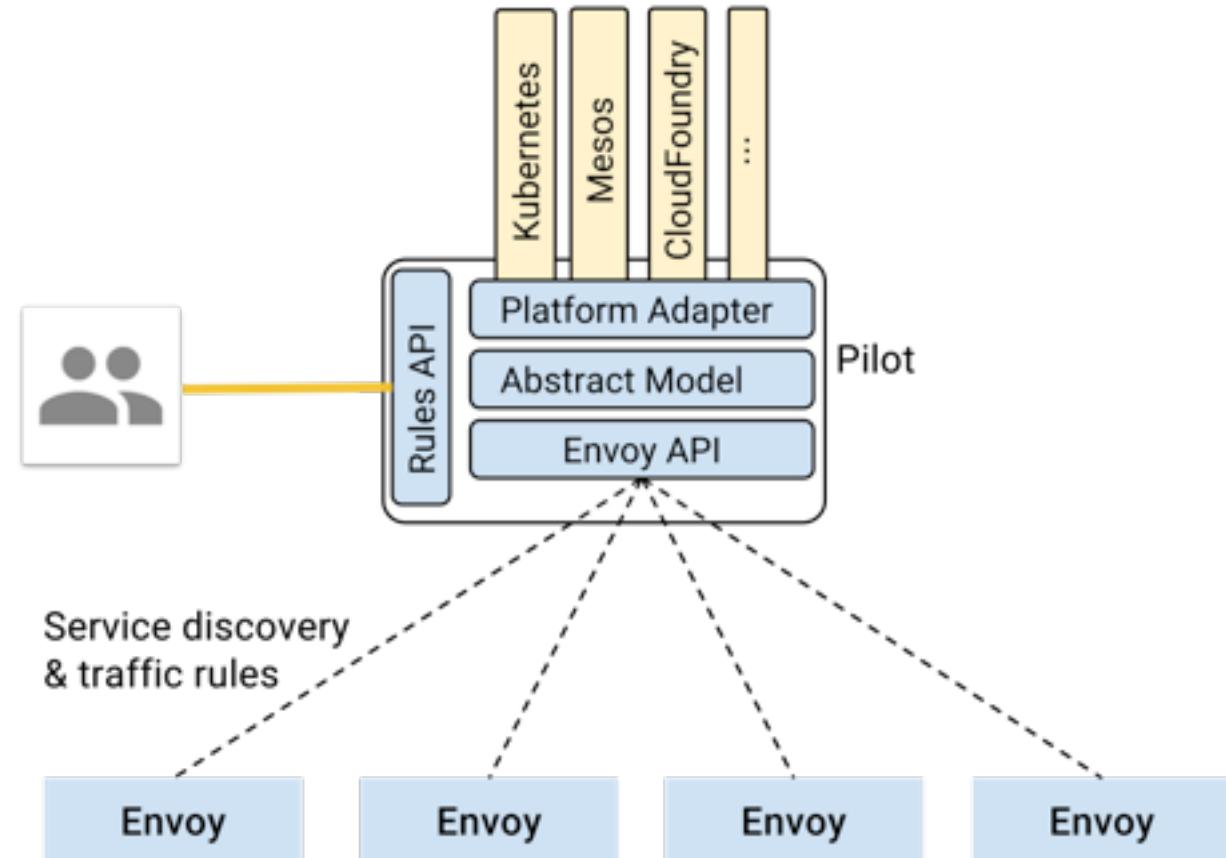
# Pilot

- Manages Envoy proxies
- Exposes APIs
  - Service discovery
  - dynamic updates (LB pools, routing tables)
- Independent of the underlying platform
- Each platform has it's own adapters



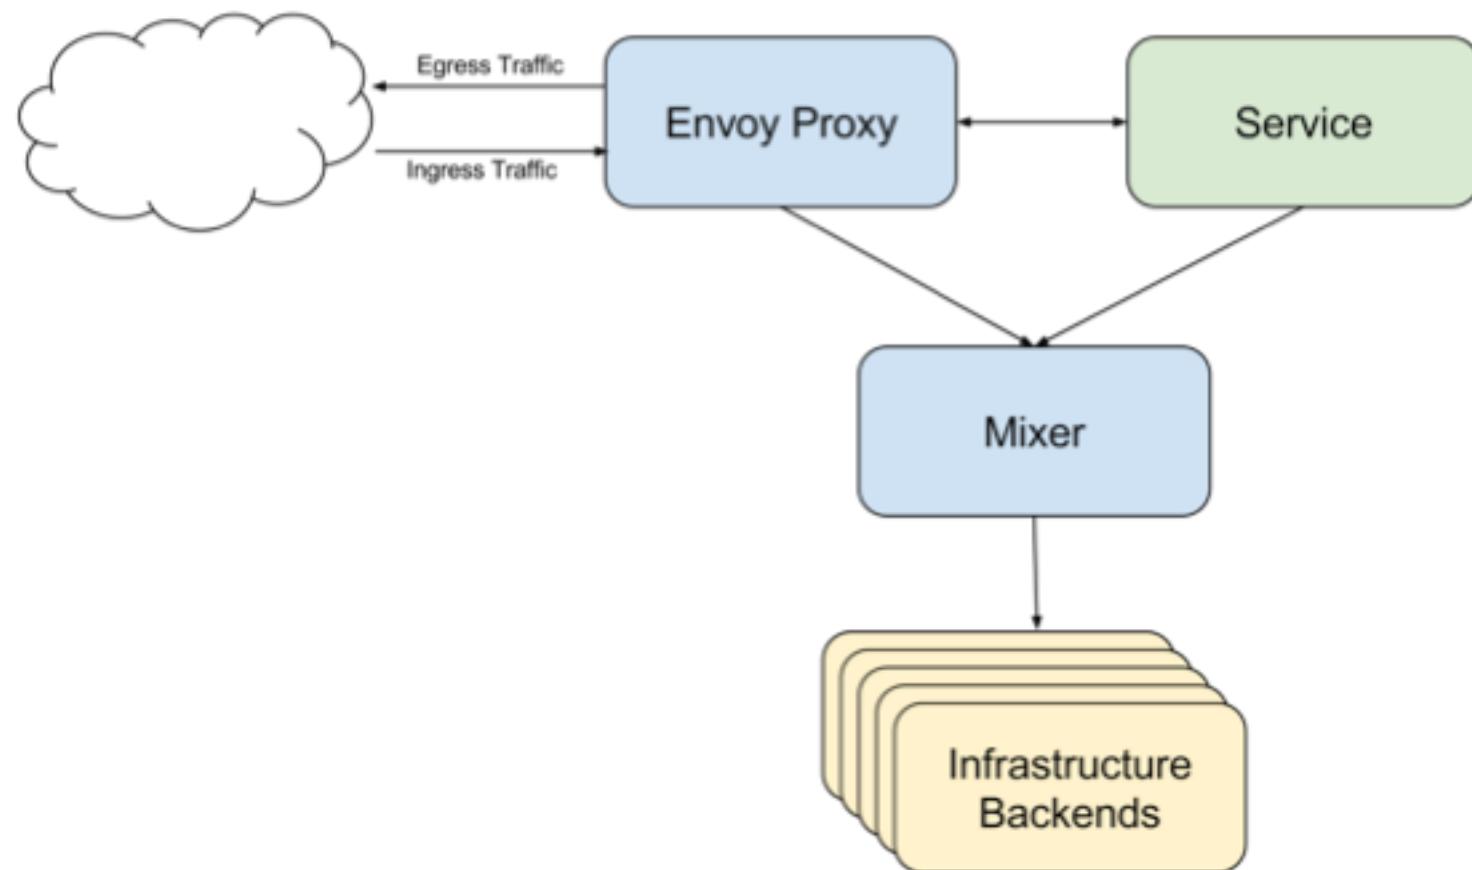
# Pilot

- Kubernetes adapter watches Kubernetes API for changes to:
  - Pod
  - Service registration
  - ingress
  - traffic management rules



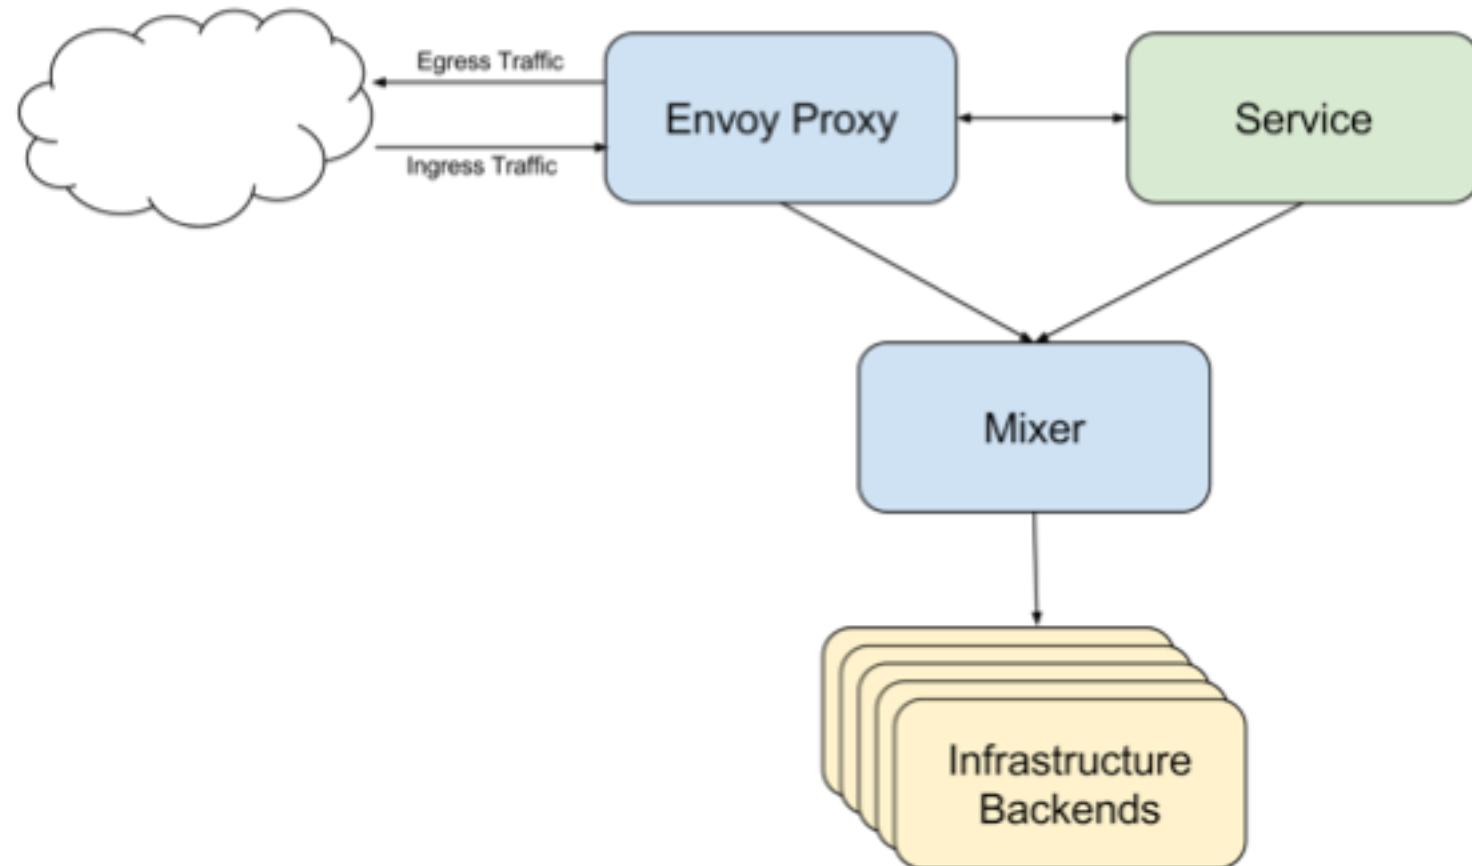
# Mixer

- Abstract services from backends
  - access control
  - telemetry
  - quotas
  - etc...



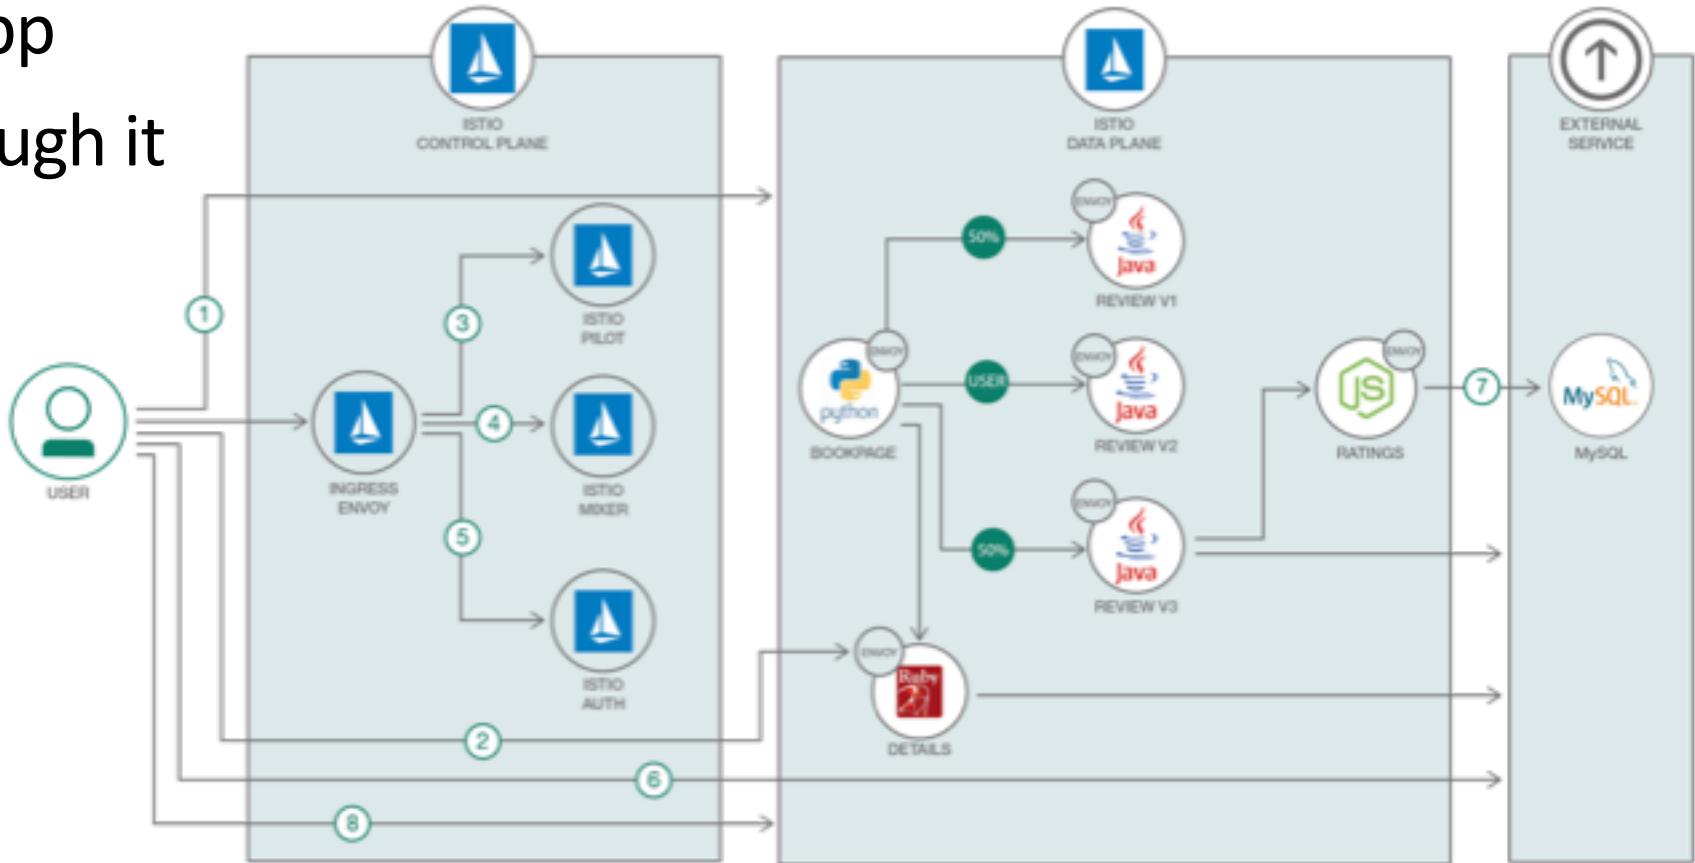
# Mixer

- Precondition checking
  - Check to see if authenticated
  - Is caller on whitelist?
  - Authorization for requested service?
- Quota management
  - API rate limits
- Telemetry reporting
  - Services can report logging and monitoring (future: billing and tracing)



# Envoy

- Deployed alongside app
- All traffic proxied through it
- Build by Lyft
  - 100+ services
  - 10,000+ VMs
  - 2M requests



# Envoy

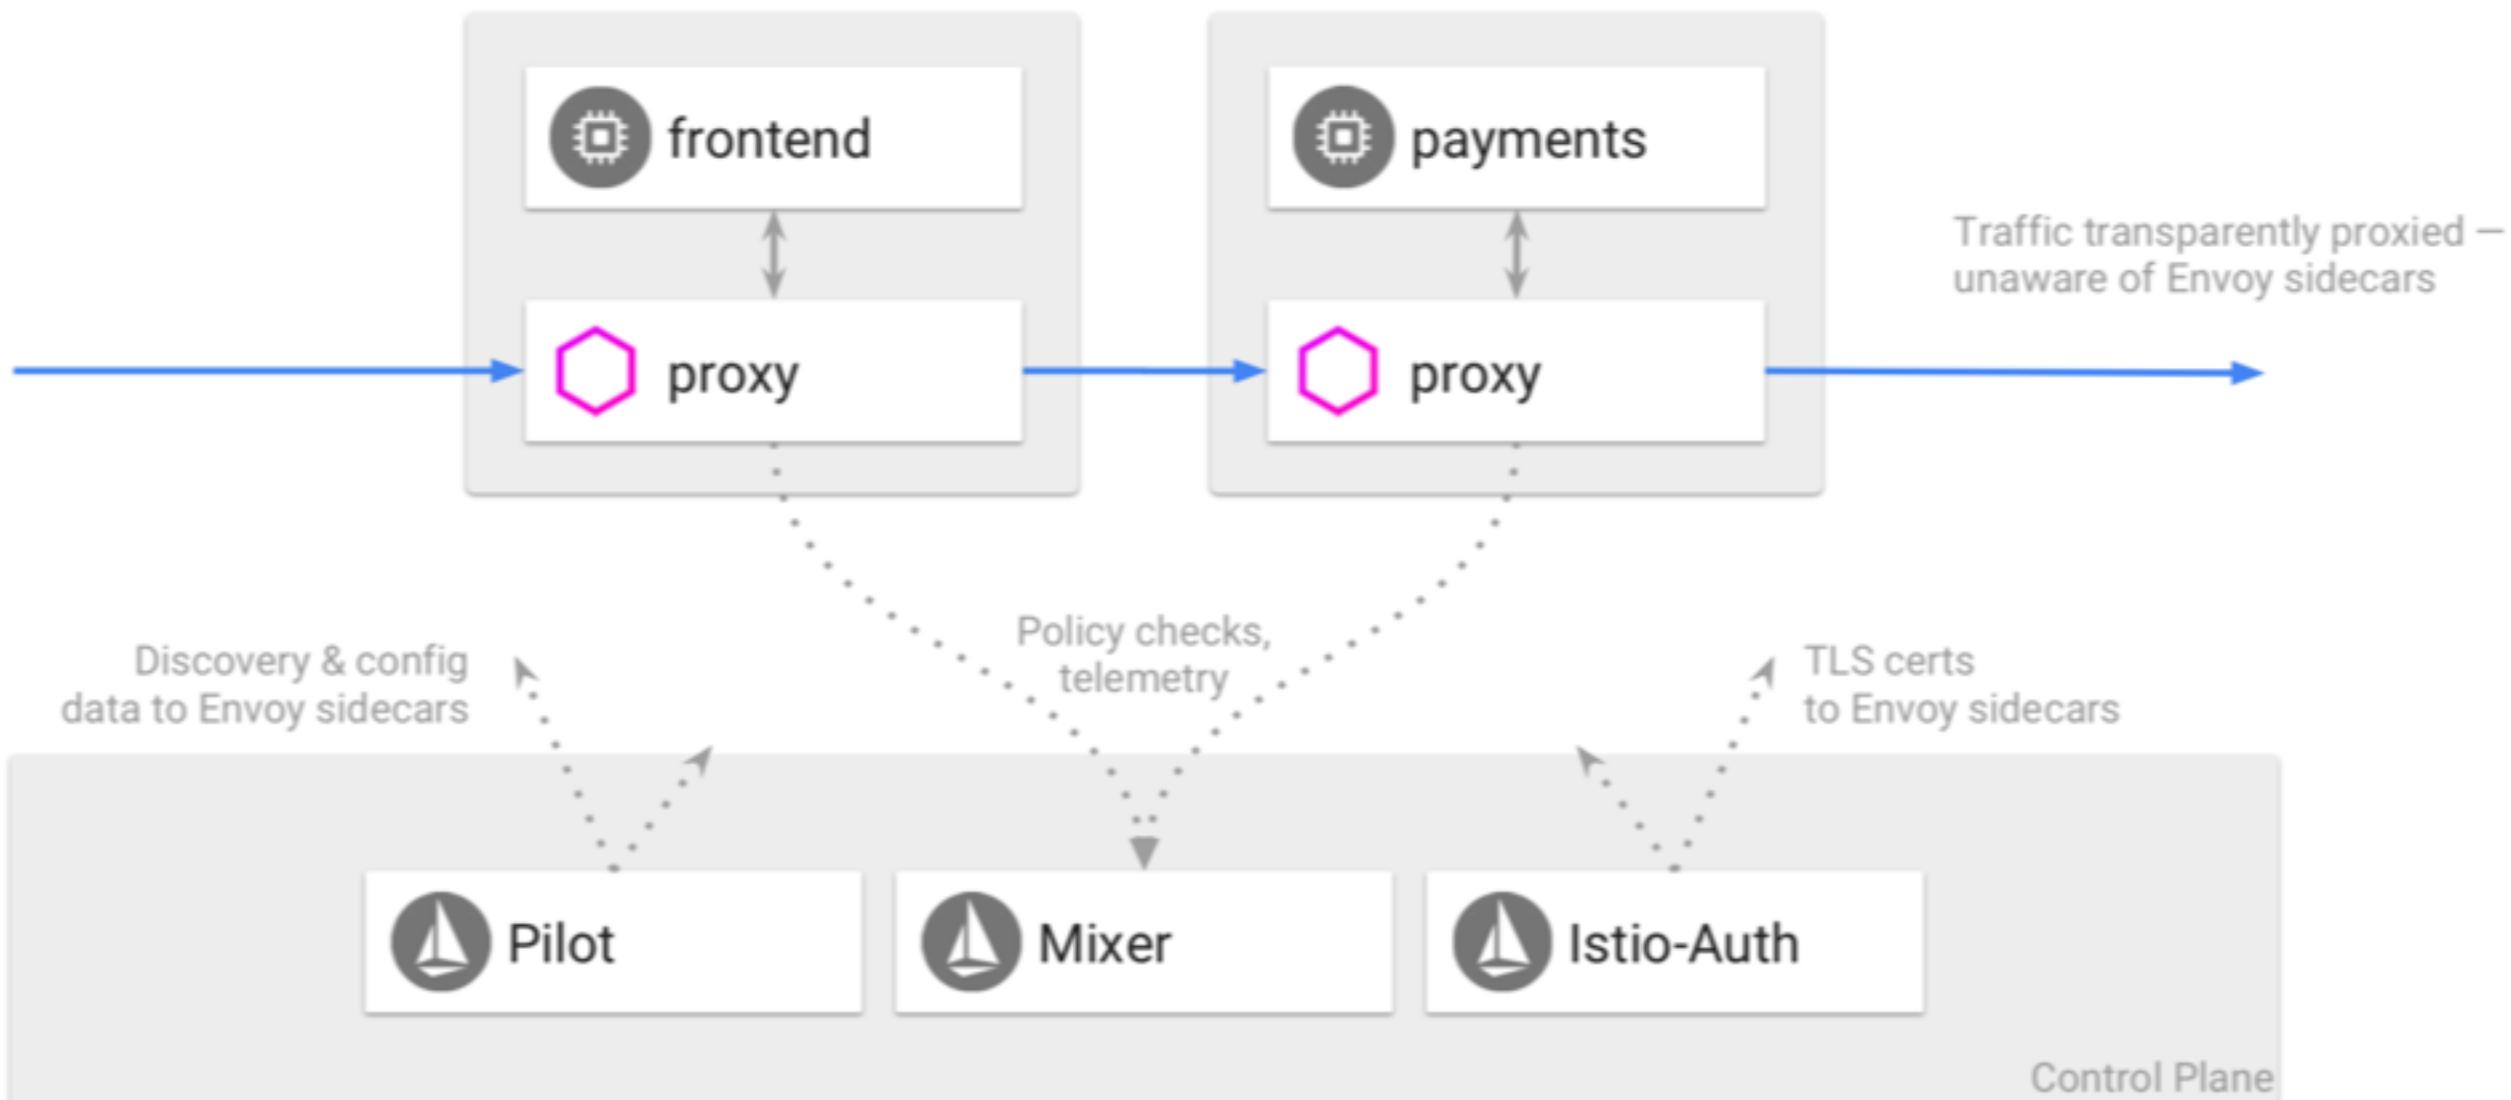
```
spec:  
  containers:  
    - image: frontend:latest
```

```
spec:  
  containers:  
    - image: frontend:latest  
    - image: istio/proxy
```



```
initImage: docker.io/istio/proxy_init  
proxyImage: docker.io/istio/proxy
```

# Envoy



Lab: Install Istio

Lab: Deploy Bookinfo application



# Traffic Management - Rules

Control how API calls & layer-4 traffic flows  
between services in application deployment

- Set service-level properties
  - timeouts
  - retries
  - circuit breakers
- Continuous deployment
  - canary
  - A/B
  - staged rollouts (% based traffic splits)

# Traffic Management - VirtualService

Defines rules that control how requests are routed within an Istio service mesh

- source
- destination
- HTTP paths
- header fields

# VirtualService

// Split traffic

```
kind: VirtualService
```

```
metadata:
```

```
  name: reviews
```

```
spec:
```

```
  hosts:
```

```
    - pictures
```

```
  http:
```

```
    - route:
```

```
      - destination:
```

```
        host: pictures
```

```
        subset: v1
```

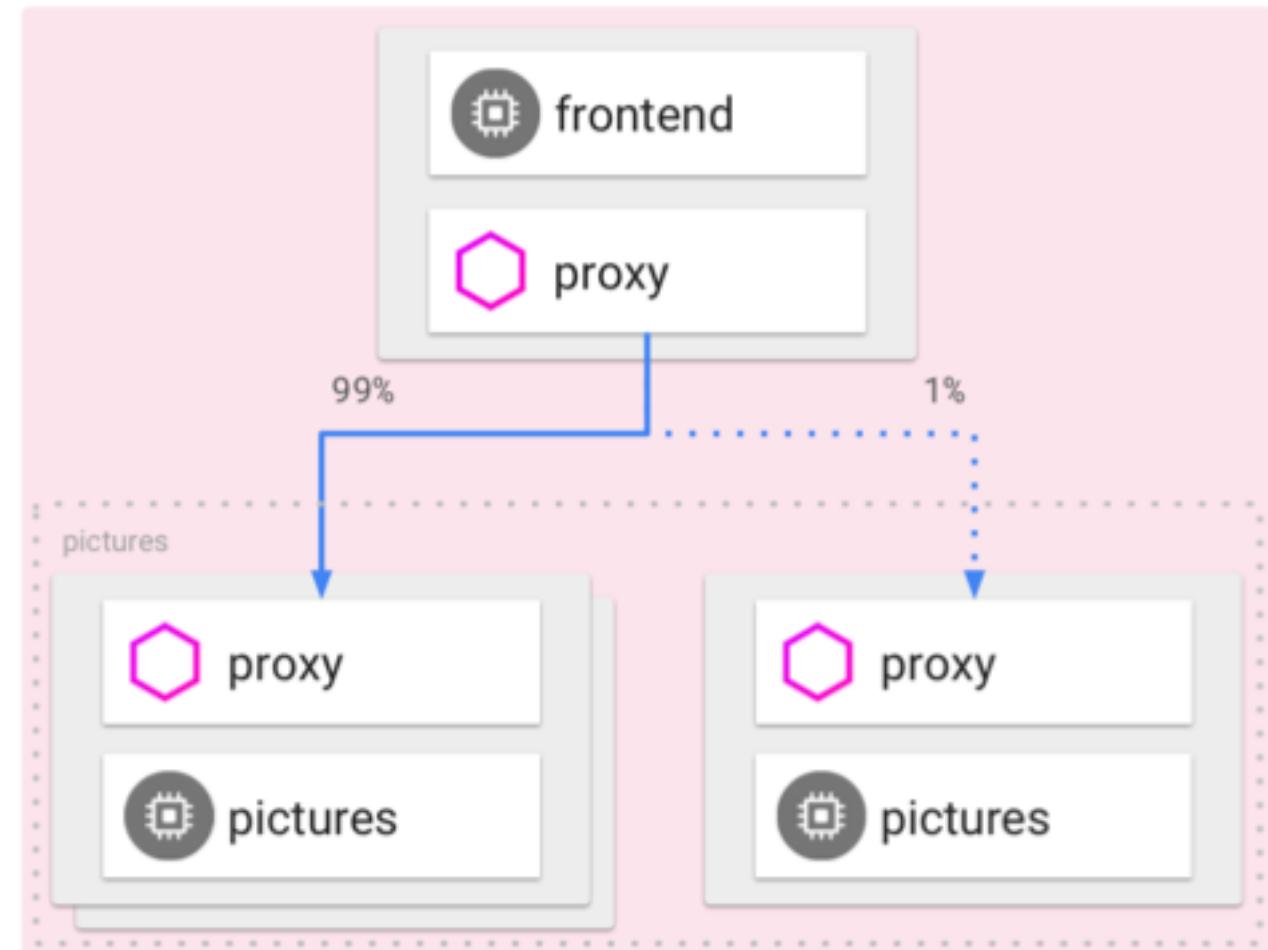
```
        weight: 99
```

```
      - destination:
```

```
        host: pictures
```

```
        subset: v2
```

```
        weight: 1
```



# VirtualService

// content based traffic steering

**kind:** VirtualService

**metadata:**

**name:** pictures

**spec:**

**hosts:**

- pictures

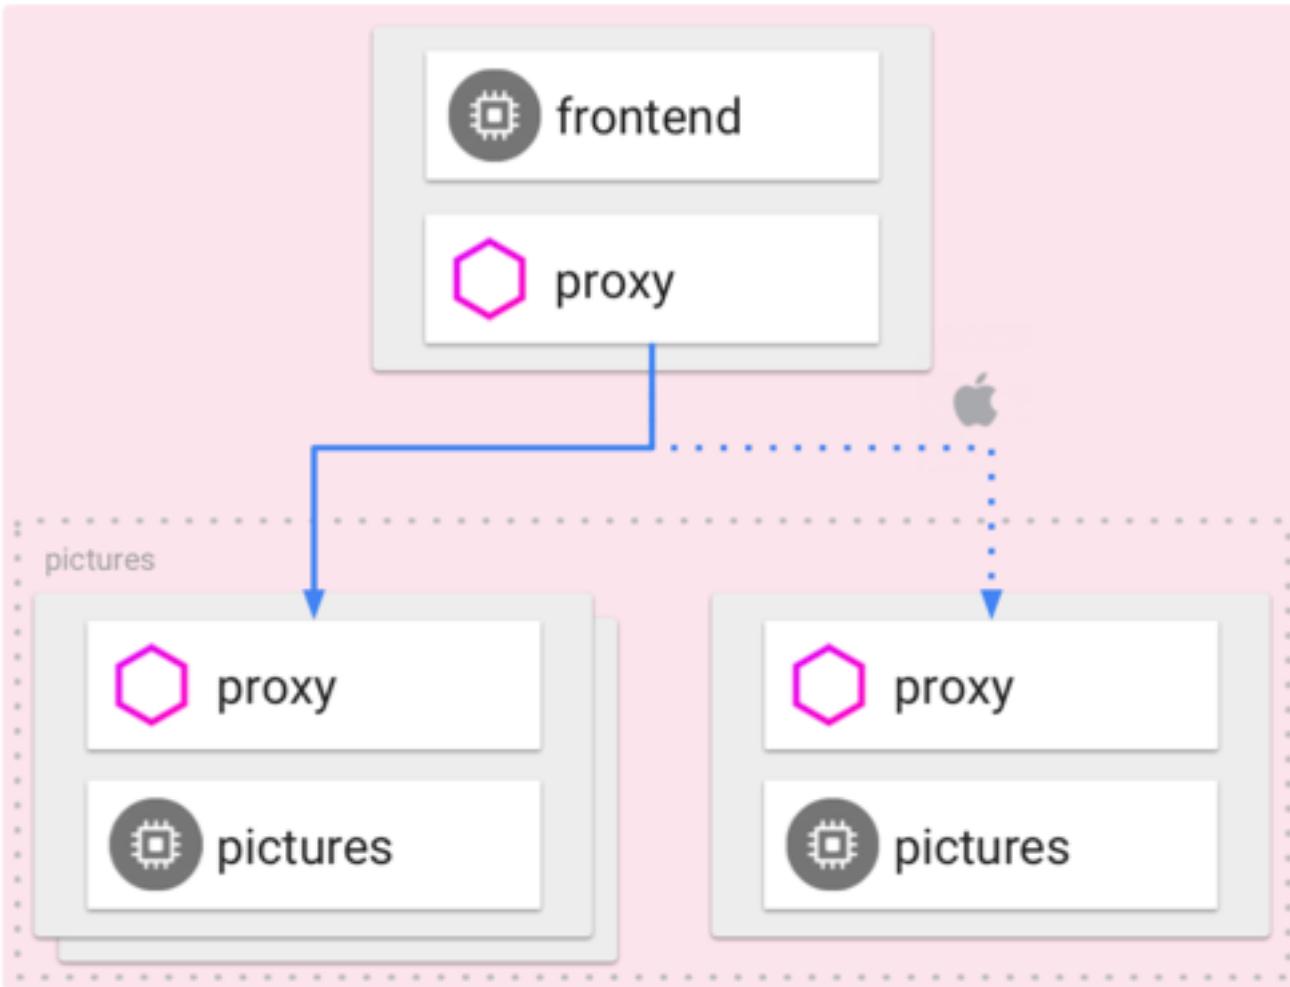
**http:**

- match:

- headers:

**cookie:**

**regex:** ^(.\*)?;(.\* iPhone)(;.\*)?\$"



# VirtualService- config

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: reviews
spec:
  hosts:
  - reviews
  http:
  - route:
    - destination:
        host: reviews
        subset: v1
```

# VirtualService- config

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: reviews
spec:
  hosts:
  - reviews
  http:
  - route:
    - destination:
        host: reviews
        subset: v1
```

The route subset specifies the name of a defined subset in a corresponding destination rule configuration:

# Traffic Management - DestinationRule

Correspond to one or more request destination hosts that are specified in a VirtualService configuration

- Hosts may or may not match actual destination workload
- Route to Service in mesh, or external
- hosts field must specify FQDN(s)

```
hosts:  
  - reviews  
  - bookinfo.com
```

# Traffic Management - DestinationRule

```
apiVersion: networking.istio.io/v1alpha3
kind: DestinationRule
metadata:
  name: reviews
spec:
  host: reviews
  subsets:
    - name: v1
      labels:
        version: v1
    - name: v2
      labels:
        version: v2
```

# Traffic Management - config

```
apiVersion: networking.istio.io/v1alpha3
kind: DestinationRule
metadata:
  name: reviews
spec:
  host: reviews
  subsets:
    - name: v1
      labels:
        version: v1
    - name: v2
      labels:
        version: v2
```

Corresponding DestinationRule  
with matching subset 'v1'

# Qualify Rules - Specific caller

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: ratings
spec:
  hosts:
  - ratings
  http:
  - match:
    - sourceLabels:
      app: reviews
    ...^
```

# Qualify Rules - Specific caller

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: ratings
spec:
  hosts:
  - ratings
  http:
  - match:
    sourceLabels:
      app: reviews
  ...
```

- `sourceLabels`: matches  
Kubernetes service selector labels

# Qualify Rules - Specific caller

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: ratings
spec:
  hosts:
  - ratings
  http:
  - match:
    sourceLabels:
      app: reviews
      version: v2
  ...
```

- Can also specify a specific version

# Qualify Rules - HTTP headers

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: reviews
spec:
  hosts:
    - reviews
  http:
    - match:
        - headers:
            cookie:
              regex: "^(.*?;)?(user=jason)(;.*?)$"
    ... ~
```

# Qualify Rules - HTTP headers

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: reviews
spec:
  hosts:
    - reviews
  http:
    - match:
        - headers:
            cookie:
              regex: "^(.*?;)?(user=jason)(;.*?)$"
    ...
...
```

- Match user=jason HTTP header

# Qualify Rules - HTTP headers

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: ratings
spec:
  hosts:
  - ratings
  http:
  - match:
    - sourceLabels:
        app: reviews
        version: v2
      headers:
        cookie:
          regex: "^(.*?;)?(user=jason)(;.*?)?$"
    ...~
```

# Qualify Rules - HTTP headers

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: ratings
spec:
  hosts:
  - ratings
  http:
  - match:
    - sourceLabels:
        app: reviews
        version: v2
      headers:
        cookie:
          regex: "^(.*?;)?(user=jason)(;.*?)?$"
    ...

```

- Multiple conditions in same  
'match' (AND)

# Qualify Rules - HTTP headers

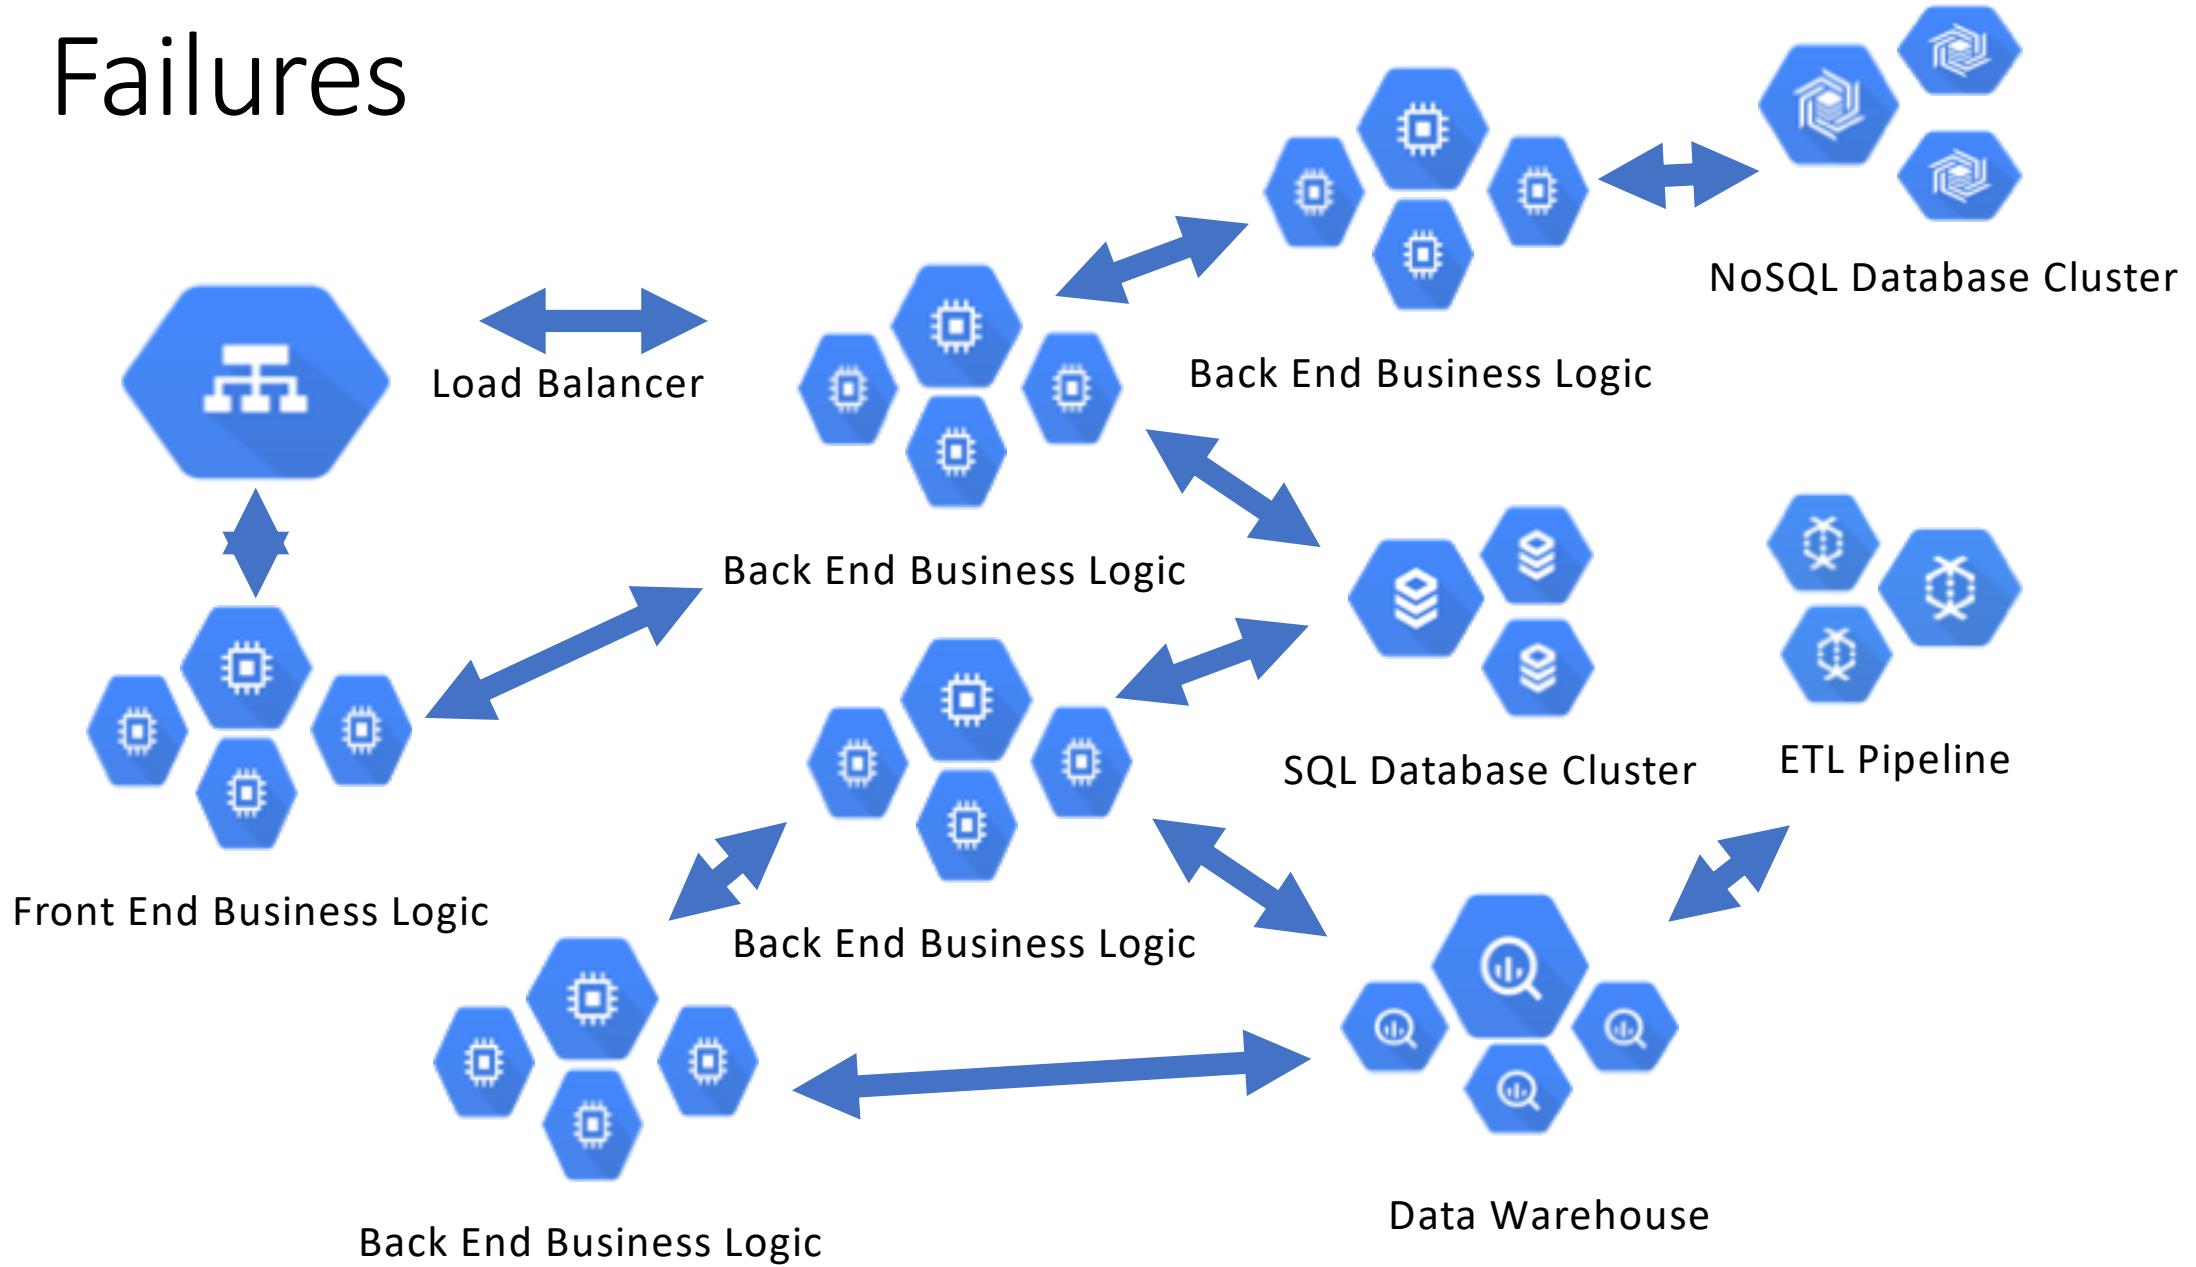
```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: ratings
spec:
  hosts:
  - ratings
  http:
  - match:
    - sourceLabels:
        app: reviews
        version: v2
  - match:
    headers:
      cookie:
        regex: "^(.*?; )?(user=jason)(; .*?)?$"
  ...~
```

- Multiple 'match' (OR)

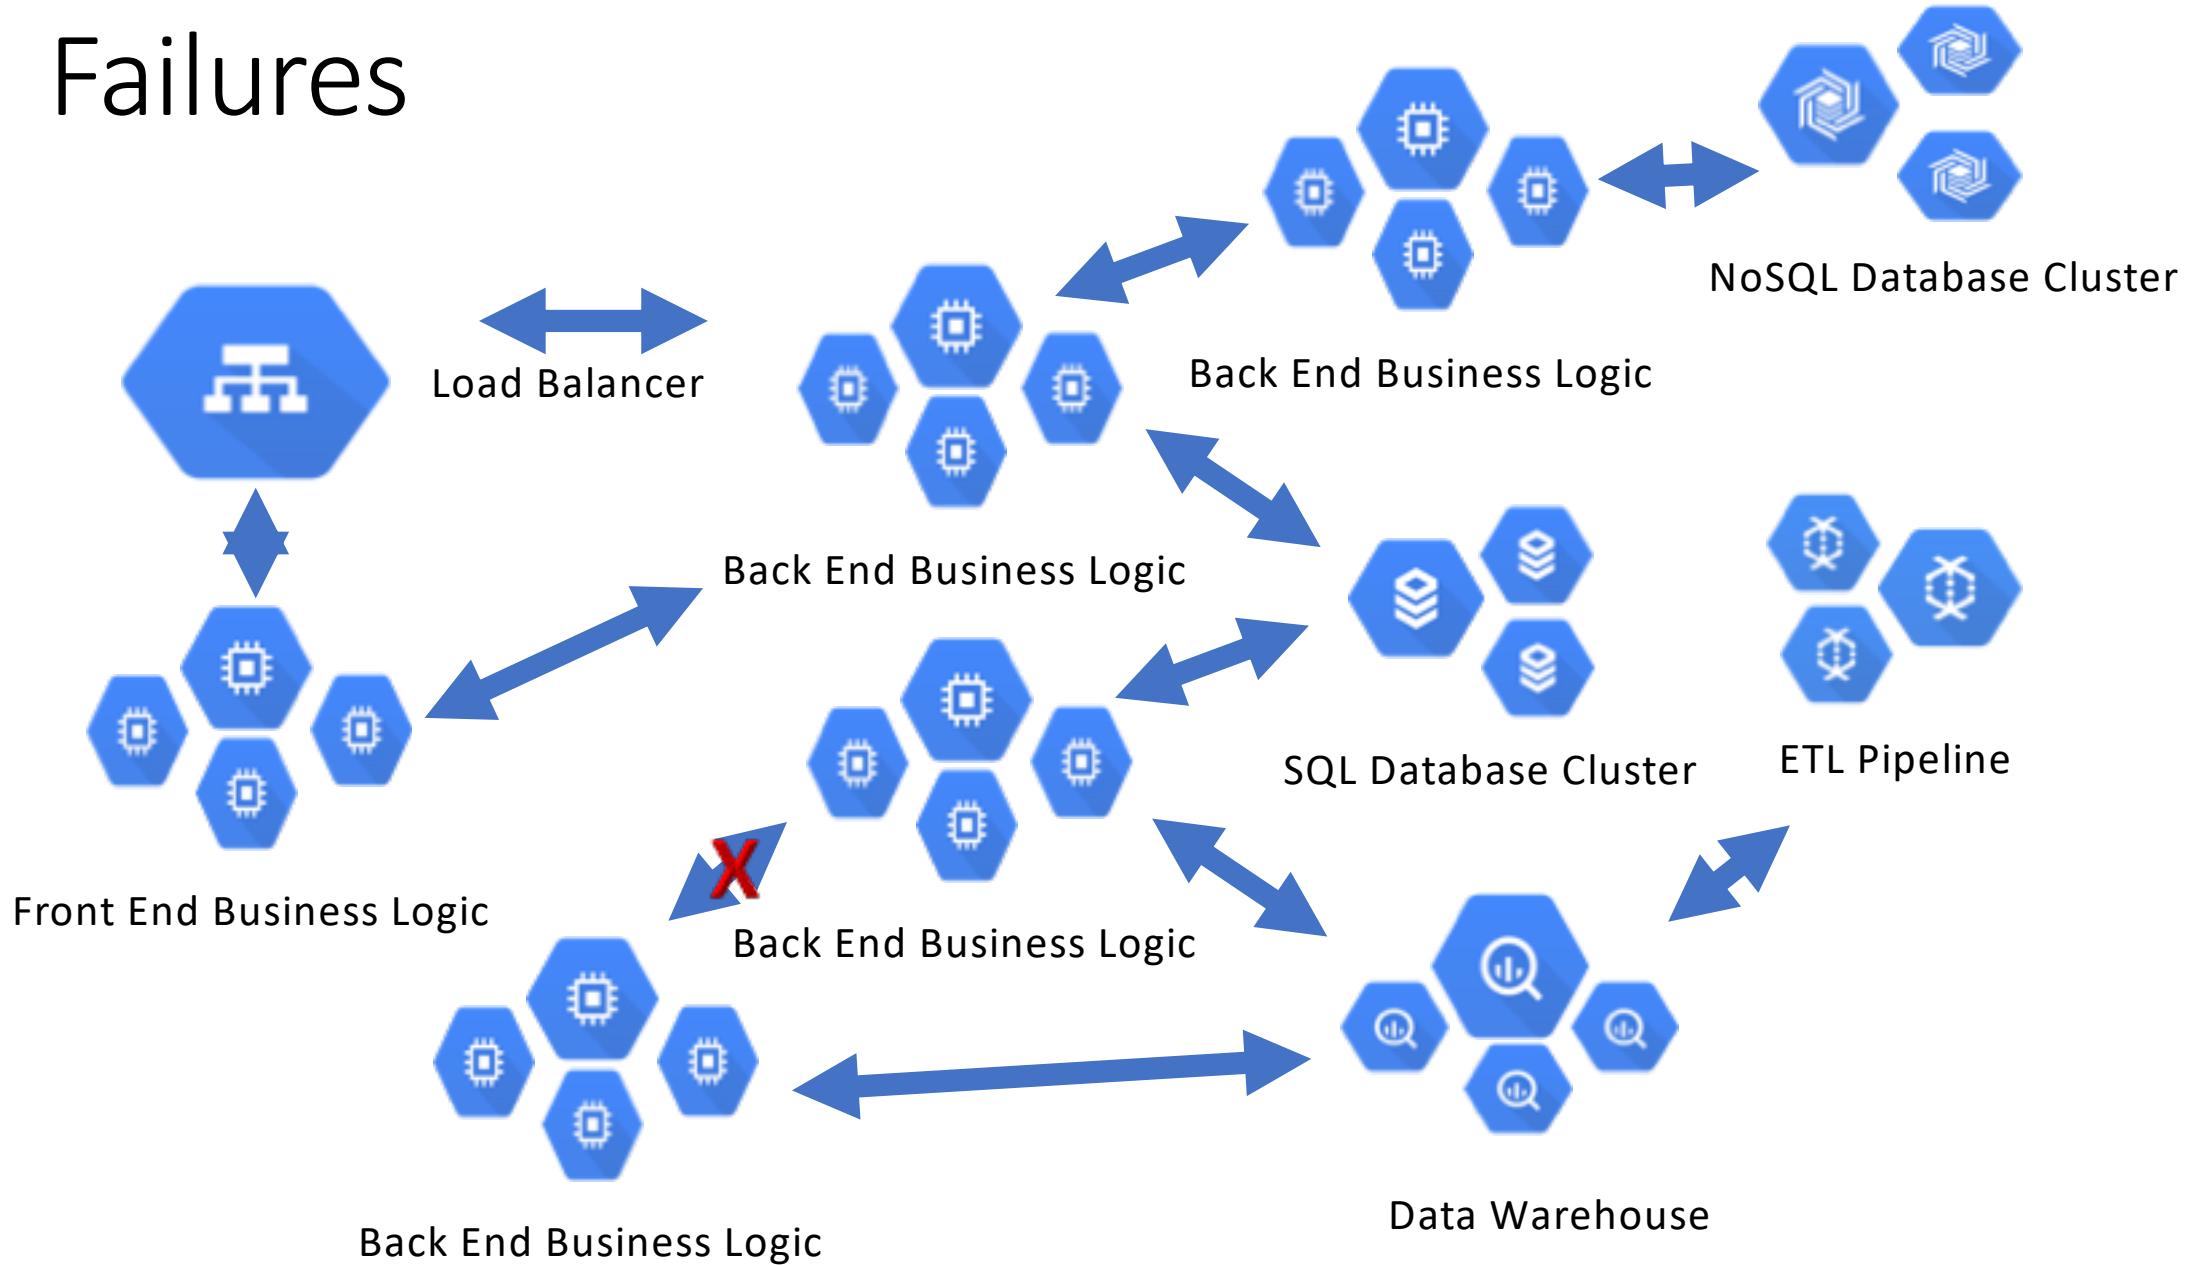
# Istio: Handling failures



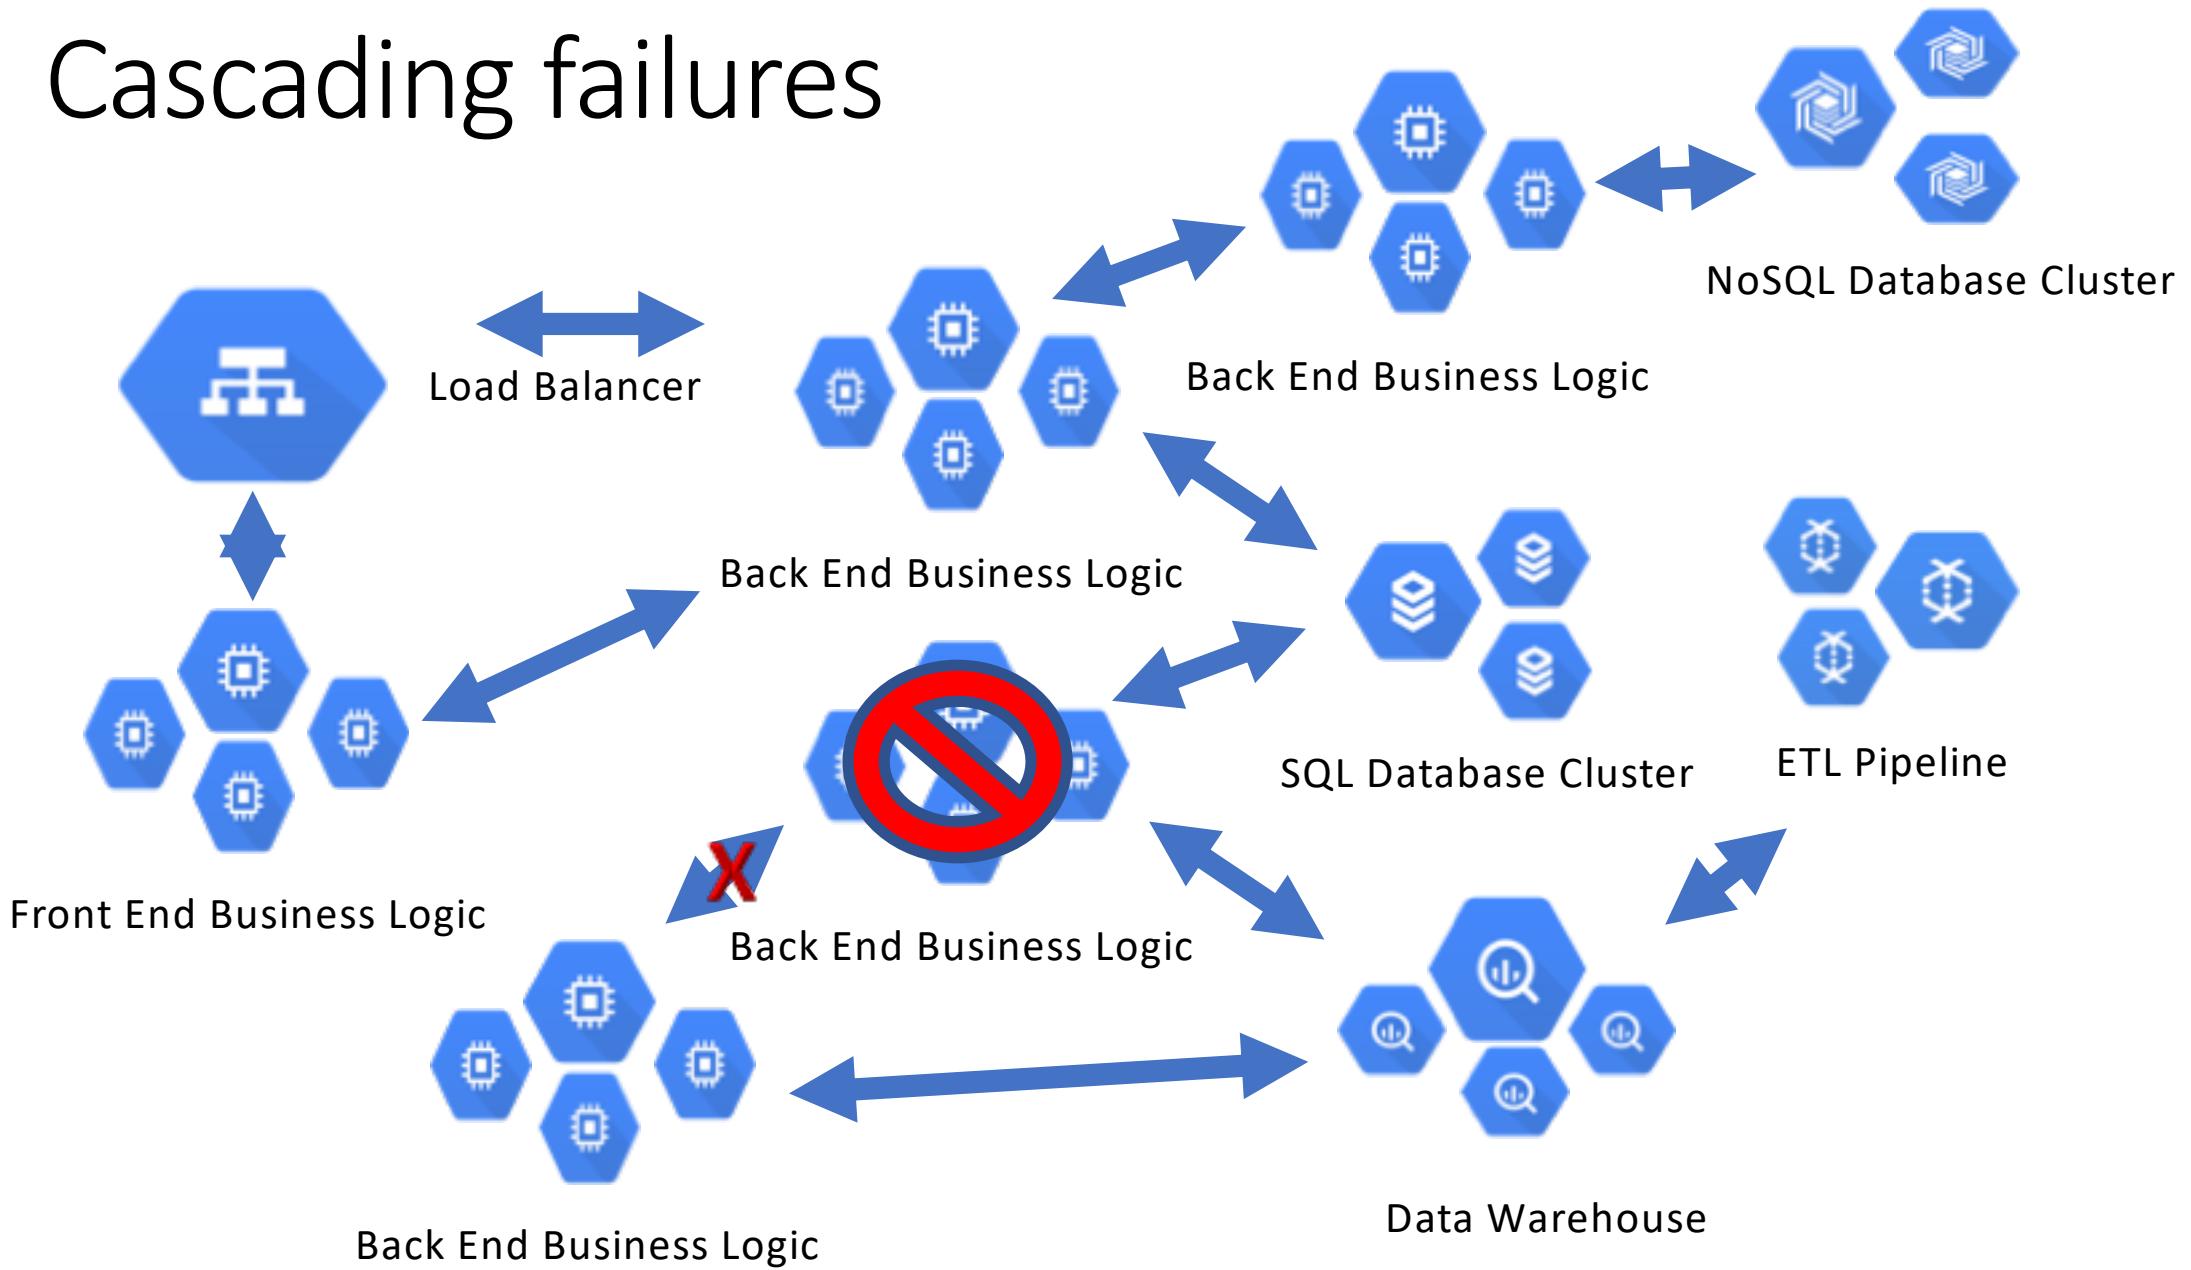
# Failures



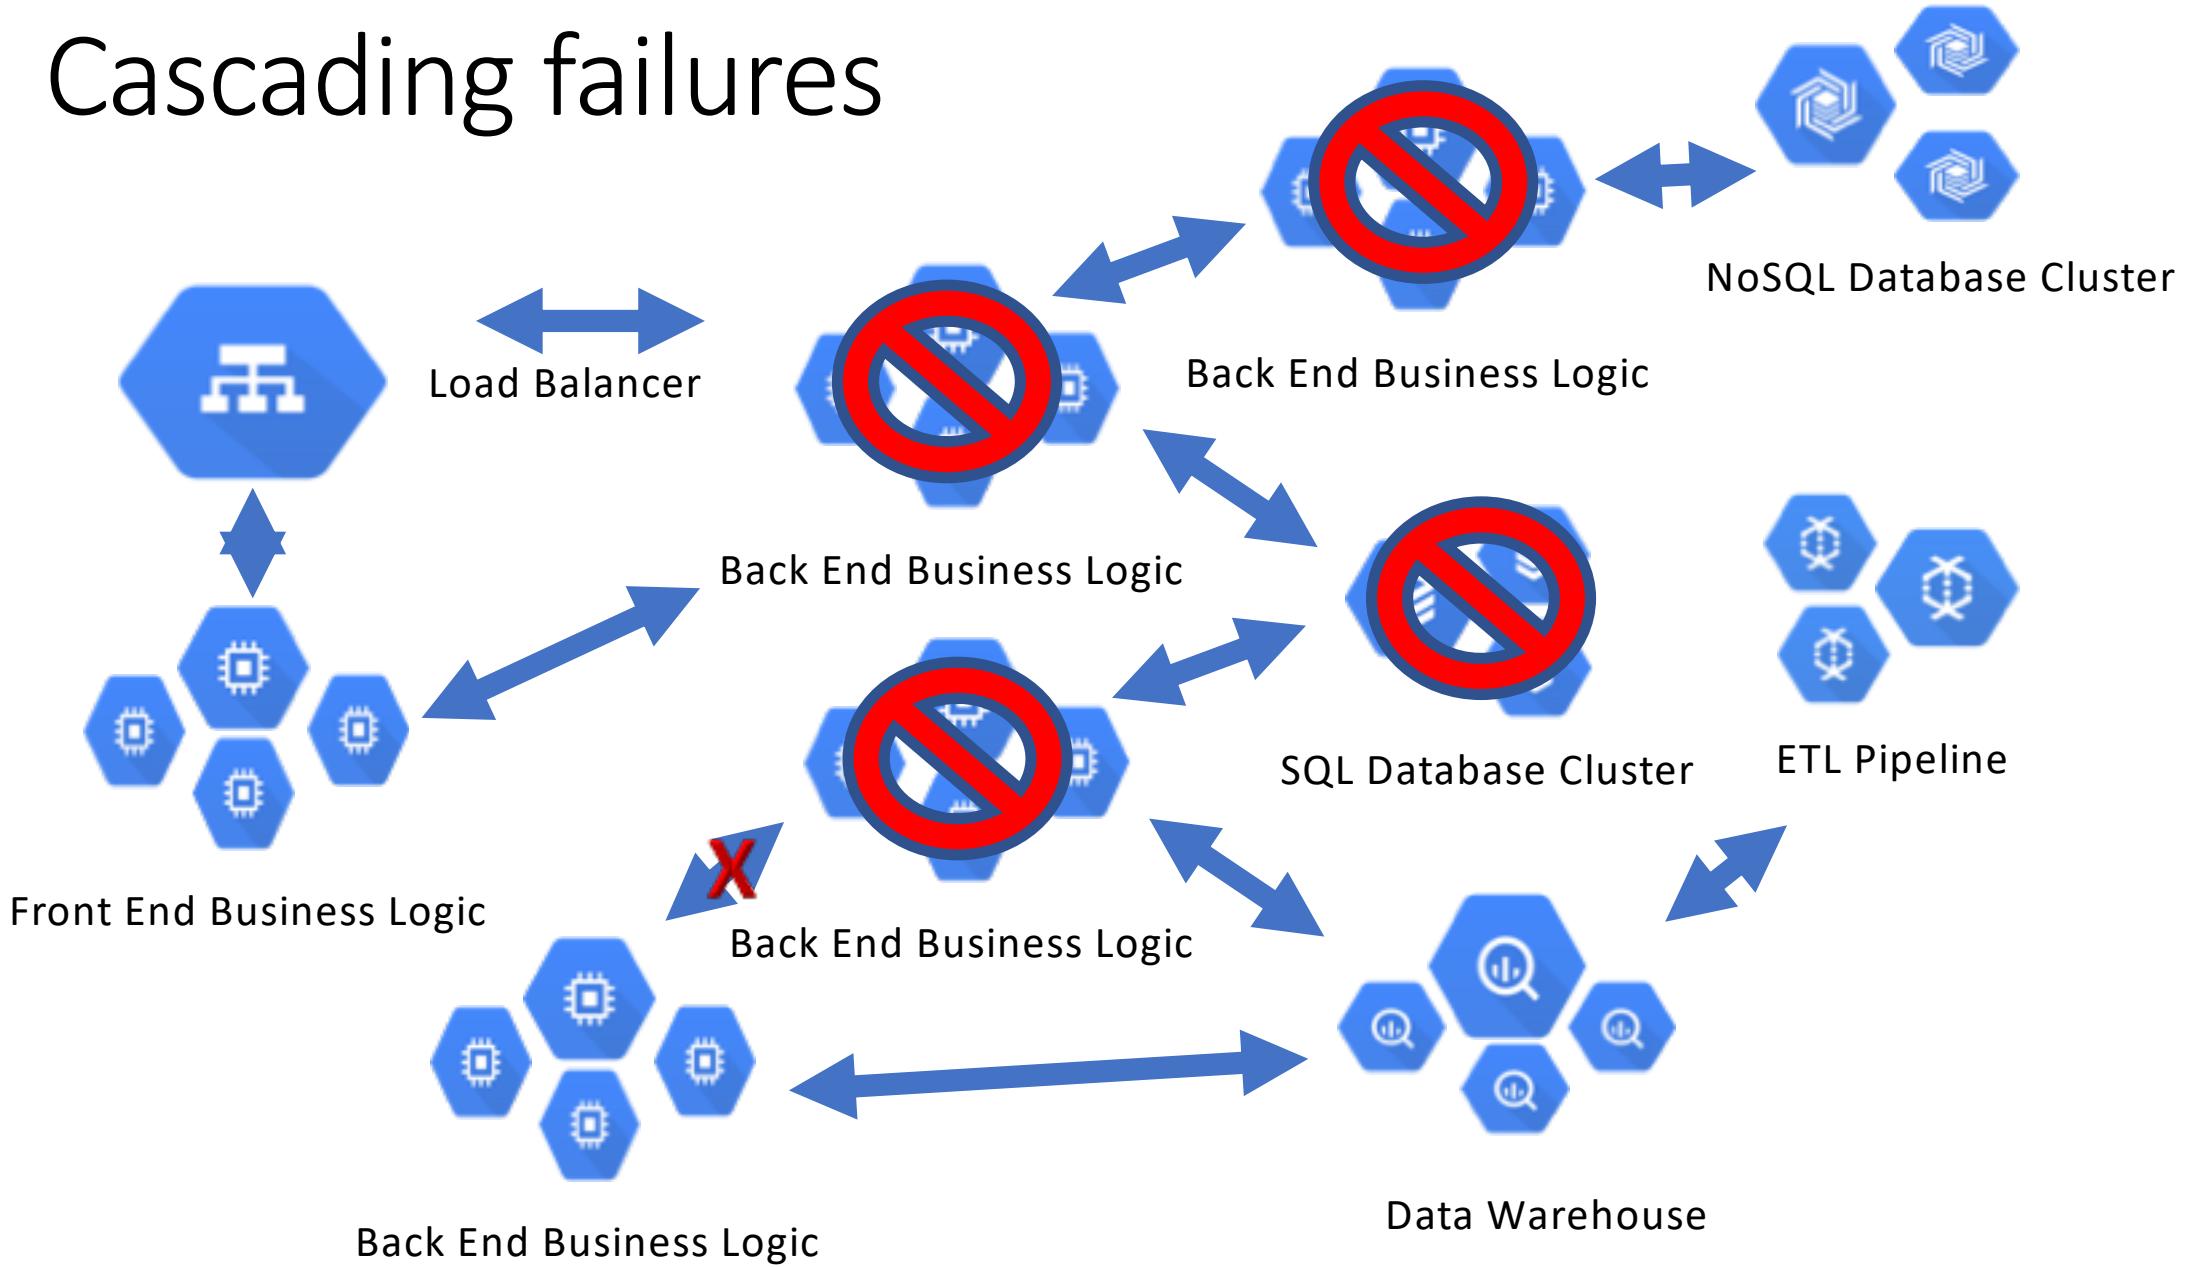
# Failures



# Cascading failures



# Cascading failures





# Timeouts & Retries

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: ratings
spec:
  hosts:
    - ratings
  http:
    - route:
        - destination:
            host: ratings
            subset: v1
        timeout: 10s
```

# Timeouts & Retries

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: ratings
spec:
  hosts:
    - ratings
  http:
    - route:
        - destination:
            host: ratings
            subset: v1
        timeout: 10s
```

- Change timeout from default 15 to 10 seconds

# Timeouts & Retries

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: ratings
spec:
  hosts:
    - ratings
  http:
    - route:
        - destination:
            host: ratings
            subset: v1
        retries:
          attempts: 3
          perTryTimeout: 2s
```

# Timeouts & Retries

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: ratings
spec:
  hosts:
    - ratings
  http:
    - route:
        - destination:
            host: ratings
            subset: v1
            retries:
              attempts: 3
              perTryTimeout: 2s
```

- Set retries

# Testing Services

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: ratings
spec:
  hosts:
  - ratings
  http:
  - fault:
      delay:
        percent: 10
        fixedDelay: 5s
    route:
    - destination:
        host: ratings
        subset: v1
```

# Testing Services

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: ratings
spec:
  hosts:
  - ratings
  http:
  - fault:
    delay:
      percent: 10
      fixedDelay: 5s
    route:
    - destination:
        host: ratings
        subset: v1
```

- Inject fault (delay 10% by 5 secs)

# Testing Services

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: ratings
spec:
  hosts:
  - ratings
  http:
  - fault:
      abort:
        percent: 10
        httpStatus: 400
    route:
    - destination:
        host: ratings
        subset: v1
```

# Testing Services

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: ratings
spec:
  hosts:
    - ratings
    http:
      - fault:
          abort:
            percent: 10
            httpStatus: 400
      route:
        - destination:
            host: ratings
            subset: v1
```

- Inject fault (400 error code 10%)

# Testing Services

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: ratings
spec:
  hosts:
  - ratings
  http:
  - match:
    - sourceLabels:
        app: reviews
        version: v2
    fault:
      delay:
        fixedDelay: 5s
      abort:
        percent: 10
        httpStatus: 400
    route:
    - destination:
        host: ratings
        subset: v1
```

# Testing Services

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: ratings
spec:
  hosts:
  - ratings
  http:
  - match:
    - sourceLabels:
        app: reviews
        version: v2
    fault:
      delay:
        fixedDelay: 5s
      abort:
        percent: 10
        httpStatus: 400
    route:
    - destination:
        host: ratings
        subset: v1
```

- Use delay and abort together to confirm services are resilient

# Rule Priority

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: reviews
spec:
  hosts:
    - reviews
  http:
    - match:
        - headers:
            Foo:
              exact: bar
        route:
          - destination:
              host: reviews
              subset: v2
    - route:
        - destination:
            host: reviews
            subset: v1
```

- Read from top to bottom
- Once match found all other rules are ignored
- Best practice: Higher priority rules first

# Rule Priority

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: reviews
spec:
  hosts:
    - reviews
  http:
    - match:
        - headers:
            Foo:
              exact: bar
        route:
          - destination:
              host: reviews
              subset: v2
    - route:
        - destination:
            host: reviews
            subset: v1
```

- Checks for  $\text{Foo} = \text{bar}$ , if match sends to v2
- If no match sends to v1

# Traffic Management - DestinationRule

configures the set of policies to be applied to a request after VirtualService routing has occurred.

- authored by service owners
- circuit breaker
- load balancer settings
- TLS settings

# DestinationRule

```
apiVersion: networking.istio.io/v1alpha3
kind: DestinationRule
metadata:
  name: reviews
spec:
  host: reviews
  trafficPolicy:
    loadBalancer:
      simple: RANDOM
  subsets:
  - name: v1
    labels:
      version: v1
  - name: v2
    labels:
      version: v2
    trafficPolicy:
      loadBalancer:
        simple: ROUND_ROBIN
  - name: v3
    labels:
      version: v3
```

# DestinationRule

```
apiVersion: networking.istio.io/v1alpha3
kind: DestinationRule
metadata:
  name: reviews
spec:
  host: reviews
  trafficPolicy:
    loadBalancer:
      simple: RANDOM
  subsets:
  - name: v1
    labels:
      version: v1
  - name: v2
    labels:
      version: v2
  trafficPolicy:
    loadBalancer:
      simple: ROUND_ROBIN
  - name: v3
    labels:
      version: v3
```

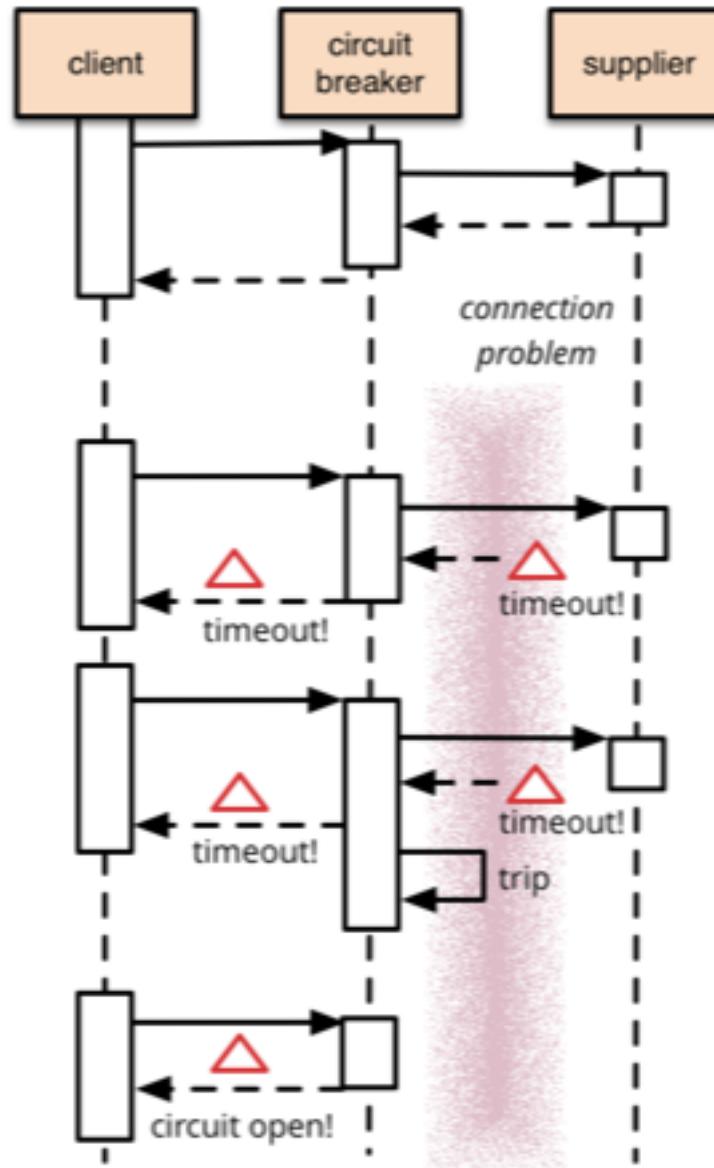
- trafficPolicy (LB RANDOM)
- trafficPolicy (LB Round\_Robin)

# Circuit breakers

critical component of distributed systems,  
which failing quickly, applying back  
pressure to downstream services

# Circuit breakers

- External service calls (HTTP timeout)
- Cascading failures



# Circuit breakers

- External service calls  
(HTTP timeout)
- Cascading failures

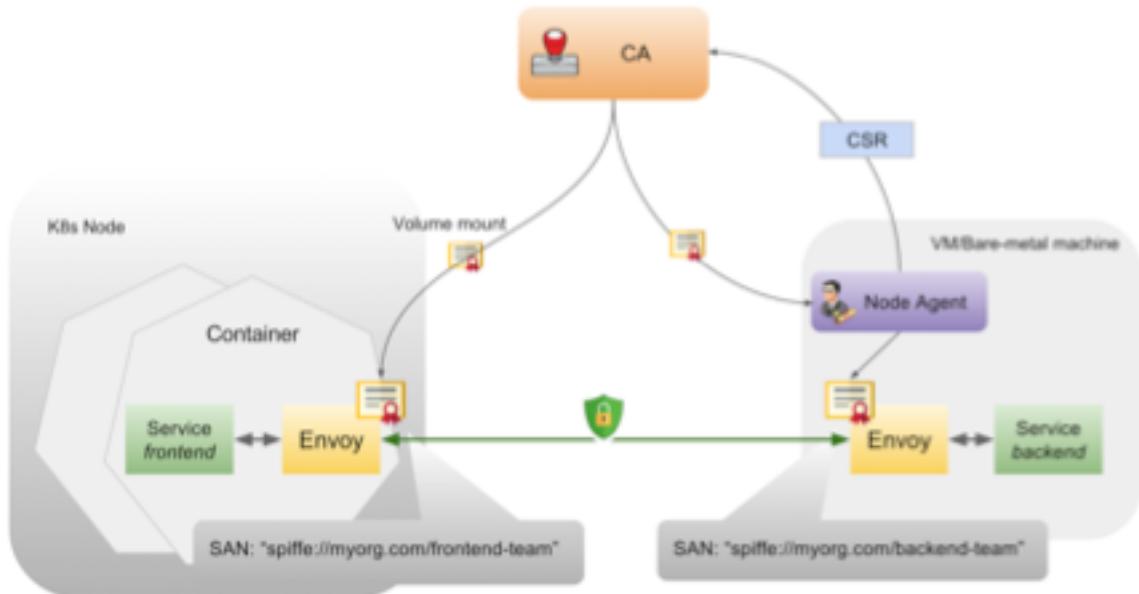
```
apiVersion: networking.istio.io/v1alpha3
kind: DestinationRule
metadata:
  name: reviews
spec:
  host: reviews
  subsets:
  - name: v1
    labels:
      version: v1
    trafficPolicy:
      connectionPool:
        tcp:
          maxConnections: 100
```

Lab: Dynamic traffic routing  
Lab: Handling failures



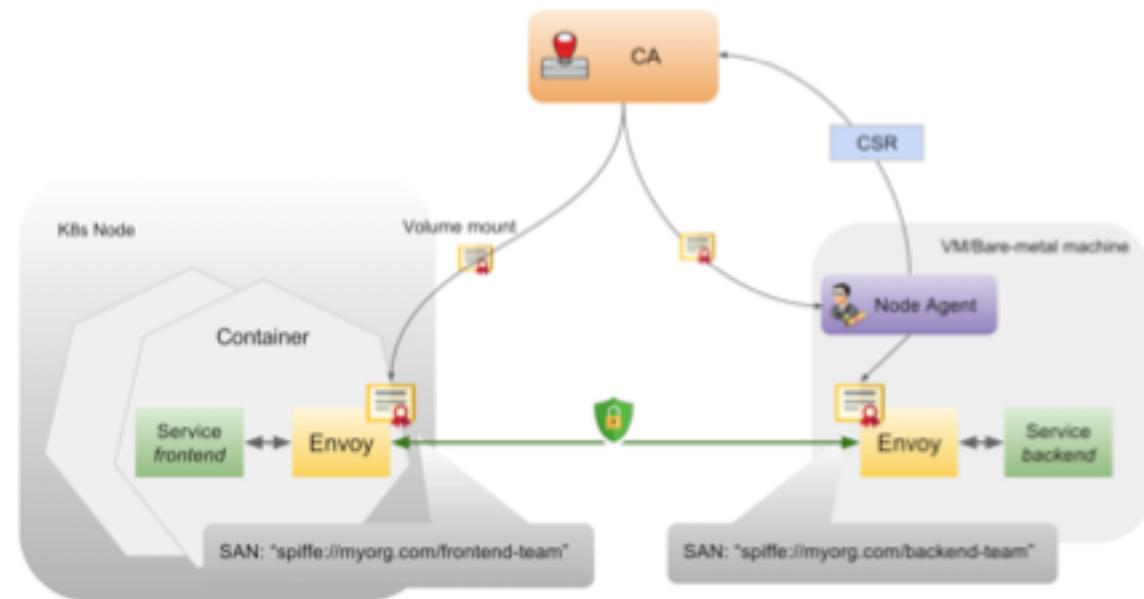
# Istio-Auth/Citadel

- Identity
- Key management
- communication security
- Leverages secret volumes to deliver keys/certs from Istio CA to Kubernetes containers.
- Optional mutual TLS encryption



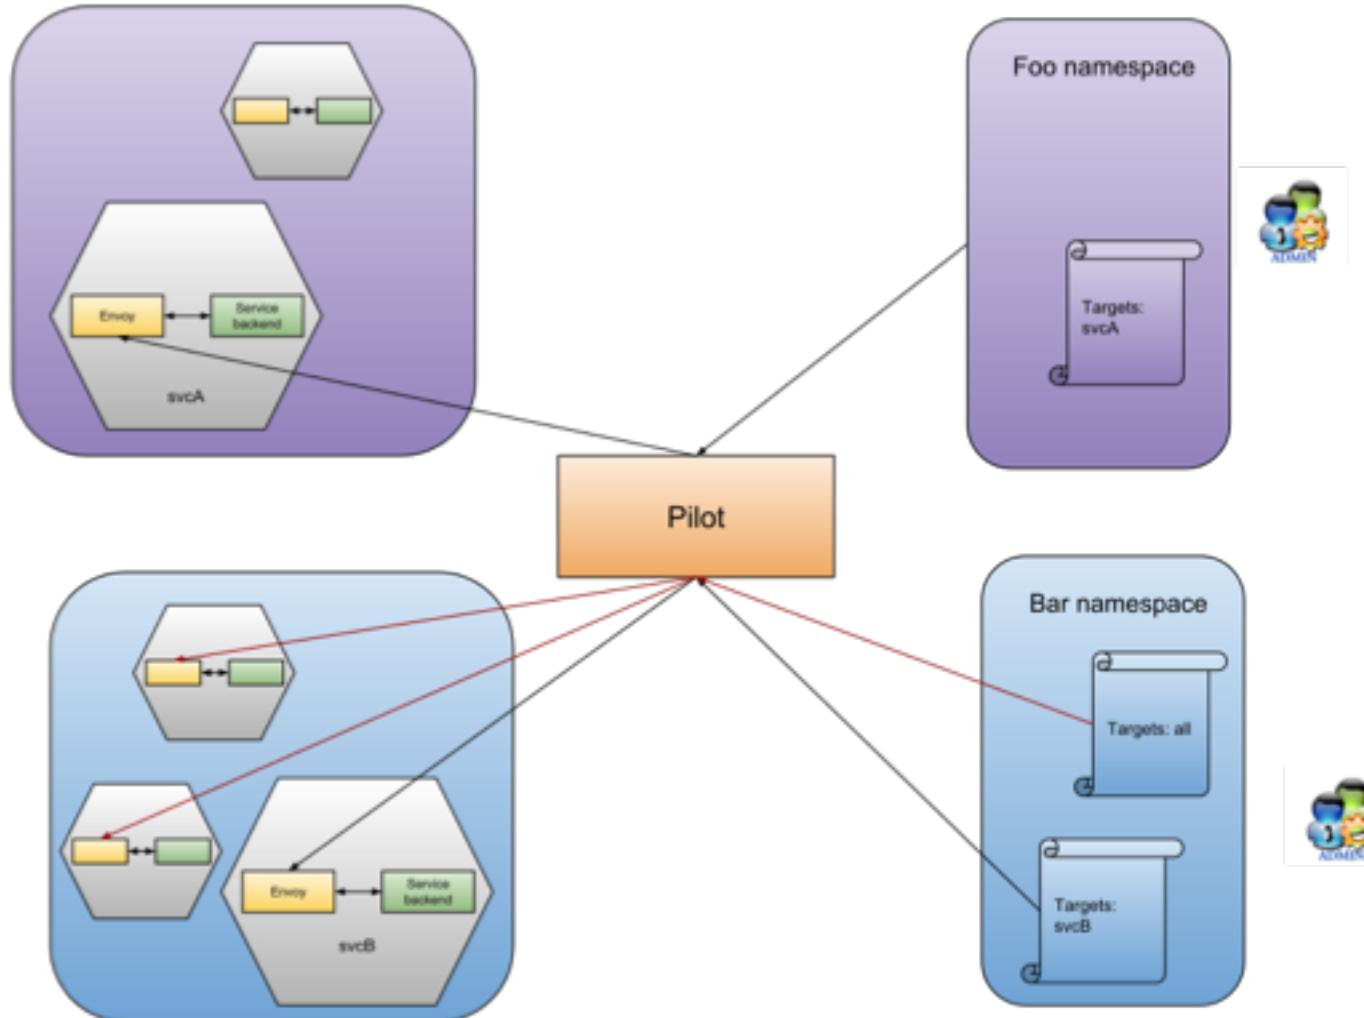
# Istio-Auth/Citadel

- Automate key & cert management process:
  - generate key/cert pair for each service account
  - distribute key/cert pair to each pod
  - rotate keys/certs periodically
  - revoke key/cert pair when needed



# Authentication Policy

- Istio authentication policy is composed of two parts:
- Peer: verifies the party, the direct client, that makes the connection. The common authentication mechanism for this is mutual TLS.
- Origin: verifies the party, the original client, that makes the request (e.g end-users, devices etc). JWT is the only supported mechanism for origin authentication at the moment.



# Anatomy of policy

```
targets:  
  - name: product-page  
  - name: reviews  
    ports:  
      - number: 9000
```

- Target selectors
  - determine which hosts/ports policies applied to.

# Enable mTLS

```
apiVersion: "authentication.istio.io/v1alpha1"
kind: "Policy"
metadata:
  name: "example-2"
spec:
  targets:
  - name: httpbin
  peers:
  - mtls:
```

- Enable mTLS

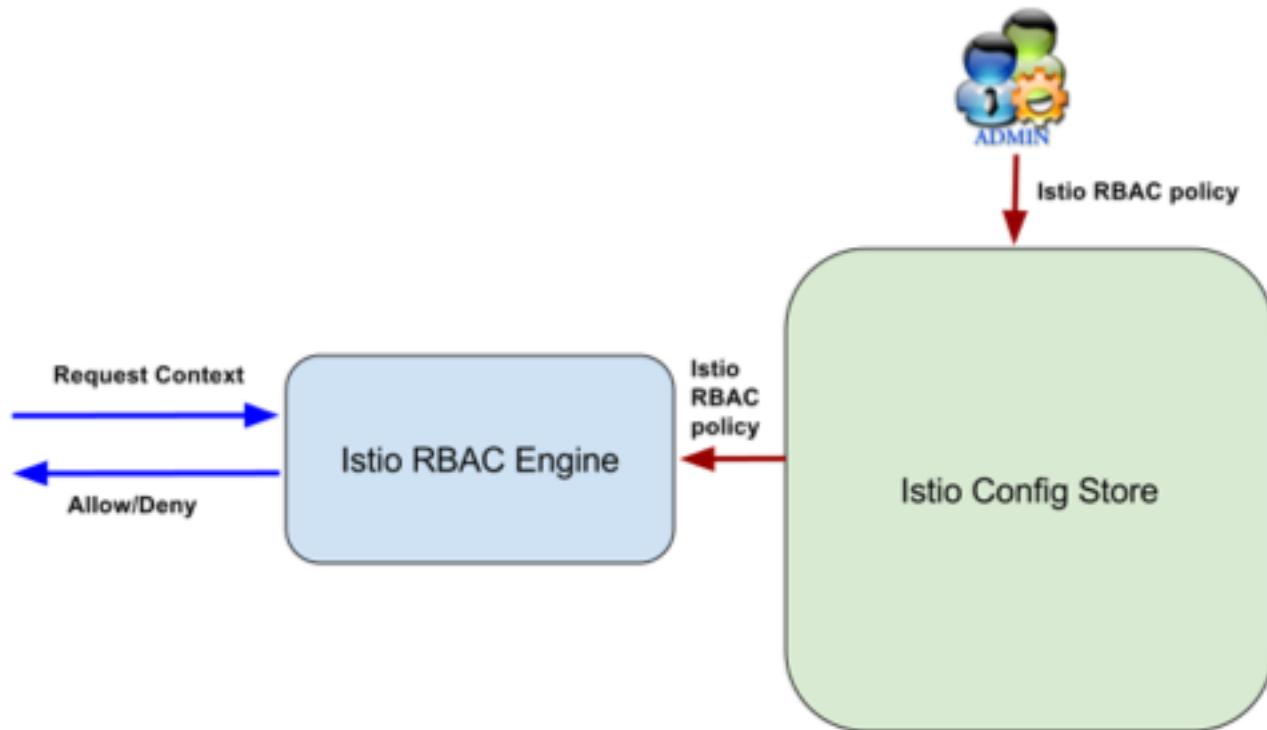
# mTLS host

```
apiVersion: "networking.istio.io/v1alpha3"
kind: "DestinationRule"
metadata:
  name: "example-2"
spec:
  host: "httpbin.bar.svc.cluster.local"
  trafficPolicy:
    tls:
      mode: ISTIO_MUTUAL
```

- Enable mTLS on httpbin host

# Istio RBAC

- More granular
- Service to service
- end user to service



# Istio RBAC

```
apiVersion: "config.istio.io/v1alpha2"
kind: authorization
metadata:
  name: requestcontext
  namespace: istio-system
spec:
  subject:
    user: source.user | ""
    groups: ""
    properties:
      service: source.service | ""
      namespace: source.namespace | ""
  action:
    namespace: destination.namespace | ""
    service: destination.service | ""
    method: request.method | ""
    path: request.path | ""
    properties:
      version: request.headers["version"] | ""
```

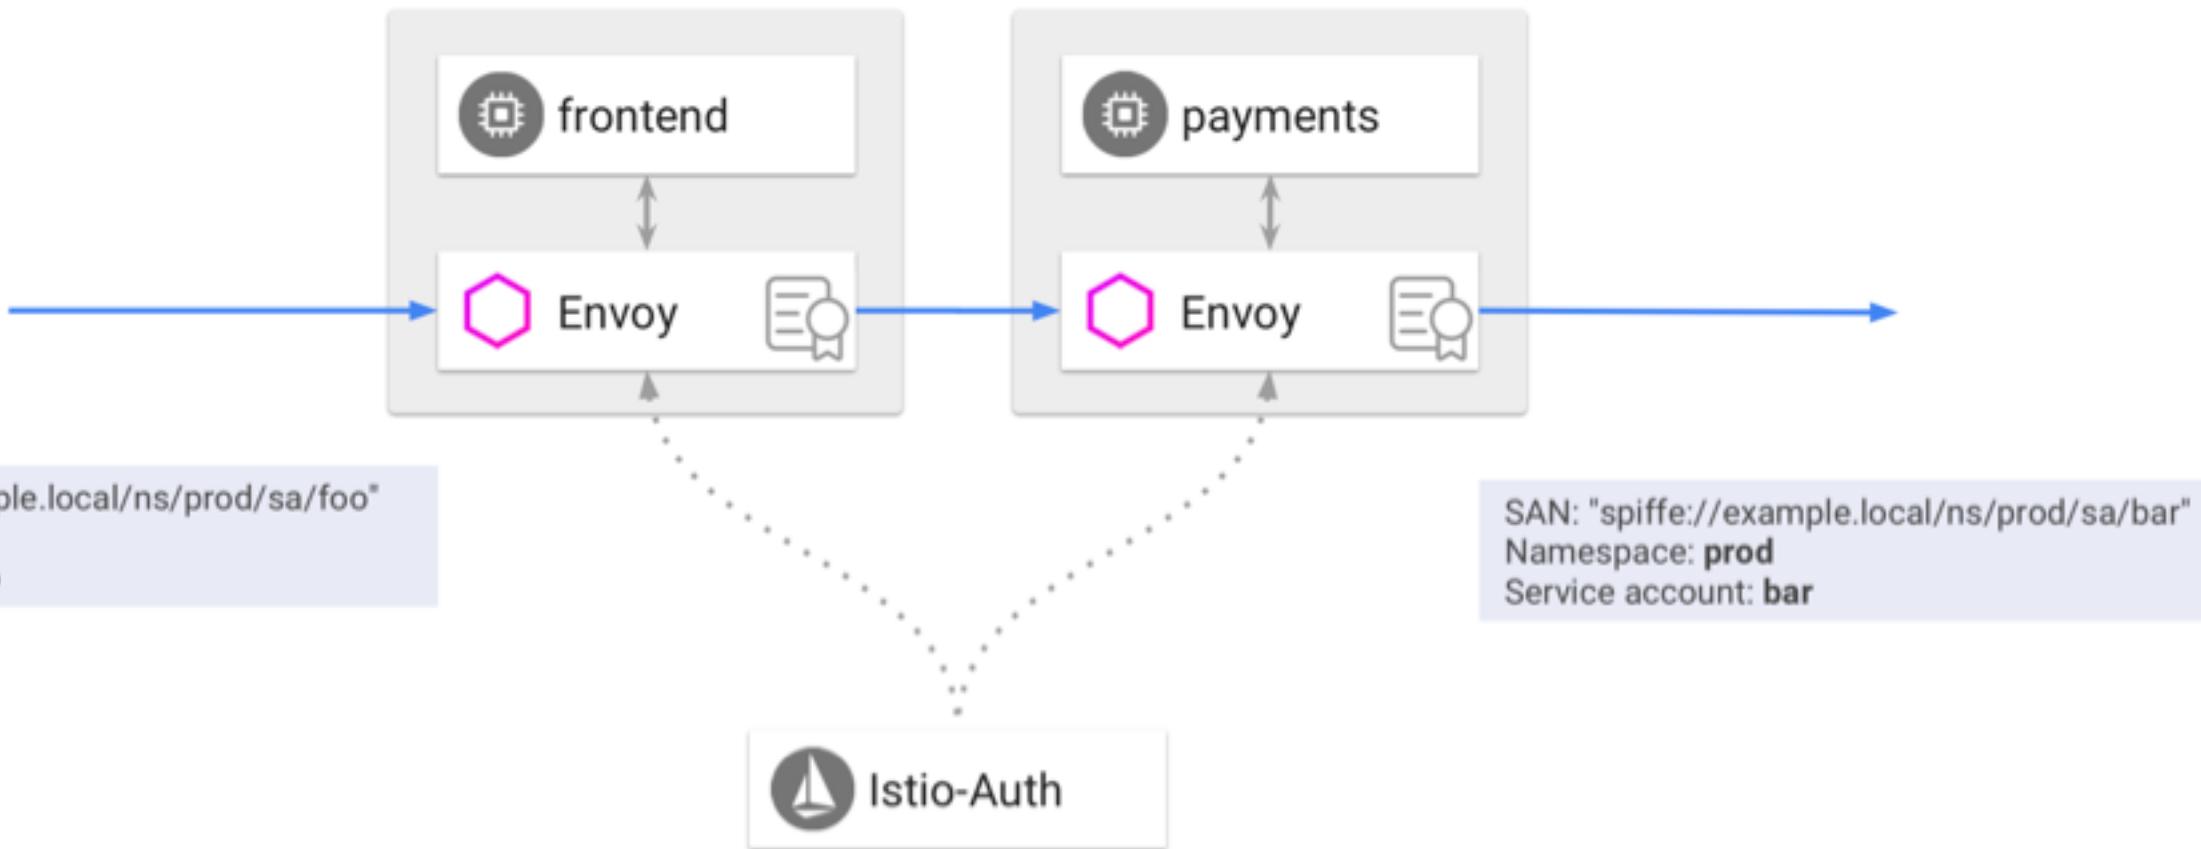
# Istio RBAC

```
apiVersion: "config.istio.io/v1alpha2"
kind: ServiceRole
metadata:
  name: products-viewer
  namespace: default
spec:
  rules:
  - services: ["products.default.svc.cluster.local"]
    methods: ["GET", "HEAD"]
```

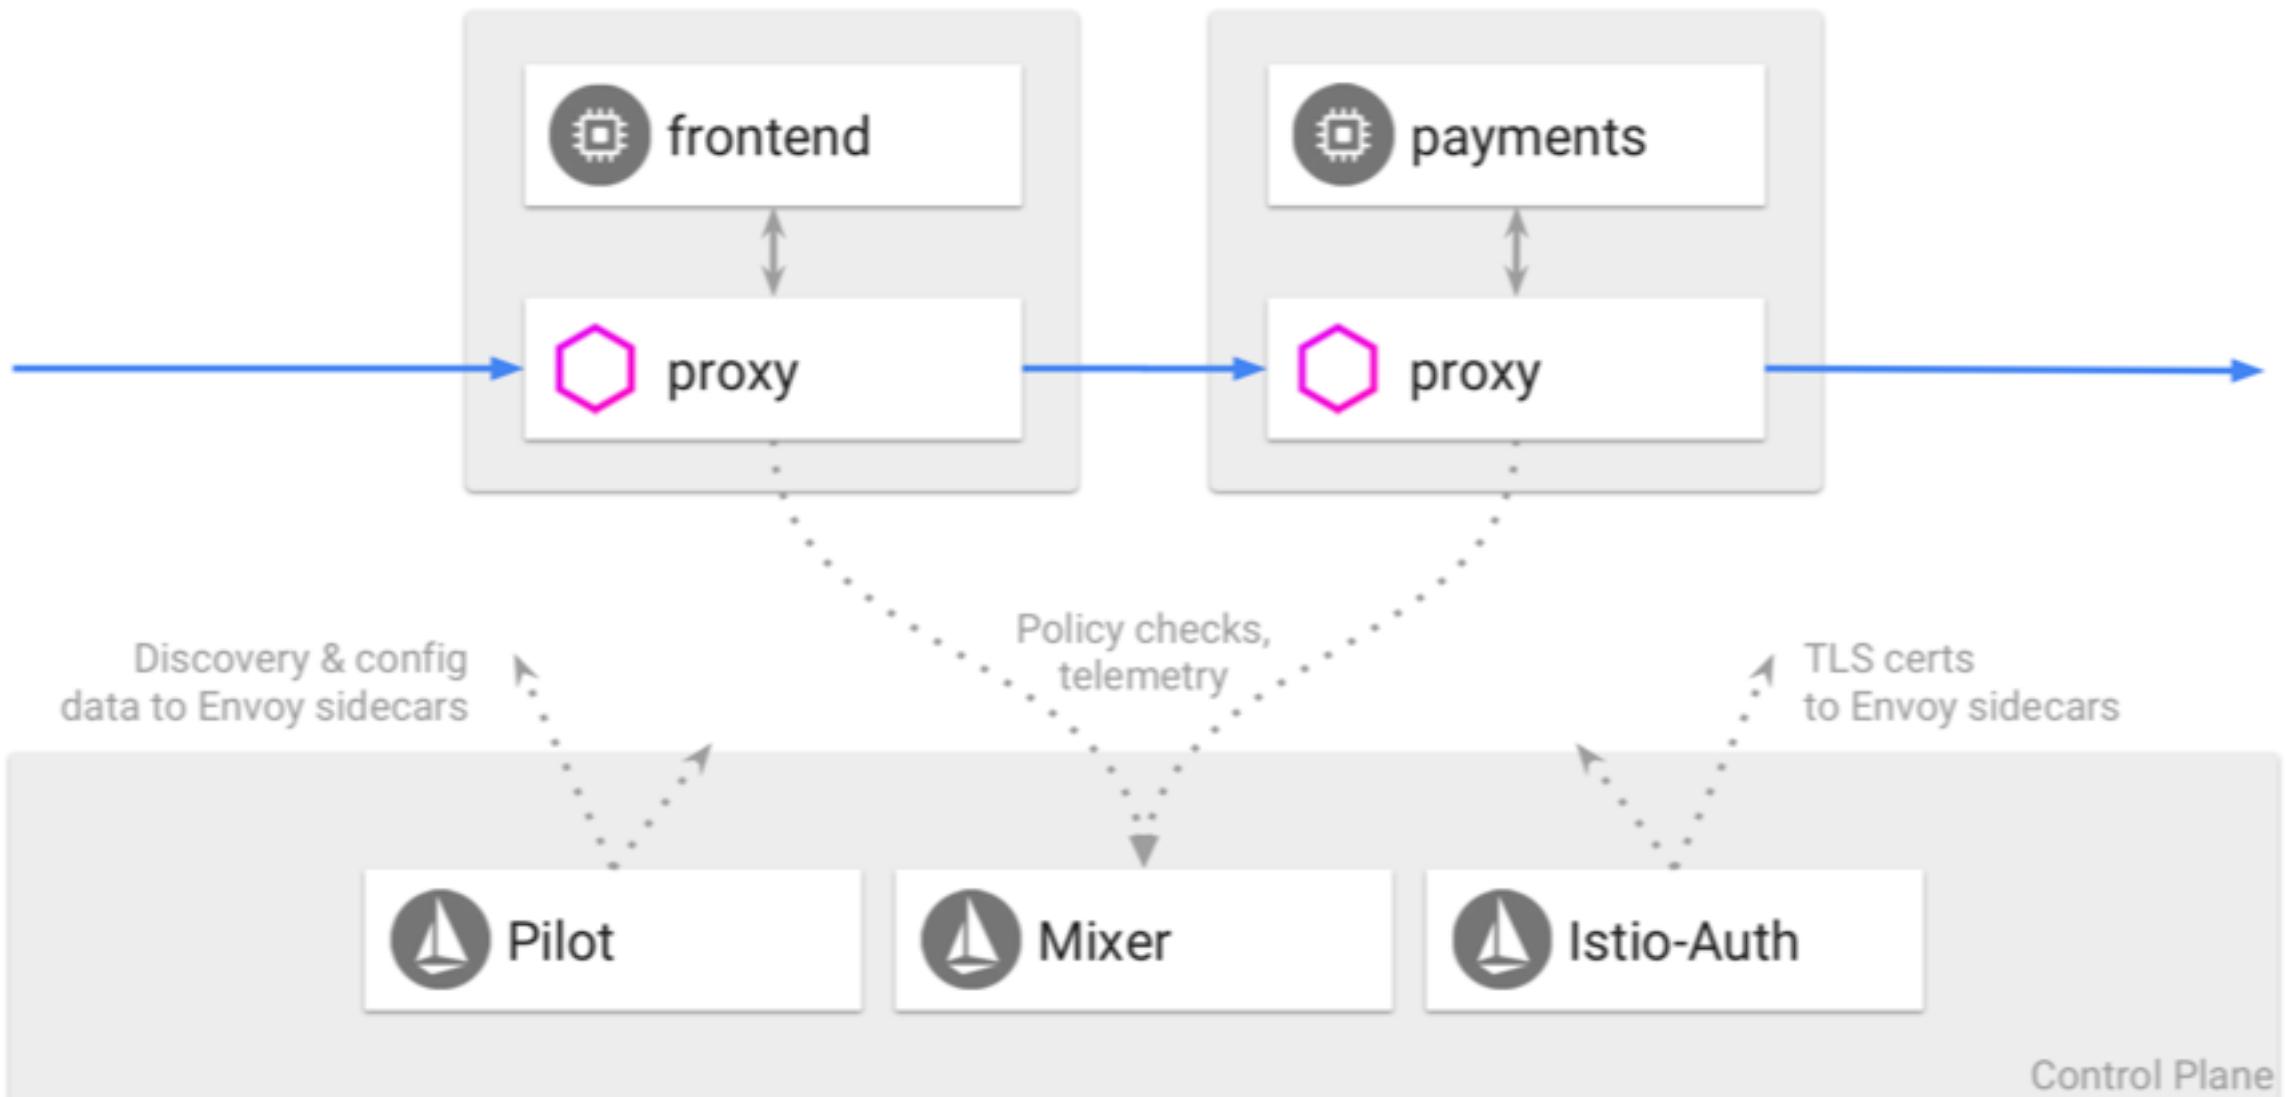
# Istio RBAC

```
apiVersion: "config.istio.io/v1alpha2"
kind: ServiceRole
metadata:
  name: products-viewer-version
  namespace: default
spec:
  rules:
    - services: ["products.default.svc.cluster.local"]
      methods: ["GET", "HEAD"]
      constraints:
        - key: "version"
          values: ["v1", "v2"]
```

# Security



# Security



# Service Entries

- Istio blocks all outgoing external requests by default
- ServiceEntry required to allow connections to outside services
- mutual TLS authentication disabled
- Policy enforcement performed on client-side

# ServiceEntry

```
apiVersion: networking.istio.io/v1alpha3
kind: ServiceEntry
metadata:
|   name: foo-ext-svc
spec:
  hosts:
  - *.foo.com
  ports:
  - number: 80
    name: http
    protocol: HTTP
  - number: 443
    name: https
    protocol: HTTPS
```

- Allow outgoing on HTTP & HTTPS to \*.foo.com

# VirtualService & ServiceEntry

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: bar-foo-ext-svc
spec:
  hosts:
    - bar.foo.com
  http:
    - route:
        - destination:
            host: bar.foo.com
        timeout: 10s
```

- Work with ServiceEntry to set a 10s timeout on external service

# Gateways

- Describes load balancer at the edge of mesh
  - exposed ports
  - protocol
  - TLS

# Gateway

```
apiVersion: networking.istio.io/v1alpha3
kind: Gateway
metadata:
  name: my-gateway
spec:
  selector:
    app: my-gatweway-controller
  servers:
    - port:
        number: 80
        name: http
        protocol: HTTP
        hosts:
          - uk.bookinfo.com
          - eu.bookinfo.com
        tls:
          httpsRedirect: true # sends 302 redirect for http requests
    - port:
        number: 443
        name: https
        protocol: HTTPS
```

# Gateway

```
hosts:
- uk.bookinfo.com
- eu.bookinfo.com
tls:
  mode: SIMPLE #enables HTTPS on this port
  serverCertificate: /etc/certs/servercert.pem
  privateKey: /etc/certs/privatekey.pem
- port:
  number: 9080
  name: http-wildcard
  protocol: HTTP
hosts:
- "*"
- port:
  number: 2379 # to expose internal service via external port 2379
  name: mongo
  protocol: MONGO
hosts:
- "*"
```

# Gateway/VirtualService

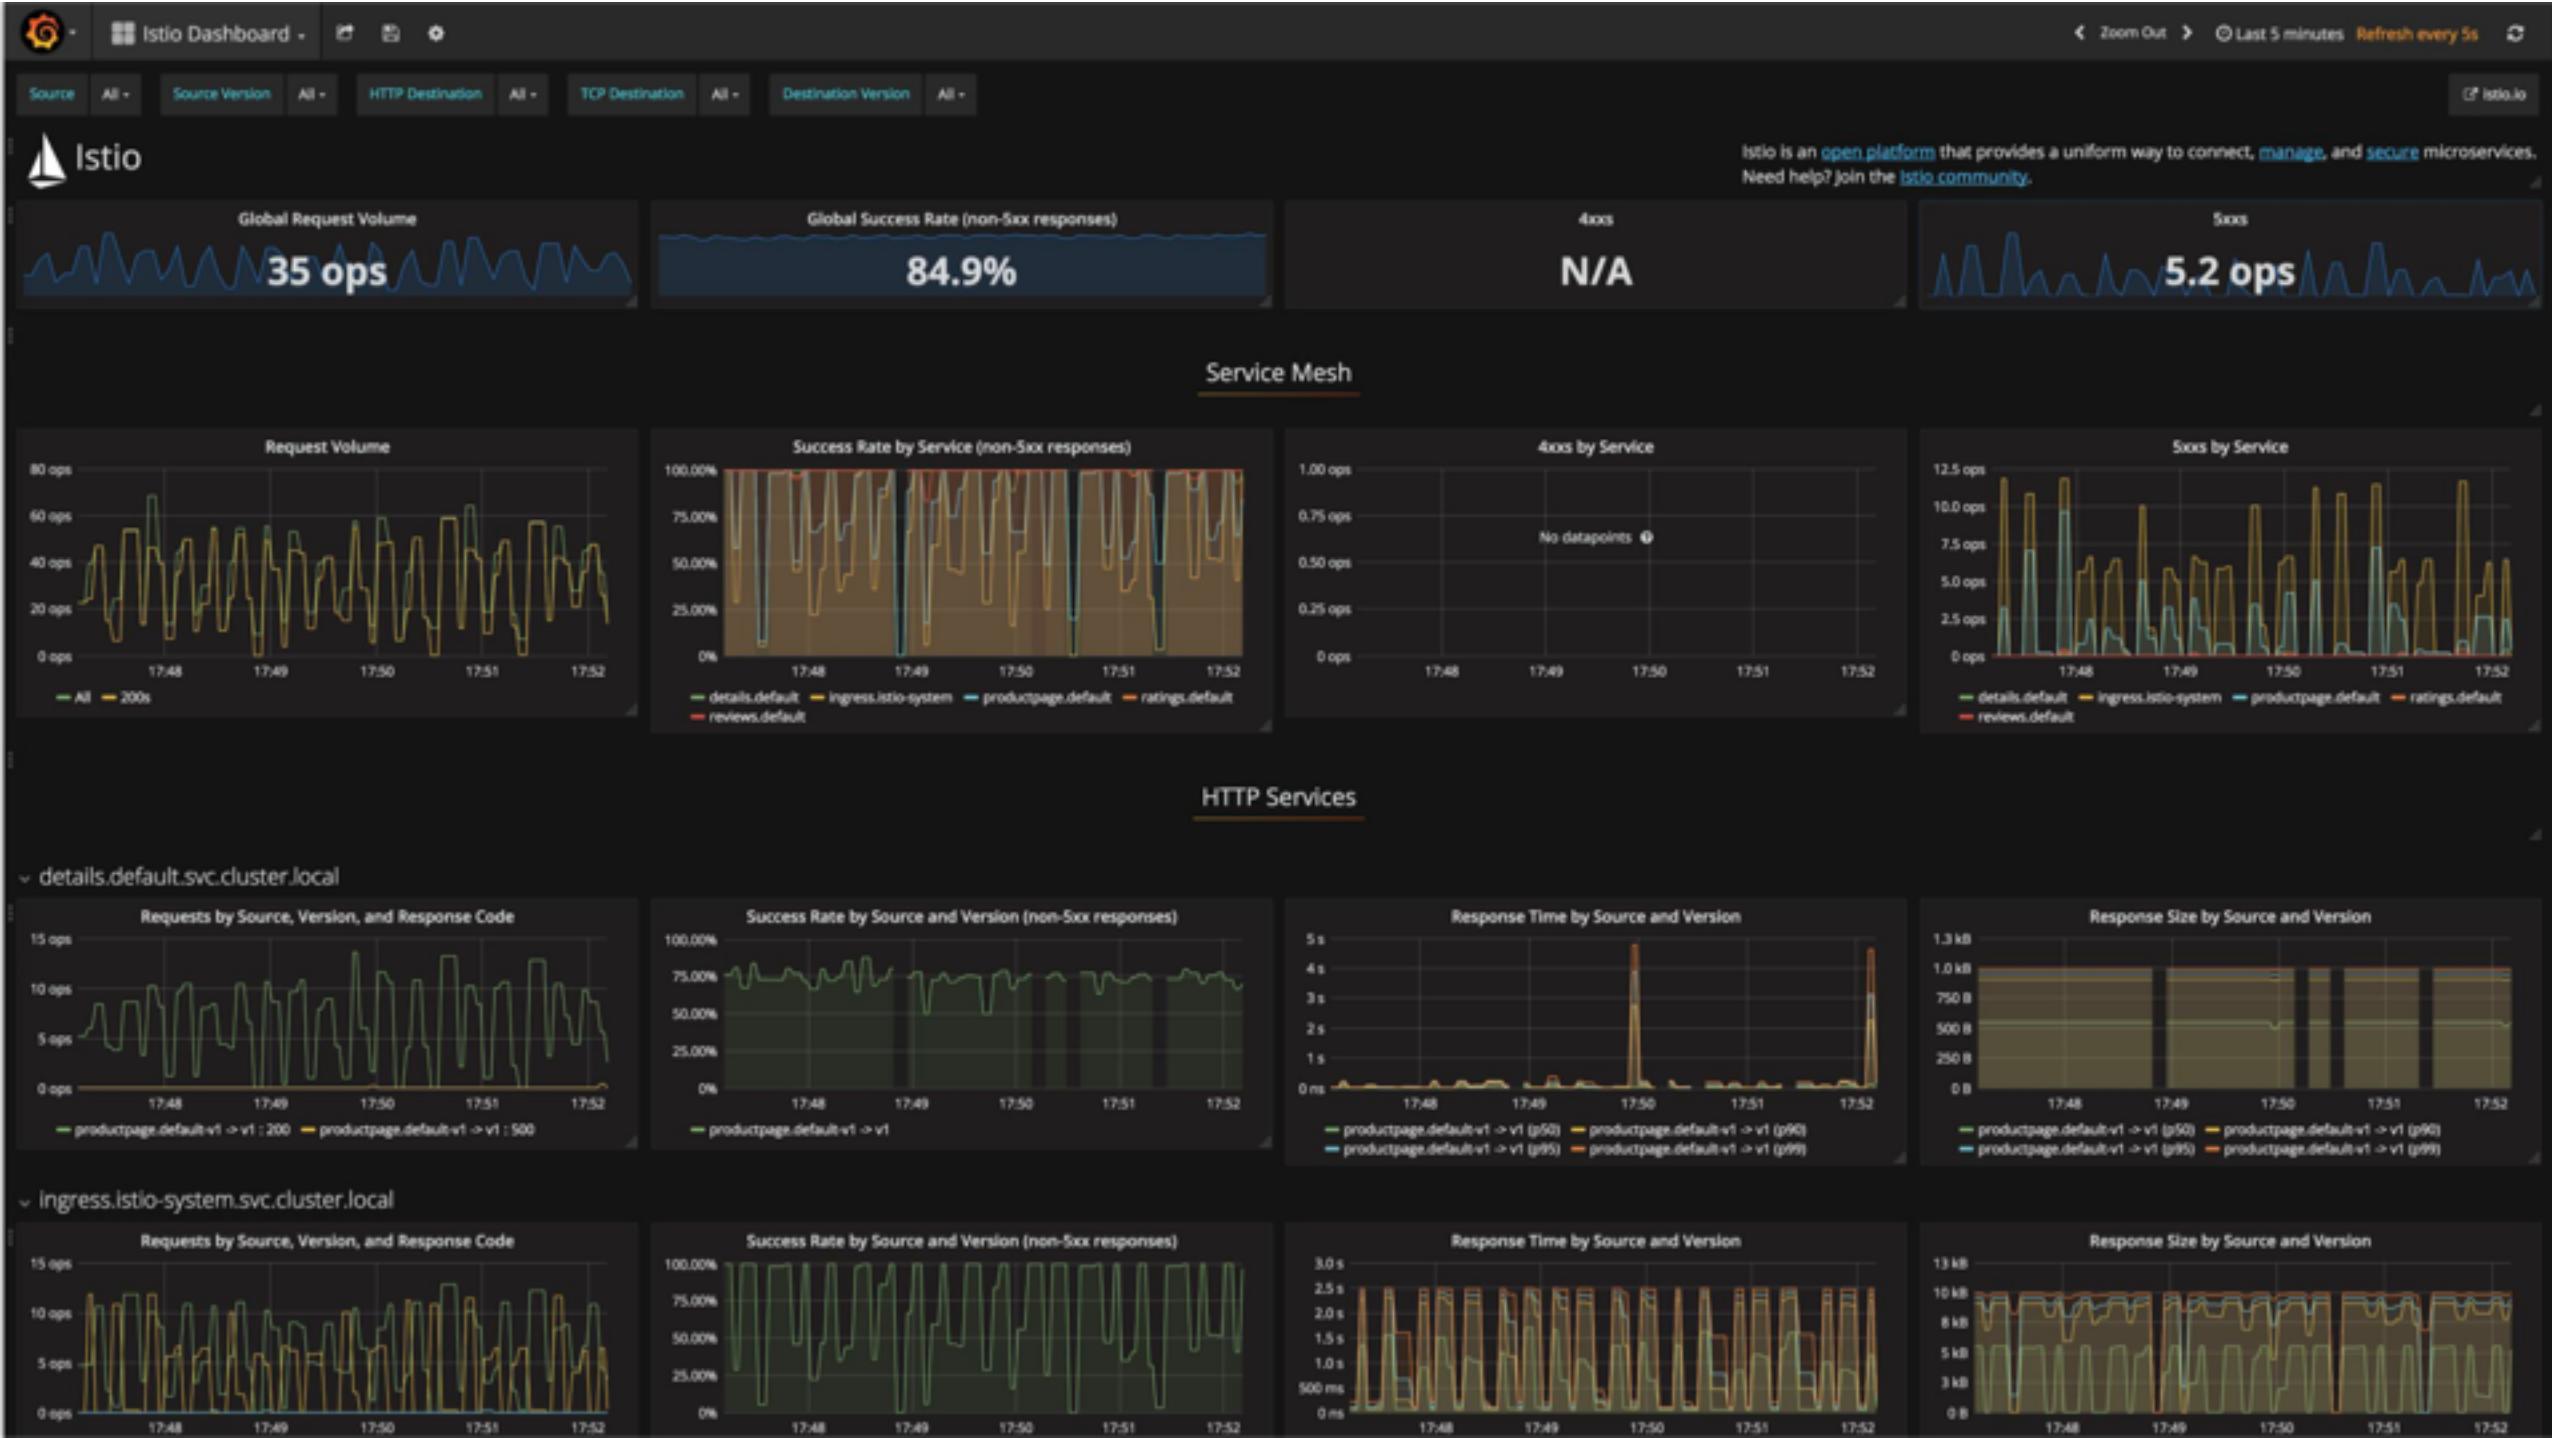
```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
| name: bookinfo
spec:
hosts:
| - bookinfo.com
gateways:
| - bookinfo-gateway # <---- bind to gateway
http:
| - match:
| | - uri:
| || prefix: /reviews
route:
```

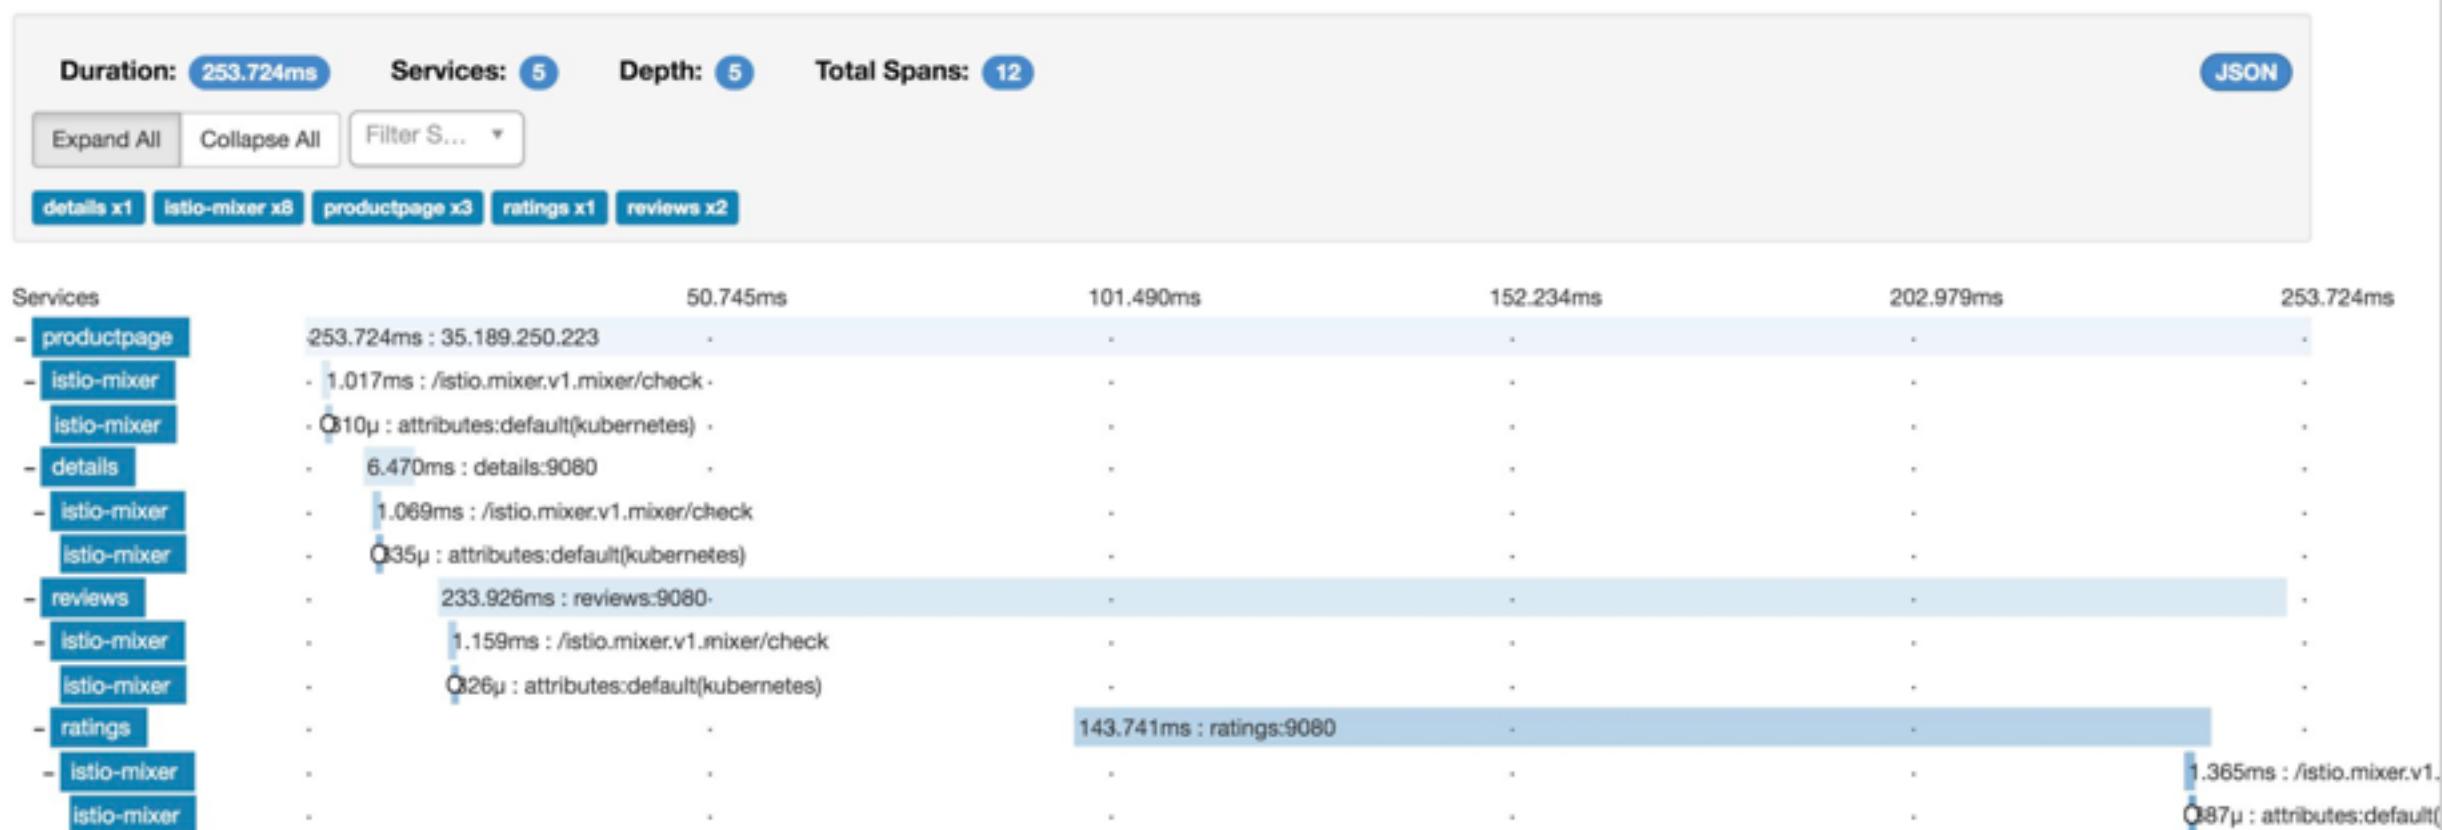
Lab: Security



# Visibility

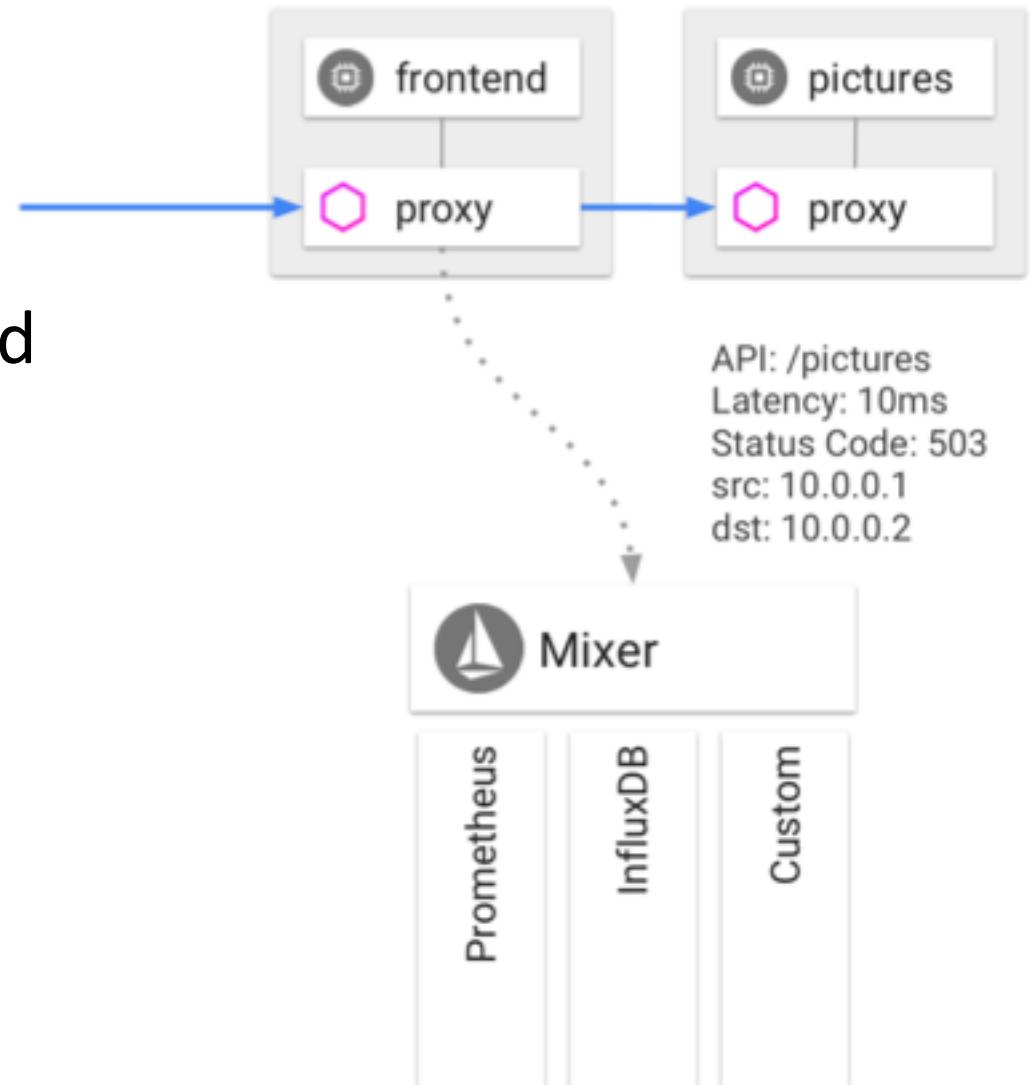
- Metrics without code changes
- Consistent metrics for every app deployed
- Trace flow of requests across services
- Portable across metric backend providers





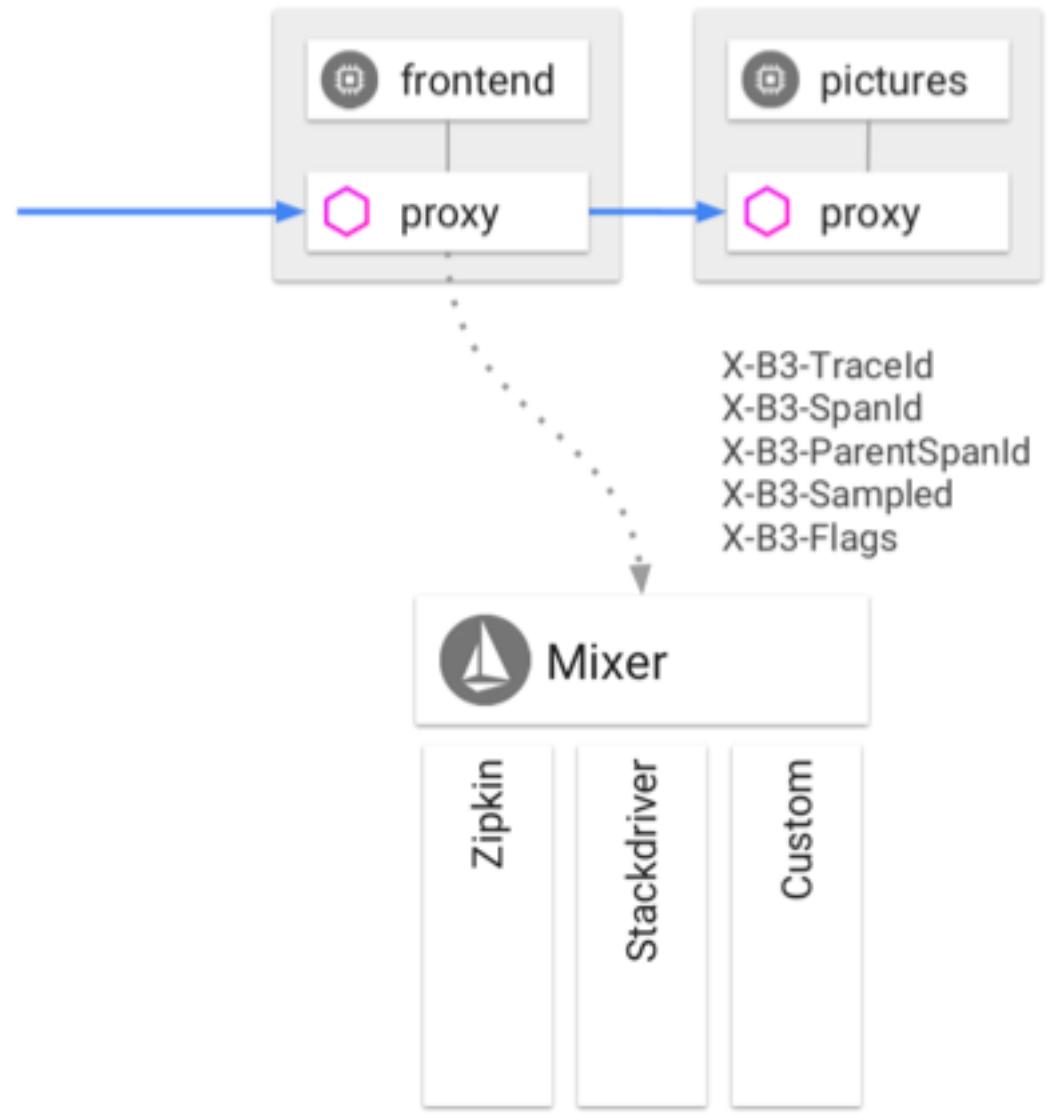
# Metrics flow

- Mixer collects metrics emitted by Envoy proxies
- Adapters in Mixer normalize and forward to monitoring backends
- Metrics backend can be swapped at runtime



# Tracing

- Apps do not have to deal with generating spans or correlating causality
- Envoy proxies generate spans
  - Apps do need to send some context heads on outbound calls
- Envoy send traces to Mixer
  - Adapters send to backends



# Tracing - headers

```
def getForwardHeaders(request):
    headers = {}

    user_cookie = request.cookies.get("user")
    if user_cookie:
        headers['Cookie'] = 'user=' + user_cookie

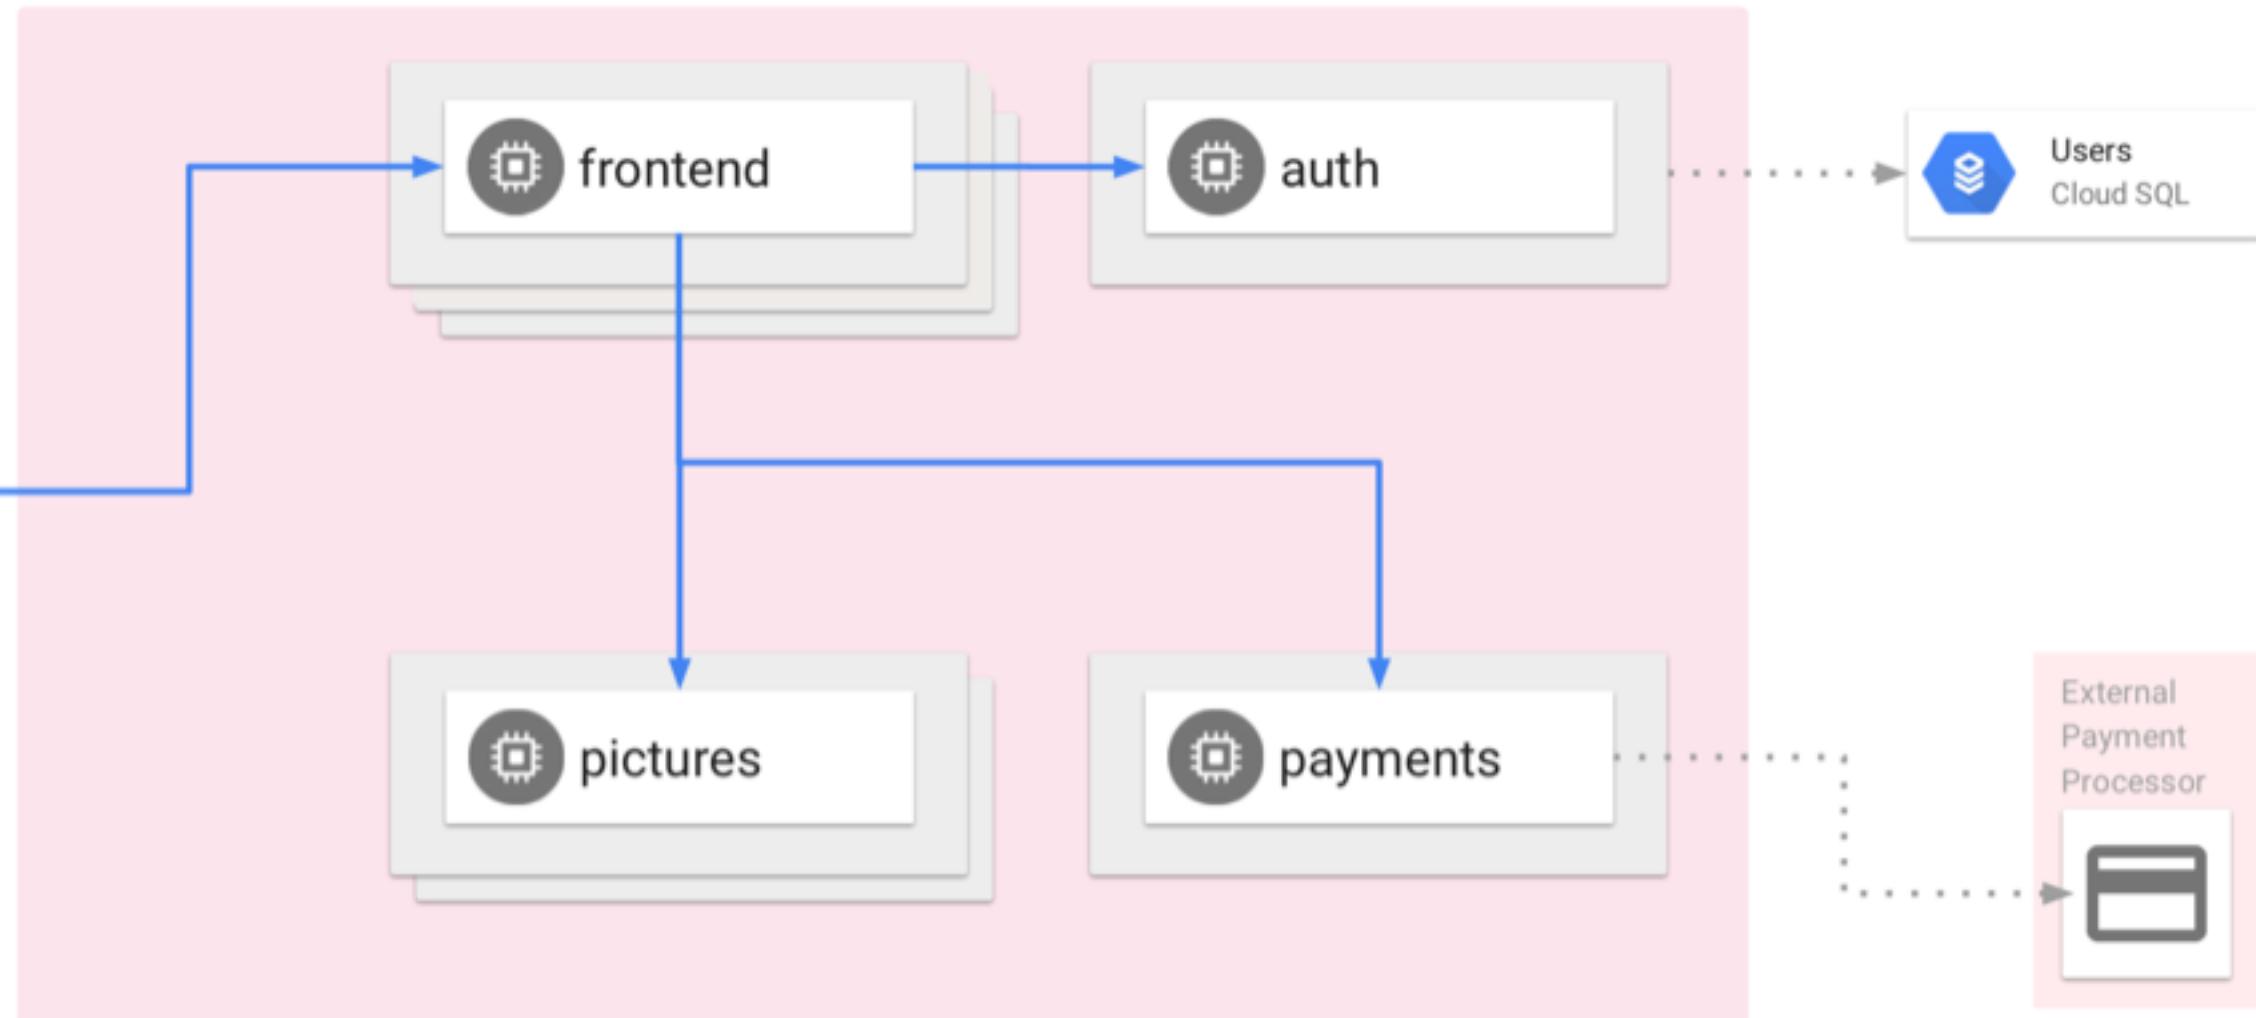
    incoming_headers = [ 'x-request-id',
                         'x-b3-traceid',
                         'x-b3-spanid',
                         'x-b3-parentspanid',
                         'x-b3-sampled',
                         'x-b3-flags',
                         'x-ot-span-context'
    ]

    for ihdr in incoming_headers:
        val = request.headers.get(ihdr)
        if val is not None:
            headers[ihdr] = val
```

# Tracing - headers

```
@GET  
@Path("/reviews")  
public Response bookReviews(@CookieParam("user") Cookie user,  
                           @HeaderParam("x-request-id") String xreq,  
                           @HeaderParam("x-b3-traceid") String xtraceid,  
                           @HeaderParam("x-b3-spanid") String xspanid,  
                           @HeaderParam("x-b3-parentspanid") String xparentspanid,  
                           @HeaderParam("x-b3-sampled") String xsampled,  
                           @HeaderParam("x-b3-flags") String xflags,  
                           @HeaderParam("x-ot-span-context") String xotspan) {  
  
    String r1 = "";  
    String r2 = "";  
  
    if(ratings_enabled){  
        JsonObject ratings = getRatings(user, xreq, xtraceid, xspanid, xparentspanid, xsampled, xflags, xotspan)
```

# Ultimately, it's just this



# Lab: Monitoring



